

Dexterous Manipulation, Single-Hand and Bimanual

Machines, simulators, and how policies are trained

Compiled 2026-09-18. Every claim traces to a note in `papers/notes/`, every note to a parsed source in `papers/md/` or `code/md/`, and every source to a hash or commit in `corpus/manifest.json`. The method is in Appendix A.

Abstract

A hand is dexterous when it can change an object's pose without putting the object down. This survey covers that problem for one hand and for two, across hands, simulators and the contact models underneath them, training methods and evaluation. The corpus holds 221 bibliography entries, 218 of which carry a structured row read from a note: 112 method papers, 33 hands, 15 simulators, 15 datasets, 14 benchmarks, 14 surveys, 8 tactile sensors and 7 evaluation protocols.

Three findings are measured rather than asserted. Papers disagree with their own released code. Sixty-two method rows released code that could be read against the paper, 38 of those record a discrepancy, and ten are contradictions where the shipped code states a different objective from the published one. `physhoi_2023` is the sharpest case. Its `compute_humanoid_reward` sets the object rotation errors to zero while the reward table in the paper weights them at 0.1 and 0.01, and its position-only success criterion could not have detected that. Nobody measures the quantity most specific to a hand. Eleven of the 96 method rows whose notes settle the question address interpenetration at all, and none reports it for its own policy's rollouts. IsaacGymEnvs already computes that depth and gates a policy update on it. Hardware and published work have come apart. 18 of the 33 hand rows appear in no method row, and 7 of those can be bought or built today.

This survey re-runs no method. It ranks nothing and publishes no leaderboard. On penetration it supplies a measurement method and a count, not a threshold. Every coverage statistic here is a floor over what this extraction captured, not a rate over what the literature reported.

Contents

1. Introduction
2. A taxonomy of the problem
3. Hands and who makes them
4. Simulators and the physics underneath
5. How policies are trained
6. Bimanual dexterous manipulation
7. Evaluation: how we would compare these methods
8. Gaps
9. Conclusion
 - 2.1 Task families
 - 2.2 What makes dexterous control hard
 - 2.3 Single hand versus two
 - Table 1. Task families against the properties that define success
 - 3.1 The design axes
 - 3.2 The hands the research literature actually runs on
 - 3.3 Open hardware and the collapse in cost
 - 3.4 Announced and unreleased hands
 - 3.5 Tactile sensing as part of the hand
 - 3.6 What is sold against what is published on
 - 4.1 What a dexterous simulation must get right
 - 4.2 Contact models and solvers, engine by engine
 - 4.3 The GPU-parallel turn
 - 4.4 Throughput, and why the reported numbers do not compare
 - 4.5 Tactile simulation and what it is calibrated against
 - 4.6 The sim-to-real gap for hands
 - 5.1 The design space
 - 5.2 Reinforcement learning
 - 5.3 Learning from human data
 - 5.4 Tracking a human reference with physics
 - 5.5 Generalist and vision-language-action policies
 - 5.6 Model-based control as the non-learning baseline
 - 5.7 Hybrids, and where the field has converged
 - 5.8 What the released code says
 - 5.9 The master table
 - 6.1 Why two hands is not twice one hand
 - 6.2 Coordination architectures
 - 6.3 Benchmarks and datasets for two hands
 - 6.4 Methods on the same axes
 - 6.5 Handover and in-hand transfer
 - 6.6 What transfers from single-hand work
 - 7.1 What the field reports, and why the numbers do not compare
 - 7.2 The axes that matter
 - 7.3 Physical plausibility as a first-class metric
 - 7.4 Statistics
 - 7.5 A proposed protocol
 - 7.6 Table 9, an empty results matrix
 - 7.7 What would have to be true
 - 8.1 What opening the repositories showed
 - 8.2 Nobody records interpenetration for their own policy's rollouts
 - 8.3 The evaluation-methodology literature contains no dexterous hand
 - 8.4 Generalist policies run on a fraction of the degrees of freedom
 - 8.5 The hands that can be bought go unused
 - 8.6 Bimanual work runs on one coordination architecture, and it has been ablated once
 - 8.7 Human data does not port across hands, and the map is usually unstated

A survey of 221 bibliography entries, 218 of which carry a structured row read from a note: 112 method papers, 33 hands, 15 simulators, 15 datasets, 14 benchmarks, 14 surveys, 8 tactile sensors and 7 evaluation protocols. "Method row" throughout means one of the 112, and every headline count here has one of those eight classes as its denominator.

1. Introduction

A hand is dexterous when it can change an object's pose without putting the object down. That property, and not the finger count, is what separates a hand from a gripper. [bicchi_grasping_chapter_2001](#) draws the same line in Sec. 1.1, between restraining an object and "manipulating objects with fingers, in contrast to manipulation with the robot arm". Restraint is a static question about whether the contacts prevent motion. In-hand manipulation is a dynamic question about whether contacts can be broken and remade while the object stays held. [an_dexil_survey_2025](#) puts the same idea in its abstract as the ability "to skillfully control, reorient, and manipulate objects through precise, coordinated finger movements and adaptive force modulation". Both definitions place the work in the fingers rather than in the arm.

The analytic theory answered the static question and stalled on the dynamic one. Form closure has a first-order test on the grasp matrix and known contact counts, four in the plane and seven in three dimensions for any polyhedron, per [bicchi_grasping_chapter_2001](#) Sec. 1.3.1. Force closure adds the wrench balance and the hand Jacobian. Neither delivers what a controller needs. The same chapter states in Sec. 1.5 that "force closure does not guarantee stability", and its Sec. 1.8 names the reason the theory could not be pushed further, which is that "the nonsmooth nature of grasp dynamics, because of the unilateral constraints on displacements and forces, has made a thorough analysis very difficult". That verdict is one architect's, on one chapter, and this survey did not survey the tradition it judges. Learned control did not solve those modelling problems. It went around them by sampling a simulator instead of solving a model, and by scoring a rollout instead of certifying a configuration. Of the 112 method papers in this corpus, 53 train with reinforcement learning and 36 run in Isaac Gym, against 6 on its successor Isaac Lab.

Three things this survey measured are worth stating before the reader commits to 27,000 words. The first is that papers disagree with their own released code. Sixty-two of the 112 method rows released code that could be read against the paper, 38 of those record a discrepancy, and ten are contradictions where the shipped code states a different objective from the published one. The sharpest case is [physhoi_2023](#). Its [compute_humanoid_reward](#) hardcodes the object rotation and rotation-velocity errors to zero, with the real computation commented out beside them, while the reward table in its own paper lists weights of 0.1 and 0.01 for exactly those terms on GRAB. Its position-only success criterion could not have caught that, and its headline 95.4 percent is cited as a baseline. Section 5.8 classifies all 38, and seven further accusations an earlier draft made were withdrawn under adversarial review and recorded beside the charge.

The second is that the quantity most specific to a hand is the one nobody records. Eleven of the 96 method rows whose notes settle the question address interpenetration at all, seven of the eleven do it outside a closed-loop policy in a grasp synthesiser, a trajectory optimiser or a contact model, and none reports a penetration number for its own trained policy's rollouts. The obstacle is not the engines. NVIDIA's own IsaacGymEnvs repository already computes a per-environment maximum interpenetration depth in Warp and gates the policy update on a 1 mm threshold. Section 7.3 has the file and the lines.

The third is that hardware and published work have come apart. 18 of the 33 hand rows in Tables 2 and 3 appear in no method row, and 7 of those can be bought today or built from published designs. Thirty-five method rows run on the Allegro, whose weight, joint torque, payload and price have no reachable source, because its product page returns HTTP 404 and everything Table 2 confirms about it comes from its ROS driver.

Fourteen corpus entries are themselves surveys or engine-comparison studies. Appendix D sets them on one set of columns in Table 10 and says what each covers. Four of the fourteen could not be obtained, or were fetched too late to read, and are entered as such. Two of the three things this survey adds are visible in that table as columns nobody else fills. The first is Table 4, which takes the engines [nine_physics_engines_review_2024](#) scored on documentation and usability and adds contact model, solver, iteration count, default timestep and penetration exposure, conditioned on what a hand does to a solver. [physics_engine_comparison_2015](#) and [contact_models_comparison_2023](#) do measure engines, on five engines and on four contact formulations, and neither surveys the field those engines are used in. The second is two hands on one object as its own problem, which none of the four field surveys gives more than a subsection.

The third is the penetration measurement, on an axis a predecessor had already named. [zhao_dexhand_survey_2026](#) states in its Sec. IV-C that the field assesses "at least two layers of performance: the quality of grasps or poses prior to execution,

and the performance of policies or generators during downstream execution", and names "physical plausibility, including penetration" as the most common criterion for the first layer. That is the reference-versus-rollout split, named before this survey measured it, and the idea that contact quality belongs on the evaluation axis is Zhao's. What Sec. IV-C does not give is a measurement method or a count, and this survey supplies those two. It does not supply a threshold. The 2 mm figure the field uses is taken from `toporetarget_2026` with no independent justification, the captured human grasps in `grab_2020` sit above it at 3.25 mm, and Section 7.7 states plainly that 2 mm is a simulator convention rather than a physical bound.

The scope is single-hand and bimanual multi-fingered manipulation. Parallel-jaw manipulation enters only as a comparison, which matters because 12 of the 112 method rows name a parallel-jaw gripper among their own embodiments. Locomotion is excluded. Prosthetics are excluded except where a hand crosses over into robot use, as the Psyonic Ability Hand does. The corpus holds 221 bibliography entries, of which 218 carry a structured row read from a note. Those 218 rows are 112 method papers, 33 hands, 15 simulators, 15 datasets, 14 benchmarks, 14 surveys, 8 tactile sensors and 7 evaluation protocols. "Method row" throughout means one of the 112, and every headline count in this survey has one of those eight classes as its denominator.


Two limits apply to every number here. This is a corpus of the learned era, which is a selection effect and not a judgement: 138 of the 221 entries are dated 2024 or later, 16 predate 2018, and two of the 112 method rows predate 2018, so any statement about a trend over time is a statement about 2022 onward. The analytic tradition is represented by one readable chapter rather than surveyed, and the planning line that took up the dynamic question directly, finger gaiting and rolling-contact manipulation and regrasp planning, is not in the corpus at all. Every coverage statistic is also a floor rather than a rate, because a value this survey failed to extract is indistinguishable from a value the paper never reported, and every such miss converts a reporting paper into a silent one.

Figure 1 puts the field on one page. Section 2 sets out the task families and what makes each hard. Section 3 covers hands, their makers, and the gap between what is sold and what is run. Section 4 covers simulators and contact models. Section 5 covers how policies are trained. Section 6 covers bimanual work as its own problem. Section 7 proposes an evaluation frame rather than a leaderboard, and is the longest section. Section 8 states the gaps as claims with their evidence. Section 9 says what to do about them, addressed to someone publishing, running experiments or buying a hand. Appendix A is the method, Appendix B and Appendix C are the full hand and reward extractions, and Appendix D is the comparison with the existing surveys.

Figure 1. The field on one page

Every node is a count over the 112 method papers in the corpus. Dashed nodes are reporting gaps, not choices.

Paradigm counts exceed 112 because a method may use several. Recomputed from corpus/rows at draw time.



The routes converge. The heaviest single path is no demonstrations, one hand, Isaac Gym, reinforcement learning, real robot. Whatever enters on the left, almost everything leaves through a GPU simulator and arrives at a real robot without contact quality ever being measured.

2. A taxonomy of the problem

2.1 Task families

The six families below are read off the rows rather than imposed on them, and Sec. 2.1's closing paragraph says what they miss. Grasping is the largest. Fifty-seven of the 112 method rows carry the grasp label, the labels are not exclusive, and one paper can sit in several families. Success is a lift that survives a hold, and the thresholds differ by more than an order of magnitude. `dexgraspvla_2025` requires the object "held 10 cm above the table for 20 s", while `omnigrasp_2024` requires it "held at least 0.5 s in simulation". What makes grasping hard at scale is the continuum of starting configurations. `unidexgrasp_pp_2023` states it in Sec. 4.3, that "we are dealing with an infinite number of tasks considering the initial object pose can change continuously".

Functional and tool use is second with 45 rows. It is the family where a stable grasp can still be the wrong answer. `dexterous_functional_grasping_2023` gives the case in Sec. 2.1, that "grabbing a hammer from the head or handle are both equally valid ways of using it", and only one of the two lets the tool be used. Success is defined against the tool's function, and the field has no shared way to state it.

In-hand reorientation has 31 rows and the most settled criteria, because the community inherited one number. `openai_dexterity_2018` declared the goal achieved below 0.4 rad of orientation error, `dextreme_2022` keeps 0.4 rad at test time against a 0.1 rad training tolerance, and `eureka_2023` counts consecutive successes at 0.1 rad. A shared tolerance is not a shared protocol, because the stopping rule differs. `visual_dexterity_2022` measures error "when the controller predicts it has reached the goal and stops", which makes the policy a judge of its own trial.

Tracking a human reference has 12 rows and inverts the problem. The target is a trajectory rather than an endpoint, so success is a per-frame error band. `maniptrans_2025` requires all four of 30 degrees of rotation, 3 cm of translation, 8 cm of joint error and 6 cm of fingertip error. The difficulty is that the reference came from a human hand and was never dynamically feasible for the robot. `dexmachina_2025` warns in App. B.4 that scoring by timesteps inside the thresholds makes "the results highly sensitive to the threshold values".

Bimanual coordination has 43 rows and handover has 10. Handover is the smallest family and the one whose failure has a single moment, because the giver must release only after the receiver has the object. `dynamic_handover_2023` reports a hit rate above its success rate and blames the gap in Sec. 5.4 on "occasional challenges encountered during the grasping phase of the catcher". Not every paper here learns both sides. In `dexterous_handover_2025` only the receiver is learned and the giver is a scripted arm.

Those six families do not cover the corpus. Twenty-four of the 112 method rows carry a label from outside them and eight carry no label from the six at all. Twenty are labelled `other`, mostly generalist policies evaluated on a task suite rather than on a dexterous task family, among them `pi0_2024`, `pi05_2025`, `pistar06_2025`, `openvla_2024` and `gemini_robotics_15_2025`. Three adjacent families are named here rather than absorbed. Locomanipulation is `humanplus_2024`, `omnih2o_2024` and `groot_n16_2025`, where the base is not fixed and the gravity argument below changes character. Piano playing is `robopianist_2023`, `rp1m_2024` and `pianomime_2024`, discussed in Section 6, where success is a per-timestep F1 against a MIDI score rather than an object pose. Deformable manipulation is `dexdeform_2023`, whose object carries its own state and its own physics. `ferrari_canny_1992` carries no task label because it is a grasp-quality measure rather than a task. Table 1's six rows are the families with enough papers to compare, and not a partition of the corpus.

2.2 What makes dexterous control hard

Contact is non-smooth, and that is a property of the problem rather than of any solver. `bicchi_grasping_chapter_2001` states in Sec. 1.2 that contact constraints are unilateral, and that losing a contact "involves an abrupt change of the structure of the model under consideration". Its Sec. 1.4 gives the sharper version, that a rod sliding on rough ground has configurations with no consistent solution and configurations with more than one. Simulators inherit the difficulty. The re-implementation

study [contact_models_comparison_2023](#) concludes in Sec. V that "there is no fully satisfactory approach at the moment, as all existing solutions compromise either accuracy, robustness, or efficiency". [pang_global_planning_2022](#) treats non-smoothness as the thing to remove, curating every contact pair in Sec. III-D so that contact points and normals change smoothly.

The hand has more actuators than the object has degrees of freedom and is still short of authority over it. The object moves only through contacts, and the contacts are what cannot be commanded. [bicchi_grasping_chapter_2001](#) Sec. 1.4 states that in multifingered grasps "the number of independent contact forces is much larger than the number of actuators. Thus, from a controllability standpoint, not all the contact forces are controllable".

Vision is occluded by the hand doing the work. [bai_unified_manip_survey_2025](#) names it in Sec. 4.3, that "occlusion hampers object tracking". [dexpoint_2022](#) keeps vision and adds imagined hand points to the point cloud. [rotating_without_seeing_2023](#) removes vision entirely and rotates objects from 16 binary touch sensors over the palm, links and fingertips.

Gravity direction changes the task rather than scaling it. [visual_dexterity_2022](#) states that reorientation with the hand below the object "is much easier", because "with a downward-facing hand, the hand must manipulate the object while simultaneously counteracting gravity". [anyrotate_2024](#) makes the direction a randomised variable by "randomly initializing hand orientations between episodes", and its Sec. 5.3 measures the cost, with performance dropping progressively from palm up and palm down, through base up and base down, to thumb up and thumb down.

2.3 Single hand versus two

Three counts describe two hands and they measure different things. Fifty-three of the 112 method rows record two hands on the robot, which is the [bimanual](#) field. Forty-three carry the [bimanual-coord](#) task label, the narrower claim that coordinating the hands is the task. Section 6 narrows again, to the 28 rows whose notes place a learned closed-loop controller on two multi-fingered hands, and its opening paragraph names every exclusion that takes the 53 down to the 28. That 28 is the denominator for every architecture count in this survey. Every bimanual claim here names which of the three it uses.

Four things genuinely change when the second hand arrives. Contact stays non-smooth, occlusion stays, and gravity stays the same problem.

Role asymmetry is the first. [asymdex_2024](#) assigns a dominant hand with full finger and wrist control and a facilitating hand with 6-DoF base pose only, so that "the facilitating hand repositions and reorients one object, while the dominant hand performs complex manipulations".

The second is that two hands on one object close a kinematic chain through the object. In [bicchi_grasping_chapter_2001](#) Sec. 1.4 the hand and object dynamics are separate and are "linked through the n rigid-body contact constraints". A second hand adds a second constraint set on the same object rather than a second independent problem. That is where physical plausibility fails first. [bimangrasp_2024](#) is the corpus row that measures it, rejecting a synthesised grasp when "total penetrations exceeded 1.5 mm" and reporting in Sec. IV-C that "penetration remains the primary cause of grasp failure".

The third is two-arm collision, which does not exist for one hand. [dydexhandover_2025](#) resets an episode when "any unintended arm contact with the environment or self-collisions" occurs. [bunny_visionpro_2024](#) pays for it in the controller, where adding a sphere-approximated self-collision cost raises motion-control time from 0.74 ms to 7.85 ms. [bidex_teleop_2024](#) solves it in hardware, mounting the hands so the human arm and the teacher arm "are perpendicular to each other and do not collide". [dexmimicgen_2024](#) states plainly that it does not handle inter-arm collision at all.

The fourth is the doubled action space. [bidexhands_2022](#) runs 20 tasks on two Shadow Hands and finds single-agent PPO beating multi-agent RL on most of them, offering in Sec. 5.2 that "PPO algorithm is able to use all observations for training the policy, while MARL can only use partial observations". [asymdex_2024](#) attacks the dimensionality from the other side, halving the observation and action dimension through its role split and a frame relative to the facilitating hand's object.

Table 1. Task families against the properties that define success

task family	what is held	what moves	what gravity does	what counts as failure	shortest honest success criterion
In-hand reorientation	one rigid object, held by the fingers for the whole episode	the object's orientation in the palm frame, while contacts break and remake	a disturbance whose direction is set by hand orientation, and a load to counteract with the palm down, visual_dexterity_2022	the object leaves the hand, e.g. the fall reset at goal distance 0.24 m in dextreme_2022	orientation error below a stated tolerance under a stated stopping rule, 0.4 rad in openai_dexterity_2018 and dextreme_2022 , 0.1 rad in eureka_2023
Grasping	nothing at the start, the object rests on a support	the hand closes, then the object is lifted clear	the test itself, since the grasp exists to resist it	no lift, or a drop inside the hold window	a stated lift height held for a stated duration, 10 cm for 20 s in dexgraspla_2025 , 10 cm to episode end in graspxl_2024 , 0.5 s in omnigrasp_2024
Functional and tool use	a tool, held in the pose its function requires rather than the most stable one	the tool's own joint, or the tool acting on a second object	secondary, the function dominates	a stable grasp in the wrong place, dexterous_functional_grasping_2023	no agreed criterion. Criteria range from never dropping the object, dexterous_functional_grasping_2023 , to a four-point rubric, dexvla_2025 . The one mechanical threshold is holding the tool joint past a commanded angle for the episode, articulated_tools_inhand_2025
Tracking a human reference	whatever the human held in the captured clip	hand and object follow a recorded trajectory under physics	full, and the reference never obeyed it, so the reference is not dynamically feasible	per-frame error leaves the band, or the object drops	all stated per-frame errors below threshold for the whole clip, 30 deg and 3 cm and 8 cm joint and 6 cm fingertip in maniptrans_2025 . Thresholds are not comparable across papers, and dexmachina_2025 calls the results "highly sensitive to the threshold values"
Bimanual coordination	one object by two hands, or one object per hand	both hands, the object, and the relative pose of the hands	can be carried by one hand while the other acts, asymdex_2024	the object drops, the arms collide, or the hands work against each other	no agreed criterion. Task-specific thresholds dominate, e.g. 0.035 m block-to-cup distance in asymdex_2024 . The one criterion with a physical-plausibility gate is 2.0 s of hold with total penetration under 1.5 mm, bimangrasp_2024
Handover and in-hand transfer	by the giver at the start, by the receiver at the end	the object crosses between hands, released only after it is regained	unopposed during the exchange, and total in a thrown transfer, dynamic_handover_2023	a drop at the moment of release, and the catcher's grasp phase, dynamic_handover_2023 Sec. 5.4	the receiver holds the object and moves it clear, more than 10 cm from the first hand in hato_visuotactile_2024 , or holds it to episode end in dydexhandover_2025

Two cells say the field has no agreed criterion, and both are in families where a second body is involved. A criterion naming a distance, a duration and a trial count can be re-run by someone else. A rubric cannot, and Section 7 takes up what follows from that.

3. Hands and who makes them

3.1 The design axes

No degree-of-freedom, force, weight or price figure in Table 2 or Table 3 was measured by anyone outside the maker. Six rows are the exception, from peer-reviewed papers with stated protocols: ILDA, Pisa/IIT, ORCA, RUKA, LEAP and BiDexHand. Table 2 dates every other claim.

The degree-of-freedom count is the first number a vendor states and the least comparable one. Shadow's specification of 4 December 2024 reads "20 actuated DOF and a further 4 under-actuated movements for a total of 24 joints" [shadow_dexterous_hand_2005](#). Table 2 prints joints first and actuated DoF second, and the gap is the informative number. Where a vendor's own DoF figure differs from the joint count, the joint count is what the table prints, as with Inspire's "Degrees of freedom 6, Numbers of joints 12" [inspire_rh56dfx_2023](#). An actuator count is neither of those columns and has its own. DexHand states no joint count at all, only 16 finger micro-servos and two wrist servos, so 18 servos is the only count quotable for it [dexhand_open_source_2023](#). Daxo states 120 actuators, no DoF figure, and a tendon structure with no rigid joints, which has far fewer kinematic degrees of freedom than actuators [daxo_muscle_v0_2025](#).

The actuation ratio is the real decision. The Pisa/IIT SoftHand drives 19 joints from one motor and its authors state the cost plainly: "No in-hand dexterous manipulation is required for this prototype" [pisa_iit_softhand_2014](#). Tendon drive moves the motors off the fingers and the mass follows them. Shadow's hand plus forearm weighs 4.3 kg, against 1.1 kg for the ILDA hand alone with every motor in its palm [shadow_dexterous_hand_2005](#) [ilda_hand_2021](#). ILDA states an 18 kg payload against Shadow's 4 kg in a power grasp, and Shadow states no fingertip force at all. What tendons cost is state estimation. The Faive Hand estimated joint angles from tendon length through a Kalman filter and its cube reorientation failed on hardware, which its authors blamed on "poor joint angle measurement from the EKFs, especially when there is contact" [faive_hand_2023](#). Ruka-v2 adds detachable AS5600 encoders to measure that problem rather than close the loop on it, and those encoders put its open-loop linear joint-to-motor map at 8.26 degrees of error over seven joints [ruka_v2_2026](#).

Direct drive trades torque density for transparency. LEAP's Dynamixel joints resist 19.5 N in a pull-out test against the Allegro Hand's 8.5 N, at 595 g for the LEAP hand [leap_hand_2023](#), and the Allegro's own weight has no reachable source. Linkage drive is the argument against the forearm. ILDA puts 15 motors, drivers and ball screws inside the palm and reports 34 N at the fingertip in the bent pose and 28 N stretched [ilda_hand_2021](#). Linkages buy coupling for free, as in BiDexHand's four-bar that drives each DIP off its PIP, whose paper claims 16 actuated DoF and 21 joints while its README describes 15 servos driving 15 joints [bidexhand_2025](#).

Fingertip force is the number a hand buyer reads first, and across these rows it is six different measurements. Table 2 puts the quantity beside the figure: pull-out resistance for LEAP's 19.5 N, pinch for RUKA's 2.74 N and BrainCo's 15 N, fingertip normal force under a 1 cm indenter for Unitree's 10 N, a calibrated five-trial fingertip force for BiDexHand's 2.14 N, a peak vendor claim for 1X's 45 N, and two unlabelled spec columns for Sharpa's "20 N | 12 N". BrainCo's whole-fist grip of 50 N is a separate figure and is no longer in the same column as a pinch. ORCA's 19.6 N was none of these. It is the newton equivalent of a 2 kg index-finger payload at a fixed 600 mA motor current, so its force cell is empty and its payload cell carries that protocol [orca_hand_2025](#). No two hands in Table 2 report payload under the same test either.

Compliance is repairability seen twice over, and both are mechanical fuses: the Pisa/IIT SoftHand's rolling-contact joints return to assembly after over-extension and ORCA's "poppable" pin joints dislocate rather than break [pisa_iit_softhand_2014](#) [orca_hand_2025](#). ORCA also reports tendon drive as maintenance: "Prolonged use requires manual re-tensioning to maintain performance". What fails is rarely the motors. Its durability run found silicone skin degrading on two fingertips after about 2,000 to 4,000 grasp cycles and sensor wires snapping on three after about 4,500 to 7,000. The sensing wore out an order of magnitude sooner than the hand.

Two columns matter before any of the above and were missing. The first is the rate a loop can be closed at, which spans 250 Hz on a Tesollo to 1 kHz on a Shadow, a Unitree or a Wuji, with the Allegro at 333 Hz on its own CAN clock and Inspire stating no rate at all. LEAP's 500 Hz is a serial query ceiling and its own sim-to-real policy runs at 20 Hz, which answer

different questions `leap_hand_2023`. A reader deploying torque control is bitten by that first. The second column says whether a URDF or MJCF exists and who ships it, which is the bridge to section 4 and to this survey's recommendation to pick a hand a simulator already carries. Shadow, Sharpa, Wuji, the Ability Hand, RUKA, Faive and the Allegro driver ship one. LEAP is the warning: its paper says the URDF is released and the released API repository contains none `leap_hand_2023`.

Table 2. Hands that can be obtained

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
shadow_dexterous_hand_2005 [1]	Shadow Robot Company	24	20	20	tendon-driven, motors in the forearm (20 Smart Motor nodes, Maxon motors, PWM)	4300			4 kg in a power grasp, vendor claim, protocol not stated	EtherCAT at 1 kHz to the host, with a 5 kHz tendon-force loop inside each motor module	MuJoCo models for Hand E and variants, plus Gazebo, in sr_common	Shadow Tactile Fingertips (STF): 17x3 DoF Hall-effect taxels, 1000 Hz, uncalibrated; up to 5 fingertips; software also supports BioTac/MST/UBIO/PST fingertip options		no	sold	datasheet	specification dated 2024-12-04	21: asymdex_2024, bidexgrasp_2026, bimangrasp_2024, ...
sharpa_wave_2026 [2]	Sharpa Robotics (Sharpa Pte Ltd)	22	22	22		1300	20 / 12	vendor spec table with two unlabelled columns, probably two configurations; grip force is listed separately as 150 N / 90 N	40 kg / 24 kg (the same two unlabelled columns)	500 Hz over 1000BASE-T Ethernet	MuJoCo and Isaac Sim models	'Dynamic Tactile Array' (DTA), camera-type array at the fingertip: resolution 240x240 / 60x60, force resolution 0.02 N / 0.05 N, spatial resolution 1 mm / 2 mm, frame rate 180 fps / 30 fps, force range 0-30 N, latency 20 ms, 6-D F/T; number of sensors per hand not stated		no	sold	vendor page	undated page, copyright 2026, fetched 2026-09-17	6: dexteleop0_2026, egoscale_2026, metis_2025, ...
proception_prohand_2026 [3]	Proception Inc (YC W25)	22			tendon-driven, motors pull cables to move the fingers							integrated skin-like sensors detecting contact, supporting grip control; same 'sensor skin' used on the		no	sold	press	launch post and press, 2026-06-29	0

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
												companion ProGlove; type/count not stated						
bidexhand_2025 [4]	Zhengyang Kris Weng, Center for Robotics and Biosystems, Northwestern University	21	16	15	cable-driven, N-configuration tendon routing (FeeTech servos, endless-loop antagonistic pull-pull); thumb CMC flexion driven via a 4-bar linkage instead of a direct tendon		2.14	fingertip normal force, calibrated sensor, five trials per finger, average peak	4.54 kg (10 lb) lifted	serial to an onboard ESP32, ROS 2 position control; rate not stated	URDF, consumed by the ROS 2 and MoveIt stack; no physics simulator	none mentioned		yes	open-source	paper	paper, 2025 (ICRA Dexterity workshop abstract)	0
wuji_hand_2025 [5]	Wuji Technology (founded 2019)	20	20	20	direct-drive rotary, backdrivable					1000 Hz across 20 axes over 100BASE-TX Ethernet, MIT force-position hybrid mode	MuJoCo and Isaac Sim models	'Multi-Axis Force/Torque Fingertip Sensing' confirmed present; no type, count or range given		no	sold	vendor page	undated page, copyright 2026, fetched 2026-09-17	2: toporetar get_2026, unidex_2026
ruka_v2_2026 [6]	Xinqi (Lucas) Liu, Ruoxi Hu, Alejandro Ojeda Olarte, Zhuoran Chen, Kenny Ma, Charles Cheng Ji, Lerrel Pinto, Raunaq Bhirangi, Irmak Guzey, New York University and NYU Shanghai	20	16	16	tendon-driven, forearm-mounted actuators; decoupled parallel 2-DoF wrist via a passive spherical ball joint; single-tendon-plus-spring abduction/adduction					Dynamixel over USB serial; no rate stated	MuJoCo XML in the shared v1 repo; not confirmed to carry the v2 wrist	optional e-flesh fingertips (magnetic touch-sensing form factor); not part of the base design	1500	yes	open-source	paper	paper, 2026	0
tesollo_dg5f_2024	Tesollo Inc. (Incheon HQ, Gwangmye	20	20	20	one integrated actuator per joint ('high-	1763			pinching 2.5 kg rated and 5 kg maximum;	250 Hz over Modbus RTU and TCP and		none standard; optional fingertip		no	sold	vendor page	undated page, copyright 2026,	0

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
	ong R&D/factory)				torque actuation', absolute encoder); page does not say direct-drive				enveloping 10 kg rated and 20 kg maximum	Ethernet TCP/IP		sensors (6-axis F/T, 3-axis force, or tactile) available; count/type not given					download links dated 2026-08	
unitree_dex5_2025	Unitree Robotics (Yushu Technology Co., Ltd.)	20	16	16	in-joint geared motor ('hollow-cup motor' + high-precision encoder + low-damping small-clearance reducer), backdrivable; page does not use the words tendon or linkage	1100	10	fingertip normal force under a 1 cm diameter cylinder pressed vertically down (vendor footnote)	3.5 kg palm down and 4.5 kg palm left, on a 5 cm round hard object	1000 Hz over USB 2.0, with per-joint stiffness and damping commands		Dex5-1: none. Dex5-1P: 94 pressure sensors per hand (2x5 palm + 2x3x5 finger pad + 2x3x5 fingertip + 2x3x4 finger root), range 10 g-2500 g.		no	sold	datasheet	undated page, copyright 2016-2025, fetched 2026-09-17	0
orca_hand_2025 [7]	Clemens C. Christoph, Maximilian Eberlein, Filippos Katsimalis, Arturo Roberti, Aristotelis Sympetheros, Michel R. Vogt, Davide Liconti, Chenyu Yang, Barnabas Gavin Cangan, Ronan J. Hinchet, Robert K. Katzschnann, Soft Robotics Lab, ETH Zurich	17	17	17	tendon-driven (antagonistic fishing-line tendon pairs per joint); wrist uses a GT2 timing belt drive	1200			10.5 kg on all four fingers and 2 kg on the index finger alone, both at a fixed 600 mA motor current	serial to Dynamixel or Feetech motors; rate not stated	URDF-derived kinematic constants only; no URDF or MJCF file in the released tree	yes: FSR-based binary tactile sensing on all 5 fingertips (RP-C7.6-ST), absolute detection threshold as low as 0.05 N		yes	open-source	paper	paper, 2025	0
leap_hand_2023	Kenneth Shaw, Ananye	16	16	16	direct-drive (Dynamixel servos, e.g.	595	19.5	pull-out resistance: the outward		up to 500 Hz querying over USB	URDF claimed in the paper;	none (future work only: 'we plan to	2000	yes	open-source	paper	paper, 2023	11: bidex_tel_eop_2024,

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
	Agarwal, Deepak Pathak, Carnegie Mellon University				XC330-M288), joint velocity ~8 rad/s			force a flexed finger resists before slipping or deviating more than 15 degrees (Table III)		serial; the paper's own sim-to-real policy runs at 20 Hz	the released LEAP_Hand_API repo contains no URDF, xacro or MJCF	develop and integrate LEAP Hand with low-cost touch sensors')						bidexhd_2024, cross_embodiment_world_mode ls_2025, ...
allegro_hand_v4_2016 [8]	Wonik Robotics Co. Ltd. (Seoul, South Korea); earlier versions 1.0/2.0 made by SimLab Co. Ltd.	16	16	16						333 Hz, CAN, the hand's own real-time clock (driver README)	URDF and xacro, left and right, in the ROS driver repo	none		no	sold	driver repo	driver repo, versions 1.0-4.0 undated, fetched 2026-09-17	35: anyrotate_2024, anyteLeop_2023, asymdex_2024, ...
ruka_2025	Anyia Zorin, Irmak Guzey, Billy Yan, Aadithya Iyer, Lisa Kondrich, Nikhil X. Bhattachali, Lerrel Pinto, New York University	15	11	11	tendon-driven (11 Dynamixel actuators in the forearm: XM430-W210T for the thumb, XL330-M288-T for the other fingers)		2.74	pinch force, best of three trials, averaged over the left and right hands (Table III)	6.0 kg: weight added to a curled cloth-bag grip until joint-angle error exceeds 15 degrees, best of three trials	Dynamixel over a USB-to-serial bridge; rate not stated, data collection ran at 15 Hz	MJCF	none (stated limitation: 'lacks tactile sensing')	1300	yes	open-source	paper	paper, 2025	0
roboteraxhand1_2024 [9]	ROBOTER A	12	12		gear-driven force-controlled joint modules per finger segment, back-drivable	1100			25 kg, tracker's 'Strength' figure, over 25 kg gripping palm-up	not stated	URDF exists but is licensed; maniptrans_2025 withholds it	tactile/force sensors on every fingertip confirmed present (senses contact, force, temperature); count and array size not stated	14000	no	sold	third-party tracker	undated tracker page, fetched 2026-09-17; the vendor page was not fetched	8: cross_embodiment_world_mode ls_2025, dexterity_2026, dexmachina_2025, ...
shadow_dex_ee_2024	Shadow Robot Company, in collaboration with Google DeepMind	12				4100				not stated; the page claims 'high bandwidth torque and position control loops'		stereo camera-based fingertip tactile sensors (hundreds of taxels each); multi-taxel		no	sold	vendor page	undated page, copyright 2026, fetched 2026-09-17	1: demostart_2024

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
												3-DoF arrays on middle and proximal phalanges; counts not stated						
inspire_rh56dfx_2023 [10]	Beijing Inspire Robots Technology Co., Ltd.	12	6			540	10	fingertip strength, vendor spec, four fingers; thumb 15 N; force resolution 0.50 N		RS485; rate not stated		none (this variant is the 'without tactile sensors' table; a tactile FTP variant exists per bib but is not on this page)		no	sold	vendor page	undated page, fetched 2026-09-17	19: ace_teleop_2024, articulated_tools_inhand_2025, being_h05_2026, ...
brainco_revo2_2025	BrainCo Inc.	11	6			383	15	pinch force, vendor minimum (≥ 15 N); the same page states a whole-fist grip of ≥ 50 N	≥ 20 kg	RS485 and CAN FD; EtherCAT on Pro and Touch; rate not stated		Touch variant only: multi-dimensional fingertip tactile module with 'Tactile Adaptive Control'; count/type not stated. Basic and Pro variants: none.		no	sold	datasheet	undated page, fetched 2026-09-17	1: bidegrasp_2026
agibot_omnihand_2025																unavailable	no source: both store.agibot.com URLs 404, fetched 2026-09-17	1: clutterdexgrasp_2025
linkerbot_l20_2025	LinkerBot (SDK author 'CHIUS INC')									USB-to-CAN at 1 Mbit/s, optional Modbus or RS485; rate not stated	URDF, with PyBullet and Isaac Gym examples in the SDK	L10/L20: per finger 4 channels (normal force, tangential force, tangential direction, proximity), 5 sites x 4 channels, 0-255 range, sensor type		no	sold	vendor page	undated page, fetched 2026-09-17	1: being_h05_2026

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
												not stated. O6 (reseller): 16 capacitive tactile regions.						
psyonic_ability_hand_2021 [11]	PSYONIC (San Diego, CA, USA)		6	6	linkage (repo file/function names indicate a four-bar finger linkage; not stated on the vendor page itself)					BLE, I2C, UART or RS485; rate not stated	URDF (left, right, large, small, no-FSR) plus MuJoCo and Isaac Sim paths in the vendor repo	30 touch sensor values streamed (FSR-based per API and URDF 'no_fsr' variant); per-finger split not stated		no	sold	vendor page	undated page, fetched 2026-09-17	7: ace_teleop_2024, asymdex_2024, bunny_visionpro_2024, ...
dexhand_open_source_2023 [12]	Rob Knight, The Robot Studio (design); electronics/firmware/ROS 2 by Trent Shumay, IoT Design Shop			18	tendon (fishing line, Suffix 832 80lb / 0.8mm kiteline); Emax ES3301/ES3302/ES3351/ES3352 micro-servos for fingers, Feetech SCS2332 or PWM servos for the wrist					Arduino SCServo or SBUS firmware, BLE; rate not stated	URDF, in the IoT Design Shop ROS 2 packages	none mentioned	300	yes	open-source	project page and GitHub README	page posts 2023-08-08 to 2023-10-01; no hardware release date	0

20 rows. 95 of 240 specification cells (39%) over the 12 specification columns are values no source stated; the key, maker, provenance and usage columns are excluded because they are never blank. Every figure here is the maker's or the authors' own claim. Nobody outside the maker has measured any DoF, force, weight or price cell in this table, except the rows sourced to a peer-reviewed paper with a stated protocol. The force column is not a ranking: read 'what the force is' first, because pull-out resistance, pinch force, a fingertip normal force under an indenter and an unlabelled vendor spec are different measurements. A blank cell means no source stated the value. No cell is truncated: a value wider than its column is wrapped at a word boundary, so a cell that runs to several rendered lines is one value and not several.

Figures with no reachable source, and other caveats on individual rows:


1. `shadow_dexterous_hand_2005` : the four distal joints are coupled and not independently controllable, so the '24 DoF' quoted in learning papers is the joint count
2. `sharpa_wave_2026` : the page footnotes 'Specifications may vary between products', and no column of the spec table is labelled
3. `proception_prohand_2026` : the evidence is a shipping announcement, not a datasheet: no weight, force, actuated-DoF count or control rate is stated anywhere
4. `bidexhand_2025` : the paper states 16 independently actuated DoF and 21 joints; the README states 15 servos driving 15 joints
5. `wuji_hand_2025` : the page footnotes that these figures 'reflect the current Beta1 product and the SPEC will continue to iterate'
6. `ruka_v2_2026` : the AS5600 encoders are attachable and detachable and are read for calibration and measurement only, not in the control loop
7. `orca_hand_2025` : the 19.6 N previously tabulated as a fingertip force is the newton equivalent of the 2 kg index-finger payload at a control-imposed current limit, and no source states it as a force
8. `allegro_hand_v4_2016` : allegrohand.com/v4 returned HTTP 404 and the Wonik wiki timed out; every cell here comes from the ROS driver repository, and weight, joint torque, payload and price have no reachable source
9. `robotera_xhand1_2024` : the 15 N in circulation comes from a survey table, not from the tracker or the vendor; the tracker states no fingertip force
10. `inspire_rh56dfx_2023` : the vendor's own page reads 'Degrees of freedom 6, Numbers of joints 12'; the joints column here prints the joint count, so the vendor's DoF figure is the actuated-DoF column
11. `psyonic_ability_hand_2021` : the four-bar finger linkage rests on repository file names and the five fingers on URDF meshes; the vendor page states neither
12. `dexhand_open_source_2023` : no joint count is given anywhere, so no DoF figure can be quoted; the servo count is 16 finger and thumb micro-servos plus 2 wrist servos, with a third optional. The \$300 is 'additional total cost of components', excluding printing and the wrist servos

3.2 The hands the research literature actually runs on

Figure 2. Which hands the surveyed methods actually run on

One count per method paper whose own experiments use the hand, from 112 method rows.

Colour: blue sold commercially, green open hardware, purple simulation-only model, orange built in-house and not sold.



announced, not sold (Table 3) | 0 — none of the 9 company-announced hands appears in any method row

research prototypes (Table 3) | 3, from 2 of the 4 lab hands in Table 3

The two oldest designs in Table 2 carry 49 of the 103 method rows that name a hand at all, and seven rows use both. Neither design's date is confirmed by its own sources, so 2005 and 2016 are the bibliography's. Figure 2 counts, per hand, the method papers whose own experiments use it. Of 112 method rows, 103 name a hand. The Allegro accounts for 35, Shadow for 21, the Inspire RH56 family for 19, a parallel-jaw gripper for 12 and LEAP for 12. Seventy-five of the 103 name an Allegro, a Shadow or Adroit model, LEAP or an Inspire.

The Allegro's position is the uncomfortable part. Its product page at allegrohand.com/v4 returned HTTP 404 while the Wonik wiki timed out allegro_hand_v4_2016. Everything Table 2 confirms about the most-used hand in the corpus comes from its ROS driver: 16 joints, four fingers, a torque interface, a 333 Hz CAN clock, 12 V, no tactile sensing. Its weight, joint torque, payload and price have no reachable source. The field's most common platform cannot be specified from a page a reader can open.

Concentration would matter less if the hand did not move the result, and it does. On the same simulated cube rotation, LEAP reaches 0.2288 rad/s against the Allegro's 0.0828 rad/s leap_hand_2023, and RUKA reports a 2.74 N pinch against the Allegro's 1.60 N under the same three-trial pinch test ruka_2025. A method compared only on Allegro hardware is compared at one point in a space where a single axis moves the headline number two or three times over. Two further patterns follow. Inspire, XHand and Sharpa take 29 of the 103 rows between them, four rows use two of the three, and none is earlier than 2024. ORCA, RUKA, Ruka-v2, BiDexHand, DexHand, the Tesollo DG-5F and the Unitree Dex5 take none at all.

3.3 Open hardware and the collapse in cost

Seven rows in Table 2 are open hardware. Five state a dollar cost in the price column, ORCA states a material cost below 2,000 CHF that the column leaves unconverted, and BiDexHand states none `orca_hand_2025` `bidexhand_2025`. The Faive Hand states neither a cost nor a licence that could be read, so it is not counted as open hardware here `faive_hand_2023`. Two of the stated bases need saying. DexHand's \$300 is "additional total cost of components", excluding the printing and the wrist servos `dexhand_open_source_2023`, and ORCA's own figure sits against Ruka-v2's table listing ORCA at about \$3.5K `ruka_v2_2026`.

The expensive end of the collapse is secondhand throughout. The only six-figure numbers on disk are RUKA's comparison table at \$100,000 for a Shadow Hand and Faive's "steep price tag of 110k GBP" `ruka_2025` `faive_hand_2023`. Shadow's own page says to discuss pricing and Table 2's price cell for it is empty. The collapse is real at the cheap end and secondhand at the expensive one, and it has barely moved the literature. Twelve of the 103 hand-naming method rows use an open-hardware hand, and all twelve are LEAP.

What the cheap hands give up is sensing. LEAP has none and names touch sensors as future work `leap_hand_2023`, RUKA states its design "lacks tactile sensing" `ruka_2025`, and BiDexHand's sources never mention touch `bidexhand_2025`. ORCA is the exception, with binary FSR fingertips whose threshold was measured as low as 0.05 N on a fresh fingertip, against the sensor's rated 0.29 N, and 6.38 N on a degraded one `orca_hand_2025`.

3.4 Announced and unreleased hands

Behind Table 3's rows sit three kinds of evidence, and conflating them is how a DoF figure with no source ends up in a survey. Vendor prose carrying numbers is the strongest, as on 1X's page of 9 July 2026 `onex_neo_hand_2026`. Video is second and carries none. Press or bibliography assertion is third, as with Tesla's V3 hand, known here only through a paraphrase of patents because the USPTO PDF parsed empty `tesla_optimus_hand_2025`. Table 3 blanks the cells resting on the third class and footnotes who did the arithmetic.

A survey can go one step past recording that a claim is unverified, which is to say which claims are implausible on their face. Daxo's 120 actuators in 750 g is about 6 g per actuator including structure, tendons, routing and skin `daxo_muscle_v0_2025`. Clone's 27 degrees of freedom under 2 pounds excludes a 500 W pump the source does not confirm is excluded `clone_robotics_hand_2024`. Figure 03's "Degrees of freedom, hands | 20" sits on the same tracker page as "Number of fingers | 10", so it is almost certainly the pair `figure_03_hand_2025`. Tesla's repeated 22 is one article's arithmetic of four DoF on each of five fingers plus two at the wrist, and Gen 2's own figure was 11 `tesla_optimus_hand_2025`.

Scepticism belongs to the evidence class, not to which table a row lands in. Sharpa's 22 of 22, Wuji's 20 of 20 and Tesollo's 20 of 20 are vendor claims about unmeasured hardware and they sit in Table 2, where Sharpa's page footnotes "Specifications may vary between products" and Wuji's says the spec "will continue to iterate" on a Beta1 product `sharpa_wave_2026` `wuji_hand_2025`. Even the best-specified vendor page in the corpus leaves its fingertip-force columns unlabelled.

Table 3. Hands announced but not purchasable

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
clone_robotics_hand_2024 [1]	Clone Robotics	27			hydraulic artificial muscle ('Myofiber' — small water-filled tubes that contract when pressurized), carbon-fiber bones and ligament-style tethers, driven by a 500 W water pump and 36 electro-hydraulic valves; each Myofiber generates up to 1 kg grip force, survived 650,000 test cycles	<910						pressure pads in the palm detect grip firmness (count not given); 70 inertial sensors track angle/speed (proprioception, not tactile)		no	internal-only	press	press, 2025-11-24, on a video of 2025-11-15	0
onex_neo_hand_2026	1X Technologies	25	25		quasi-direct-drive tendons via the '1X Tendon Drive' at low gear ratios (~5:1-15:1); motors in the forearm, tendons pulled through the wrist		45	distal flexion force, peak vendor claim		not stated		high-resolution tactile sensing across fingertips and finger surfaces (normal force, contact location, shear); sensor count/technology not stated		no	announced	vendor page	vendor page, 2026-07-09	0
sanctuary_phoenix_hand_2024	Sanctuary AI (Sanctuary Cognitive Systems Corporation)	21			hydraulic, 'unique miniaturized hydraulic valves'							not in the press release text; sidebar-only headlines mention new		no	internal-only	press	undated press release; trade reprint 2024-12-17	0

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
												touch/tactile sensors as a separate later announcement, not confirmed as part of this hand						
leap_hand_v2_adv_2025 [2]	Kenneth Shaw, Deepak Pathak, Carnegie Mellon University	21	17	17	tendon (PIP/DIP coupled by a single tendon per finger); motor type not stated					not stated	URDF, loaded into PyBullet by the released teleoperation node	none mentioned	3000	yes	announced	vendor page	undated project page; API repo at commit a0936196, fetched 2026-09-17	2: bidex_tel_eop_2024, dexwild_2025
ilda_hand_2021	Ajou University, Korea Institute of Machinery & Materials, Korea University	20	15	15	linkage-driven, direct linear drive (3 Maxon DCX8M motors per finger with GPX8 16:1 gearboxes driving ball screws via parallel+serial four-bar linkages)	1100	34	fingertip normal force in the bent pose; 28 N stretched, and 25 N per finger measured while crushing an aluminium can	18 kg	CAN to a desktop; loop rate not stated		6-axis F/T sensor per fingertip (5 total), force resolution 62 mN, range +/-35 N; not a distributed taxel skin		no	prototype	paper	paper, accepted 2021-11-05	0
pisa_iit_softhand_2014	Centro E. Piaggio, University of Pisa and IIT (Catalano, Grioli, Farnioli, Serio, Piazza, Bicchi)	19	1	1	tendon-driven adaptive synergy, single motor via differential gears (6 W Maxon RE-max21, 84:1 reduction)				holding force about 20 N along z and holding torque 2 N m, 28 N and 3.5 N m with a stronger motor	not stated	URDF and Gazebo models in the ROS repo (BSD 3-Clause)	none (motor encoder only; a padded work glove supplied contact compliance during experiments)		no	prototype	paper	paper, IJRR 2014	0
faive_hand_2023 [3]	Yasunori Toshimitsu, Benedek Forrai, Barnabas Gavin Cangan, Ulrich Steger, Manuel Knecht, Stefan Weirich,	16	11	16	tendon-driven, rolling-contact joints (16 Dynamixel XC330-T288-T servos, 6 antagonistic pairs)	1100			10 kg, whole-hand downward power grasp on a dumbbell	Dynamixel over serial; joint angles estimated from tendon length by EKF; rate not stated	MJCF, used directly by the released Isaac Gym training repo	none on the physical hand; a simulated 'fingertip force' (15-dim, critic-only/privileged) exists only in the RL training observation, not real			prototype	paper	paper, 2023	1: graspxl_2024

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
	Robert K. Katzschmann, Soft Robotics Lab, ETH Zurich											tactile hardware						
paxini_dexh13_2024	PaXini Tech											1,140 ITPU multidimensional tactile processing units (press-release claim; unclear if PX-6AX-GEN3 sensor family)			announced	press	press, CES 2026-01-07	0
tesla_optimus_hand_2025 [4]	Tesla				V3: tendon-driven from forearm actuators, three tendons per finger through a crosstalk-managing wrist. Gen 2: in-hand actuators/sensors (unlocated in the article).							Gen 2: tactile sensing on all fingers, demonstrated matching contact location and force in an egg-pickup video; type/count not stated. V3: not stated.		no	prototype	press	Gen 2 video 2023-12-13; V3 patents filed 2024-10, relayed by press 2026-04	0
boston_dynamics_atlas_hand_2026	Boston Dynamics				electric (robot is fully electric); specific hand mechanism not stated							tactile sensing in the fingers and palms, confirmed present; type and count not stated		no	prototype	press	trade press, 2026-01-05 (CES); the Boston Dynamics blog of the same date has no hand content	0
figure_03_hand_2025 [5]	Figure AI											not stated on page for the hand itself; palm cameras in each hand are vision, not tactile		no	prototype	third-party tracker	undated tracker page, fetched 2026-09-17; figure.ai launch page returned HTTP 404	0

hand	maker	joints	act. DoF	actuators	actuation	weight g	force N	what the force is	payload	control rate	URDF or MJCF	tactile	price USD	open HW	status	source	claim date	corpus methods using it
daxo_muscle_v0_2025 [6]	Daxo Robotics			120	ultra-redundant tendon-driven; compliant structure with no rigid joints, flexible materials and tendon routing	750									prototype	third-party tracker	third-party catalogue post, 2026-09-16; daxo-robotics.com returned HTTP 404	0
xiaomi_cyberone_hand_2026	Xiaomi				motors located in the hand (compact motors generate heat requiring liquid cooling); no tendon/linkage statement, though contrast with tendon-driven hands' cycle life implies a non-tendon design							full-palm tactile sensing, area ~8,200 sq mm, detects pressure and contact across the whole palm not just fingertips; taxel count and type not stated		no	prototype	press	press, 2026-03-30	0

13 rows. 96 of 156 specification cells (61%) over the 12 specification columns are values no source stated; the key, maker, provenance and usage columns are excluded because they are never blank. Every figure here is the maker's or the authors' own claim. Nobody outside the maker has measured any DoF, force, weight or price cell in this table, except the rows sourced to a peer-reviewed paper with a stated protocol. The force column is not a ranking: read 'what the force is' first, because pull-out resistance, pinch force, a fingertip normal force under an indenter and an unlabelled vendor spec are different measurements. A blank cell means no source stated the value. No cell is truncated: a value wider than its column is wrapped at a word boundary, so a cell that runs to several rendered lines is one value and not several.

Figures with no reachable source, and other caveats on individual rows:

1. `clone_robotics_hand_2024`: 'under 2 pounds' is an upper bound, and the source does not say whether the 500 W pump and 36 valves are inside it
2. `leap_hand_v2_adv_2025`: the page claims 21 DoF with 17 powered motors and the README calls the same hand 17-DOF
3. `faive_hand_2023`: the row's open-hardware status is unconfirmed: the repo has LICENSE and LICENSE-NVIDIA files whose contents were not captured, and no open-hardware claim appears in the sources
4. `tesla_optimus_hand_2025`: the 22 DoF in circulation is a Teslarati article's arithmetic of four DoF on each of five fingers plus two at the wrist, not a Tesla statement; Gen 2's own figure is 11 DoF. This row's actuation cell describes V3 and its tactile cell describes Gen 2
5. `figure_03_hand_2025`: the tracker's 'Degrees of freedom, hands | 20' sits beside 'Number of fingers | 10', so the 20 is most likely the pair and about 10 per hand; the 16-per-hand figure in circulation has no reachable source
6. `daxo_muscle_v0_2025`: 120 actuators in 750 g is about 6 g per actuator including structure, tendons, routing and skin; no DoF number is stated, and a redundant tendon architecture with no rigid joints has fewer kinematic DoF than actuators

3.5 Tactile sensing as part of the hand

DIGIT set the cost floor. It is 20 by 27 by 18 mm, weighs about 20 g, streams 640x480 at 60 fps, and its paper states a "total estimated manufacturing cost is approximately 15 USD per sensor ... when manufactured in a batch of 1000" `digit_2020`. Durability was as much the contribution as price: its gel degraded 0.3 percent over 15 abrasion passes, against 805 and 918 percent for the two gels compared with it.

Digit 360 is the same lineage at a different operating point, claiming about 8.3 million taxels, spatial features to 7 um, normal and shear force resolution of 1.01 mN and 1.27 mN, and on-device processing that cuts event-to-action latency from 6 ms to 1.2 ms `digit360_2024`. It is also not for sale, and no corpus paper's parsed text names it.

XELA's uSkin is the non-camera alternative and the one that bolts onto hands the field already uses. Its curved fingertip kit for the Allegro V4 and LEAP carries 30 three-axis sensing points, a full Allegro integration reaches 368, resolution is 0.1 gram-force, and the kit runs at 275 Hz `xela_uskin_2020`. The 500 Hz figure often quoted belongs to the product family, not the curved fingertip. One corpus paper names it, `sparsh_2024`, whose self-supervised encoders are pre-trained on 462.7k tactile images and beat matched end-to-end models by 95.1 percent when both see 33 to 50 percent of the labels. Its own bead-maze policies never complete the maze on the real robot.

The usage number is the one to keep, and it needs its inclusion rule stated: a sensor physically on the hand, read by the deployed policy. Eight of 112 method rows meet it: `anyrotate_2024`, `articulated_tools_inhand_2025`, `dexteleop0_2026`, `dexumi_2025`, `hato_visuotactile_2024`, `robot_synesthesia_2023`, `rotateit_2023` and `rotating_without_seeing_2023`. The rule excludes `penspin_2024`, whose 20 binary contacts are simulated on an Allegro that has no tactile hardware and whose released config sets `enable_tactile: False`, and it excludes `dexndm_2025` and `dexplore_2025` for the same reason. Sixty-five of the 112 method papers mention tactile sensing somewhere, which is the gap worth quoting. Thirty-five is the count over this survey's notes.

Almost none of the eight uses a high-resolution sensor. `rotating_without_seeing_2023` removes vision entirely and rotates objects from 16 binary touch sensors over the palm, links and fingertips, deployed zero-shot to a real Allegro, and

Robot Synesthesia uses 16 force-sensing resistors read as binary [robot_synesthesia_2023](#). Two papers exclude touch deliberately, HORA reporting rotation "even without the usage of vision and tactile sensing" [hora_2022](#), and ORCA's authors dropping it from their reinforcement learning "due to the additional complexity involved in accurately modeling them" [orca_hand_2025](#). Taxel counts have risen by three orders of magnitude while the policies consuming them have stayed at binary contact.

3.6 What is sold against what is published on

The bottom rows of Figure 2 carry the finding. None of the nine company-announced hands in Table 3 appears in a single method row whose own experiments use it. Tesla, Figure, 1X, Sanctuary, Boston Dynamics, Xiaomi, Clone, Daxo and PaXini account for zero of the 112 method papers' experiments. The nearest thing to a counterexample is [helix_2025](#), a Figure blog post claiming a 35-DoF whole-upper-body action space at 200 Hz that includes individual finger control. It never names the hand, gives no per-hand DoF count, and reports no success rate or trial count for any task. It is the maker describing its own unreleased hand, which is the evidence class the finding is about. The four research prototypes in Table 3 are in a different position, since the Faive Hand and LEAP Hand v2 Advanced account for three method rows between them [graspxl_2024](#) [bidex_teleop_2024](#).

Table 3's emptiness is measurable and part of the same finding. Over the twelve specification columns, 61 percent of its cells are values no source stated, against 39 percent for the hands that can be bought. The hands with the highest advertised DoF counts have the least specification behind them.

The one hand that crosses the gap crosses it because its maker published rather than announced. ByteDexter reaches a method row through a ByteDance technical report stating 21 DoF per hand, 16 piezoresistive fingertip channels in the action vector and real-robot success rates [gr_dexter_2025](#), not through a product page. That mechanism is available to every vendor in Table 3 and none has used it.

The gap runs the other way too, and the last column of Table 2 shows it. Seven hands that are sold and documented take zero method rows each: the Unitree Dex5 with 94 pressure sensors on its P variant [unitree_dex5_2025](#), the fully actuated Tesollo DG-5F [tesollo_dg5f_2024](#), ORCA, RUKA, Ruka-v2, BiDexHand and DexHand. Cheapness is established only for RUKA, Ruka-v2 and DexHand, because Table 2's price cell is empty for Unitree, Tesollo, ORCA and BiDexHand, and the only Dex5 price on disk is Ruka-v2's secondhand "\$25K" [ruka_v2_2026](#). A reader choosing a hand on this corpus's evidence has two well-precedented options, an Allegro or an Inspire, and a third in LEAP if eleven papers is enough. Everything else is a press release or a hand nobody has published on.

4. Simulators and the physics underneath

4.1 What a dexterous simulation must get right

The clearest demonstration that hands are the hard case came from a benchmark that was not about hands. Erez et al. built a 35-DOF arm modelled on the Shadow Hand, closed it around a capsule with fixed spring-dampers, and asked four of the five engines for the largest timestep at which the object was still in the hand. MuJoCo held the grasp at 16 ms, PhysX at 2 ms, ODE at 0.25 ms and Bullet at 0.03 ms, a spread of a factor of 500 ([physics_engine_comparison_2015](#), Sec. IV-D). Havok is the fifth and was excluded for want of a working PD controller. Three caveats travel with that spread. The engines run a deliberately restricted common model, hinge joints with sphere and capsule geometry, no boxes and no meshes (Sec. II). The authors call it "an open question whether this model system can be tuned to work better in the gaming engines" (Sec. IV-D). The timesteps are log-spaced and good only to a factor of two, and the authors wrote MuJoCo and disclose it.

A grasp is hard for nameable reasons. It is many contacts at once, all persistent, all near stiction, on an object much lighter than the mechanism holding it. Persistence and multiplicity make the problem hyperstatic, and Le Lidec et al. show that per-contact solvers of the projected Gauss-Seidel family, and RaiSim's, then inject spurious jamming forces at stiction that vanish only once the object slides. The mass ratio makes it ill-conditioned, and in their stacked-cube test at a 10^3 to 10^{-3} kg ratio those same methods fail to converge. Global methods with proximal regularisation stay robust in both cases ([contact_models_comparison_2023](#), Sec. IV-A). What degrades a truncated solve is therefore conditioning and redundancy, not stiffness, and a grasp supplies both by construction. Locomotion is forgiving by comparison. On flat ground the contact model and solver choice "hardly affects" their quadruped's tracked base velocity, and only on rough, slippery terrain do RaiSim and CCP deviate. A walking robot makes and breaks a few contacts against ground far heavier than itself. A hand does neither.

Speed enters by the same door. MuJoCo Playground reports that contact time scales with the number of possible contacts rather than the active ones, because JAX requires static shapes, which is why its tasks carry hand-tuned `max_contact_points` and `max_geom_pairs` overrides ([mujoco_playground_2025](#), Sec. VI). Every override printed in that paper is a locomotion port. A hand's bound would have to be set the same way, for the worst case, but no hand configuration is printed.

Figure 3 sets out the stages of one simulation step. Three of them make overlap, and they do not answer to the same knob. In an engine that enforces non-penetration at the velocity level, integration turns any residual approach velocity into overlap of order v times Δt ; an engine that enforces the gap at the next configuration carries no such term, which is why Dojo's hard-contact NCP keeps its feet above the floor at every timestep it was tested at. Constraint assembly fixes the compliance a loaded contact then rests at. A truncated solver leaves a residual that grows with conditioning. Which of the three dominates in a grasp is not measured anywhere in this corpus, and the three are not ordered here. A fourth item on the figure is not a source of overlap at all. It is a mismatch between the geometry the solver uses and the geometry the renderer draws, and it runs in both directions.

4.2 Contact models and solvers, engine by engine

Figure 3. One simulation step, and where engines differ

The stages every rigid-body engine runs. Under each, the choices that separate the engines used for hands. Dashed red marks the three stages that make overlap. They do not answer to the same knob.




Table 4 is the engine-by-engine comparison, one row per simulator, built only from what a note confirmed. Le Lidec et al. supply the taxonomy that organises it, checking each formulation against the Signorini condition, Coulomb's law, and the maximum dissipation principle. Linear complementarity, the family of Bullet, ODE, PhysX and DART, satisfies Signorini alone, because linearising the friction cone to a pyramid biases friction toward its corners. The cone complementarity problem satisfies the other two but relaxes Signorini, so contact acts at a distance of size $\Delta t \cdot \mu \cdot l_{v_tl}$. The full nonlinear problem satisfies all three and is non-convex ([contact_models_comparison_2023](#), Table II). A contact model is not the algorithm that solves it. PhysX linearises the cone, which places its model in the LCP family, and it does not solve an LCP. Its Temporal Gauss-Seidel scheme folds substepping into the Gauss-Seidel sweep and exposes position and velocity iteration counts separately ([isaacgym_2021](#), Sec. 3).

MuJoCo is inside that taxonomy rather than outside it. Le Lidec et al. classify it as CCP-MuJoCo, a convex relaxation solved by a Newton method on the primal QCQP, in the same family as CCP-Drake, the SAP-style scheme (Table III). The relaxation makes two errors of opposite sign in one engine. A loaded contact carries a violation, and a sliding contact carries force at a positive gap. Drake's SAP inherits the same pair, and its gliding effect at distance $\phi \approx \delta t \cdot \mu \cdot l_{v_tl}$ "unfortunately does not go away as $\delta t \rightarrow 0$ " ([castro_sap_contact_2021](#), Sec. V-A). Le Lidec et al. call that compliance a "numerical trick designed to circumvent the issues due to hyper-staticity or ill-conditioning at the cost of impairing the simulation".

The depth a loaded contact carries is a chosen number. In a regularised formulation the steady-state violation at a contact is the normal load times the compliance. It is zero at zero load and it grows with the load carried. MuJoCo drives that violation coordinate with a critically damped stabiliser parameterised by ε and κ , and for an object resting under gravity the steady-state depth has a closed form independent of the object's mass ([mujoco_convex_contact_2014](#), Sec. V). The closed form did not survive the parse of that paper, so the cancellation is quoted and the algebra is not. Neither MuJoCo paper states why mass cancels, and the explanation this survey offers is its own inference rather than a cited one: the regulariser is scaled by the inverse effective inertia at the contact, so compliance falls as $1/m$ exactly as the gravity load rises as m . The nearest support in a parsed source is for a different solver of the same engine, MuJoCo's diagonal solver, a "mass-aware spring-damper" that uses the diagonal of the A matrix to keep contacts critically damped ([mujoco_2012](#), Sec. II-E). This is not a

penalty spring, and depth is not always non-zero. The impulse solves a regularised convex program over the whole contact set rather than a per-contact function of the gap, both MuJoCo papers reject spring-dampers by name, and the 2012 ball-drop figure is captioned "there is no penetration" ([mujoco_2012](#), Fig. 2).

The first of two concrete measurements comes from the other end. Dojo solves a hard-contact nonlinear complementarity problem with an exact second-order friction cone, by an interior-point method converging within 15 iterations on the three robots of its convergence study. Its Table II drops an Atlas humanoid and reports foot-floor penetration against the timestep. MuJoCo penetrates -28 mm at $\Delta t = 0.01$ s and -46 mm at $\Delta t = 0.001$ s, while Dojo stays above the floor at every step tested ([dojo_2022](#), Sec. V-A). The MuJoCo column is not a trend. A ten-times-smaller step produces more overlap, which no timestep-independent stabiliser does. Either that configuration ties the compliance to Δt , or the quantity is an impact transient on a drop. Neither reading is a steady-state grasp depth. Drake's SAP quotes 2.5×10^{-5} m at $\delta t = 10^{-2}$ s and 2.5×10^{-7} m at $\delta t = 10^{-3}$ s, three and five orders of magnitude below Dojo's two MuJoCo cells at the same steps, on an engine that is also compliant ([castro_sap_contact_2021](#), Sec. V-B). Those two figures are analytical bounds for a single point mass at rest on a plane under that paper's near-rigid stiffness rule, not measured depths, and they are not a Drake grasp-penetration number: a point-mass bound and a humanoid drop transient differ in load, effective inertia and regime, so the distance between them is not a measurement of anything. What the pair of engines does show is that in a compliant formulation the depth follows from a stiffness that someone chose.

Dojo's Table V times 1000 steps of forward simulation with gradients, at a matched $\Delta t = 0.01$ s, for engines that are not all computing gradients. MuJoCo is fastest on every system, 0.335 ± 0.001 s against Dojo's 1.159 ± 0.077 s on a Franka Panda, and the authors call the comparison difficult because Dojo is stable at five times the step size (Sec. VI-B). What the regularisation buys, on Le Lidec's reading quoted above, is conditioning on a hyperstatic problem, and the depth it costs is tuned separately. Whether that is also what holds Erez's grasp at 16 ms is a question his data do not answer. His planar chain is contact-free, so its numbers speak to the coordinate formulation and not to contact: MuJoCo runs at 243.2 kHz there against Bullet's 22.8 and PhysX's 6.4, and Bullet's articulated Featherstone mode at 81.4 kHz beats every Cartesian-coordinate engine ([physics_engine_comparison_2015](#), Sec. IV-B). Those are throughput in evaluations per second, not a largest-stable-timestep result, and the articulated Bullet mode was never run on the grasp test at all, being usable only in tests without contact (Appendix). Joint coordinates explain the contact-free speed advantage. The grasp timestep is a contact result, and the paper attributes it to nothing: it reports that the other engines go unstable and "effectively simulate a different physics model which can no longer hold the object" (Sec. IV-D), without an experiment that separates the coordinate formulation from the soft absorption of penetration.

The GPU era moved the compliance knob rather than removing it. ComFree-Sim resolves contact in closed form in the dual cone of the friction cone, so penetration becomes an explicit tuning parameter and the paper reports it. On a drop test of convex primitives about 5 cm across at $dt = 0.002$ s, MuJoCo Warp penetrates 1.7 ± 4.9 mm and ComFree-Sim runs from 3.9 ± 6.9 mm at $k_{\text{user}} = 0.1$ down to 0.9 ± 1.5 mm at $k_{\text{user}} = 0.5$, though its prose describes the opposite direction from its own table ([comfree_sim_2026](#), Sec. IV-A). Its impedance acts on the signed gap through the "identical" default hyperparameter API as MuJoCo's solver, so the knob it is credited with exposing is one MuJoCo users already set. Those are 5 cm primitives under their own weight, not a fingertip loaded by a grasp, and depth scales with the normal load. Nobody has published the fingertip number. The baseline it measures against is MuJoCo Warp, which Newton builds on and which MuJoCo Playground names as its intended replacement for JAX ([mujoco_playground_2025](#), Sec. VI).

That brings Table 4's most important column, and the claim it does not support. The column records whether a parsed source reported a penetration depth, not what an engine can compute. Two of the 15 engines report one, Dojo and ComFree-Sim, and both are engines whose paper is about contact accuracy. Four are recorded as not reporting it: Brax, Isaac Gym, Isaac Lab and Orbit. Nine rows are blank. The word "penetration" appears nowhere in the Isaac Gym paper, and Isaac Lab's contact sensor reports force, duration and an average contact point with no contact-quality metric. Those two simulators carry 42 of the 112 method papers in this corpus.

The tooling exists and the number is still not recorded. NVIDIA's own IsaacGymEnvs repository computes interpenetration depth in simulation. Its IndustReal tasks load plug and socket meshes into Warp, sample points on one, query them against the other, and reduce to a per-environment maximum interpenetration distance ([code/md/isaacgym_2021.md](#),

lines 8809 to 8862). The policy update is gated on that number. Environments are split on `max_interpen_dists <= interpen_thresh` and the surviving reward is scaled by `1 - tanh(max_interpen_dists / interpen_thresh)`, with `interpen_thresh: 0.001` commented as the maximum allowed interpenetration between plug and socket (lines 3630 and 3753). That is a shipped Isaac Gym task measuring simulated interpenetration per environment at a millimetre threshold, during RL, and acting on it. Table 4 records Isaac Gym as not exposing penetration and `corpus/rows/isaacgym_2021.json` carries `penetration_exposed: false`, which the code parse in the same corpus contradicts. Tactile Genesis makes the point from the other side, shipping penetration depth as a sensor on an analytic SDF backend and a BVH backend (`tactile_genesis_2026`, App. A.1), while Table 4 leaves the Genesis cell blank. The depth is computable from the poses and the meshes in a few lines of Warp, and the field's own benchmark repository already does it. Interpenetration in a dexterous rollout is a setting nobody records and a measurement nobody takes.

A policy will exploit what nobody looks at. DexTrack's configs carry PhysX's `max_depenetration_velocity` at 10.0 or 1000.0 depending on the task variant, with no explanation in the paper or in a config comment (`dextrack_2025`). The same parameter appears across unrelated stock IsaacGymEnvs tasks at 5.0, 10.0, 100.0 and 1000.0 (`code/md/isaacgym_2021.md`, lines 897, 1924, 2083, 2448 and 3113), so DexTrack inherited a template rather than choosing per variant. The parameter caps the rate at which the solver pushes overlapping bodies apart, so it sets how long an overlap persists and how violently it is undone, not how deep the overlap gets. The one knob here that governs interpenetration behaviour is being copied without being read. DexTrack's paper defines a maximum hand-object penetration depth, applies it only to its input kinematic references, and presents tolerance of "severe hand-object penetrations" as evidence of robustness (App. B). `toporetaret_2026` is the one corpus method that reports the number carefully, and it reports it on retargeted references rather than on a rollout, which section 7 takes up.

The rest of Table 4 is largely empty, and the emptiness is a result. Sixty-eight of its 165 cells are values no parsed source stated, which is 68 of the 150 cells outside the engine-key column, or 45 percent. No engine paper states a default physics timestep. Three report one for a named experiment, and the timestep column reports those experiment settings. Isaac Gym's cell is its Shadow Hand step, from the only per-task timestep table any engine paper here publishes, which runs 1/120 s for Shadow Hand and Allegro, 1/200 s for ANYmal and TriFinger and 1/60 s for Franka (`isaacgym_2021`, Table 4). MuJoCo's 0.01 s is the 27-DoF humanoid test's step and ComFree-Sim's 0.002 s is its benchmark step, against a stated stability limit near 0.02 s. Four engines state a solver iteration count and seven ship any dexterous hand at all. The licence column answers a question a reader choosing an engine actually has, and thirteen of fifteen rows do not answer it.

Table 4. Simulators and physics engines

engine	contact model	solver	iters	diff.	GPU	dt s	penetration exposed	throughput	hands shipped	licence
brax_2021	rigid bodies in maximal coordinates; naive/quadratic-scaling pairwise collision detection (no broad-phase acceleration structure); no convex-decomposition or mesh-vs-primitive distinction described	not LCP-based; velocity-level collision updates with Baumgarte stabilization (inspired by the Tiny Differentiable Simulator); joints modeled as spring constraints rather than Featherstone-style articulated-body methods		yes	yes		no	millions of simulation steps/sec on a single accelerator (e.g. MuJoCo Ant equivalent on a single chip); scales to hundreds of millions of steps/sec distributed across a 4x2 TPUv3 and other TPU topologies; Ant on a TPUv3 8x8 estimated at ~hundreds of millions of steps/sec vs ~thousands of steps/sec for single-threaded CPU MuJoCo Gym (all figures qualitative/order-of-magnitude, no exact numeric table)	synthetic 4-fingered claw hand (Grasp environment, non-commercial, from-scratch morphology; not confirmed present in current repo snapshot)	
comfree_sim_2026	complementarity-free, analytical (closed-form) contact resolution in the dual cone of the Coulomb friction cone via a prediction-correction/impedance-style update, extended to a unified 6D model (normal, tangential, torsional, rolling friction)	no complementarity solve; a 4-kernel GPU pipeline (Algorithm 1: smooth-velocity prediction, per-contact/per-face dual-cone solve, generalized-impulse accumulation, velocity correction) computes the impulse in closed form per contact facet	none (closed-form, no per-step iterative solve)		yes	0.002	yes	AMD 32-core CPU + NVIDIA RTX 4090 GPU: ~3x faster (near-linear vs. MJWarp's superlinear) step time vs. contact count at 512 parallel envs; ~2x parallel-environment throughput vs. MJWarp at 256-4096 envs (Allegro-hand cube-grasping benchmark); on real LEAP-hand MPC (AMD 32-core CPU + RTX 4090), MPPI compute 13.9-28.2 ms vs. MJWarp's 34.0-68.8 ms (~2.4x average speedup)		
dojo_2022	hard-contact nonlinear complementarity problem (NCP) with an exact nonlinear (second-order-cone) friction cone, avoiding pyramidal/linearized approximation, to avoid interpenetration and creep artifacts	custom primal-dual interior-point solver (Algorithm 2), based on Mehrotra's predictor-corrector algorithm, extended for non-Euclidean quaternion variables, with cone-handling borrowed from CVXOPT; gradients via implicit differentiation (implicit function theorem) of the KKT/complementarity system, with central-	converges within 15 iterations for all three robots tested in the convergence study (Sec. V-A); default tolerances $r_{tol} = \kappa_{appa_tol} = 1e-5$	yes	no		yes	Intel Core i9-10885H, 32GB RAM, CPU only (no GPU support; described as future work). Table V, 1000-step forward+gradient wall-clock time at $dt=0.01s$: Humanoid - Dojo 1.750+/-0.135s vs MuJoCo 1.512+/-0.045s vs Drake 5.463+/-0.078s vs Brax 6.975+/-0.485s; Unitree A1 - Dojo 5.235+/-0.071s vs MuJoCo		

engine	contact model	solver	iters	diff.	GPU	dt s	penetration exposed	throughput	hands shipped	licence
		path parameter kappa trading gradient smoothness for physical accuracy						1.114+/-0.007s vs Drake 3.870+/-0.024s vs Brax 11.064+/-0.521s; Franka Panda - Dojo 1.159+/-0.077s vs MuJoCo 0.335+/-0.001s vs Drake 2.352+/-0.055s vs Brax 10.954+/-0.395s; Skydio X2 - Dojo 0.807+/-0.003s vs MuJoCo 0.047+/-0.002s vs Drake 0.571+/-0.011s vs Brax 7.953+/-0.484s. MuJoCo is fastest on every system.		
genesis_2024	Constraint-based rigid solver (equality/inequality constraints, plus an explicit noslip() post- pass, architecturally similar to MuJoCo's solver+noslip design) combined with multi- material coupling: FEM, MPM, particle- based PBD and SPH, a libuipc Incremental Potential Contact solver, an explicit inter-solver coupler, and a SAP (semi- analytic-primal-style) compliant-contact path used for grasp- coupling scenes.	ConstraintSolver with Newton-style line- search iterative resolve (linereach.py) and constraint-island decomposition for parallelism (island.py); multiple pluggable physics solvers (Rigid, FEM, MPM, PBD, SPH, libuipc, SAP) compiled via the Quadrants (Taichi- derived) backend.		yes	yes				Shadow Hand	
isaacgym_2021	Rigid-body contacts via PhysX; contact geometry can be primitive shapes or meshes loaded from URDF/MJCF (Sec. 2.2); net contact forces exposed per- rigid-body via the Tensor API (Table 1, Get-only). No explicit statement of convex- vs-mesh contact generation or per-pair contact-point count. Shadow Hand tendons simulated via PhysX Fixed Tendon mechanics:	Temporal Gauss- Seidel (TGS) solver (Sec. 3, ref. 18), not classic PGS/LCP: folds sub-stepping into a per-body accumulated velocity- delta buffer projected onto the constraint Jacobians. Table 3 exposes two user- tunable iteration knobs, Position iterations (biased, velocity+positional- error-correcting) and Velocity iterations (unbiased, velocity- error-only-correcting);		no	yes	0.0083333333333333 33	no	Shadow Hand: 150,000 parallel environment steps/sec at 16,384 environments, single NVIDIA A100 GPU (Sec. 5.3); Shadow Hand OpenAI reproduction (feed- forward): >20 consecutive successes in <1 hour on 1x A100 vs. 30 hours on OpenAI's cluster of 384x16-core CPUs (6144 cores total) + 8x NVIDIA V100 GPUs (Sec. 6.4.1).	Shadow Dexterous Hand, Allegro Hand, TriFinger (3-finger, 9- DoF manipulator; the paper itself says this is not a hand)	

engine	contact model	solver	iters	diff.	GPU	dt s	penetration exposed	throughput	hands shipped	licence
	spring+damping force proportional to deviation from rest length, propagated through a tendon-joint tree (Sec. 3, Appendix A.1).	no default numeric values for these are given in the main text.								
isaacsim_2025	Rigid contacts by default via PhysX 5 (SDF collisions, Featherstone articulation solver, Sec. 2.2); Factory assembly envs specifically use 'SDF-based contact generation, a contact reduction technique, and a Gauss-Seidel solver' (Sec. 6.4). ContactSensor exposes net normal force per body, optional filtered pairs, contact duration, average contact point, and a short history (Sec. 3.3.1); no contact-quality/penetration metric is defined (survey note).	NVIDIA PhysX 5's internal rigid-body/Featherstone-articulation solver, used as-is; Factory assembly envs specifically named as using a Gauss-Seidel solver with SDF-based contact generation and contact reduction (Sec. 6.4). No timestep, decimation, or solver-iteration values are stated anywhere in the paper for the general (dexterous) suite; the only note is qualitative, that the Digit humanoid 'requires a higher solver iteration count for stable simulation' (Fig. 14).		no	yes		no	DextrAH teacher task (state-based grasp-and-lift): over 900,000 FPS training throughput with 8 GPUs (RTX Pro 6000, the paper's only multi-GPU platform) and 16,384 environments (Sec. 4.1.1); Franka cabinet drawer task: over 1.6 million FPS at the same 8-GPU/16,384-environment setting. Single-GPU dexterous-suite numbers are not stated.	KUKA Allegro hand (first-party dexterous suite, Sec. 7.2.3), ShadowHand (third-party GraspQP grasp evaluation only, Sec. 6.4), AbilityHand (third-party GraspQP grasp evaluation only, Sec. 6.4), Fourier GR1-T2 dex hands, Unitree Inspire hand, Unitree trihand (XR teleop retargeters in code only)	
maniskill3_2024	PhysX (via SAPIEN) rigid-body contacts on GPU; no statement of convex-vs-mesh contact generation or contact-point count in the note. AllegroHandRightTouch reads FSR contact impulse (get_fsr_impulse) for touch sensing (App. VII-H, code).	PhysX (via SAPIEN), used as-is; no LCP/PGS/TGS/Newton solver-algorithm name given. Benchmark (cartpole) settings state 'Solver Position Iterations: 4 / Solver Velocity Iterations: 0' (App. XI-A), but per-task _default_sim_config values are not exported in the parsed source.	4 position / 0 velocity iterations (benchmarked cartpole configuration only, App. XI-A); per-task defaults not exported in the parsed source		yes			Up to 30,000+ FPS (RGBD+segmentation) on a single RTX 4090 GPU, environment count for that headline figure not stated; at 128 parallel environments with cameras, ManiSkill3 uses 3.5GB GPU memory vs. Isaac Lab's 14.1GB (Sec. III-B); PPO uses up to 4096 parallel environments for state-based training and 256-1024 for RGB training (App. IX-A), all on RTX 4090 (App. XI-A benchmark hardware).	Allegro Hand (incl. touch-sensor variant, AllegroHandRightTouch), Ability Hand, Inspire Hand, Delta 3-finger hand, TriFinger	Apache-2.0
mujoco_2012	soft, convex velocity-stepping contact; three interchangeable solvers replace the standard LCP-with-friction-pyramid	implicit-complementarity solver (customized non-smooth Newton method, most accurate); convex		no	no			up to ~400,000 dynamics evaluations/sec on a 12-physical-core machine (2x 6-core Intel X5860 3.33GHz,		

engine	contact model	solver	iters	diff.	GPU	dt s	penetration exposed	throughput	hands shipped	licence
	approach: implicit-complementarity, convex (kinetic-energy minimization with a soft non-penetration cost), and diagonal (mass-aware spring-damper)	solver (interior-point method used for reported timings; projected Newton/CG/Gauss-Seidel also mentioned as unverified-faster alternatives); diagonal solver (per-contact critically-damped spring-damper using the diagonal of the A matrix, fastest/least accurate)						24 threads via hyper-threading), 3D humanoid with 18 DOF and 6 active contacts (Abstract, Table 3)		
mujoco_convex_contact_2014	soft, convex, complementarity-free; a unified impulse vector covers joint dry friction, joint/tendon/distance limits, and frictional contacts together, with an elliptical friction cone whose dimensionality per contact is 2 (sliding), 3, or 5 (sliding+torsional+rolling)	GPJS (generalized projected Gauss-Seidel, described as the authors' own unpublished-at-the-time method handling cone and pyramid constraints) for the contact/limit impulse; a separate 'soft Gauss principle' derivation for equality/holonomic constraints; an analytically-invertible closed-form solver for a restricted circular-cone special case	5 and 50 (both tested, Fig. 3, Sec. VI-B)	no	no	0.01		single-core Intel i7-3930K (Windows 7): forward dynamics of a 27-dof humanoid with 10 contacts evaluated in 0.1 ms (100x real-time at a 10 ms timestep); inverse dynamics of the same humanoid under 100 impulses in 30 microseconds (10 microseconds with no impulses); no parallel/multi-environment throughput reported		
mujoco_warp_2025	MuJoCo's native solref/solimp-parameterized soft-constraint contact model ported to Warp — per-contact Jacobian construction (dense and sparse variants, with flex variants) and equality/limit/friction constraint rows; not a re-derived contact model	MuJoCo's default (implicit CCP-style) solver and Newton constraint solver (via a 'newton' flag threaded through equality/limit/friction kernels) are supported; legacy PGS solver and noslip post-processing pass are 'not yet supported' per the README		no	yes					
newton_2025	Solver-dependent: MuJoCo-style soft-constraint (solref/solimp-style) contact via SolverMuJoCo/MJWarp; Proximal-ADMM / Dual-Variational-Inequality contact resolution via SolverKamino; position-based-dynamics-style constraint projection	Multi-solver architecture: SolverMuJoCo (wraps MJWarp), SolverKamino (Proximal-ADMM + DVI), XPBD (rigid), VBD (cloth/deformables), Style3D (cloth), plus a separate IKSolver (LM and L-BFGS backends).			yes				Allegro Hand	

engine	contact model	solver	iters	diff.	GPU	dt s	penetration exposed	throughput	hands shipped	licence
	for the XPBD and VBD solvers.									
orbit_2023	PhysX SDK 5 signed-distance-field (SDF) collision checking for rigid bodies (handles non-convex geometry such as screw threads); FEM-based solver using a constraint-based stable neo-hookean material formulation for deformable bodies; PBD (position-based dynamics) solver for cloth	PhysX SDK 5's internal rigid-body and FEM solvers, used as-is; no LCP/PGS/TGS, Newton, or CG iteration counts given		no	yes		no	125,000 FPS physics-only ceiling (no env count stated); ~10x rigid-body throughput vs. CPU-vectorized frameworks (robosuite, ManiSkill2, IsaacGymEnvs) swept 256-4096 envs; ~3x deformable/cloth throughput vs. DEDO; 270 FPS aggregate for 10 cameras at 640x480; hardware: 16-core AMD Ryzen 5950X, 64GB RAM, NVIDIA RTX 3090 (Sec. V-E-b, VII)	Allegro hand	
pybullet_2016					yes					
raisim_2018										requires a valid license and activation key from the RaiSim Tech website; the repo README also states RaiSim itself is no longer supported in favor of a RaiSim2 repo
sapien_2020	Rigid-body contact with convex-decomposed collision meshes (PhysX 4.1); three joint systems offered (kinematic, dynamic, and PhysX articulation) trading control accuracy against speed.	PhysX 4.1's internal rigid-body solver, used as-is; no LCP/PGS/TGS name or iteration count given in the paper.			no			~5000 Hz engine, ~700 Hz OpenGL render; single-instance CPU-physics figure with no parallel environments, on a laptop with a 2.2 GHz Intel i7-8750 CPU and an Nvidia GeForce RTX 2070 GPU.		

15 rows; 69 of 165 cells (41%) are values no source stated. No cell is truncated. A value wider than its column is wrapped at a word boundary, so a cell that runs to several rendered lines is one value and not several.

4.3 The GPU-parallel turn

Isaac Gym set the pattern. Physics, observations, rewards and actions stay on the GPU, and PhysX resolves contacts with the Temporal Gauss-Seidel sweep described above. Its per-task timesteps are published, which is rare: the Shadow Hand runs a 1/120 s physics step under a 1/60 s control step, or 1/20 s in the OpenAI variant. The result that reorganised the field is that reproducing OpenAI's Shadow Hand cube reorientation took under an hour on one A100, against 30 hours on 6144 CPU cores and 8 V100s ([isaacgym_2021](#), Sec. 6.4.1). Thirty-six of the 112 method papers in this corpus run on it.

Orbit and Isaac Lab moved the stack to PhysX 5, and the dexterous offering is thinner than the predecessor's: the first-party suite is lifting, grasping and reorienting with the KUKA Allegro hand, while Shadow and Ability hands appear only through third-party grasp evaluation ([isaacgym_2025](#), Sec. 7.2.3). MJX and MuJoCo Playground took the other route, putting MuJoCo's own solver on the GPU through JAX at the cost described above. MuJoCo Warp ports the same solref and solimp soft-constraint rows to NVIDIA Warp, and its README states that the legacy PGS solver and the noslip pass are unsupported and that differentiability "is not yet available" ([mujoco_warp_2025](#)). GPU MuJoCo is not a differentiable MuJoCo. Genesis is the broadest, pairing a rigid constraint solver with FEM, MPM, PBD and SPH solvers and a differentiable backward path, and it ships a Shadow Hand. Its README quotes no throughput number at all, so the large FPS figures that circulate for Genesis trace to nothing in [genesis_2024](#).

Newton changes how these names should be read. It is explicitly multi-solver: SolverMuJoCo wraps MuJoCo Warp, SolverKamino is a proximal-ADMM and dual-variational-inequality contact solver, and XPBD and VBD do position-based constraint projection ([newton_2025](#)). Its contact model is solver-dependent, not a property of the engine, which is how Table 4 records it. Isaac Lab's code already exposes a `--physics newton_mjwarp` backend switch in its hands demo, and its roadmap announces Newton integration. An experiment reported as Isaac Lab may be running PhysX 5 with TGS, or MuJoCo's soft constraint rows under Warp, and those two make different contact errors. The name also fails to fix the physics inside one engine. MuJoCo's convex solver is a family, interior point or projected Newton, conjugate gradient or Gauss-Seidel ([mujoco_2012](#), Sec. II-D), and MuJoCo Warp supports neither PGS nor the noslip pass. PhysX 4 and PhysX 5 differ in whether non-convex rigid bodies get SDF collision, which is a contact-geometry difference between Isaac Gym and Isaac Lab under one vendor name ([orbit_2023](#), [isaacgym_2025](#)). Naming a simulator no longer names its physics. Papers should report the backend and the solver beside the framework.

4.4 Throughput, and why the reported numbers do not compare

Four headline figures measure four different quantities. Isaac Gym reports parallel environment steps per second with physics, observations and rewards on device: 150,000 for the Shadow Hand at 16,384 environments on one A100. Isaac Lab reports frames per second in training, which includes the learning update: over 900,000 for the DextrAH teacher task at 16,384 environments on eight RTX Pro 6000 GPUs, with no single-GPU dexterous number anywhere in the paper. ManiSkill3 reports "up to 30,000+ FPS" with RGBD and segmentation on one RTX 4090, measured over 1000 random actions with reward and termination removed from the timing, and does not state the environment count behind the figure ([maniskill3_2024](#)). Orbit reports a 125,000 FPS physics-only ceiling on an RTX 3090, also with no environment count ([orbit_2023](#), Sec. VII).

The one number reported cleanly enough to reuse is MuJoCo Playground's, in PPO steps per second on a single A100 over five seeds. LeapCubeReorient runs at $76,354 \pm 143$ and PandaRobotiqPushCube at $487,341 \pm 4,346$. Same hardware, same measurement, same codebase, a factor of 6.4 between two environments. It is not a factor from the task in isolation. Playground tunes solver iterations, line-search iterations, timestep and contact bounds per environment, with values as far apart as one and four solver iterations ([mujoco_playground_2025](#), Table III), and it does not print the two configurations side by side. That omission is the reporting failure this subsection is about. The same confound sits inside Isaac Gym's paper on one A100: 700,000 environment steps per second for Ant, 200,000 for Humanoid, 150,000 for the Shadow Hand. A training-loop FPS is also mostly not a measurement of physics. Playground breaks the fractional cost down on an RTX 4090:

for CartpoleBalance, physics is 0.02, rendering 0.06, inference 0.01 and the policy update 0.91, and the policy update still dominates on the Franka task.

Cross-framework comparisons add further free parameters. ManiSkill2's PickCube table takes the best result over 16 to 512 environments for each system (maniskill2_2023, Table 1a), giving ManiSkill2 with a render server 2487 ± 24 FPS at its optimum of 64 environments against Isaac Gym's 865 ± 35 at its optimum of 512. The env-count tuning is the second problem, not the first. The two systems are not the same kind of thing. ManiSkill2 runs rigid-body physics on CPU worker processes behind a shared GPU render server, and Isaac Gym runs physics on the GPU, a difference its own paper states plainly. The measured quantity includes 128×128 rendering at 500 Hz simulation and 20 Hz control for both, so this is a visual sample-collection loop rather than two physics engines. Its authors add the fidelity caveat themselves, and Playground is equally explicit that its cross-simulator plot borrows its Isaac Lab and ManiSkill3 numbers from the ManiSkill3 paper. Several sources give no number at all: MuJoCo Warp's README points to an external nightly dashboard, and Newton's and Genesis's READMEs contain no FPS or speedup anywhere.

The failure reaches past the engine papers into the methods. Of the 47 corpus method papers that name a GPU-batched simulator, 19 state no environment count anywhere, and Isaac Lab's own paper states no physics timestep, no decimation and no solver iteration count for any task. A throughput figure is interpretable only with the environment count, the GPU, the timesteps, the solver iteration budget, and a statement of whether rendering and the learning update sit inside the measurement. Almost nobody reports all five.

4.5 Tactile simulation and what it is calibrated against

The three tactile simulators in this corpus calibrate against three different things, and none of them is a manipulation outcome. TACTO is a rendering layer over a host engine, by default PyBullet's rigid contact model. It reads post-solve link poses and the engine's reported normal force and maps that force to gel-mesh deformation at the rendering level, so it contributes no contact physics of its own. Its only sim-to-real number is a tactile pose-estimation task, at 1.66 ± 0.16 mm with colour-jitter augmentation against 0.76 ± 0.07 mm for a model trained on 128 real datapoints (tacto_2020, Table II).

Taxim is example-based rather than simulated, with an optical model calibrated from 50 real indentations, and it beats TACTO on every optical-similarity metric against real images. Its marker-motion model is calibrated against an ANSYS FEM, and the errors are 1.00×10^{-2} mm real against FEM, 1.02×10^{-2} mm real against Taxim, and 3.96×10^{-3} mm FEM against Taxim. Taxim tracks its own reference more tightly than that reference tracks reality, so its residual is bounded below by the calibration target. The authors also state that they simulate quasi-static contact only, with slip left as future work (taxim_2021, Sec. V). Slip is what a hand needs.

Tactile Genesis inverts the survey's thread. It implements two penetration-depth backends on the collision geometry, an analytic SDF query and a BVH raycast, then binarises the result with Schmitt hysteresis at 5×10^{-4} m on and 2×10^{-4} m off. Here the penetration depth is the sensor signal, so the quantity every other engine treats as an error sets this sensor's operating point. Against a real GelSight it reports relative marker RMSE of 0.329 in dilation and 0.174 in shear, against HydroShear's 0.403 and 0.217, with each simulator tuned to match the real image first (tactile_genesis_2026, Fig. 2). Its own sim-to-real check is a matched success count rather than a fidelity measurement. The real XHand1 SDK reports a per-taxel raw pressure field as well as an aggregate contact pressure, but documents no taxel positions or response characteristics, so the raw field cannot be registered to the simulated probe layout and only the aggregate is comparable (App. C). The undocumented calibration is the barrier, not a missing signal.

4.6 The sim-to-real gap for hands

Most of what is written about this gap is attribution without measurement. PenSpin asserts that the pure physics gap "cannot be bridged by extensive domain randomization alone" while reporting no experiment that isolates it (penspin_2024). MuJoCo Playground attributes its LEAP hand failures to physical flex in low-cost hardware and names more accurate collision geometries as the fix, without measuring either. DeXtreme lists four candidate causes for its shortfall, including a malfunctioning Allegro thumb used in most trials, and disambiguates none by experiment (dextreme_2022).

What has been measured is narrower, and section 7 collects the transfer numbers. The one result that links a quantified model error to transfer is a humanoid recipe paper, which ranks its system identification runs by dynamics-model MSE and pairs each with real success over 10 trials: the lowest-MSE model grasps 8 out of 10, the median-MSE model 3 out of 10, the highest-MSE model 0 out of 10 ([humanoid_sim2real_recipe_2025](#), Table 1). Where a cause has been pinned down elsewhere it is usually perception or actuation, not contact. OpenAI's pose estimator has 3.12 mm error on rendered images and 9.27 ± 4.02 mm on 992 real ones, and PDDM reports a camera tracker with 5 mm average error and 20 ms latency as the unmodelled source in its real numbers ([pddm_2019](#), App. C).

The contact side stays unmeasured, and the one paper that looks at it is usually read backwards. DeXtreme's real-to-sim replay interpenetrated because the replayed poses carried the pose estimator's error, not because the contact model failed. The paper says so: "there is still some sim-to-real gap in pose estimation. This is manifested when we played back the real states in sim (real-to-sim) with physics enabled, which sometimes resulted in interpenetrations. Therefore, we were not able to easily calibrate physics parameters of the cube" ([dextreme_2022](#), Sec. 5). A replayed trajectory is a placement, so the overlap it shows bounds the state estimate rather than the physics. That is why it could not calibrate the cube, and it is why calibrating contact against hardware still has no worked example for a hand. The only direct measurement of simulator fidelity against hardware in this corpus is Dojo's, an average final-position gap of about 0.5 cm over 5 box-pushing trials. It is a parallel-jaw arm pushing a box, and there is no equivalent number for a hand.

5. How policies are trained

5.1 The design space

A dexterous policy is defined by what supervises it, and four sources are in use across the corpus's 112 method rows: a reward function supervises reinforcement learning, a human demonstration supervises imitation, a human reference trajectory supervises a physics-based tracker, which is imitation with a simulator in the loop, and nothing supervises a model-based planner, which is handed a cost and a model instead.


The labels do not partition the corpus. Of the 112 method rows, 53 carry the tag `RL`, 27 `BC`, 23 `distillation`, 18 `VLA`, 14 `teleop-system`, 14 `diffusion`, 10 `data-collection`, 8 `flow`, 6 `RL+demo`, 6 `trajopt`, 5 `MPC`, 4 `grasp-synthesis` and 2 `world-model`. The tags sum to far more than 112 because most methods published since 2024 sit on two branches at once. Figure 4 draws the tree and the cross-links. Table 7 is the row-by-row version of the same thing, and is the table to scan when looking for work comparable to your own.

What the deployed policy looks like once training is done matters more than the name a paper gives itself, and on that axis the field has converged hard. Section 5.7 shows why.

Figure 4. What supervises a dexterous policy

Four kinds of supervision, subdivided to the level at which the 112 method papers actually differ.

Every leaf count is the rows its stated predicate selects, recomputed from corpus/rows; a paper may appear more than once.



The branches are not exclusive, and the dashed curves on the right are the overlaps. Most work after 2024 trains with a reward in simulation and ships something shaped like an imitation policy, so it belongs to the first branch and the second at once. Eight further rows carry a trajectory-optimisation or MPC tag inside a learned pipeline and are counted on the branch that learns. The double-dashed smoothing edge is the one overlap that is a theorem rather than a pipeline: `pang\_global\_planning\_2022` proves that a policy gradient and an analytic log-barrier relaxation compute the same local model of contact.

5.2 Reinforcement learning

5.2.1 The standard recipe

Fifty-nine of the 112 method rows learn from a reward. Forty-eight name PPO as the algorithm, and the next most frequent, DAPG, appears four times. Forty-three run their own experiments in a GPU-parallel simulator from the Isaac family or Genesis, and 36 do both. Only 31 rows state an environment count, with a median of 8192 and a maximum of 64000. Twenty-three rows distil a privileged teacher into a deployable student, and 16 combine all three of PPO, a GPU simulator and distillation. That 16-row intersection is the recipe as it is actually practised.

The hand is almost always an Allegro. It appears in 35 of the 112 method rows and in 17 of the 31 reorientation rows, ahead of Shadow at 21 and 6 and LEAP at 12 and 4.

Privilege is handled two ways. `openai_dexterity_2018` gives the value network object and target orientation, joint angles, joint velocities and object velocities the policy never sees, with a footnote recording that current object orientation was left out of the policy inputs by accident, and `dextreme_2022` uses the same asymmetric critic with a 2048-unit LSTM against the actor's 1024. `hora_2022` instead compresses nine privileged object properties into an eight-dimensional vector regressed from 30 steps of proprioception, and `visual_dexterity_2022` distils twice, from a state teacher to a synthetic point cloud and then to a rendered one, for a fivefold speedup.

Domain randomisation is the part of the recipe with the least discipline. `openai_rubiks_cube_2019` made it adaptive, pushing each range boundary out when performance at that boundary exceeds 20 successes and pulling it in below 10, and `dextreme_2022` reproduced the mechanism with a 256-sample queue and 40 percent of environments dedicated to boundary evaluation. Both publish the discovered ranges. Below that standard, two shipped configurations disagree with their own papers about what was randomised: `hora_2022` states a joint-noise range of $U(0, 0.005)$ against a shipped `jointNoiseScale` of 0.02, in a commit its own README says is not the one that reproduces the paper, and `penspin_2024` zeroes the disturbance force its appendix describes. `dexpbt_2023`'s `randomize: False` is not a third case: it is the IsaacGymEnvs default and agrees with the paper's statement that randomisation was not used here.

5.2.2 Reward engineering

Table 5 puts the 21 in-hand reorientation methods against nine recurring term families and marks each cell by where the term was found: in the paper, in the code, in both, or in the code with every shipped configuration setting its weight to zero.

Table 5. Reward terms across in-hand reorientation methods

method	yr	goal or rotation tracking	object velocity	finger-object distance	hand-pose deviation	action rate or magnitude	torque, work or joint velocity	drop or failure	contact or force	success bonus	terms	code	paper/code mismatch
openai_dexterity_2018	2018	paper						paper		paper	3	no	
openai_rubiks_cube_2019	2019	paper						paper		paper	3	no	
dextreme_2022	2022	both				both	both	code (0)		both	6	yes	yes
hora_2022	2022	both	both		both		both				5	yes	yes
visual_dexterity_2022	2022	both		both	paper		both	both	paper	both	7	yes	yes
dexpbt_2023	2023	both		both		code	paper	code (0)		both	4	yes	yes
eureka_2023	2023	paper	paper	paper								yes	yes
robot_synesthesia_2023	2023	paper	paper	paper		paper	paper	paper			6	no	
rotateit_2023	2023	paper	paper		paper		paper				6		
rotating_with_out_seeing_2023	2023	paper	paper	paper			paper	paper			6		
anyrotate_2024	2024	paper	paper		paper		paper	paper	paper	paper	10	no	
demostart_2024	2024									paper	1	no	
dreureka_2024	2024	paper	paper		paper			paper			4	yes	yes
penspin_2024	2024	both	both		both	code (0)	both				7	yes	yes
dexndm_2025	2025	paper	paper		paper		paper			paper	7	no	
dexremoe_2025	2025	paper	paper			paper				paper	5	no	
dexteritygen_2025	2025											no	
force_grasp_sim2real_2026	2026	paper			paper	paper	paper	paper	paper	paper	7	no	
poise_2026	2026	paper					paper	paper	paper	paper	4	no	
teledexter_2026	2026	paper								paper		no	
viserdex_2026	2026	paper	paper			paper	paper	paper		paper	10	no	

21 rows; 97 of 189 term cells (51%) are families the method does not use. `paper` means the term is in the paper and either no code was released or it is absent from the released reward code; `code` means it is in the released code and not in the paper's stated reward; `both` means it is in both. Marks are read from papers/notes/, term by term; the per-method source section is in corpus/reward_matrix.json.

The families are not equally popular. Nineteen of 21 methods have a goal or rotation tracking term, which is the task. Thirteen penalise effort as torque, work or joint velocity, 13 pay a sparse success bonus, 11 penalise object velocity and 10 penalise a drop. Eight penalise deviation of the hand from a canonical grasp pose, a family the plan for this table did not anticipate and which had to be added. Six penalise action rate or magnitude, five reward closing the distance from fingertips to the object, and four carry a contact or force mark at all.

Three cells an earlier draft marked `code` print `code (0)` instead, because in each the term is in the released code with every shipped configuration setting its weight to zero: `dextreme_2022`'s `timeout_rew` and `dexpbt_2023`'s fall penalty through `fallPenalty: 0.0`, and `penspin_2024`'s `action_penalty_scale: 0.0`. Marking them `code` would tell a reader the code optimises something the paper does not state, and leaving them blank would hide a term that is in the file. The same mark carries the config key and value in `corpus/reward_matrix.json`. Section 5.8 says why the distinction matters.

Three of the four contact marks are the finding, not four. `anyrotate_2024` scores good and bad fingertip contacts, `poise_2026` rewards a friction-cone wrench margin, and `force_grasp_sim2real_2026` tracks a commanded grasp force. The fourth, `visual_dexterity_2022`, penalises the *object* touching the table, a task-shaping term against using the table as a third finger rather than a hand-object term, and its `pen_tb_contact` flag defaults to `False` with no shipped config setting it true. Three of 21 in-hand reorientation methods, then, put a hand-object contact or force quantity in the reward, and none puts interpenetration in it. `teledexter_2026` penalises interpenetration with a differentiable signed-distance term, but during offline reference construction, not in the policy's reward. Across all 112 method rows, 85 do not address penetration at all, 5 constrain it, 3 measure it, 3 penalise it and 16 say nothing either way. The physical quality of the contact is not something this literature optimises.

Term counts range from one to ten. `demostart_2024` gives a single binary terminal reward of 1 and lets a curriculum over demonstration states do the shaping. `anyrotate_2024` and `viserdex_2026` each publish ten weighted terms. `viserdex_2026` is the cleanest specification in the table, with every term, weight and equation in one appendix table, and a statement that the same weights are used for every object.

Weights are frequently unrecoverable. Three of the 21 rows give no numeric weight that survives PDF conversion, because the equation blocks are images, which is a limit of this survey and is recorded as such. `robot_synesthesia_2023` is a worse case. Its six coefficients are described only as "tuned hyper-parameters" and no number is printed anywhere in the paper, so the reward it reports cannot be reconstructed by anyone. `dexremoe_2025` names an angular-velocity penalty weight in prose with no entry in its hyperparameter table, and `dextrack_2025` names an affinity reward that is missing from its weight table.

5.2.3 Curricula, populations and machine-written rewards

Curricula in this literature relax physics or tighten tolerances. `dexpbt_2023` tightens a success tolerance from 0.075 to 0.01 by a factor of 0.9 every 3000 environment steps once three successes are logged, and `maniptrans_2025` starts at zero gravity and high friction and restores both while narrowing a fingertip threshold from 6 cm to 4 cm. `dexmachina_2025` attaches virtual spring-damper controllers to the object, lets them do the task at first, then decays their gains to zero. The two that ablate the curriculum report it is load-bearing: `rotateit_2023` sets its rotation-penalty weight to zero and ramps it to -0.1, because applying it from the start makes the policy learn only to hold the object still, and `viserdex_2026`, whose performance-based curriculum runs over regularisation weight, action latency and time between successes, fails completely without the regularisation component.

Population-based training appears once. `dexpbt_2023` runs populations of 8, 16 and 32 agents, splits them 30/40/30, mutates the middle and replaces the bottom with mutated copies of the top, each float hyperparameter multiplied or divided by a

factor drawn from $U(1.1, 1.5)$ with probability 0.2. It reports 30 hours on a single V100 for a five-billion-transition single-arm run, and 0.32 trillion environment steps for its largest population.

[eureka_2023](#) has GPT-4 write the reward function directly, constrained to return a total and a dictionary of named components, and feeds per-component statistics back between iterations. Its appendix gives an example whose reward correlates at -0.26 with the human-written one and still scores 1.45 on the human-normalised metric, the strongest argument in the corpus that hand-tuned weights are not a ceiling. [dreureka_2024](#) moves the safety requirement into the prompt instead of a penalty term, asking in words for a cube rotating at about 0.25 radians per second with fingers penalised for leaving their initial pose. Neither can be checked against a file. Eureka's generated rewards are written at run time to a gitignored directory and never committed, so the reward behind any reported number cannot be recovered from the repository, though its README documents the paper's budget of five iterations and sixteen samples as the default and marks the shipped one-iteration block a fast-test preset. DrEureka's repository contains only the locomotion and globe-walking trees, so its LEAP-hand reward has no code at all.

5.2.4 Where reinforcement learning stalls

Reinforcement learning stalls on objectives that are not learnable as stated. [anyrotate_2024](#) found the angular-velocity objective unlearnable in the multi-axis setting and replaced it with a moving keypoint target. [rotateit_2023](#) reports that its multi-axis policy "does not converge when training with only reinforcement learning" and adds an imitation loss against single-axis oracles. [eureka_2023](#) cannot spin a pen from scratch and needs a pretrain-then-finetune split, with the from-scratch ablation failing to complete one cycle. Section 5.4's contact-graph term exists because of the general shape of the problem: contact moves the object off the reference, so return goes down, so the policy learns not to touch the object.

It stalls on geometry too. [hora_2022](#) fails on objects under 4 cm across because the fingers collide with each other. What it does not stall on is publication: the real-trial counts behind these headlines are small enough that section 7.1 treats them as the reporting problem they are.

5.3 Learning from human data

5.3.1 Teleoperation and retargeting

Table 6 lists the 30 corpus rows that describe a system for getting human motion onto a robot hand, with the retargeting objective compressed to one clause each.

Table 6. Teleoperation and human-data systems

system	operator interface	robot hand	retargeting objective	latency	rig USD	data collected
dexpilot_2020	four RealSense D415 RGB-D cameras tracking a bare hand, markerless; a coloured glove is worn only offline to train the neural hand-pose priors	Wonik Robotics Allegro hand (16-DoF, 4x4-joint fingers), retrofitted with 4 SynTouch BioTac tactile sensors at the fingertips	minimise a weighted squared distance between human and Allegro fingertip task-space vectors, with switching weights that close primary-finger-to-thumb gaps and enforce a minimum finger separation, plus a regulariser toward the open hand, solved online with SLSQP	about one second, end to end (Sec. IV)		
dexmv_2021	human demonstrator recorded on video (cubic capture rig, two RealSense D435 cameras), no live teleoperation; hand pose (MANO) and object pose extracted offline from the recorded video	Adroit Hand	match human (MANO forward-kinematics) and robot task-space vectors (palm-to-fingertip and palm-to-mid-phalanx pairs) per frame via SLSQP with a temporal-consistency term, then fit a minimum-jerk joint trajectory and compute torques via inverse dynamics			700 traj, 7 h
dexvip_2022	no live operator; a consensus hand pose is mined offline from curated HowTo100M video frames via FrankMocap 3D hand-pose estimation and k-medoid clustering	Adroit Hand	deterministic 4-stage geometric mapping (not an optimization) from FrankMocap's 21-joint hand pose to Adroit's 30-DoF joint space, applied once offline to produce a single static consensus target pose per object category, used only as a reward target			
dime_2022	vision-based, single RGB camera, MediaPipe hand detector, no glove/headset/exoskeleton	Allegro Hand	map human fingertip 2D image locations directly to fixed-height robot fingertip 3D targets (no depth estimation), then solve per-finger inverse kinematics for joint angles	not stated; targets at 30 Hz, PD loop at 300 Hz		30 traj
holo_dex_2022	VR headset (Meta Quest 2), built-in 4-camera hand tracker, no glove/exoskeleton	Allegro Hand	direct joint-angle copy for index/middle/ring fingers; thumb fingertip position matched via inverse kinematics; pinky ignored (no cost function or optimization)	under 100 ms on a local network (Sec. IV-D)	399	
videodex_2022		LEAP Hand	hand poses mapped via a distilled MLP replicating a Robotic-Telekinesis-style fingertip/palm keypoint-vector energy function; wrist trajectory computed via a PnP + monocular SLAM + heuristic gravity-alignment + workspace-rescaling pipeline, both applied offline to build pretraining trajectories, not run online at deployment			965 traj
anyteleop_2023	vision-based, camera only (RGB or RGB-D, single or multi-camera), MediaPipe hand-keypoint detection, no glove/headset/exoskeleton	Allegro Hand	minimize the difference between human and robot fingertip keypoint vectors via forward kinematics, with joint-limit constraints and a temporal-smoothness penalty	26-35 ms hand pose, 9-10 ms retargeting (Table II); no end-to-end figure		
pgdm_2023		ShadowHand	for pre-grasps only: inverse kinematics fits the robot hand's fingertip positions to a single human hand pose frame near first contact, used solely to initialize RL exploration; full object trajectories			40 traj

system	operator interface	robot hand	retargeting objective	latency	rig USD	data collected
			are used directly as the RL goal with no hand-trajectory retargeting			
ace_teleop_2024	visual exoskeleton: 3D-printed bimanual 6-DoF-per-arm exoskeleton with wrist-mounted cameras for MediaPipe finger tracking; not a glove or VR headset	Ability Hand / Inspire Hand / parallel-jaw gripper (embodiment-dependent)	wrist/end-effector pose mapped via a scale-and-recenter IK transform (Normal/Mirror/Bimanual modes) from exoskeleton forward kinematics; finger joints retargeted via AnyTeleop's fingertip-keypoint-vector optimization with joint limits and temporal smoothing	27 Hz hand tracking; no end-to-end figure	600	
bidex_teleop_2024	Manus Meta motion-capture gloves for the fingers plus GELLO-style teacher arms for wrist and arm pose; no headset	LEAP Hand (16 DoF); LEAP Hand V2 (21 DoF) used for 'extreme dexterity' experiments	inverse kinematics from glove fingertip positions to LEAP fingertip targets, with the arm commanded directly in joint space from the teacher arm rather than in end-effector space	claimed low-latency, no number given	6000	50 traj
bunny_visionpro_2024	VR headset (Apple Vision Pro) hand/wrist tracking, plus a custom ERM haptic feedback device driven by fingertip FSR tactile sensors; no glove or exoskeleton	Ability Hand	minimize scaled human-robot fingertip keypoint-vector distance via forward kinematics with joint limits and temporal smoothness (hand), solved online via SQP with a reduced-dimension reformulation for loop joints; separately, a unified IK + singularity-avoidance + sphere-approximated self-collision objective for arm motion	3.43 ms retargeting, 15.93 ms motion control, >60 Hz overall (Table 1)		
cyberdemo_2024	vision-based teleoperation (single RealSense camera observing operator hand motion, real-time hand detection and retargeting via a cited third-party system), not a glove or headset	Allegro Hand	delegated entirely to the cited third-party teleoperation system; CyberDemo does not describe or modify the human-to-robot mapping itself			
dexcap_2024	wearable backpack rig: Rokoko EMF gloves, one SLAM camera per wrist and a chest-mounted RGB-D LiDAR; no robot in the loop during collection	LEAP Hand	fingertip-position inverse kinematics into the 16-dim LEAP joint space anchored by the mocap wrist pose, with the little finger discarded		4000	787 traj, 4.5 h
egomimic_2024	human data: head-worn Project Aria glasses (egocentric RGB + onboard SLAM/hand-tracking, no glove or exoskeleton); robot demonstrations: standard bimanual ViperX teleoperation	parallel-jaw gripper (not a dexterous hand; see note scope flag)	no per-finger kinematic retargeting (end effector is a 2-finger gripper); domain alignment via a shared per-timestep camera-centered reference frame, per-embodiment Gaussian normalization of proprioception/actions, and SAM2-based visual masking with a directional line overlay applied identically to human and robot streams		1000	2,150 traj, 4 h
hudor_2024	VR headset (Meta Quest 3) built-in hand tracking for fingertip 3D positions, ArUco-marker rigid-body calibration into the robot frame; demonstrator is in the same scene as the robot, not internet video	Allegro Hand	direct Cartesian-space correspondence: human fingertip 3D positions are treated as target Cartesian positions for the corresponding robot fingertips with no kinematic-chain optimization, converted to joint commands via a custom Jacobian-based gradient-descent IK			4 traj

system	operator interface	robot hand	retargeting objective	latency	rig USD	data collected
okami_2024	single human demonstrator recorded on a static RGB-D camera (Intel RealSense D435i), no live teleoperation and no glove/headset; single video per task	Inspire Hand (x2)	factorized retargeting run once per demonstration video: arm/shoulder-elbow-wrist IK (via the Pink library, wrist position dominating the objective) plus hand joint angles computed via dex-retargeting from SMPL-H hand keypoints; the resulting trajectory is affinely SE(3)-warped at test time to match newly localized object start/end poses, with fingers retargeted independently of arm warping			
open_television_2024	VR headset (Apple Vision Pro or Meta Quest 3) with active stereo video streamed back from a 2-3 DoF camera gimbal on the robot's neck	Inspire Robots hand (6 actuated DoF, 12 total); Fourier GR-1 embodiment instead uses a 1-DoF parallel-jaw gripper	dex-retargeting keypoint-vector optimisation over five wrist-to-fingertip and two thumb-to-finger vectors, solved with SLSQP, with end-effector position taken relative to the head and orientation absolute	60 Hz stereo round trip and 60 Hz control		
dexteritygen_2025	human teleoperator tracked via a Manus Glove (retargeted to the Allegro hand at 300Hz via an unreleased, confidential fast retargeting method) plus a Vive tracker for 6-DoF wrist pose driving the arm	Allegro Hand				
dexumi_2025	hand-specific 3D-printed wearable exoskeleton (per-robot-hand optimized joint-to-fingertip kinematics), joint encoders reading exoskeleton angles directly (no visual fingertip retargeting), a wrist-mounted camera (ARKit-tracked wrist pose), and matching tactile sensors installed on both exoskeleton and target robot hand	Inspire Hand (12 DoF, 6 active) and XHand (12 active DoF)	no visual/kinematic-chain retargeting at inference; a learned per-joint regression model maps exoskeleton encoder values directly to robot motor values (calibrated by visually overlaying exoskeleton and robot at matched motor values), with the exoskeleton's mechanical linkage itself optimized offline so its fingertip workspace matches the target robot hand's forward-kinematics workspace	per-sensor latency measured and compensated offline, no figure		1,355 traj
dexwild_2025	wearable, calibration-free rig: motion-capture glove + two palm-mounted stereo cameras (pinky-side and thumb-side) + wrist ArUco marker tracking, worn by untrained operators; no VR headset	LEAP Hand / LEAP Hand V2 Advanced	robot hand kinematics optimized to match observed human fingertip positions (fixed hyperparameters, no per-user tuning); a relative (delta) end-effector state-action representation removes the need for a global coordinate frame, and identical palm-camera mounting on human and robot hands makes RGB observations look the same across embodiments			9,290 traj, 46.2 h
doglove_2025	haptic force-feedback exoskeleton glove (21-DoF encoder-based motion capture, 5-DoF cable-driven force feedback, per-fingertip LRA haptic feedback), plus an HTC Vive Tracker for wrist localization; not vision-based	LEAP Hand	map glove-FK fingertip positions to robot fingertip targets via the third-party Mink differential-1K solver, with a single scalar hand-size scaling factor; no paper-original cost function	30 Hz minimum system rate (120 Hz mocap, 30 Hz haptics)	600	
egozero_2025	Project Aria smart glasses (egocentric fisheye RGB + SLAM	Franka Panda parallel-jaw gripper (not a dexterous hand; see note	no kinematic hand-to-robot retargeting; both human and robot			700 traj

system	operator interface	robot hand	retargeting objective	latency	rig USD	data collected
	cameras + onboard hand/camera pose tracking), no glove, no live teleoperation of a robot during data collection; at inference an iPhone provides the egocentric view instead of Aria	scope flag)	are represented in a shared morphology-agnostic 3D-point space so the policy transfers zero-shot without any cross-embodiment pose-mapping step			
geometric_retartargeting_2025	Manus glove for fingertip keypoints + HTC Vive tracker for wrist pose (real-world evaluation); the retargeting method itself is tracking-source-agnostic	Allegro Hand	learned per-finger neural network trained offline (3-5 minutes) to jointly satisfy motion-direction preservation, C-space coverage, flatness, pinch correspondence, and self-collision avoidance, replacing per-frame online optimization with a single 1kHz forward pass at inference	1 kHz retargeting inference, against 60-100 Hz for optimisation baselines		
h_rdt_2025	for pretraining: none (consumes EgoDex's pre-existing released 3D hand-pose annotations, no live capture by this paper); for fine-tuning: standard teleoperation of each robot platform, not described in detail	parallel-jaw / 2-jaw grippers at deployment (Aloha-Agilex, ARX5, Franka-Panda, UR5+UMI); no dexterous hand deployed, though pretraining action space includes bimanual fingertip positions from EgoDex	no explicit kinematic retargeting; the 48-D bimanual wrist-fingertip human action space is treated as a superset of most end-effector action spaces, and the human-to-robot correspondence is instead learned implicitly by reinitializing and fine-tuning the state/action adaptors and action decoder on paired robot data while keeping the pretrained vision/language backbone			338,000 traj, 829 h
humanoid_policy_human_policy_2025	consumer VR headsets (Apple Vision Pro built-in camera with ARKit hand/head tracking, or Meta Quest 3/Vision Pro with a 3D-printed ZED Mini stereo mount via an OpenTelevision web app), <\$700; robot-side demonstrations collected via standard humanoid teleoperation	Inspire Hand (x2, 5-fingered)	no offline kinematic retargeting network; human and robot share an identical 54-D state-action space via a bijective fingertip-to-fingertip correspondence (both are 5-fingered hands), bridged only by action-speed interpolation (human trajectories slowed ~4x) and instructing operators to minimize torso movement		700	
wm_dex_human_videos_2025		Allegro Hand	no explicit robot retargeting network; both human hands (MANO 21-keypoint) and robot end-effectors (Allegro fingertips via forward kinematics; parallel-jaw grippers approximated with dummy circular keypoints) are mapped into the same fingertip-keypoint-difference action representation, with the missing pinky filled by reusing the ring-finger-equivalent Allegro keypoints			829 h
dexteleop0_2026	VR headset (Meta Quest 3), egocentric hand/wrist tracking (26 tracked hand-joint transforms per hand), no external motion capture	Sharpa Wave (x2)	DexPilot-style vector matching of critical inter-joint/finger-to-palm target vectors for the hand, plus analytical IK on wrist pose for the arm, producing a raw joint command that a real-time QP then corrects to keep fingertip contact forces in a safe window and balance net object force/torque	30 Hz QP control cycle		
egoscale_2026	large-scale pretraining: in-the-wild egocentric video with off-the-shelf	Sharpa Wave (22-DoF); cross-embodiment target Unitree G1 (7-	human 21-keypoint (MANO-style) hand pose retargeted via an			20854 h

system	operator interface	robot hand	retargeting objective	latency	rig USD	data collected
	SLAM + hand-pose estimation, no dedicated capture rig; aligned mid-training: Vive trackers for 3D wrist pose + Manus gloves for 25-joint in-hand pose, same camera rig as the robot	DoF tri-finger hand)	optimization-based procedure enforcing joint limits and kinematic constraints into the 22-DoF Sharpa Wave hand joint space; wrist motion is represented as a shared relative SE(3) transform identical across human and robot data			
teledexter_2026	NOKOV motion-capture system tracking operator hand pose and object 6D pose in real time (not vision-based, not a glove)	SharpaWave (22-DoF, headline); also LeapHand (16-DoF)	two-stage geometry-aware retargeting: Stage 1 fingertip-vector alignment to human hand geometry; Stage 2 object-mesh-aware refinement combining a surface-contact term, a differentiable-SDF interpenetration penalty, a self-collision sphere term, and a temporal-smoothness term, producing the fingertip+object-pose subgoals the RL controller tracks			
unidex_2026	human-video side: no live capture rig, converts existing egocentric RGB-D datasets (H2O, HOI4D, HOT3D, TACO) via offline retargeting; fine-tuning demonstrations collected via OpenTeleVision + dex-retargeting on an Apple Vision Pro	Inspire, Leap, Shadow, Allegro, Ability, Oymotion, XHand, Wuji (8 hands in UniDex-Dataset); real-robot eval uses Inspire, Wuji, and Oymotion hands	two-stage human-in-the-loop kinematic retargeting: an automatic PyBullet multi-end-effector IK stage aligns a 6-DoF offset between a virtual human-hand-fixed base and the real robot base to match fingertip positions (respecting joint limits, with iterative mimic-joint correction), followed by an interactive GUI stage where a human manually corrects alignment; a separate visual-alignment step masks the human hand from the pointcloud and re-renders the retargeted robot-hand mesh into the scene			52,000 traj

30 rows; 60 of 210 cells (28%) are values no source stated. No cell is truncated. A value wider than its column is wrapped at a word boundary, so a cell that runs to several rendered lines is one value and not several.

Retargeting splits four ways. The dominant family matches task-space vectors between the human and robot hands. `dexpilot_2020` set the pattern with a weighted squared distance between fingertip vectors, switching weights of 1, 200 and 400 by whether a pair is within threshold, a minimum-separation constant of 3 cm between primary fingers, and a regulariser toward the open hand, solved online with SLSQP; `anyteleop_2023`, `open_television_2024`, `bunny_visionpro_2024` and `ace_teleop_2024` all use the same formulation. The second copies joint angles where the kinematics allow it, as `holo_dex_2022` does for index, middle and ring while solving the thumb by inverse kinematics and discarding the pinky the Allegro does not have. The third learns the map offline: `geometric_retargeting_2025` trains a per-finger network in three to five minutes against motion-direction, coverage, flatness, pinch and collision criteria, then runs at 1 kHz against the 60 to 100 Hz of online optimisation, raising configuration-space coverage from 38 to 90 percent and one-time grasp success from 55 to 87.5 percent. The fourth removes retargeting from the loop: `dexumi_2025` builds an exoskeleton whose fingertip workspace is optimised to match the target hand, then regresses encoder values straight to motor values.

The table's empty cells are the point. Only 12 of the 30 rows carry a latency cell at all, three of those give no number, and only 7 state a rig cost. Where costs are stated they are low and falling, from `holo_dex_2022`'s \$399 headset through the \$600 exoskeleton of `ace_teleop_2024` and the \$600 haptic glove of `doglove_2025` to the \$4,000 backpack of `dexcap_2024` and the \$6,000 rig of `bidex_teleop_2024`. Latency is not one quantity: `dexpilot_2020`'s "about one second" and `holo_dex_2022`'s "under 100 milliseconds" are end to end, while `geometric_retargeting_2025`'s 1 kHz and `anyteleop_2023`'s 9 to 10 ms retargeting are single stages and several cells give a control rate. The column mixes the two, so its spread is not a range. Only 13 rows report how many trajectories were collected, and only 7 report hours.

Where throughput is measured, human collection beats teleoperation by a wide margin. `humanoid_policy_human_policy_2025` times a grasp demonstration at 3.79 seconds for a bare human and 19.72 seconds through a teleoperated humanoid, and a pour at 4.81 against 37.31 seconds. It attributes the gap to retargeting latency and to the workspace of a 7-DoF arm rather than to anything about the hand.

5.3.2 Imitation architectures

Four architectures cover the corpus. Action chunking came from `aloha_act_2023`, which predicts a chunk of joint targets with a CVAE and combines overlapping chunks by exponentially weighted ensembling, reaching 80 to 90 percent on fine bimanual tasks from about 50 demonstrations each. Diffusion came from `diffusion_policy_2023`, which denoises an action sequence and reports a 46.9 percent average success improvement over LSTM-GMM, IBC and BET across 15 tasks. `dp3_2024` swaps the image encoder for a small point-cloud encoder and reports 74.4 percent against 59.8 percent across 72 simulated tasks, and 85.0 against 35.0 percent on four real tasks from 40 demonstrations each. Flow matching is the 2024-onward default for large models and carries 8 of the 112 rows. Discrete action tokens are the fourth, used by `openvla_2024` and `metis_2025`, and are the only one of the four that needs no continuous head at all.

Data generation is a separate lever and is underrated. `dexmimicgen_2024` turns 60 human source demonstrations into 21,000 simulated ones across 9 tasks and 3 embodiments, and a real Fourier GR1 with two Inspire hands reaches 90 percent from 40 generated demonstrations against 0 percent from the 4 source ones. `dex1b_2025` iterates optimisation, a CVAE proposal model and a simulator filter into roughly a billion synthetic grasps, and its CVAE baseline beats the prior best by 22 points.

5.3.3 Human video without a robot

Seven rows train a dexterous-hand policy from video with no teleoperation at any stage. `dexmv_2021` retargets 700 self-recorded demonstrations by matching palm-to-fingertip task-space vectors and estimating actions by inverse dynamics, `videodex_2022` mines 965 trajectories from EpicKitchens to pretrain a neural dynamic policy, and `dexvip_2022` goes further into the reward, clustering 715 curated HowTo100M frames into one consensus grasp pose per object category and adding it as a reward term. `okami_2024`, `human2sim2robot_2025` and `hudor_2024` work from a single video each, and `wm_dex_human_videos_2025` pretrains a world model on 829 hours of EgoDex.

5.3.4 The embodiment gap, and what closing it is worth

Three papers measure the gap directly rather than asserting it. `okami_2024` runs the same reference-plan and warping pipeline with and without human body and hand retargeting, and the version without it loses 75.0, 42.1, 82.0 and 74.0 points across four task and setting pairs. That is the largest clean embodiment-gap number in the corpus.

`humanoid_policy_human_policy_2025` ablates two mitigations separately on the same vertical-grasp task over 10 trials. The full method scores 4 out of 10. Removing the interpolation that slows human trajectories by a factor of four drops it to 1 out of 10. Removing the unified 54-dimensional state space while keeping the slowdown drops it to 0 out of 10, because the policy is handed a shortcut for telling the two embodiments apart. `egozero_2025` reports the sharpest version of the same effect. Removing 3D augmentation drops every one of its seven tasks to 0 out of 15, and substituting monocular depth for triangulated depth does the same, because the best metric depth models it tested carry more than 5 cm of error. It also quantifies what its hand-pose pipeline costs, at 1 to 2 cm of action-label error.

`objdex_2024` supplies the counterexample. It retargets only the wrist and lets reinforcement learning discover finger motion, and its ablation that adds a fingertip-matching reward "does not yield benefits and even leads to lower performance". Richer human correspondence is not uniformly better, and this is the one result in the corpus that says so with an ablation.

5.4 Tracking a human reference with physics

Twelve method rows carry the `track-human-ref` task family. Eight track a human hand on an object and are covered here; the other four sit at the edges, `human2sim2robot_2025` tracking only the object's trajectory, `dexman_2025` retargeting bimanual video onto a full humanoid, and `humanplus_2024` and `omnih2o_2024` tracking whole-body motion. A reference is retargeted onto the robot and a policy trained to make the simulated body follow it: the reference supplies the shaping reward engineering would otherwise have to invent, and the simulator the physical consistency pure imitation lacks.

`physhoi_2023` is the origin of the reward form. It multiplies a body term, an object term, an interaction-graph term and a contact-graph term, and reaches 95.4 percent success on GRAB against 27.0 percent for a DeepMimic baseline. The contact-graph term exists to stop the policy learning not to touch the object. `omnigrasp_2024` replaces the action space with a pretrained latent motion prior and reports 94.6 percent grasp success and 84.8 percent trajectory success on GRAB, while stating outright that it omits penetration metrics.

`dextrack_2025` adds an imitation loss to PPO and mines its own demonstrations through a homotopy search over easier neighbouring trajectories, reaching 46.70 and 65.48 percent on GRAB at loose and strict thresholds against 38.58 and 54.82 for the best baseline. It has a penetration-depth formula and applies it only to input references, never to its own rollouts, and presents tolerance of "severe hand-object penetrations" as robustness. `maniptrans_2025` freezes a generalist hand imitator and trains a per-task residual on top, reaching 58.1 percent single-hand and 39.5 percent bimanual success on OakInk-V2. Its stance on contact is to raise the friction coefficient above the real value rather than model skin deformation, and its real deployment is open-loop replay.

`dexmachina_2025` moves to Genesis with 12,000 environments and six hand URDFs, and drives the object with virtual controllers that decay to zero. Its re-implementation of `maniptrans_2025`'s curriculum does not beat no curriculum on its own setup, which is a direct disagreement between two papers worth quoting in both directions. `dexplore_2025` drops the explicit retargeting stage and treats the reference as soft guidance, with termination thresholds derived from the failure rate, reaching 87.7 percent on GRAB with an Inspire hand.

`toporetaret_2026` is the only corpus method that quantifies penetration as an evaluation quantity, reporting maximum depth and the fraction of frames past 2 mm, and constraining it during retargeting with a 1 mm soft tolerance and a 30 mm hard bound. It reaches 87.5 percent on its own pen-spinning set against 46.9 for the best baseline. It does not re-measure penetration after the tracking policy runs, so the property it constrains is a property of the reference and not of the behaviour.

That is the pattern across all eight, and it is the reference-versus-rollout split in its sharpest form. Penetration is handled at the reference, if at all, and never at the rollout. `objdex_2024` completes the set with the only real-robot numbers among them, from 100 percent on a microwave and a laptop down to 41.2 percent on a ketchup bottle over 20 trials each.

5.5 Generalist and vision-language-action policies

The finding is the size of the hand. Eighteen method rows carry the `VLA` tag. Fourteen of them settle whether the reported evaluation ran on a multi-fingered hand. Eight of those fourteen did, six report no hand result at all, and the remaining four never say, `groot_n16_2025` naming no end effector anywhere on its page. Section 8.4 counts the same eighteen the same way.

Five of the eight state a hand size, and that comparison has to be made in actuated degrees of freedom rather than joints, for the reason section 3 opens with. Four are six, the Inspire RH56DFX among them actuating six of its twelve joints, and one is twelve, `dexora_2026`'s XHAND. Their median is 6. Two of the four rows that never settle the hand question do state a size, and both are large: 21 for `gr_dexter_2025`'s ByteDexter V2 and 22 for `egoscale_2026`'s Sharpa Wave. Of the 59 reward-learning rows, 48 state a count, 38 of those are 16 or above, and their median is 16: an Allegro and a LEAP actuate 16, and a Shadow actuates 20 of its 24 joints. So no row that settles the question reaches the band the reinforcement-learning literature of section 5.2 works in, and the two rows that reach it on paper are the two that never say whether the hand was in the evaluation. An earlier version of this claim counted eleven rows as evaluating on a hand, which mixed the settled eight with rows the notes leave open.

The six with no hand result are parallel-jaw throughout (`pi0_2024`, whose released code encodes each ALOHA gripper as one scalar, `pi05_2025`, `pistar06_2025`, `openvla_2024`, `helix_2025`, which claims individual-finger control of a 35-DoF upper body with no success rate or trial count for any task, and `gemini_robotics_2025`, whose Apollo hand appears in a figure caption with no number attached). Its successor `gemini_robotics_15_2025` closed that gap with Apollo progress scores, not success rates, of 0.74 down to 0.62 across five generalisation axes. `dexvla_2025` is the borderline case the other way: a hand in one real-robot task family, grippers in the other five and in all 91 pretraining tasks.

The released code also drifts. `groot_n1_2025`'s repository is a later N1.7 generation with a different backbone, a different action horizon and embodiments the paper does not describe, and `pi05_2025`'s supports only the flow-matching head, not the hybrid recipe the paper describes.

5.6 Model-based control as the non-learning baseline

Three corpus methods solve dexterous tasks without learning a policy, and they are the control group the rest of this section lacks.

`pddm_2019` learns an ensemble of dynamics models and plans through it with an MPPI-style optimiser, replanning every step. It needs 1 to 2 hours of data in simulation and 2 to 4 hours on a real 24-DoF Shadow Hand. `mjpc_2022` removes the learned model too and samples ten rollouts per step through MuJoCo itself, planning in 1 to 20 milliseconds. It reorients a cube with a Shadow Hand in real time from scratch, and reports no success criterion, no trial count and no real-robot result.

`pang_global_planning_2022` is the most substantial of the three. It proves that the randomised smoothing implicit in reinforcement learning and an analytic log-barrier relaxation compute the same local linear model of contact, then uses the analytic version inside a trajectory optimiser and an RRT. Allegro in-hand rotation takes 19.59 seconds to optimise and its plate-pickup task 117.16 seconds of planning on a 16-core CPU, against the GPU-days of section 5.2, and it is the only method here that imposes non-penetration as a hard constraint. Its own limitation is the honest part: its 3D systems transfer to hardware far worse than its 2D ones, because the quasi-dynamic assumption breaks and planned grasps miss contacts under a second-order solver.

5.7 Hybrids, and where the field has converged

The recurring shape is reinforcement learning in simulation distilled into a policy that looks like an imitation policy, and it appears in four variants. The first distils a privileged teacher into a vision student inside one paper, which is `hora_2022`, `visual_dexterity_2022`, `rotateit_2023`, `robot_synesthesia_2023` and `viserdex_2026`. The second uses reinforcement learning as a demonstration factory: `dextrack_2025` mines demonstrations with per-trajectory RL and trains a generalist on them, `maniptrans_2025` does the same to build a 3.3K-episode dataset, and `dexteritygen_2025` trains a diffusion controller on ten billion simulated grasp-to-grasp transitions and projects a human teleoperator's coarse command

onto it, taking four reorientation tasks from 0 out of 20 under raw teleoperation to 12, 13, 10 and 9 out of 20. The third generates data without reinforcement learning at all, which is `dexmimicgen_2024`, `dex1b_2025` and `deximit_2026`; `dex1b_2025`'s row is classed as a dataset, so figure 4's leaf, which counts method rows, holds the other two. The fourth wraps an RL policy in a residual, which is `resdex_2024` at 88.8 percent over 3,200 objects in 12 GPU-hours.

The consequence is worth stating plainly. In the variants that dominate 2025 and 2026 work, the reward is no longer the objective of the deployed policy. It is the objective of the process that produced the deployed policy's training data. Reward engineering has moved upstream into data curation, and every reward-shaping pathology in section 5.2 now reaches the shipped policy through a dataset rather than a gradient.

5.8 What the released code says

Thirty-eight of the 112 method rows record a discrepancy between a paper and the code it released, and all 38 released code, so they sit inside the 62 rows that released anything. They are not one kind of thing. Ten are contradictions, where paper and code state different values or different terms. Thirteen are limits of this survey's own parse, where the body or config that would settle the question was never recovered and the row says so. Eight released code without the described component in it, four are version skew against a later repository, and three are a paper disagreeing with itself. The fourth version skew is `groot_n16_2025`, which ships a main branch one generation later than the checkpoint its paper describes.

Ten is the number to quote, eight at high confidence and two, `penspin_2024` and `omnih2o_2024`, held at medium pending a direct read of the code. Ten of 62 is 16 percent, bounded on both sides: a floor, because the census covers method rows only and `robopianist_2023`, whose row is a benchmark, sums five reward terms against the three its Table 2 documents; a ceiling, because seven accusations an earlier draft of this section made were withdrawn under adversarial review, each with its reason recorded in the accused row's `mismatch_review` field. Among the 21 reorientation methods of Table 5, seven released a repository: four disagree with their paper, one (`dreureka_2024`) ships no cube-rotation environment at all, one (`hora_2022`) is a later generation its own README flags, and one (`eureka_2023`) was a default-value question the same README settles.

The most consequential case is `physhoi_2023`. Its `compute_humanoid_reward` hardcodes the body position-velocity error and both object rotation errors to zero, with the real computation commented out beside them, and does so unconditionally rather than per dataset, while its Table 4 lists non-zero weights of 0.1 and 0.01 for those rotation terms on GRAB. The reward that produced the paper's numbers never tracked object orientation: a method presented as tracking a 6-DoF reference was, in the code that ran, tracking the object in position only, with body rotation and body rotation-velocity still live. `omnigrasp_2024`, in the same file family, keeps its object rotation term live and fails the other way, internally: `compute_pregrasp_reward_time` takes weights as arguments and then hardcodes `w_pos, w_rot = 0.9, 0.1`.

Zeroed terms recur, and are not the same failure. A term present and zeroed is worse than a term missing, because it survives a reader's check of the file, but only when the paper claims it. `dexpbt_2023` sums eight components against the paper's four and multiplies `hand_delta_penalty` by zero with the comment "currently disabled", a term the paper never claims, and its five named weights are present at their stated values. `pianomime_2024`'s Table 3 states two weighted terms while its environment sums roughly five unweighted ones, two of them inherited stubs returning zero and a third, forearm collision, the paper never lists. `penspin_2024` ships `forceScale: 0.0` against the disturbance force in its appendix, and its 96-dimensional observation carries no tactile channel, which is also what its proprioception-only student should carry, hence the medium confidence.

Weights drift. `dextreme_2022` states an action-delta penalty of -0.25 in Table 2 and ships -0.2 and -0.01 in its two DR yamls, neither matching. `visual_dexterity_2022`'s Eq. 8 penultimate-joint penalty is absent from `dexenv/envs/rewards.py`, and its two configs disagree about the fall distance. `unidexgrasp_2023` and `dexpoint_2022` ship rewards structured differently from their equations, and `pddm_2019`'s Baoding reward carries a -10 wrist-height term Table 2 omits.

In 13 rows the repository does not settle the question, which is this survey's limit and not an accusation. The reward code is C++ and outside the parse in `graspxl_2024`, whose configs expose four velocity coefficients where the paper prints two,

and in `artigrasp_2023`, whose two weight sets are the two phases of a curriculum the paper documents. `dextrack_2025` ships several unreconciled coefficient sets and which produced its headline table cannot be identified, and `maniptrans_2025`'s learning rate and environment count match its own config, the differing values being a README example's override and an unused fallback. `hora_2022` is the one repository that discloses its own gap, telling the reader to check out tag v0.0.1. Elsewhere a paper disagrees with itself: `aloha_act_2023`'s Algorithm 1 says MSE and its Section IV.C says L1, `dp3_2024`'s prose says the network predicts noise while its config sets `prediction_type: sample`, and `dexmachina_2025`'s multiplicative task reward, charged to its code in an earlier draft, is printed in the paper.

A reward table is a claim about a training run and the code is a claim about a repository. Here the two contradict each other in ten cases, in the other 28 the released artefacts do not settle the question, and in exactly one, `hora_2022`, the repository says so itself. Read the reward function before the reward table, and treat a printed weight as a hypothesis about the code.

5.9 The master table

Table 7's emptiest columns are the ones a reader most needs: only 31 rows state an environment count, only 70 state how many real trials are behind the headline number, and only 39 state how many unseen objects were tested. The trial and unseen-object figures are the audited ones, after section 7.1 recovered 15 trial counts and 7 unseen-object counts that the notes carried and the extraction had dropped. A mostly empty row is not a weak method, but it is one that cannot be compared with any other row here.

Table 7. Methods

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
ferrari_canny_1992	1992												no
dapg_2017	2017	grasp, reorient, functional/tool, other	RL+demo	NPG (Natural Policy Gradient) + BC pretraining (DAPG)	ADROIT	24	no	MuJoCo	no			not addressed	yes
openai_dexterity_2018	2018	reorient	RL	PPO	ShadowRobot Dexterous Hand (EDC electric-motor version)	24	no	MuJoCo	yes	10		not addressed	no
openai_rubiks_cube_2019	2019	reorient	RL, distillation	PPO with Automatic Domain Randomization (ADR); Policy Cloning (DAGGER/distillation-style) for architecture changes	Shadow Dexterous E Series Hand (E3M5R)	20	no	MuJoCo (+ ORRB/Unity3D for rendered vision training images)	yes	10		not addressed	no
pddm_2019	2019	reorient, other	MPC	PDDM: ensemble of feed-forward dynamics models + MPPI-style filtering/reward-weighted-refinement trajectory optimizer (no policy network, no policy-gradient training)	Shadow Hand (24-DoF, real and sim); D'Claw (9-DoF, sim only, valve turning)		no	MuJoCo	yes			not addressed	yes
dexpilot_2020	2020	reorient, grasp, functional/tool, other	teleop-system	DART model-based hand tracker bootstrapped by two learned neural priors (GloveNet ResNet-50 keypoint regressor; PointNet++-based keypoint network + JointNet FC mapping), online fingertip-vector kinematic retargeting solved via SLSQP (NLopt) warm-started each frame, Riemannian	Wonik Robotics Allegro hand (16-DoF, 4x4-joint fingers), retrofitted with 4 SynTouch BioTac tactile sensors at the fingertips	16	no		yes	150		not addressed	no

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				Motion Policies for arm motion generation									
castro_sap_contact_2021	2021	bimanual-coord		SAP (Semi-Analytic Primal) solver: a custom Newton-type solver on an unconstrained convex primal reformulation of compliant contact, with analytic gradients, line search, and sparsity exploitation	Allegro hand (16 DoF each), two	16	yes	Drake	no			measured	yes
dexmv_2021	2021	grasp, functional/tool	RL+demo, data-collection	DAPG / SOIL / GAIL+ on top of TRPO, using demonstrations translated from human video via TSV retargeting + inverse-dynamics action estimation	Adroit Hand	30	no	MuJoCo	no		500	not addressed	yes
dexpoint_2022	2022	grasp, other	RL	PPO	Allegro Hand	16	no	SAPIEN	yes	260	40	not addressed	yes
dextreme_2022	2022	reorient	RL	PPO (rl-games implementation, LSTM actor/critic, asymmetric/state-privileged critic)	Allegro Hand	16	no	Isaac Gym (PhysX)	yes	10		not addressed	yes
dexvip_2022	2022	grasp	RL	PPO with a video-mined consensus grasp-pose auxiliary reward (R_pose) combined with an affordance reward (R_aff) and a lift-success reward (R_succ)	Adroit Hand	30	no	MuJoCo	no			not addressed	no
dime_2022	2022	reorient, grasp	teleop-system, BC, RL+demo	INN/VINN (nearest-neighbor imitation) vs. BC; DAPG/PPO/BCRL (simulation RL-finetuning)	Allegro Hand		no	MuJoCo	yes	10		not addressed	yes
holo_dex_2022	2022	reorient, grasp, functional/tool	teleop-system, BC	VINN (BYOL nearest-neighbor) vs.	Allegro Hand	16	no		yes	10	10	not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				Behavior Cloning / BC-Rep									
hora_2022	2022	reorient	RL, distillation	PPO	Allegro Hand (Wonik Robotics)	16	no	IsaacGym	yes	240	30	not addressed	yes
mjpc_2022	2022	reorient	MPC	Predictive Sampling: zero-order, derivative-free sampling-based shooting MPC over a spline-parameterised control sequence; the framework (MJPC) also implements iLQG and Gradient Descent as alternative planners, but the Hand result is attributed to Predictive Sampling	Shadow Hand (MuJoCo Menagerie model)		no	MuJoCo	no			not addressed	yes
pang_global_planning_2022	2022	reorient, grasp, bimanual-coord, other	trajopt	iMPC (iLQR-style trajectory optimizer over a locally-smoothed convex quasi-dynamic contact model, CQDC) for local case studies; an RRT variant (RRT-through-contact) using a smoothed Mahalanobis distance metric, single-step dynamically-consistent extension, and contact sampling for global planning on harder tasks	Allegro Hand (full 3D model, in-hand rotation / pen placement / plate pickup / door opening); a 2-DoF-per-finger planar hand (2 fingers) for 2D reorientation tasks		yes	Custom Convex Quasi-Dynamic Differentiable Contact (CQDC) model implemented in Drake (MultibodyPlant/SceneGraph), validated against Drake's own high-fidelity second-order contact solver	yes			constrained	yes
videodex_2022	2022	grasp, reorient, functional/tool	BC, data-collection	two-stream Neural Dynamic Policy (NDP) trajectory regression (L1 loss); pretrained on human-video-retargeted trajectories, fine-tuned on real	LEAP Hand	16	no		yes			not addressed	no

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				teleoperated demonstrations; R3M visual encoder									
visual_dexterity_2022	2022	reorient	RL, distillation	PPO (teacher, privileged state) + DAgger imitation distillation to a depth-point-cloud student (two-stage: synthetic-PC pretrain then rendered-occluded-PC fine-tune)	D'Claw (open-source, low-cost; 3-finger and modified 4-finger versions)		no	Isaac Gym (PhysX)	yes	20	12	not addressed	yes
aloha_act_2023	2023	bimanual-coord	BC	ACT (Action Chunking with Transformers, CVAE)	parallel-jaw gripper (not dexterous), two		yes	MuJoCo	yes	25		not addressed	yes
anyteleop_2023	2023	grasp, functional/tool, handover	teleop-system, RL+demo	DAPG (downstream IL); keypoint-vector retargeting optimization (dex_retargeting)	Allegro Hand		no	SAPIEN	yes	100		not addressed	yes
artigrasp_2023	2023	grasp, functional/tool, bimanual-coord	RL	PPO (D-Grasp-based implementation)	MANO hand model, two	51	yes	RaiSim	no		1	not addressed	yes
dexdeform_2023	2023	bimanual-coord, other	BC, trajopt	latent skill model (VAE over demonstrations) for skill-planned trajectory initialization, refined by gradient-based trajectory optimization through a differentiable MPM simulator, with refined rollouts fed back as new demonstrations	Shadow Dexterous Hand (simulated)	28	yes	PlasticineLab (MLS-MPM, CUDA), built on Hu et al. 2018's Moving Least Squares Material Point Method	no			not addressed	yes
dexpbt_2023	2023	grasp, reorient, bimanual-coord	RL	PPO (rl_games) with decentralized Population-Based Training (PBT) over RL hyperparameters and reward-shaping coefficients	Allegro Hand	16	yes	Isaac Gym (PhysX)	no			not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
dexterous_functional_grasping_2023	2023	functional/tool	RL	PPO	LEAP hand	16	no	IsaacGym	yes	70	5		no
diffusion_policy_2023	2023	grasp, functional/tool, bimanual-coord, other	BC, diffusion	conditional DDPM over action sequences (CNN or transformer noise-prediction network), DDIM at inference			yes	task-specific: Robomimic (MuJoCo-backed), Push-T (custom 2D physics per IBC), BlockPush (per BET), Franka Kitchen (per Relay Policy Learning); no single engine/timestep stated for the combined benchmark	yes	20		not addressed	yes
dynamic_handover_2023	2023	handover, bimanual-coord	RL	MAPPO	Allegro Hand	16	yes	IsaacGym	yes	15	14	not addressed	no
eureka_2023	2023	reorient, bimanual-coord, other	RL	PPO with LLM-authored (GPT-4) evolutionary reward search	Shadow Hand (also Allegro Hand; bimanual pairs of Shadow Hands for Dexterity tasks)	20	yes	IsaacGym	yes			not addressed	yes
pgdm_2023	2023	grasp, functional/tool	RL	PPO (Stable-Baselines3) with an object-trajectory-tracking reward plus pre-grasp-based exploration initialization (PGDM)	ShadowHand	24	no	MuJoCo	yes			not addressed	no
physhoi_2023	2023	track-human-ref	RL	PPO	SMPL-X humanoid (simulated, not a robot hand)	90	no	Isaac Gym	no			not addressed	yes
robot_synesthesia_2023	2023	reorient	RL, distillation	PPO	Allegro Hand	16	no	Isaac Gym	yes	5		not addressed	no
rotateit_2023	2023	reorient	RL, distillation	PPO	AllegroHand (Wonik Robotics)	16	no	IsaacGym	yes		15	not addressed	
rotating_with_out_seeing_2023	2023	reorient	RL	PPO	Allegro Hand	16	no	IsaacGym	yes	30	5	not addressed	
unidexgrasp_2023	2023	grasp	grasp-synthesis, RL, distillation	PPO (teacher) + DAgger (student distillation)	ShadowHand	26	no	Isaac Gym	no		241	penalised	yes
unidexgrasp_p_2023	2023	grasp	RL, distillation	PPO + DAgger (with critic distillation)	Shadow Hand	24	no	Isaac Gym	no			not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
ace_teleop_2024	2024	grasp, functional/tool	teleop-system, BC	3D Diffusion Policy (DP3) for xArm platforms; ACT for H1 humanoid (downstream IL); IK-based scale/recenter end-effector mapping + AnyTeleop fingertip-vector hand retargeting	Ability Hand / Inspire Hand / parallel-jaw gripper (embodiment-dependent)		yes		yes	10	18	not addressed	no
anyrotate_2024	2024	reorient	RL, distillation	PPO	Allegro Hand	16	no	IsaacGym	yes		10	not addressed	no
asymdex_2024	2024	bimanual-coord, grasp	RL	PPO	Shadow Hand (30-DoF: 24-DoF hand + 6-DoF floating wrist) in simulation; 16-DoF Allegro hand + 6-DoF Ability Hand (mismatched, hardware-availability artifact) in real-world experiments	30	yes	NVIDIA Isaac Gym	yes	20		not addressed	yes
bidex_teleop_2024	2024	bimanual-coord, handover, functional/tool	teleop-system, BC	ACT (Action Chunking Transformer)	LEAP Hand (16 DoF); LEAP Hand V2 (21 DoF) used for 'extreme dexterity' experiments	16	yes	none	yes	20		not addressed	no
bidexhd_2024	2024	functional/tool, bimanual-coord	RL, distillation	IPPO (teacher) + DAgger (student distillation)	LEAP Hand, left and right	16	yes	Isaac Gym (PhysX)	no			not addressed	yes
bimangrasp_2024	2024	grasp, bimanual-coord	grasp-synthesis, diffusion	MALA (Metropolis-adjusted Langevin) energy optimisation + DDPM (BimanGrasp-DDPM)	Shadow Hand, pair	22	yes	Isaac Gym	no		225	penalised	no
bunny_visionpro_2024	2024	grasp, functional/tool, bimanual-coord	teleop-system, BC	ACT, Diffusion Policy, DP3 (downstream IL); SQP fingertip-vector hand retargeting + unified IK/singularity/coll	Ability Hand	24	yes		yes	10		not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				ision arm-motion optimization									
cyberdemo_2024	2024	grasp, reorient, functional/tool	BC, teleop-system, data-collection	ACT (Action Chunking with Transformers) BC policy, trained on sim teleoperation demonstrations augmented via Sensitivity-Aware Kinematics Augmentation, automatic curriculum learning, and action aggregation, fine-tuned on real demonstrations	Allegro Hand	16	no	SAPIEN	yes	20	2	not addressed	no
demostart_2024	2024	grasp, functional/tool, reorient, other	RL, distillation	MPO (teacher, auto-curriculum) + BC distillation (student, Perceiver-Actor-Critic)	DEX-EE Hand	12	no	MuJoCo	yes	100		not addressed	no
dexcap_2024	2024	grasp, bimanual-coord, functional/tool	BC, diffusion, data-collection	Diffusion Policy with Perceiver point-cloud encoder (DP-perc)	LEAP Hand	16	yes	none	yes	60	9		yes
dexmimicgen_2024	2024	bimanual-coord, grasp, functional/tool	data-collection, BC	MimicGen-style per-arm subtask transform/replay generation; downstream BC-RNN / BC-RNN-GMM / Diffusion Policy	Inspire dexterous hand (6-DoF, real-world GR1); unnamed 'dexterous hands' in the two simulated dexterous-hand embodiments	6	yes	RoboSuite (MuJoCo)	yes	20			yes
dp3_2024	2024	reorient, grasp, functional/tool, bimanual-coord, other	BC, diffusion	Diffusion Policy backbone (DDPM training / DDIM inference) conditioned on a compact 3D point-cloud feature (small MLP + max-pool encoder) instead of image features	Shadow Hand (Adroit, Bi-DexHands, DexDeform, DexMV domains); Allegro Hand (DexArt, HORA domains; also real Roll-Up/Dumpling/Drill); parallel gripper (MetaWorld; real Pour)		yes	MuJoCo (Adroit, MetaWorld), IsaacGym (Bi-DexHands, HORA), Sapien (DexArt, DexMV), PlasticineLab (DexDeform)	yes	40	5	not addressed	yes
dreureka_2024	2024	reorient	RL	PPO with LLM-authored (GPT-	LEAP hand	16	no	Isaac Gym	yes			not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				4) reward search + reward-aware physics prior domain randomization									
egomimic_2024	2024	grasp, functional/tool, bimanual-coord	BC	shared-backbone ACT variant co-trained on paired egocentric human video and teleoperated bimanual robot demonstrations, with dual pose-space and joint-space action heads	parallel-jaw gripper (not a dexterous hand; see note scope flag)	1	yes		yes	135		not addressed	yes
graspxl_2024	2024	grasp	RL	PPO	MANO (also Shadow Hand, Allegro Hand, Faive Hand)	45	no	RaiSim	no		503409	not addressed	yes
hato_visuotactile_2024	2024	handover, bimanual-coord	diffusion	Diffusion Policy (DDPM, CNN-based)	Psyonic Ability Hand, two	6	yes		yes	10		not addressed	yes
hudor_2024	2024	grasp, functional/tool	RL	open-loop IK replay of a single retargeted human fingertip trajectory plus online residual RL (DrQv2) trained with an object-centric point-tracking trajectory-matching reward	Allegro Hand	16	no		yes	10		not addressed	no
humanplus_2024	2024	track-human-ref, locomanipulation	RL, BC	PPO (HST shadowing policy) + decoder-only Transformer BC (HIT)	Inspire-Robots RH56DFX, two	6	yes	Isaac Gym-based legged_gym/rsi_rl	yes	10		not addressed	yes
objdex_2024	2024	track-human-ref, bimanual-coord	BC, RL, distillation	BC (Transformer high-level planner) + PPO (low-level controller)	Shadow Hand		yes		yes	80	1	not addressed	no
okami_2024	2024	grasp, functional/tool, bimanual-coord	trajopt, BC	single-video imitation: SLAHMR-extended SMPL-H body+hand reconstruction, GPT-4V/Grounded-SAM/Cutie	Inspire Hand (x2)	6	yes	RoboSuite (2 of 6 tasks only)	yes	12		not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				object-centric reference-plan construction, factorized arm IK (Pink) + dex-retargeting hand-joint mapping, SE(3) trajectory warping to new object poses at test time; optional downstream ACT BC policy trained on OKAMI-generated rollouts									
omnigrasp_2024	2024	grasp, track-human-ref	RL, distillation	PPO (hierarchical: PHC-X finger imitator → PULSE-X distilled latent prior → PPO over latent action space)	SMPL-X whole-body humanoid (simulated, not a robot hand)	90	yes	Isaac Gym	no		5	not addressed	yes
omnih2o_2024	2024	track-human-ref, locomanipulation	RL, distillation	PPO (teacher) + DAgger (student) + diffusion policy (autonomy layer)	Inspire hands (open-loop, DoF not stated)		yes		yes	20			yes
open_televisi on_2024	2024	bimanual-coord, functional/tool	teleop-system, BC	ACT (Action Chunking Transformer) with DinoV2 ViT backbone	Inspire Robots hand (6 actuated DoF, 12 total); Fourier GR-1 embodiment instead uses a 1-DoF parallel-jaw gripper	6	yes	none	yes	20		not addressed	yes
openvla_2024	2024	other	VLA	autoregressive next-token prediction over discretized (256-bin per-dimension, quantile-based) action tokens, single-step (no action chunking)			no	LIBERO benchmark (simulated Franka) for fine-tuning experiments; engine/version not stated	yes	230		not addressed	yes
penspin_2024	2024	reorient	RL, distillation	PPO + BC	Allegro Hand	16	no	Isaac Gym	yes	50	7	not addressed	yes
pi0_2024	2024	other	VLA, flow	conditional flow matching (linear-Gaussian/optimal-transport probability path), forward-Euler integration at inference (10 steps), shifted-			yes		yes	10		not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				Beta noisy-timestep sampling									
pianomime_2024	2024	music	RL+demo, distillation	PPO (per-song specialists); DDPM diffusion (generalist)	Shadow Hand E3M5	23	yes	ROBOPIANIST (MuJoCo)	no			not addressed	yes
rdt1b_2024	2024	bimanual-coord	diffusion	RDT (DiT + cross-attention, QKNorm/RMSNorm, MLP decoder, Alternating Condition Injection)	parallel-jaw gripper (ALOHA-style, Cobot Mobile ALOHA), two		yes		yes	139	2	not addressed	yes
resdex_2024	2024	grasp	RL, distillation	PPO + DAgger	ShadowHand	18	no	IsaacGym	no		241	not addressed	yes
twisting_lids_2024	2024	functional/tool, bimanual-coord	RL	PPO (asymmetric actor-critic)	Allegro Hand, two	16	yes	Isaac Gym	yes	20	15	not addressed	no
umi_2024	2024	functional/tool, bimanual-coord, other	BC, diffusion, data-collection, teleop-system	Diffusion Policy (unmodified loss/decoder); contribution is the hand-held data-collection hardware and the observation/action interface (relative-trajectory representation, inference-time latency matching)			yes		yes	260			yes
articulated_tools_inhand_2025	2025	functional/tool	RL, BC, distillation	PPO (oracle) + BC (student + CATFA online adaptation)	Inspire Hand	6	no	IsaacLab	yes	50		not addressed	no
being_h0_2025	2025	grasp, functional/tool	VLA, BC	autoregressive next-token prediction over discretized MANO motion tokens (Grouped Residual Quantization) for pretraining; L1-loss MLP regression head over learnable query tokens for post-training imitation learning	6-DoF Inspire hand mounted on a 7-DoF Franka Research 3 arm (real robot); MANO parametric hand model used as the pretraining representation (no physical hardware)	6	no		yes	20		not addressed	yes
clutterdexgrasp_2025	2025	grasp	RL, diffusion, distillation	PPO (teacher) + DP3 diffusion	AgiBot dexterous hand	6	no	Isaac Gym	yes	167	2029	constrained	no

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				policy (student)									
cross_embodiment_world_models_2025	2025	other	world-model, MPC	DPI-Net (graph neural network particle-dynamics world model) + sampling-based receding-horizon MPC	PSYONIC Ability Hand (6-DoF) and Robot Era XHand (12-DoF) for real-world deployment; trained across 6 simulated hands (Ability, Allegro, XHand, Leap, Shadow, forearm-less Shadow variant) for the embodiment-scaling study		no	SAPIEN (rigid-body Object Pushing data) and Rewarped (differentiable multiphysics, deformable Plasticine Reshaping data)	yes	20		not addressed	no
dexgraspvla_2025	2025	grasp, other	VLA, diffusion	hierarchical VLA: frozen VLM high-level planner (bounding-box affordance) + DIT diffusion-policy controller conditioned on frozen DINOv2 features and a SAM+Cutie object mask, trained end-to-end from scratch by imitation learning (MSE noise-prediction loss); Immiscible Diffusion noise assignment; DDIM sampling (16 of 50 training steps)	PsiBot G0-R, 6-DoF, single (right) hand	6	no		yes	1287	360	not addressed	yes
dexmachina_2025	2025	track-human-ref, bimanual-coord	RL	PPO (rl-games)	six URDFs: Inspire, Allegro, XHand, Schunk (main), Ability, DexRobot Dex Hand		yes	Genesis	no			constrained	yes
dexman_2025	2025	track-human-ref, bimanual-coord	RL, trajopt	PPO (RL_GAMES) residual policy on top of IK-retargeted human motion	Shadow Dexterous Hand		yes	NVIDIA Isaac Gym	no				no
dexndm_2025	2025	reorient	RL, distillation	PPO (oracle) + BC (generalist) + supervised residual policy	LEAP hand		no	Isaac Gym	yes			not addressed	no
dexplore_2025	2025	track-human-ref	RL, distillation	PPO (teacher) + DAgger-style	Inspire hand (also Allegro)	6	no	Isaac Gym	yes		2	not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				VAE distillation (student)	hand)								
dexremoe_2025	2025	reorient	RL	PPO	GX11 three-fingered dexterous hand (custom)	11	no	IsaacGym	no		50	not addressed	no
dexteritygen_2025	2025	reorient, functional/tool, grasp	RL, diffusion	RL-generated Anygrasp-to-Anygrasp transitions used to train a UNet DDPM diffusion foundation controller (DexGen) with gradient-guided diffusion sampling that projects an external (teleoperated) command onto a safe, high-likelihood learned action	Allegro Hand	16	no		yes	240			no
dexterous_handover_2025	2025	handover, grasp	RL	PPO	Allegro Hand		no	IsaacLab	no		3	not addressed	
dextrack_2025	2025	track-human-ref	RL+demo	PPO + IL action-supervision loss	Allegro hand (sim); LEAP hand (real)	16	no	Isaac Gym	yes			measured	yes
dexumi_2025	2025	grasp, functional/tool	BC, data-collection	diffusion policy (DDPM-style action-chunking) trained on exoskeleton-collected demonstrations converted to robot-hand-looking video via a SAM2 segment/inpaint/composite pipeline	Inspire Hand (12 DoF, 6 active) and XHand (12 active DoF)		no		yes	20			yes
dexvla_2025	2025	grasp, functional/tool, bimanual-coord	VLA, diffusion	ScaleDP (Scale Diffusion Policy, transformer-based diffusion action head, up to 1B parameters, multi-head per embodiment) plugged into a Qwen2-VL-2B VLA backbone, trained with a 3-stage 'Embodied	Robotiq parallel-jaw gripper (Franka rig a, Bimanual UR5e rig c); Inspire multi-fingered dexterous hand, 6-DoF hand joint space (Franka rig b); AgileX bimanual arm end-effector unnamed (rig d)			LIBERO (secondary, App. A.2 only; main results are real-robot)	yes	10	30	not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				Curriculum Learning' recipe; loss $L = L_{diff} + \alpha L_{ntp}$ ($\alpha=1$)									
dexwild_2025	2025	grasp, functional/tool, bimanual-coord	BC, data-collection	diffusion U-Net policy (also compared against ACT) co-trained on human wearable-rig demonstrations and teleoperated robot demonstrations via fixed-ratio batch sampling	LEAP Hand / LEAP Hand V2 Advanced	17	yes		yes		11		yes
dydexhandover_2025	2025	handover, bimanual-coord	RL	MAPPO (human-regularized, CTDE, with hybrid advantage estimation)	11-DoF (6 actuated) dexterous hand, per side (vendor not stated)	11	yes	NVIDIA Isaac Sim / Isaac Lab	no		9	not addressed	no
egozero_2025	2025	grasp, functional/tool	BC	closed-loop Transformer policy (BC, Gaussian NLL loss) over a morphology-agnostic 3D-point state-action space (triangulated object keypoints + thumb/index 3D coordinates + a thresholded grasp scalar)	Franka Panda parallel-jaw gripper (not a dexterous hand; see note scope flag)	1	no		yes	15		not addressed	yes
gemini_robotics_15_2025	2025	other	VLA		ALOHA parallel gripper (by platform convention, not re-specified); Bi-arm Franka parallel gripper; Apollo humanoid multi-finger dexterous hand (grasp types shown: single-finger push, cylindrical side grasp, medium wrap, five-finger power grasp, bimanual grasp/rotate; no DoF or vendor model given for		yes	MuJoCo (used to generate evaluation scenes at scale, not for training; engine timestep/contact model not stated)	yes			not addressed	no

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
					any embodiment)								
gemini_robotics_2025	2025	other, handover	VLA, distillation		ALOHA 2 parallel gripper, two fingers (primary embodiment); bi-arm Franka parallel gripper; Apollo humanoid five-fingered dexterous hand (no model/DoF given, evaluated only qualitatively)		yes		yes	20		not addressed	no
geometric_retargeting_2025	2025	grasp	teleop-system	per-finger MLP retargeting network (GeoRT) trained offline against five geometric losses: motion-direction preservation, C-space coverage (Chamfer), flatness, pinch correspondence, and a learned self-collision classifier	Allegro Hand		no		yes			not addressed	yes
gr_dexter_2025	2025	bimanual-coord, grasp	VLA	GR-Dexter (Mixture-of-Transformer VLA, flow matching + next-token prediction, following GR-3)	ByteDexter V2, two	21	yes		yes	125	23	not addressed	no
groot_n16_2025	2025	grasp, functional/tool, bimanual-coord, locomanipulation	VLA	flow-matching DiT action head (32 layers, per page and cross-confirmed by code changelog) fed by a Cosmos-2B VLM backbone (per page); predicts state-relative action chunks for most embodiments; no algorithm name beyond "DiT" is given			yes	Galaxea R1 Pro simulated on the BEHAVIOR suite; physics engine, timestep, and env count not stated	yes			not addressed	yes
groot_n1_2025	2025	grasp, functional/tool, bimanual-coord, handover	VLA, flow	dual-system architecture: Eagle-2 VLM (System 2,	Fourier dexterous hands, on the Fourier GR-1 humanoid;		yes	RoboCasa (Kitchen and GR-1 Tabletop benchmarks);	yes	10	5	not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				reasoning) feeding a DiT flow-matching action head (System 1, acting) via alternating cross-/self- attention; K=4- step forward- Euler denoising at inference; latent-action VQ- VAE for labeling action-less video	DoF/finger count not stated			DexMimicGen (cross- embodiment suite and sim pretraining-data generation); underlying physics engine and timestep not stated					
h_rdt_2025	2025	grasp, functional/tool, bimanual-coord	flow	2B-parameter flow-matching diffusion transformer (H- RDT), pretrained on EgoDex human hand- pose trajectories then fine-tuned per robot embodiment by reinitializing the state/action adaptors and action decoder	parallel-jaw / 2- jaw grippers at deployment (Aloha-Agilex, ARX5, Franka- Panda, UR5+UMI); no dexterous hand deployed, though pretraining action space includes bimanual fingertip positions from EgoDex		yes	RoboTwin 2.0	yes	25		not addressed	no
helix_2025	2025	grasp, functional/tool, handover	VLA	dual-network "System 1, System 2" VLA: S2 (7B VLM, 7- 9Hz) produces a continuous latent conditioning vector consumed by S1 (80M- param cross- attention encoder-decoder transformer, 200Hz); trained end-to-end with a "standard regression loss" (no diffusion/flow- matching, no formula given)			yes		yes			not addressed	no
human2sim2robot_2025	2025	track-human-ref	RL	PPO	Allegro Hand	16	no	IsaacGym	yes	70		not addressed	yes
humanoid_policy_human_policy_2025	2025	grasp, functional/tool	BC	Human Action Transformer (HAT), an ACT- style action- chunking transformer co-	Inspire Hand (x2, 5-fingered)	6	yes		yes	230		not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				trained jointly on unified human (PH2D) and teleoperated-robot demonstrations in a shared 54-D wrist-pose + fingertip state-action space									
humanoid_sim2 real_recipe_2025	2025	grasp, bimanual-coord, handover	RL, distillation	PPO (asymmetric actor-critic) for per-task specialist policies, distilled into a generalist via Diffusion Policy trained on filtered successful rollouts	Fourier GR1 hand (6 actuated + 5 underactuated DoF); cross-embodiment check also uses Inspire hand (6 actuated + 6 underactuated DoF)	6	yes	NVIDIA Isaac Gym	yes	10			no
maniptrans_2025	2025	track-human-ref, bimanual-coord	RL, BC, diffusion	PPO (two-stage: frozen generalist imitator + per-task residual policy)	Inspire Hand (12-DoF sim); also Shadow(22), MANO(22), Allegro(16)	12	yes	Isaac Gym	yes			not addressed	yes
metis_2025	2025	grasp, functional/tool, bimanual-coord	VLA	Autoregressive next-token cross-entropy over discretized 'motion-aware dynamics' tokens (VQ-VAE visual dynamics + RQ-VAE motion dynamics), plus a supervised continuous-action regression head (Action Decoder); combined loss $L = L_{ar} + \lambda L_{action}$	Inspire dexterous hand (paired, main real-robot eval); cross-embodiment eval uses 22-DoF SharpaWave Dexterous Hand	6	yes		yes	20		not addressed	no
pi05_2025	2025	other	VLA, flow	hierarchical high-level discrete-subtask prediction (FAST tokenizer, autoregressive) + low-level flow-matching action expert; two-stage schedule pretrain ($\alpha=0$, 280k steps) then post-train			yes		yes	40		not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				(alpha=10.0, 80k steps)									
pistar06_2025	2025	other	VLA, RL, flow	RECAP: advantage-conditioned policy extraction on a flow-matching VLA, using a distributional value function (cross-entropy over 201 discretized return bins) to compute an N-step advantage and a binarized improvement indicator, with classifier-free-guidance-style advantage-conditioning dropout (30%)			yes		yes	750		not addressed	yes
wm_dex_human_videos_2025	2025	grasp	world-model, MPC	DexWM: a deterministic latent-space (DINOv2) world model with a CDiT-based predictor conditioned on MANO fingertip-keypoint-difference actions, trained with a state-prediction loss plus an auxiliary hand-consistency heatmap loss; controlled via CEM/MPC planning in latent space	Allegro Hand		no	RoboCasa	yes	12			no
being_h05_2026	2026	bimanual-coord, other	VLA, flow	Rectified Flow (continuous action velocity-field prediction) + Masked Motion Token Prediction (discrete cross-entropy channel) inside a Mixture-of-Transformers / Mixture-of-Flow (MoF) action	Inspire Hand (Franka FR3); LinkerBot O6 (Unitree G1); unnamed 6-DoF dexterous hands on PND Adam-U and BeingBeyond D1; parallel gripper on	6		LIBERO, RoboCasa (engine/timestep not stated)	yes	20		not addressed	yes

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				expert with shared 'Foundation Expert' layers and routed Top-K specialized experts	LeRobot SO-101 (not dexterous)								
bidexgrasp_2026	2026	grasp, bimanual-coord	grasp-synthesis, diffusion	bi-level QP-ADMM optimisation (region-pair init + decoupled per-hand force closure) + DDPM generator	Shadow Hand pair (sim); Inspire and BrainCo hands (real)		yes	MuJoCo	yes	260	30	measured	no
deximit_2026	2026	bimanual-coord, grasp, functional/tool	data-collection, BC, grasp-synthesis, trajopt	4D hand-object reconstruction (ST2+FPose) + LLM-based (Qwen3-VL) subtask decomposition with Action-Centric Scheduling + BODex-style force-closure grasp synthesis and keyframe motion planning, feeding a 3D Diffusion Policy (DP3)	XHands (real-world deployment); exact simulation hand not explicitly named		yes		yes			constrained	no
dexora_2026	2026	grasp, functional/tool, bimanual-coord	VLA, diffusion, teleop-system	Decoder-only diffusion transformer (28 layers, hidden size 1024, 16 attention heads) with T5 (language) and SigLIP (vision) encoders feeding conditional tokens; DDPM training, DPMSolver++ sampling for inference; training clips reweighted by a discriminator trained with a positive-unlabeled objective (DWBC-style score-to-weight mapping) that	XHAND dexterous hand (paired), 12 fully actuated joints per hand, thumb and index additionally support lateral ab/adduction	12	yes	MuJoCo (digital twin of the real platform, used to generate the synthetic corpus and mirrored in real-time teleoperation)	yes	20		not addressed	

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				scores demonstration quality									
dexteleop0_2026	2026	grasp, functional/tool, bimanual-coord	teleop-system, MPC	box-constrained QP shared-autonomy residual controller (force tracking + multi-contact force-torque balance + nominal stabilization terms) on top of DexPilot-style vector retargeting	Sharpa Wave (x2)	22	yes	NVIDIA IsaacSim 4.5	yes	35			no
egoscale_2026	2026	grasp, functional/tool, bimanual-coord	flow, VLA	flow-based VLA (GR00T N1-style): VLM backbone + DiT action expert trained with a flow-matching objective, in a three-stage recipe (large-scale human-video pretraining, aligned human-robot mid-training, task-specific post-training)	Sharpa Wave (22-DoF); cross-embodiment target Unittree G1 (7-DoF tri-finger hand)	22	yes		yes	10			no
force_grasp_sim2real_2026	2026	grasp, reorient	RL	PPO (asymmetric actor-critic)	xHand	12	no	IsaacLab	yes	70	2		no
poise_2026	2026	reorient	RL	asymmetric PPO	Sharpa Wave hand	22	no	Isaac Lab	yes	10	1	not addressed	no
simtoolreal_2026	2026	functional/tool	RL	SAPG (PPO variant)	Sharpa five-fingered hand	22	no	IsaacGym	yes	120	12	not addressed	no
teledexter_2026	2026	reorient, functional/tool	RL, teleop-system	single-stage RL (SAPG) with consecutive subgoal co-tracking reward	SharpaWave (22-DoF, headline); also LeapHand (16-DoF)	22	no	Isaac Gym	yes	15		penalised	no
toporetaret_2026	2026	track-human-ref	trajopt, RL	constrained Laplacian-optimization retargeting + PPO tracking	Wuji Hand		no		yes			constrained	no
unidex_2026	2026	grasp, functional/tool	flow, VLA, data-collection	UniDex-VLA: a pi0-style flow-matching VLA with a Uni3D pointcloud	Inspire, Leap, Shadow, Allegro, Ability, Oymotion, XHand, Wuji (8 hands in UniDex-		no		yes	20	1		no

method	yr	task	paradigm	algorithm	hand	DoF	bi	sim	real	trials	unseen obj	penetration	code
				encoder and Gemma-based backbone, predicting actions in a unified Function-Actuator-Aligned Space (FAAS) across 8 dexterous hands	Dataset); real-robot eval uses Inspire, Wuji, and Oymotion hands								
viserdex_2026	2026	reorient	RL, distillation	PPO (RSL-RL)	Allegro Hand	16	no	Isaac Lab	yes	50		not addressed	no

112 rows; 224 of 1568 cells (14%) are values no source stated. No cell is truncated. A value wider than its column is wrapped at a word boundary, so a cell that runs to several rendered lines is one value and not several.

6. Bimanual dexterous manipulation

A warning first. Fifty-three corpus method rows record `bimanual: true`, and the flag says only that the robot has two end effectors. Thirteen of the 53 put no dexterous hand on the robot at all. `aloha_act_2023`, `rdt1b_2024`, `egomimic_2024`, `h_rdt_2025` and `umi_2024` are parallel-jaw throughout; `pi0_2024`, `pi05_2025`, `pistar06_2025` and `diffusion_policy_2023` name no hand and run grippers on every embodiment; `gemini_robotics_2025` and `gemini_robotics_15_2025` report every bimanual number on grippers and show five-fingered hands only qualitatively; and `helix_2025` and `groot_n16_2025` name no end effector anywhere on the page. Three more run a gripper on one embodiment and a hand on another: `ace_teleop_2024`, `dp3_2024` and `open_television_2024`, whose Unitree H1 carries 6-DoF Inspire hands and whose Fourier GR-1 carries a 1-DoF jaw.

The denominator for this section is 28: the method rows whose notes place a learned closed-loop controller on two multi-fingered hands. It is the 53 less those 16; less two static grasp-synthesis methods, `bimangrasp_2024` and `bidexgrasp_2026`; less three systems with no learned policy, `castro_sap_contact_2021`, `dexteleop0_2026` and `pang_global_planning_2022`; and less four rows where no learned policy holds both hands: `omnih2o_2024`, whose fingers are mapped open-loop from the Vision Pro and sit outside its 19-DoF policy; `okami_2024`, whose headline pipeline is open-loop retargeting with a learned policy only in a side experiment; `dexdeform_2023`, a skill model refined by trajectory optimisation rather than a closed-loop controller; and `omnigrasp_2024`, a simulated human body with no bimanual task. `dexterous_handover_2025` never enters, because its row records `bimanual: no`. Benchmarks and datasets are outside the 28 by class, `bidexhands_2022`, `bench2dex_2026` and `robopianist_2023` being benchmarks and `rplm_2024` and `humanoidgen_2025` datasets; they are quoted here as evidence and never counted. Every paper below runs two multi-fingered hands unless said otherwise.

6.1 Why two hands is not twice one hand

When both hands hold the same object, the object closes a kinematic loop between them. Neither hand can move without changing what the other must do. `bidexgrasp_2026` reports that the coupled objective "often yields imbalanced solutions, where one hand dominates stability while the other contributes marginally." `artigrasp_2023` reports simulation speed scaling "roughly quadratically with the number of contacts," and trains each hand alone before pairing them.

Role asymmetry is an assumption almost everyone makes silently. `bidexhd_2024` states it outright: "we assume the robot to be right-handed by default, i.e., the left hand handles the target object and the right hand handles the tool," and `twisting_lids_2024` bakes the same split into its reward, putting reference contact keypoints on the bottle base for the left fingertips and on the lid for the right. The bias is in the data first: `taco_2024` recruited 14 right-handed subjects and measures the right hand moving consistently faster.


The arms collide, and the corpus handles this by construction rather than by control. `robopianist_2023` ships a forearm-forearm collision term that its own reward table never lists. Two papers name arm collision as unsolved. `dexmimicgen_2024` "does not explicitly handle collisions," and `dexman_2025` says its control parameterisation "overlooks full arm posture," which it needs to avoid them.

The action space doubles and the observation grows faster. `bidexhands_2022` uses a 52-dimensional action for 18 of its 20 tasks against observations of 398 to 446 dimensions. `gr_dexter_2025` emits an 88-dimensional chunk of arm joints, end-effector poses, hand joints and fingertips. What the larger vector does not carry is any statement of who does what. The usual answer is one scalar reward summed over both sides, and the cost is a compounding failure rate. `maniptrans_2025` scores success only if both hands succeed, and its OakInk-V2 success rate falls from 58.1 percent on single-hand sequences to 39.5 percent on bimanual ones.

6.2 Coordination architectures

Figure 5. Four ways to control two dexterous hands

Counts are the 25 corpus papers that put a learned controller on two dexterous hands. Two are counted twice.



The field is concentrated in the first panel. The two on the right are where the structure of the problem is actually used, and asymdex_2024 is the only corpus paper that puts either one inside a policy.

Figure 5 sets the four architectures side by side. Every row of the 28 is assigned to exactly one of them, read from its note.

Twenty-one of the 28 put one policy over both hands, which is three quarters of the set: the trackers and tracking-adjacent methods `dexmachina_2025`, `dexman_2025`, `maniptrans_2025`, `objdex_2024` and `humanplus_2024`; the reinforcement-learning tasks `dexpbt_2023`, `twisting_lids_2024`, `eureka_2023`, `pianomime_2024` and `humanoid_sim2real_recipe_2025`; the demonstration pipelines `bidex_teleop_2024`, `dexcap_2024`, `dexwild_2025`, `dexmimicgen_2024`, `hato_visuotactile_2024` and `humanoid_policy_human_policy_2025`; and the generalist policies `gr_dexter_2025`, `groot_n1_2025`, `dexora_2026`, `metis_2025` and `egoscale_2026`. In every one of them a single network takes a concatenated two-hand observation and emits a two-hand action. Two rows do not say which they are, `bunny_visionpro_2024` and `deximit_2026`, whose notes describe the rig and the data pipeline but never the policy's own decomposition. The concentration is not the outcome of a comparison that was won.

Four of the 28 give each hand its own network, and the two papers that compare the choice disagree. `bidexhands_2022`, a benchmark row and so outside the 28, ships the MARL baselines and finds PPO over the full observation beats HAPPO and MAPPO "in most cases," with the gap narrowing on tasks that need both hands, because PPO "can use all observations" where MARL sees only part. `bidexhd_2024` concludes the opposite. Its independent PPO teachers score 74.59 percent stage-two tracking rate on trained tasks against 53.88 for a centralised policy over both observations. Both are simulation only, on different tasks and hands, so neither settles it. `artigrasp_2023` trains one PPO policy per hand, and `dynamic_handover_2023` and `dydexhandover_2025` use MAPPO with one agent per arm-hand system.

One of the 28 assigns explicit leader and follower roles, and the same one is the only policy expressed in a relative frame. `asymdex_2024` is that paper, and the only corpus paper that puts either mechanism inside a policy. It gives the dominant hand full finger and wrist control, restricts the facilitating hand to a 6-DoF base pose, and writes the dominant hand and the object in a frame attached to the object the facilitating hand holds. That cuts the observation from 176 dimensions to 88 and the action from 52 to 26. It ablates the two mechanisms separately, which no other corpus paper does. On Block in cup the full method scores 0.7701 over five seeds, against 0.1086 for relative frames without asymmetry and 0.0164 for asymmetry without them. Twist Lid transfers zero-shot at 18 of 20 real trials. The real rig pairs a 16-DoF Allegro with a 6-DoF Ability Hand because only one Allegro was available, so the roles are confounded with the hardware. Two other rows use a relative quantity without giving either hand a role: `dexwild_2025` appends the inter-hand pose to its observation, and `dexmimicgen_2024` applies one shared SE(3) transform to both arms' source segments when generating data, which is the mechanism offline rather than in a policy.

6.3 Benchmarks and datasets for two hands

Two purpose-built bimanual dexterous suites exist in the corpus, four years apart.

`bidexhands_2022` is 20 tasks on two Shadow Hands in Isaac Gym, ordered by the infant age at which humans acquire the skill, at 2048 environments and a reported 30,000-plus FPS. Its measurement discipline is weaker than its coverage. It reports reward and normalised score, never a success rate. The only success flag in the code is `goal_dist < 0.03`, a 3 cm object-to-goal test that ignores orientation and exists only in the four catching tasks. Any success rate later work attributes to Bi-DexHands comes from that flag or its own definition.

`bench2dex_2026` is the more instrumented of the two. It runs 26 long-horizon tasks in Isaac Lab across 12 arm-and-hand embodiments, with roughly 1.3K teleoperated demonstrations in eight modalities, including a ray-cast tactile image that maps all 12 hand geometries into one representation. Success needs the terminal predicate held for 0.5 s. Matched, GR00T N1.5 leads at 631 of 1300 rollouts, 48.5 percent, against ACT at 29.5, $\pi 0.5$ at 27.3 and Diffusion Policy at 12.9. Under combined perturbation those first two compress to 19.8 and 19.7. Tasks are not crossed with embodiments, so no contrast isolates the hand, and the authors state the hand and tactile models are not calibrated against matching physical hardware.

Ten corpus rows are bimanual datasets, and five are two-hand interaction capture. `arctic_2022` gives 339 mocap sequences of 11 one-DoF articulated objects held in both hands, and deliberately labels over-shooting vertices as in contact, so its ground truth records interpenetration as contact. `taco_2024` gives 2.5K bimanual tool-use sequences fitted markerlessly under an attraction loss and a penetration loss the paper says trade against each other. `oakink2_2024` is the only corpus dataset publishing a penetration number for its own annotations, a mean depth of 0.25 cm over its 627 long-horizon sequences. `hot3d_2024` gives 833 egocentric minutes with mocap-grade poses for both hands and up to six objects, and `gigahands_2024` gives 2,034 markerless minutes from 56 subjects, neither with a penetration metric. These five feed the tracking work here: `dexmachina_2025` and `artigrasp_2023` from ARCTIC, `bidexhd_2024` from TACO, `maniptrans_2025` from OakInk-V2 and `dexman_2025` from both.

6.4 Methods on the same axes

Table 7 carries the bimanual papers on the same columns as every other method here, and three of those columns are worth reading together.

Real trials. `asymdex_2024` reports 20 per task, `hato_visuotactile_2024` 10 per condition, `dexmimicgen_2024` 20 on a Fourier GR1 and `bidexgrasp_2026` 260 across 30 objects. `maniptrans_2025` reaches real Inspire hands by open-loop replay with no trial count and no success rate. Seven of the 28 have no real robot at all: `artigrasp_2023`, `bidexhd_2024`, `dexmachina_2025`, `dexman_2025`, `dexpbt_2023`, `dydexhandover_2025` and `pianomime_2024`. The two piano rows outside the denominator, `robopianist_2023` and `rp1m_2024`, are simulation-only as well.

Penetration. Most bimanual rows read "not addressed." `bimangrasp_2024` gates on it, failing any grasp whose total penetration across hand-object, self and inter-hand exceeds 1.5 mm, and reports that "penetration remains the primary cause of grasp failure." `bidexgrasp_2026` reports it as a number, 0.15 to 0.20 cm maximum depth for its own grasps. `deximit_2026` carries a named hand-hand penetration term in its grasp-synthesis objective. `dexmachina_2025` resolves it once, replaying retargeted joints against a fixed object "to eliminate object penetrations" and measuring nothing during rollout. `artigrasp_2023` declines the metric outright, because its baselines "include a physics simulation which exhibits no interpenetration."

Comparability. `bimangrasp_2024` reports 54.03 percent success in Isaac Gym at friction 3. `bidexgrasp_2026` re-runs the same grasps in MuJoCo at friction 0.6 and gets 26.80 percent, at 1.52 cm penetration depth. Neither number transfers, and section 5.4's re-implementation result is the same lesson on the training side.

6.5 Handover and in-hand transfer

Handover is the one bimanual task where the hands are unambiguously asymmetric, because one gives and one receives. All three corpus handover papers use a single shared reward across giver and receiver. None defines separate objectives for the two roles.

`dynamic_handover_2023` states one formula, $r = r_{\text{dis}} + r_{\text{linvel}} + r_{\text{torque}}$, and never says whether thrower and catcher receive different decompositions of it. `dydexhandover_2025` is explicit that they do not: "Both hands share aligned

objectives, forming a fully cooperative relationship," with one weighted sum feeding both agents. Only the thrower gets an extra KL regulariser toward a policy pretrained on human throws, which constrains style and not role.

The third case is weaker still. In `dexterous_handover_2025` only the receiver is learned. The giver is a UR5e with an Allegro hand that "holds the object without moving during the whole episode." There is one agent, one reward and no second policy, its row in Table 7 accordingly records `bimanual: no`, and it is one of the rows the 28 excludes. The 94 percent often attached to this paper needs its conditions. It is Total Success, which counts "Indetermination" cases where the simulator failed to resolve collisions and the object clipped through the giver's hand, on the short prism, in simulation, over 100 episodes, with no real robot in the paper.

Nothing in the handover literature measures contact quality. Contact appears only as a positive signal, a boolean per-phalange touch in `dexterous_handover_2025` and a boolean contact reward in `dydexhandover_2025`. The two systems that do handover well sidestep the problem. `hato_visuotactile_2024` reaches 10 of 10 on a real slippery handover by imitating teleoperation with fingertip touch sensing and no physics model. `humanoid_sim2real_recipe_2025` stages its reward with a discrete variable switching which hand's fingertips are scored, and still calls handover its hardest task at 52.5 percent real success.

6.6 What transfers from single-hand work

The training recipe transfers intact. PPO with thousands of parallel environments, an asymmetric critic on privileged state, domain randomisation and distillation into a vision policy is the same in `twisting_lids_2024` and `humanoid_sim2real_recipe_2025` as in section 5.2. Reward forms transfer literally, and the contact-goal reward $r_{\text{contact}} = \sum 1/(1+\alpha d)$ appears in both.

Multi-task learning does not transfer, and the bimanual benchmarks show it clearest. `bidexhands_2022` finds multi-task PPO and ProMP fail on MT4, MT20, ML4 and ML20. `robopianist_2023` finds F1 on the shared training song drops "from roughly 0.7 F1 for 1 song to almost 0 F1 for 16 songs," with no positive transfer "regardless of the size of the pre-training tasks."

Single-hand evaluation does not transfer either. Success criteria compound, because `maniptrans_2025` fails a trajectory "if either hand fails to meet these conditions." And object-centric metrics cannot say which hand caused the failure. `bidexhd_2024` scores the steps where both objects track their references, and `dexmachina_2025` averages ADD across object parts before an AUC. Both say whether the task worked and nothing about which hand lost it.

One failure mode has no single-hand counterpart. Two hands can penetrate each other, and the corpus almost never looks. Three papers carry an inter-hand penetration term, and all three are static grasp synthesis: `bimangrasp_2024`, `bidexgrasp_2026` and `deximit_2026`. Not one learned bimanual controller in the corpus measures or penalises hand-hand penetration during a rollout. `pianomime_2024` ships a flag that turns inter-hand collision off in the physics.


7. Evaluation: how we would compare these methods

7.1 What the field reports, and why the numbers do not compare

Of the 112 method rows in the corpus, 89 report a real-robot experiment, 22 do not and one row is unsettled: 79 percent of all 112 rows, or 80 percent of the 111 the note settled. Among those 89, 70 state how many real trials produced the headline number, which is 79 percent of them and 62 percent of all 112 method rows. The 89 is the denominator that belongs to this statistic: the 22 rows with no real robot cannot state a real trial count, and counting them as silent turns a definitional impossibility into a reporting failure. Thirty-nine rows state a count of unseen test objects, 35 percent. Ninety-eight state how a rollout is scored, 88 percent. Sixty-two released code and 46 did not, with four rows unsettled: 55 percent of all rows, or 57 percent of the 108 the note settled. Figure 6 draws these six shares, each against the denominator that belongs to it.

Figure 6. What the method literature reports

Share of the 112 surveyed method papers whose row records each item, each against its own denominator: the pale track is the rows the item can apply to, the grey tail the rows whose note did not settle it. A bar is evidence the item was stated, not that the work did it well.



Red marks an item fewer than a third of the rows it applies to state. Every bar is a floor: a hand audit of the nulls recovered 15 trial counts, 19 criteria and 7 object counts (§7.1).

These are counts of what this survey captured, not of what papers reported, and every one is a floor. A row in `corpus/rows` holds a scalar. A paper that reports ten trials on each of nine tasks, or a scoring rubric instead of a threshold, or a count spread over four tables, has nothing the extraction can reduce to one integer, so it produces a null, and a null is then indistinguishable from a paper that said nothing. The bias runs one way: every miss converts a reporting paper into a silent one, and the survey's argument is that the field reports badly, so the artefact flatters the argument. Section 5.8 makes the same disclosure about the paper/code count, subtracting the 13 disagreements that are limitations of this survey's own parsing before declaring which number to quote, and the coverage statistics above need it more.

So the nulls were audited by hand against the notes they came from, and the numbers above are the audited ones. Of the 34 method rows with a real robot and no trial count, 15 had the count written in their own note. `pi0_2024` at ten trials per task, `rdt1b_2024` at 139 across seven tasks, `umi_2024` at 260, `pistar06_2025` at 750, `gemini_robotics_2025` at twenty per task, and ten more. That is 44 percent of the audited nulls, and it moved the headline from 55 rows to 70, from 62 percent of real-robot papers to 79 percent, and the "never says" figure from 34 of 89 down to 19. The success criterion moved further: of 33 null rows, 19 do state a criterion, so the count rose from 79 to 98. The unseen-object count moved least, 7 recovered from 80 audited nulls, 32 to 39. The audit only counted cases where the note itself carried the number, so a count that the note also missed is still uncounted, and 19 trial counts, 8 unseen-object evaluations whose object count the note never gives, and 4 criteria remain genuinely unsettled. Those nulls are now a defensible claim rather than an artefact. Every bar in Figure 6 should still be read as a lower bound.

The remaining gap is the one that matters. Nineteen of the 89 papers with a real robot never say how many times they ran it. A percentage with no denominator cannot be given an interval, so it cannot be compared with anything.

Where the denominator is stated it is small, and it is not one quantity. Some stored counts are per-cell, meaning ten trials on each task, or twenty per condition, or five per object. Others are grand totals over every cell. The rows now carry a `real_trials_kind` beside every value, and the two distributions are quoted separately. The 39 per-cell counts run from 5 to

100 with a median of 15 and quartiles at 10 and 20. The modal cell is 10 trials, in 16 rows, then 20, in 12. The 24 grand totals run from 12 to 750 with a median of 110. Pooling the two gives a median of 20 and a range of 5 to 1287, and that pooled figure is the one an earlier draft of this section quoted. It describes nothing, because `hora_2022`'s 240 and `visual_dexterity_2022`'s 20 are experiments of comparable size recorded on different bases. Seven further counts could not be assigned a basis at all.

The per-cell count is the one a Wilson interval attaches to. At 20 trials a reported 60 percent carries a 95 percent Wilson interval of 39 to 78 percent, and a reported 80 percent an interval of 58 to 92 percent. Overlapping intervals are not a test, so the point is made with one: 12 of 20 against 16 of 20 is $z = 1.38$, $p = 0.17$, and the difference is not established. Two methods separated by 20 points at 20 trials, the larger of the two modal cell sizes, are not separated at all.

The share of papers stating a count has risen, from 39 of the 65 rows before 2025, 60 percent, to 31 of the 47 rows from 2025 and 2026, 66 percent. The counts themselves have not. The median stated count is 20 in 2024, 22.5 in 2025 and 20 in 2026. Before 2024 the per-year medians rest on one to seven observations and should not be read as a trend.

The denominators also sit on different hardware. The 103 method rows that name their own hand give 78 distinct hand strings between them, and this survey applies no normalisation to those strings, so 78 is a count of strings and not of hand designs. Matching on the string, Allegro appears in 35, Shadow in 21, Inspire in 19 and LEAP in 12. A success rate on a 16-degree-of-freedom Allegro and a success rate on a 6-actuator Inspire hand are not measurements of the same thing.

Nor are the criteria, and the audit changed what that sentence can claim. `dexverse_2026` counts PickCube a success when the cube is "lifted at least 0.20 m above its resetting height". `bench2dex_2026` requires its terminal predicate to hold for a continuous dwell time of 0.5 s, to reject transient contacts. `colosseum_2024` counts an episode successful "if the model completes the task fully". `dextrack_2025` reports every success rate as a pair under two threshold sets, which on GRAB gives 46.70 and 65.48 percent for the same rollouts. Of the 14 rows that still state no criterion, ten have no success predicate at all. They report radians rotated or time-to-fall and never define a success, which is a fact about the paper rather than a gap in this survey. And four are unsettled by the note. What the audit found in the other 19 was mostly not a threshold: eleven score by rubric or staged partial credit, five judge binary completion against a task description by eye, two defer to a benchmark's own definition, and exactly one, `pistar06_2025`, states a verbatim numeric threshold. A rubric is a milder failure than silence and a worse one than a threshold, because it is reproducible inside a lab and not across two.

7.2 The axes that matter

Seven quantities dissociate in the published data, so they have to be reported separately.

Task success. Binary success discards the difference between near-misses and inaction. `beyond_binary_success_2026` puts it plainly: "a policy that completes 90% of the task is clearly better than a policy that is frozen the whole time, yet their success rates would be identically 0%". In `kress_gazit_policy_eval_2024` policy C scores 17 percent overall on the pancake task while picking up the spatula and flipping the pancake in 23 of 23 attempts.

Robustness to perturbation. `colosseum_2024` measures a 30 to 50 percent success drop under single perturbation factors and at least 75 percent under all 14 together. In `bench2dex_2026` GR00T N1.5 leads the matched condition at 48.5 percent and falls to 19.8 percent under combined shift, while $\pi 0.5$ goes from 27.3 to 19.7 and retains the most at 72.1 percent. The ranking at the anchor is not the ranking under shift.

Generalisation to unseen objects. Only 39 rows state a count and the median is 11 objects.

Physical plausibility of the contact. Eleven of the 96 rows whose contact handling the note settled address it, 11 percent, with 16 rows unknown. Section 7.3 takes them apart.

Sample and wall-clock cost. Thirty-one of 112 rows state a parallel environment count and 18 a simulated episode count. `robopianist_2023` is the exception, at 5 million samples per song and roughly 5 hours per run on four Tesla K80 GPUs.

Real-robot transfer. Simulated rank order is not real rank order. `autoeval_2025` scores Open- $\pi 0$ on put-eggplant-in-sink at 6 of 50 in SIMPLER and 47 of 50 on the real WidowX. `suresim_2025` states the limit directly: "the simulation-to-real gap

precludes rigorous statistical inferences about real-world outcomes from simulation results alone".

Reproducibility. 62 rows released code that could be parsed against the paper, and 38 of the 112 rows record a disagreement of some kind between the paper and that code. All 38 released code, so the raw rate among code-releasing rows is 61 percent. That raw rate is not the finding, because the 38 are not one thing. Section 8.1 classifies them: 10 contradictions, 13 limitations of this survey's own parsing, 8 components never released, 4 version skews and 3 inconsistencies internal to a paper. Only the contradictions are a finding about the work rather than about this survey, so 10 of 62 code-releasing rows, which is 16 percent, is the figure this section and Table 8 use. `physhoi_2023` is the clearest of the 10. It lists a non-zero object-orientation weight for GRAB in Table 4, and its released `compute_humanoid_reward` hard-sets that orientation error to zero, so the reward that produced the published numbers tracked the object in position only.

7.3 Physical plausibility as a first-class metric

Eleven method rows handle interpenetration in any form: three penalise it, three measure it, five constrain it. The denominator is 96, not 112, because the `penetration` field is null for 16 rows, and a null there means the note did not settle the question, not that the paper ignored penetration. Eleven of 96 is 11 percent.

Where in the pipeline those eleven act is the reference-versus-rollout split that section 1 takes from `zhao_dexhand_survey_2026`. A reference is a pose or a trajectory scored before execution, and a rollout is what the trained policy actually did. What this section supplies on that axis is a measurement method and a count, and it does not supply a threshold. The count is the eleven of 96 above, with four closed-loop policies inside it and none reporting a number for its own rollouts. The method is the plausibility row of Table 8: maximum and mean penetration depth over the evaluation rollouts, on a dense surface sample, computed by code that never entered the reward or the termination rule. The threshold is borrowed, and section 7.7 says from where and why it does not bind. Six of the eleven are grasp synthesisers or trajectory optimisers, namely `bidexgrasp_2026`, `bimangrasp_2024`, `deximit_2026`, `pang_global_planning_2022`, `toporetarget_2026` and `unidexgrasp_2023`, and `castro_sap_contact_2021` is a contact model rather than a controller. That leaves four closed-loop policies in the whole corpus: `clutterdexgrasp_2025`, `dexmachina_2025`, `dextrack_2025` and `teledexter_2026`.

The reason the number is four is a measurement trap. A quantity a policy optimises cannot also judge it, because the policy learns the measure rather than the property the measure stands for. `physhoi_2023` documents the failure in the clean direction. Its kinematic imitation reward was maximised by not touching the object at all, because contact perturbed the reference trajectory: "the policy may learn not to touch the object and falls into a local optimal". A contact-graph reward closed the hole, and that reward reads a force threshold rather than geometry, so it constrains contact presence and says nothing about penetration depth.

`dextrack_2025` shows the trap in the other direction. It defines a maximum hand-object penetration depth over all frames in its Appendix B, and applies it only to the input kinematic references, as one component of a reference-quality score. No penetration number appears for any of its own rollouts, in simulation or on the LEAP hand. Tolerance of the failure is then reported as a result: "Despite severe hand-object penetrations in Figure 4c and Figure 4a, the hand still interacts effectively with the object, highlighting the resilience of our tracking controller".

`toporetarget_2026` is the strongest case in the corpus and still stops one step short on the same reference-versus-rollout line. It constrains penetration during retargeting with a 1 mm soft tolerance and a 30 mm hard bound, and it reports two numbers on 25 ContactPose grasps: a maximum penetration of 1.07 mm and 0.00 percent of frames above 2 mm, against 22.22 mm and 96 percent of frames for its GeoRT baseline. Then a PPO controller tracks those references, and its four reward terms and its five termination criteria govern object pose, link position, joint error and action smoothness, never penetration. The constrained quantity is the reference, and the rollout is not re-measured.

Definitions are not shared either. `grab_2020` estimates contact by proximity, because "contact cannot be directly observed", with a 4.5 mm tolerance, and reports that "'Use' grasps have 3.25 ± 0.68 mm average penetration", without saying whether 3.25 mm is a maximum, a mean or a median. `oakink_2022` supplies the fullest published vocabulary: penetration depth, solid intersection volume and simulation displacement. Its Table 3 scores the GRAB GrabNet split at 2.53 cm penetration

depth, against GRAB's own 3.25 mm. The two differ by a factor of about eight, and neither source states its distance function precisely enough to reconcile them, and the two are not scored over the same grasps either, since GRAB's figure is over its own captured "use" grasps and OakInk's is a model's output on the GrabNet split, so a difference of population and a difference of definition are confounded in the same ratio.

The analytic tradition scored a grasp without simulating it, and `ferrari_canny_1992` and `roa_suarez_grasp_quality_2015` are its wrench-space reference points. Neither could be obtained. The first DOI fetch returned HTTP 202 and the second a Springer JavaScript interstitial, so both are cited by metadata only and no definition here rests on their contents. The learned literature has not replaced that tradition with anything it measures on its own rollouts. Penetration is a quantity between meshes, so it needs a simulator or a mesh reconstruction, and the protocol below treats it as a simulation-only axis.

7.4 Statistics

`kress_gazit_policy_eval_2024` is the field's reference protocol and it prescribes process, not numbers. Write the success criteria before the run and have someone other than their author score the runs. Match initial conditions with image overlays. Interleave the policies blind within one session. Report counts rather than percentages, alongside the initial conditions and the failure modes. Use a posterior over the Bernoulli parameter instead of a point estimate. It prescribes no minimum trial count anywhere and no frequentist confidence-interval width anywhere. Its own example report uses 10 initial conditions with two runs each, 20 evaluations per policy, and its worked case shows what 20 buys. Pancake success of 15 of 18 against 11 of 17, nominally 83 against 65 percent, leaves a 0.11 posterior probability that the worse policy is actually better. At 150 of 180 against 110 of 170 the same rates separate.

`lbm_careful_examination_2025` supplies the missing numbers by fiat rather than derivation: 50 rollouts per task per policy per condition on hardware, 200 in simulation, blind, with randomised policy order inside per-initial-condition bundles. It replaces confidence intervals with Beta-posterior violins and corrects its pairwise tests under Bonferroni, "unless otherwise noted". Its own warning is the strongest sentence in this literature: "there is significant risk that many robotics papers are measuring statistical noise due to insufficient statistical power".

The rest fall short of their own advice. `roboarena_2025` runs 612 double-blind pairwise comparisons across seven institutions and 4284 rollouts, and reports no confidence intervals and no per-policy trial counts. `colosseum_2024` evaluates 235 test sets at 25 episodes each with "one training seed and one evaluation seed" per baseline and gives no intervals in simulation. `autoeval_2025` runs 50 trials per policy per task and calls ± 10 percent "the natural variance of robot evaluations", without giving the formula behind the intervals it plots.

Two papers do give usable numbers. `suresim_2025` pairs real and simulated trials and de-biases the simulation with a rectifier, saving 20 to 25 percent of hardware trials at paired correlations of roughly 0.6 to 0.7, and nothing at all at a correlation near zero. The 25 percent figure is the better of its two reported settings, the DP case at $p = 0.702$. Its decision rule is exact: combining helps only when the rectifier variance is below the variance of the real evaluations. `beyond_binary_success_2026` gives the largest saving. On the LBM rubrics its sequential test on graded scores cuts simulated evaluation by about 70 percent and hardware by about 45 percent, 286 trials against a nominal 500, with per-task decisions landing in 12 to 36 paired trials. On RoboArena's data a 30-point gap on continuous progress scores reaches significance in 18 trials, while a 20-point gap on binary success needs about 80.

7.5 A proposed protocol

Every count below is printed by `tools/make_eval_tables.py --derive`, and each axis is derived for the statistic that axis actually reports: a single rate takes a Wilson half-width, a matched comparison takes McNemar, a ratio takes the standard error of the log ratio, a correlation takes the Fisher-z interval.

Fix the width first, then read off the count. Take a 95 percent Wilson interval on a single reported rate, at the worst case of $p = 0.5$. A half-width of 20 points needs 21 trials, 15 points needs 39, 10 points needs 93 and 5 points needs 381. Ten points is the coarsest width at which a single rate is worth printing, so the absolute-rate minimum is 93, rounded to 100. At 100 trials a

reported 80 percent has an interval of 71 to 87 percent, and a reported 50 percent has 40 to 60. A comparison is a different question and a harder one: two rates each carrying ± 10 points do not resolve a 10-point difference between them, because the difference's standard error is larger by a factor of $\sqrt{2}$, so the width argument sets a floor on what is worth reporting and not on what can be compared.

For the A/B comparison the relevant calculation is power, and the design is paired. Table 8 matches initial conditions by image overlay and interleaves the two policies in one session, so the unit is a matched pair and the count follows McNemar, which depends on the discordance rate. The share of initial conditions on which the two policies disagree, and not on the two rates alone. To separate 50 from 70 percent at $\alpha = 0.05$ with 80 percent power: 37 pairs per arm at a discordance of 0.2, 57 at 0.3, 77 at 0.4 and 96 at 0.5. The protocol assumes 0.3 and asks for 57, and states the sensitivity rather than hiding it, because 0.5 is the discordance the same two rates produce when the pairing buys nothing, and at that value the paired count returns to the 93 per arm an unpaired test would need. The saving from pairing is real but smaller than the pair counts suggest, since a pair costs two rollouts: 57 pairs is 114 rollouts against 186. An earlier version of this section quoted 93, 169 and 387 per arm for gaps of 20, 15 and 10 points, which are correct for independent arms and are the wrong test for this protocol. That 93 was also the same integer as the half-width calculation in the paragraph above, which is a coincidence of the worst-case arithmetic and not a second derivation of the same number.

One hundred is a cap and not a bill, because on a graded score a sequential test reached its decision in 12 to 36 paired hardware trials in `beyond_binary_success_2026`. At 30 rollouts a continuous score already carries a half-width of ± 0.36 standard deviations, which is why a graded score can stop where a binary one cannot. A cell that stops early does not report a Wilson interval. Optional stopping breaks the coverage of a fixed-n interval, which is the reason `beyond_binary_success_2026` and `suresim_2025` use anytime-valid betting intervals rather than Wilson, so Table 8 asks a cell run to a fixed 100 for a Wilson interval and a cell stopped early for a confidence sequence, and never for both. Intervals are also marginal rather than simultaneous. Table 8 has seven axes and Table 9 twelve methods. At 84 independent 95 percent intervals, four excursions are expected by construction, so a paper comparing k policies on m tasks corrects its $k(k-1)/2$ pairwise tests to a global 95 percent level, as `lbm_careful_examination_2025` does, or says its intervals are not simultaneous.

Simulation is cheap, so simulated cells take 200 episodes, giving a 6.9-point half-width. Perturbation axes are screened rather than certified, and 40 per axis buys a 15-point half-width on each axis's own absolute rate, which is enough to rank the axes and pick the two worst for hardware. It is not enough for the ratio to the anchor that an earlier draft asked each cell to report. At 40 trials in each arm, a fall from a 0.50 anchor to 0.30 is a ratio of 0.60 with a 95 percent interval of 0.34 to 1.06, which contains 1: at the screening count you cannot establish that the perturbation hurt at all. Certifying that same drop takes 101 per arm by the log-ratio standard error, so Table 8 now asks for absolute rates with their own intervals at 40, reports the ratio without an interval, and prescribes 101 before any claim that a named axis hurt.

For unseen objects the resampling unit is the object and not the trial, so 20 objects at 5 trials each gives 100 trials and an object-level half-width near 20 points. That 20 points is the Wilson width at $n = 20$ and it treats each object's outcome as a single Bernoulli draw, which the five within-object trials are not. It is the right order of magnitude and the assumption belongs in the cell. A 10-point claim about an object distribution needs about 93 objects. Seven of the 39 rows that state an unseen count reach that: 225 in `bimangrasp_2024`, 241 in `resdex_2024` and `unidexgrasp_2023`, 360 in `dexgraspvla_2025`, 500 in `dexmv_2021`, 2029 in `clutterdexgrasp_2025` and 503409 in `graspxl_2024`. For a continuous score the half-width is 1.96 standard deviations over the square root of the count, so 100 rollouts give ± 0.20 standard deviations, and the unit is the rollout because frames within one are correlated.

The transfer axis is the one where 100 is least defensible. On 100 matched pairs a measured correlation of 0.70 carries a Fisher-z interval of 0.58 to 0.79, a half-width of about 0.10 that 130 pairs would be needed to guarantee. That is enough to establish that a simulator tracks reality at all, and it is not enough to separate `suresim_2025`'s useful regime from its marginal one, since those differ by about 0.11 in correlation. Both limits come within 0.05 of the estimate only at about 457 pairs. Table 8 states which of the two decisions each count supports rather than leaving a reader to assume the larger one.

Table 8. The proposed evaluation protocol

axis	what is measured	how	minimum trial count	reported alongside it	why
Task success	fraction of episodes meeting a criterion, and a graded rubric score in [0,1] over equally weighted milestones	criterion written before the run by the designer and scored by someone else; terminal predicate held for a 0.5 s dwell; initial conditions matched by image overlay; policies interleaved blind in one session	100 real and 200 sim per policy per task per condition for an absolute rate at a 10-point interval; a comparison alone needs 57 matched pairs per arm (McNemar, 50 against 70 percent, discordance 0.3)	raw counts, the criterion verbatim, the rubric, the initial-condition protocol, and an interval whose kind is named: Wilson for a cell run to a fixed n, an anytime-valid confidence sequence (WSR, as in <code>beyond_binary_success_2026</code>) for a cell stopped early by a sequential test, because a Wilson interval computed at a data-dependent stopping time is not a 95 percent interval. Intervals are marginal, not simultaneous: a paper comparing k policies corrects its $k(k-1)/2$ pairwise tests to a global 95 percent level, as <code>lbm_careful_examination_2025</code> does, or says it has not	a rate without a denominator and an interval cannot be compared; the rubric separates a near-miss from inaction; the design is paired, so the count that governs it is paired
Robustness	success under each perturbation axis as an absolute rate; the ratio to the unperturbed anchor is reported descriptively and is not tested at the screening count	axes that must not change the action and axes that must be scored separately; ranges stated; anchor replayed exactly	40 per axis per policy in sim to rank the axes; 101 per axis and 101 at the anchor before any claim that a named axis hurt	both absolute rates with their Wilson intervals, the ratio without an interval unless the certifying count was run, the axis ranges, and which split each axis is in	the ranking at the anchor is not the ranking under shift; at 40 per arm a fall from 0.50 to 0.30 is a ratio of 0.60 with a 95 percent interval of 0.34 to 1.06, which contains 1
Unseen objects	mean over held-out objects of per-object success	objects drawn from a stated distribution disjoint from training; bootstrap over objects, not trials	20 objects at 5 trials for a screening claim; 93 objects for a 10-point claim	the object list, per-object rates, the bootstrap interval, the object-count caveat, and the within-object trial count, since the 93 is a Wilson width that treats one object as one Bernoulli draw	the resampling unit is the object, so trial counts overstate the precision
Physical plausibility	maximum and mean hand-object penetration depth per rollout, and the fraction of frames above 2 mm	dense surface sample against the object mesh or SDF, computed on the policy's own rollouts by code that never entered the reward or the termination rule	100 rollouts, simulation only	interval over rollouts, sample density, threshold, solver depenetration settings, and at least one rendered rollout	a measure the policy optimised is not evidence about the policy
Sample and wall-clock cost	environment steps and wall-clock to the checkpoint that produced the headline number	counted to that checkpoint, not to the end of training; GPU model and count stated	3 training seeds	the range over seeds, and the statement that 3 seeds is a range and not an interval	training variance is not rollout variance and the two are routinely conflated
Real-robot transfer	paired real and simulated outcomes on matched initial conditions, their correlation, and the rectifier variance against the variance of the real evaluation	the same 100 initial conditions run in both; correlation on the pairs; variance ratio bootstrapped over the same pairs	the 100 real trials paired to 100 sim, which fixes the correlation to about ± 0.10 ; about 457 pairs for a claim that separates two correlations 0.1 apart	the correlation with its Fisher-z interval, the variance ratio with its bootstrap interval, and which of the two decisions the count supports	a simulator earns a real-world claim only through its measured paired correlation; at 100 pairs $\hat{\rho} = 0.70$ carries 0.58 to 0.79, enough to say a simulator tracks reality at all and not enough to separate <code>suressim_2025</code> 's 0.70 regime from its 0.59 one
Reproducibility	code, checkpoints, and the exact config that produced the headline run	diff the paper's stated objective against the released config and name the file and line of every disagreement	not a trial count	the named file and line for each disagreement, or an explicit statement that none was found	16 of the 62 code-releasing method rows carry a true contradiction between the paper's stated objective and the released code, and 38 carry a recorded disagreement of any kind (§5.8)

7.6 Table 9, an empty results matrix

The rows are the 12 most-mentioned dexterous-hand policy methods in the corpus, and the rule is the one `tools/make_eval_tables.py` implements, stated here in the same words. A candidate is a method row with a non-null `hand`; its `hand` string must not contain "parallel" or "gripper"; it must carry at least one paradigm tag that produces a closed-loop policy and must not carry `teleop-system`, because an interface is scored on latency and operator effort rather than on a policy's success rate. And its name must be at least four characters, so that a short string does not match everything. Candidates are then scored by the number of other corpus papers whose parsed text in `papers/md` contains the name, and the top 12 by count, ties broken by key, are the rows.

Two corrections to that ranking are worth stating, because both changed it. The match is on a whole word. Under the bare substring test an earlier version used, "UniDex" matched inside "UniDexGrasp" and "UniDexGrasp++", and `unidex_2026`. A 2026 paper. Sat sixth in a ranking over a corpus written mostly before it, on 34 mentions that belonged to a different work. As a whole word it has 3 and it is not in the table. And the interface rule is now applied to every row that carries the tag rather than only to rows that carry nothing else, which drops `anyteleop_2023` at 34 mentions, `dime_2022` at 28 and `holo_dex_2022` at 23, along with `dexpilot_2020`, which the earlier prose already excluded by hand. The mention counts are printed under the table so a reader can audit them.

The ranking is not one quantity even so. A method's name is taken from the first line of its note, which yields an acronym for some works and a full title for others, and a title is matched mostly inside reference lists while an acronym is matched in running text. Those have different base rates, so the table marks which kind each row was matched on and the two kinds are not comparable with each other. Mention counts are counts of mentions and not of use, as `METHOD.md` records.

Every cell is empty. This survey re-ran nothing, and no cell can be filled at the denominator Table 8 asks for. `dextreme_2022` comes closest and is the reason the claim is stated that narrowly: it reports a criterion, a trial count and an interval. Object orientation within 0.4 rad of target, 27.8 ± 19.0 average consecutive successes with the $\pm$ a 90 percent confidence interval. On 10 trials. Table 7 is not a counter-example either, though it looks like one: it carries `trials`, `unseen obj`, `penetration` and `code` columns for all 112 method rows, including all 12 of these. Table 7 records what each method reported. Table 9 asks for what Table 8 defines. A value with an interval, a stated denominator and a criterion written before the run. And none of Table 7's values meets that. The first row of Table 9 is a worked example so that the format of a cell is unambiguous. Every number in it is fabricated and labelled as such.

Table 9. The matrix, for someone else to fill

method	task success (rate $\pm$ Wilson 95, n)	robustness (per-axis rate $\pm$ Wilson 95, n; ratio to anchor)	unseen objects (mean per-object rate, bootstrap 95, k objects)	plausibility (max / mean mm, frac frames > 2 mm)	cost (env steps, GPU-h, 3-seed range)	transfer (ρ [lo, hi], rectifier var / real var)	reproducibility (file:line of each disagreement, or none)
<i>worked example — every number fabricated</i>	0.72 $\pm$ 0.09 (100)	lighting 0.41 $\pm$ 0.15 (40), ratio 0.57; anchor 0.72 $\pm$ 0.09 (100)	0.55, [0.41, 0.68], 20 objects	3.1 / 0.8 mm, 0.12	1.2e9 steps, 46 GPU-h, 0.68–0.74	0.61 [0.47, 0.72], 0.43	cfg/train.yaml:88 orient weight 0.0 vs paper 0.5
openai_dexterity_2018*							
dexmv_2021							
dapg_2017*							
dextreme_2022							
dexcap_2024							
hora_2022*							
visual_dexterity_2022*							
rotating_without_seeing_2023*							
unidexgrasp_2023							
pddm_2019*							
hato_visuotactile_2024*							
dexpoint_2022							

12 rows, 84 cells, all 84 empty; the first row is a worked example and every number in it is fabricated. Rows are the 12 most-mentioned dexterous-hand policy methods in the corpus, by the rule in §7.6, scored on whole-word matches over papers/md: `openai_dexterity_2018` 64; `dexmv_2021` 50; `dapg_2017` 49; `dextreme_2022` 44; `dexcap_2024` 40; `hora_2022` 34; `visual_dexterity_2022` 32; `rotating_without_seeing_2023` 27; `unidexgrasp_2023` 27; `pddm_2019` 25; `hato_visuotactile_2024` 23; `dexpoint_2022` 22. Mention counts are counts of mentions, not of use. * marks a work matched on its title rather than a short name; a title is matched mostly inside reference lists and a short name in running text, so the two kinds of count are not comparable with each other.

7.7 What would have to be true

The bill comes first. Per policy and per task the protocol asks for 57 matched trials on the anchor set, which is the paired count from §7.5, and 100 on the unseen-object set, at 20 objects and 5 trials each, with robustness and plausibility absorbed by simulation. A two-policy comparison on three tasks is then 342 real rollouts on the matched set and 600 on the unseen-object set, 942 in all, and at one minute per rollout including the reset that is about 16 hours of robot time, before failed resets, repairs and scoring. The same bill computed from the independent-arm count, which an earlier draft used, was 1200 rollouts and 20 hours. The pairing removes about a fifth of it rather than the four fifths the pair counts suggest, because a pair is two rollouts and only the anchor set is paired. `autoeval_2025` ran about 850 episodes in 24 hours on a WidowX with three human interventions, and had to pause 20 minutes every 6 hours once the motors overheated. A tendon-driven multi-finger hand is more fragile, so 16 hours of rollouts is most of a week of calendar time.

That bill is large but not unprecedented. `autoeval_2025` records that evaluating OpenVLA against its baselines took more than 2500 rollouts and more than 100 hours of human labour across three institutions, and `lbm_careful_examination_2025` analysed about 1800 real rollouts across nine hardware stations. What is unprecedented is paying it for a single dexterous-hand paper, where the modal per-cell count is 10 trials and the median per-cell count 15, and where the median stated count has been 20 or 22.5 in each of the last three years.

Four things would have to change. Reviewers would have to reward 57 matched trials on one task over 20 unmatched trials on five, and nothing in the corpus suggests that is happening. The loop would have to be automated, and `autoeval_2025` shows it is buildable for a gripper at one to three hours of setup, while stating that it supports binary success only and no robustness axes. Somebody would have to run the penetration measure on their own rollouts, which is a choice rather than a capability: section 4.2 shows the depth is computable from the poses and the meshes in a few lines of Warp, and IsaacGymEnvs already ships a task that does it every step. An earlier draft of this section made the engine the barrier, and that claim is withdrawn, because the released code refutes it. And the comparison would have to be sequential, because the savings in `beyond_binary_success_2026` are the only reason 100 is a cap rather than a cost.

Four limits apply to the proposal itself. This survey re-ran no method, so every count in Table 8 is derived from an interval width, a power calculation or another paper's measurement, and Table 9 is empty because we filled no cell. The counts are worst-case at $p = 0.5$, so a method near 90 percent needs fewer trials for the same width and a method near 50 percent needs all 100, and the paired count additionally rests on an assumed discordance of 0.3, which no paper in this corpus reports. The perturbation axes are borrowed from `colosseum_2024` and `simpler_2024`, which run parallel-jaw grippers on rigid objects, where `simpler_2024` found physical parameters moved success rates by at most 15 percent. That is the sensitivity expected to grow with multi-finger contact, and nobody has measured it. The 2 mm penetration threshold is taken from `toporetargt_2026` with no independent justification, and the captured human grasps in `grab_2020` sit above it at 3.25 mm, which makes 2 mm a simulator convention rather than a physical bound. So the three things this survey adds to the reference-versus-rollout axis are a count, a measurement method and a protocol slot for them, and a threshold is not among them. It is borrowed from one paper and reported as borrowed, and it will stay a convention until somebody measures penetration on rollouts across hands and engines and finds a value that separates behaviour a physicist would accept from behaviour they would not.

8. Gaps

8.1 What opening the repositories showed

This is a result rather than a gap, and a result about publishing practice in robot learning: the dexterous corpus is its sample, not its subject.

Sixty-two method rows released code that could be parsed against the paper, and 38 carry a recorded discrepancy. The classification is the finding: ten contradictions, where the paper states one value and the shipped code demonstrably states another; thirteen limits of this survey's own parse, which captured signatures or a truncated body rather than the component; eight cases where the code was never released; four version skew; three inconsistencies inside a paper with no code involved. Ten is the number to quote, eight at high confidence, with `penspin_2024` and `omnih2o_2024` held at medium against innocent readings a direct code read would settle.

The first count was sixteen. An adversarial re-reading withdrew seven accusations. `maniptrans_2025`, `eureka_2023`, `open_television_2024`, `dexmachina_2025`, `artigrasp_2023`, `graspxl_2024`, and the domain-randomisation half of the charge against `dexpbt_2023`. Two refuted by the accused repository's own README, two resting on reward code never in the parse, one charging the code with structure the paper prints, one against a paper with no reward function. Each is recorded in its row in a `mismatch_review` field: a survey that names people should carry its retractions beside its accusations.

`physhoi_2023` survived every attempt to break it. Table 4 weights object rotation at 0.1 for GRAB, and the released `compute_humanoid_reward` sets the object rotation and rotation-velocity errors to `torch.zeros_like`, unconditionally, so the dataset exemption the paper states does not cover it. The reward behind its 95.4 percent tracked the object in position only, and its own position-only success criterion could not have caught that. `robot_synesthesia_2023`, at the other end, prints its six weights as symbols and released nothing, so that objective exists in no machine-readable form. `hora_2022` alone discloses its own gap, in its README.

Ten is a floor, since forty-six method rows released nothing to check and four more are unsettled. What would close it: publish the reward table generated from the released config at a named commit, so a reviewer diffs two artefacts instead of reading two documents.

8.2 Nobody records interpenetration for their own policy's rollouts

Eleven method rows handle interpenetration at all, of the ninety-six whose notes settle the question, and only four of them are closed-loop policies. Every one of the eleven sits on the reference side of the reference-versus-rollout split, scoring a pose or a trajectory before execution rather than the behaviour that followed. The engines are not the obstacle: IsaacGymEnvs ships a task that computes a per-environment maximum interpenetration depth in Warp and gates its policy update on a 1 mm threshold, and `tactile_genesis_2026` offers two penetration-depth backends on Genesis geometry as sensors. The depth is computable by anyone from poses and meshes; nobody reports it for a dexterous rollout. What would close it: maximum and mean penetration depth over the evaluation rollouts, on a dense surface sample, from a measure the policy never optimised.

8.3 The evaluation-methodology literature contains no dexterous hand

Seven corpus rows are evaluation protocols, a small denominator, and not one uses a dexterous hand: `suressim_2025` runs a parallel-jaw gripper and the other six state no hand. `kress_gazit_policy_eval_2024` sets no minimum trial count and no confidence-interval width. What would close it: run the Table 8 protocol once, on one in-hand reorientation task with one 16-DoF hand, and release the rollouts as the first row of Table 9.

8.4 Generalist policies run on a fraction of the degrees of freedom

The gap is size, not absence. Eighteen rows carry the generalist tag. Fourteen of them settle whether the reported evaluation ran on a multi-fingered hand, and eight of those fourteen did. The remaining four do not say, which is itself the smaller half of

this gap. Of the eight, five state the hand's degrees of freedom: four at 6 and one at 12. The median is 6, against 16 over the 48 reinforcement-learning rows that state one. So the generalist policies that do touch a hand run it at roughly a third of the actuation the reinforcement-learning literature assumes. The stronger claim, that no generalist reaches the 16-to-24 band, holds over those five stated counts and no further, which is what section 5.5 says. `gr_dexter_2025` and `egoscale_2026` name hands at 21 and 22 actuated degrees of freedom, so they do reach the band on paper, and neither settles whether the reported evaluation ran on that hand. What would close it: one dexterous-hand task in the standard generalist suite, with the hand's DoF and vendor beside the number.

8.5 The hands that can be bought go unused

7 documented hands that can be bought or built from published designs take zero method rows between them: Unitree Dex5, Tesollo DG-5F, Proception ProHand, ORCA, RUKA, RUKA v2, BiDexHand. Only three of the fifteen simulator rows name a real hand at all, and the ones they do name are the field's defaults. 18 of the 33 hands in Tables 2 and 3 appear in no method row, but 11 are neither sold nor open and appear in none for that reason, so 33 is not the denominator for a software-lag claim. The rule that decides used from unused is one regular expression per hand against the method rows' own hand field, in `tools/hand_usage.py`, so the partition can be recomputed rather than argued about. `bench2dex_2026` compares 12 hands and `dexverse_2026` six without stating a DoF count for any. What would close it: a conformance suite for hand models, one URDF or MJCF per hand, with fixed joint-limit, mass and collision checks and a published pass or fail per engine.

8.6 Bimanual work runs on one coordination architecture, and it has been ablated once

Twenty-one of the 28 rows that Section 6 counts as putting a learned closed-loop controller on two multi-fingered hands run one policy over a concatenated two-hand observation. That choice has been compared twice, with opposite outcomes, and ablated once: `asymdex_2024` scores 0.7701 on Block in cup over five seeds, against 0.1086 with relative frames but no role asymmetry and 0.0164 with asymmetry but no relative frames. All three handover papers share one reward across giver and receiver. What would close it: one handover task, three architectures, the same hand and the same seeds, with separate giver and receiver returns.

8.7 Human data does not port across hands, and the map is usually unstated

Fifty-three method rows use human data and name 45 distinct hand strings between them. Twenty appear in Table 6 with a stated retargeting objective. The other 33 do not say how the human motion reached the hand, and the stated objectives do not converge, from a fingertip keypoint-vector energy in `anyteleop_2023` to a per-joint regression from exoskeleton encoders in `dexumi_2025` to no retargeting at all in `egozero_2025`. More correspondence is not better: `objdex_2024` retargets the wrist only and beats both finger-joint and fingertip mapping, and its added fingertip-matching reward "does not yield benefits and even leads to lower performance". What would close it: a retargeting benchmark in TopoRetarget's form, reporting contact precision, contact alignment, maximum penetration and share of frames past 2 mm, per hand and per dataset.

9. Conclusion

The binding constraint on this field is not ideas. It is verification. Sixty-two method papers released code that could be read against the paper, 38 of those record a discrepancy, and ten are contradictions where the shipped code states a different objective from the published one. The first count was sixteen, and an adversarial re-reading withdrew seven accusations, each withdrawal recorded in the row beside the charge. Ten is a floor, because forty-six method rows released nothing to check. A reward table in a paper is a claim about a document, not about a run. `physhoi_2023` is the case to remember, because the term its table weights at 0.1 is set to zero in the code, and its own success criterion could not have detected that.

The second finding is that the quantity most specific to dexterous manipulation is the one nobody records. Contact is what separates a hand from a gripper. Eleven of the 96 method rows whose notes settle the question address interpenetration at all, only four inside a closed-loop policy, and not one reports a penetration number for its own policy's rollouts. The obstacle is not the engines. IsaacGymEnv ships a task that computes a per-environment maximum interpenetration depth against meshes and gates the policy update on a 1 mm threshold, and `tactile_genesis_2026` offers penetration depth on Genesis geometry as a sensor. The tooling sits in the field's own benchmark repository and the number is still not reported. `dextrack_2025` has the formula and points it at its inputs.

The third is that hardware and software have come apart, on a narrower claim than the hand count first suggests. Tables 2 and 3 hold 33 hands and 18 appear in no method row. 11 of those 18 are neither sold nor open and appear in none for that reason, which leaves 7 hands that can be bought today or built from published designs and that take zero method rows between them. Eight of the fourteen generalist policies that settle the question do evaluate on a dexterous hand, at a median of 6 degrees of freedom against 16 across the reinforcement-learning rows.

What this survey cannot establish is which method is better than which. It re-runs nothing, and Section 7 argues that the published numbers do not compare. Five works are cited by metadata only and no claim rests on them. Every coverage statistic here counts what this survey's extraction captured rather than what the literature reported. Each is a floor and not a rate, because every miss converts a reporting paper into a silent one.

Three things to do next week, cheapest first.

If you are publishing, generate the reward table from the config that trained the reported run, print it, and cite the commit. State the trial count, state the success predicate, and release the per-trial outcomes. None of that needs a GPU.

If you are running experiments, measure penetration depth over your evaluation rollouts on a dense surface sample, and never let that measure become a reward. `toporetaret_2026` shows what the number looks like when someone takes it seriously, and what the widely used retargeters look like when nobody does.

If you are choosing hardware, the corpus names four hands and no more. The Allegro carries 35 method rows, the Shadow 21, the Inspire 19 and LEAP 12, and every other hand in Tables 2 and 3 carries eight rows or fewer. An announced hand has no URDF, no datasheet that can be checked, and no paper in this corpus that used it.

Appendix A. Method

The survey is written from a corpus that lives on disk. No claim about a paper is made from memory. Each claim traces to a note, each note to a parsed markdown file, and each markdown file to a PDF or repository with a recorded hash or commit.

Selection

Six topic-specific bibliographies were assembled in parallel (simulators and physics, hands and vendors, reinforcement learning, imitation and human data, bimanual, benchmarks and evaluation), seeded with the canonical works in each area and extended by search up to 2026-09-17. Every arXiv identifier was checked by fetching the abstract page and matching the title, and entries that could not be checked that way are marked. The six lists were merged with deduplication on arXiv identifier and normalised title, giving 221 entries.

Acquisition

PDFs were downloaded from arXiv or, for work without a preprint, from the publisher or vendor page recorded in the bibliography. Vendor pages for hands without any paper were fetched as HTML and converted to markdown, so that a specification quoted in this survey is quoted from a stored copy of the page and dated. Repositories were shallow-cloned and parsed, then deleted. What remains is one markdown per repository holding the README, a pruned file tree, the task and reward configuration files, and the bodies of reward and observation functions.

Parsing

PDFs were converted with pymupdf4llm, falling back to raw text extraction when the layout parse returned too little. The corpus manifest records, per entry, the source URL, page count, sha256 of the PDF and the converter used. Two consequences matter for reading this survey. Equations rendered as images do not survive conversion, so where a reward weight exists only inside a figure it is recorded as unreadable rather than guessed. Tables are flattened, so a number taken from a table is quoted with the table it came from and, where the flattening is ambiguous, the ambiguity is stated.

Notes

Each entry was read into a structured note under a fixed template: embodiment, learning method, objective, contact handling, evaluation, reproducibility, stated limitations, and quotable claims. Notes were written only from the parsed files. Where the source is silent the note says "not stated", and nothing was inferred from the reviewer's prior knowledge of the work. For method papers the reward or loss was quoted from the paper and, separately, from the released code, so that disagreements between the two are visible rather than smoothed over. Those disagreements turned out to be common enough to become a finding in their own right.

Sources that could not be obtained

Six works are cited by metadata only and no claim in this survey rests on their contents. Five are behind publisher paywalls with no author-hosted copy found on 2026-09-18: Okamura et al. 2000, Bicchi 2000 (IEEE T-RO), Piazza et al. 2019, Roa and Suarez 2015, and Butterfass et al. 2001 on DLR-Hand II. The sixth is Hwangbo et al. 2018 on RaiSim, whose hosted PDF returned no body. The freely circulating PDF often taken for Bicchi 2000 is a different work, a book chapter, and is listed separately.

Ma and Dollar 2011 is a seventh case with a different cause. The fetch that failed when its note was written succeeded on 2026-09-18, so a seven-page PDF and its hash are in the manifest, but no note has been read from it. It is cited by metadata only for that reason and not because the source is unavailable.

Three of the 221 bibliography entries carry no structured row. Two are the paywalled Bicchi 2000 and DLR-Hand II entries. The third is [bicchi_grasping_chapter_2001](#), which was read into a note and quoted throughout but is a book chapter rather than a work with an embodiment, a method or a result to record in a row.

What this method cannot do

Mention counts over the corpus are counts of mentions, not of use: a related-work sentence counts the same as an experiment. Vendor specifications are manufacturer claims and are labelled as such throughout, and where a page has since gone offline the note says so. The corpus is large but not exhaustive, and selection by search favours work that is indexed, in English, and posted as a preprint. It also favours recent work. Of the 221 bibliography entries, 138 are dated 2024 or later and 16 predate 2018, so this is a corpus of the learned era and any claim here about a trend over time is a claim about 2022 onward.

Every coverage statistic in this survey measures what this survey's extraction captured, not what the literature reported. A structured row holds a scalar. A paper that reports a per-task count, a rubric, or a total spread across several tables produces a null, and a null is then counted as silence. The bias runs one way. Every miss converts a reporting paper into a silent one, so the field is made to look worse at reporting than it is. The size of the effect was measured on the statistic the survey leads with. Of the 34 method rows that had a real robot and no recorded trial count, 15 carried a count in plain text in their own

note, dropped because the paper reports it per task and the field takes a single integer. Those 15 have since been re-extracted, which moved the stated-trial-count row from 55 to 70. The success criterion and the unseen-object count were audited the same way, rising from 79 to 98 and from 32 to 39. The same mechanism reaches the penetration, code-release and failure-mode fields, none of which has been audited that way. Read every coverage statistic in this survey as a floor rather than as a rate, and read the bars in Figure 6 the same way.

Appendix B. The full hand table with its sources

Table 2 and Table 3 split these rows by whether the hand can be bought, and truncate their cells to stay readable. This appendix prints every hand row once and at full width. Nothing here is a measurement by this survey. Every specification is a manufacturer claim, a figure in a paper, or a figure in a press article, and the `source quality` column in B.2 says which. Cells reproduce the text stored in the row, with the source's own punctuation and wording.

B.1 Specifications as the source states them

hand	maker	DoF	act. DoF	actuation	weight g	tip force N	payload	tactile	control rate	price USD	BOM USD	open HW	licence	sim model	status
agibot_omni_hand_2025															
onex_neo_hand_2026	1X Technologies	25	25	quasi-direct-drive tendons via the '1X Tendon Drive' at low gear ratios (~5:1-15:1); motors in the forearm, tendons pulled through the wrist		45		high-resolution tactile sensing across fingertips and finger surfaces (normal force, contact location, shear); sensor count/technology not stated	not stated			no			announced
leap_hand_v2_adv_2025	Kenneth Shaw, Deepak Pathak, Carnegie Mellon University	21	17	tendon (PIP/DIP coupled by a single tendon per finger); motor type not stated				none mentioned	not stated	3000	3000	yes		URDF, loaded into PyBullet by the released teleoperation node	announced
paxini_dexh13_2024	PaXini Tech							1,140 ITPU multidimensional tactile processing units (press-release claim; unclear if PX-6AX-GEN3 sensor family)							announced
clone_robotics_hand_2024	Clone Robotics	27		hydraulic artificial muscle ('Myofiber' — small water-filled tubes that contract when pressurized), carbon-fiber bones and ligament-style tethers, driven by a 500 W water pump and 36 electro-hydraulic valves; each Myofiber generates up to 1 kg grip force, survived 650,000 test cycles	<910			pressure pads in the palm detect grip firmness (count not given); 70 inertial sensors track angle/speed (proprioception, not tactile)				no			internal-only
sanctuary_phoenix_hand_2024	Sanctuary AI (Sanctuary Cognitive)	21		hydraulic, 'unique miniaturized				not in the press release text; sidebar-				no			internal-only

hand	maker	DoF	act. DoF	actuation	weight g	tip force N	payload	tactile	control rate	price USD	BOM USD	open HW	licence	sim model	status
	Systems Corporation)			hydraulic valves'				only headlines mention new touch/tactile sensors as a separate later announcement , not confirmed as part of this hand							
bidexhand_2025	Zhengyang Kris Weng, Center for Robotics and Biosystems, Northwestern University	21	16	cable-driven, N-configuration tendon routing (FeeTech servos, endless-loop antagonistic pull-pull); thumb CMC flexion driven via a 4-bar linkage instead of a direct tendon		2.14	4.54 kg (10 lb) lifted	none mentioned	serial to an onboard ESP32, ROS 2 position control; rate not stated			yes	MIT (code, per README); CAD/hardware licence unconfirmed (bib's claimed CC BY-NC-SA is not verified in the parsed README or file tree)	URDF, consumed by the ROS 2 and MoveIt stack; no physics simulator	open-source
ruka_v2_2026	Xinqi (Lucas) Liu, Ruoxi Hu, Alejandro Ojeda Olarte, Zhuoran Chen, Kenny Ma, Charles Cheng Ji, Lerrel Pinto, Raunaq Bhirangi, Irmak Guzey, New York University and NYU Shanghai	20	16	tendon-driven, forearm-mounted actuators; decoupled parallel 2-DoF wrist via a passive spherical ball joint; single-tendon-plus-spring abduction/adduction				optional e-flesh fingertips (magnetic touch-sensing form factor); not part of the base design	Dynamixel over USB serial; no rate stated	1500	1500	yes		MuJoCo XML in the shared v1 repo; not confirmed to carry the v2 wrist	open-source
orca_hand_2025	Clemens C. Christoph, Maximilian Eberlein, Filippas Katsimalis, Arturo Roberti, Aristotelis Sympetheros, Michel R. Vogt, Davide Liconti, Chenyu Yang, Barnabas Gavin Cangan, Ronan J. Hinchet, Robert K. Katzschmann,	17	17	tendon-driven (antagonistic fishing-line tendon pairs per joint); wrist uses a GT2 timing belt drive	1200		10.5 kg on all four fingers and 2 kg on the index finger alone, both at a fixed 600 mA motor current	yes: FSR-based binary tactile sensing on all 5 fingertips (RP-C7.6-ST), absolute detection threshold as low as 0.05 N	serial to Dynamixel or Feetech motors; rate not stated			yes		URDF-derived kinematic constants only; no URDF or MJCF file in the released tree	open-source

hand	maker	DoF	act. DoF	actuation	weight g	tip force N	payload	tactile	control rate	price USD	BOM USD	open HW	licence	sim model	status
	Soft Robotics Lab, ETH Zurich														
leap_hand_2023	Kenneth Shaw, Ananye Agarwal, Deepak Pathak, Carnegie Mellon University	16	16	direct-drive (Dynamixel servos, e.g. XC330-M288), joint velocity ~8 rad/s	595	19.5		none (future work only: 'we plan to develop and integrate LEAP Hand with low-cost touch sensors')	up to 500 Hz querying over USB serial; the paper's own sim-to-real policy runs at 20 Hz	2000	2000	yes	Code: MIT License; CAD: CC BY-NC-SA (non-commercial use with attribution)	URDF claimed in the paper; the released LEAP_Hand_API repo contains no URDF, xacro or MJCF	open-source
ruka_2025	Anya Zorin, Irmak Guzey, Billy Yan, Aadithya Iyer, Lisa Kondrich, Nikhil X. Bhattasali, Lerrel Pinto, New York University	15	11	tendon-driven (11 Dynamixel actuators in the forearm: XM430-W210T for the thumb, XL330-M288-T for the other fingers)		2.74	6.0 kg: weight added to a curled cloth bag grip until joint-angle error exceeds 15 degrees, best of three trials	none (stated limitation: 'lacks tactile sensing')	Dynamixel over a USB-to-serial bridge; rate not stated, data collection ran at 15 Hz	1300	1300	yes		MJCF	open-source
dexhand_open_source_2023	Rob Knight, The Robot Studio (design); electronics/firmware/ROS 2 by Trent Shumay, IoT Design Shop			tendon (fishing line, Sufix 832 80lb / 0.8mm kiteline); Emax ES3301/ES3302/ES3351/ES3352 micro-servos for fingers, Feetech SCS2332 or PWM servos for the wrist				none mentioned	Arduino SCServo or SBUS firmware, BLE; rate not stated	300	300	yes	CC BY-NC-SA 4.0 (README); Onshape CAD freely available	URDF, in the IoT Design Shop ROS 2 packages	open-source
ilda_hand_2021	Ajou University, Korea Institute of Machinery & Materials, Korea University	20	15	linkage-driven, direct linear drive (3 Maxon DCX8M motors per finger with GPX8 16:1 gearboxes driving ball screws via parallel+serial four-bar linkages)	1100	34	18 kg	6-axis F/T sensor per fingertip (5 total), force resolution 62 mN, range +/-35 N; not a distributed taxel skin	CAN to a desktop; loop rate not stated			no			prototype
pisa_iit_softhand_2014	Centro E. Piaggio, University of Pisa and IIT (Catalano, Grioli, Farnioli, Serio, Piazza, Bicchi)	19	1	tendon-driven adaptive synergy, single motor via differential gears (6 W Maxon RE-max21, 84:1 reduction)			holding force about 20 N along z and holding torque 2 N m, 28 N and 3.5 N m with a stronger motor	none (motor encoder only; a padded work glove supplied contact compliance during experiments)	not stated			no		URDF and Gazebo models in the ROS repo (BSD 3-Clause)	prototype

hand	maker	DoF	act. DoF	actuation	weight g	tip force N	payload	tactile	control rate	price USD	BOM USD	open HW	licence	sim model	status
faive_hand_2023	Yasunori Toshimitsu, Benedek Forrai, Barnabas Gavin Cangan, Ulrich Steger, Manuel Knecht, Stefan Weirich, Robert K. Katzschmann, Soft Robotics Lab, ETH Zurich	16	11	tendon-driven, rolling-contact joints (16 Dynamixel XC330-T288-T servos, 6 antagonistic pairs)	1100		10 kg, whole-hand downward power grasp on a dumbbell	none on the physical hand; a simulated 'fingertip force' (15-dim, critic-only/privileged) exists only in the RL training observation, not real tactile hardware	Dynamixel over serial; joint angles estimated from tendon length by EKF; rate not stated					MJCF, used directly by the released Isaac Gym training repo	prototype
boston_dynamics_atlas_hand_2026	Boston Dynamics			electric (robot is fully electric); specific hand mechanism not stated				tactile sensing in the fingers and palms, confirmed present; type and count not stated				no			prototype
daxo_muscle_v0_2025	Daxo Robotics			ultra-redundant tendon-driven; compliant structure with no rigid joints, flexible materials and tendon routing	750										prototype
figure_03_hand_2025	Figure AI							not stated on page for the hand itself; palm cameras in each hand are vision, not tactile				no			prototype
tesla_optimus_hand_2025	Tesla			V3: tendon-driven from forearm actuators, three tendons per finger through a crosstalk-managing wrist. Gen 2: in-hand actuators/sensors (unlocated in the article).				Gen 2: tactile sensing on all fingers, demonstrated matching contact location and force in an egg-pickup video; type/count not stated. V3: not stated.				no			prototype
xiaomi_cyberone_hand_2026	Xiaomi			motors located in the hand (compact motors generate heat				full-palm tactile sensing, area ~8,200 sq mm, detects				no			prototype

hand	maker	DoF	act. DoF	actuation	weight g	tip force N	payload	tactile	control rate	price USD	BOM USD	open HW	licence	sim model	status
				requiring liquid cooling); no tendon/linkage statement, though contrast with tendon-driven hands' cycle life implies a non-tendon design				pressure and contact across the whole palm not just fingertips; taxel count and type not stated							
shadow_dexterous_hand_2005	Shadow Robot Company	24	20	tendon-driven, motors in the forearm (20 Smart Motor nodes, Maxon motors, PWM)	4300		4 kg in a power grasp, vendor claim, protocol not stated	Shadow Tactile Fingertips (STF): 17x3 DoF Hall-effect taxels, 1000 Hz, uncalibrated; up to 5 fingertips; software also supports BioTac/MST/U BIO/PST fingertip options	EitherCAT at 1 kHz to the host, with a 5 kHz tendon-force loop inside each motor module			no		MuJoCo models for Hand E and variants, plus Gazebo, in sr_common	sold
proception_prohand_2026	Proception Inc (YC W25)	22		tendon-driven, motors pull cables to move the fingers				integrated skin-like sensors detecting contact, supporting grip control; same 'sensor skin' used on the companion ProGlove; type/count not stated				no			sold
sharpa_wave_2026	Sharpa Robotics (Sharpa Pte Ltd)	22	22		1300	20 / 12	40 kg / 24 kg (the same two unlabelled columns)	'Dynamic Tactile Array' (DTA), camera-type array at the fingertip: resolution 240x240 / 60x60, force resolution 0.02 N / 0.05 N, spatial resolution 1 mm / 2 mm, frame rate 180 fps / 30 fps, force range 0-30 N, latency 20 ms, 6-D	500 Hz over 1000BASE-T Ethernet			no		MuJoCo and Isaac Sim models	sold

hand	maker	DoF	act. DoF	actuation	weight g	tip force N	payload	tactile	control rate	price USD	BOM USD	open HW	licence	sim model	status
								F/T; number of sensors per hand not stated							
tesollo_dg5f_2024	Tesollo Inc. (Incheon HQ, Gwangmyeong R&D/factory)	20	20	one integrated actuator per joint ('high-torque actuation', absolute encoder); page does not say direct-drive	1763		pinching 2.5 kg rated and 5 kg maximum; enveloping 10 kg rated and 20 kg maximum	none standard; optional fingertip sensors (6-axis F/T, 3-axis force, or tactile) available; count/type not given	250 Hz over Modbus RTU or TCP and Ethernet TCP/IP			no			sold
unitree_dex5_2025	Unitree Robotics (Yushu Technology Co., Ltd.)	20	16	in-joint geared motor ('hollow-cup motor' + high-precision encoder + low-damping small-clearance reducer), backdrivable; page does not use the words tendon or linkage	1100	10	3.5 kg palm down and 4.5 kg palm left, on a 5 cm round hard object	Dex5-1: none. Dex5-1P: 94 pressure sensors per hand (2x5 palm + 2x3x5 finger pad + 2x3x5 fingertip + 2x3x4 finger root), range 10 g-2500 g.	1000 Hz over USB 2.0, with per-joint stiffness and damping commands			no			sold
wuji_hand_2025	Wuji Technology (founded 2019)	20	20	direct-drive rotary, backdrivable				'Multi-Axis Force/Torque Fingertip Sensing' confirmed present; no type, count or range given	1000 Hz across 20 axes over 100BASE-TX Ethernet, MIT force-position hybrid mode			no		MuJoCo and Isaac Sim models	sold
allegro_hand_v4_2016	Wonik Robotics Co. Ltd. (Seoul, South Korea); earlier versions 1.0/2.0 made by SimLab Co. Ltd.	16	16					none	333 Hz, CAN, the hand's own real-time clock (driver README)			no		URDF and xacro, left and right, in the ROS driver repo	sold
inspire_rh56dfx_2023	Beijing Inspire Robots Technology Co., Ltd.	12	6		540	10		none (this variant is the 'without tactile sensors' table; a tactile FTP variant exists per bib but is not on this page)	RS485; rate not stated			no			sold
robotera_xhand1_2024	ROBOTERA	12	12	gear-driven force-controlled joint	1100		25 kg, tracker's 'Strength'	tactile/force sensors on every fingertip	not stated	14000		no		URDF exists but is licensed; maniptrans_	sold

hand	maker	DoF	act. DoF	actuation	weight g	tip force N	payload	tactile	control rate	price USD	BOM USD	open HW	licence	sim model	status
				modules per finger segment, back-drivable			figure, over 25 kg gripping palm-up	confirmed present (senses contact, force, temperature); count and array size not stated						2025 withholds it	
shadow_dex_ee_2024	Shadow Robot Company, in collaboration with Google DeepMind	12			4100			stereo camera-based fingertip tactile sensors (hundreds of taxels each); multi-taxel 3-DoF arrays on middle and proximal phalanges; counts not stated	not stated; the page claims 'high bandwidth torque and position control loops'			no			sold
brainco_rev_o2_2025	BrainCo Inc.	11	6		383	15	>=20 kg	Touch variant only: multi-dimensional fingertip tactile module with 'Tactile Adaptive Control'; count/type not stated. Basic and Pro variants: none.	RS485 and CAN FD; EtherCAT on Pro and Touch; rate not stated			no			sold
linkerbot_l20_2025	LinkerBot (SDK author 'CHIUS INC')							L10/L20: per finger 4 channels (normal force, tangential force, tangential direction, proximity), 5 sites x 4 channels, 0-255 range, sensor type not stated. O6 (reseller): 16 capacitive tactile regions.	USB-to-CAN at 1 Mbit/s, optional Modbus or RS485; rate not stated			no		URDF, with PyBullet and Isaac Gym examples in the SDK	sold
psyonic_ability_hand_2021	PSYONIC (San Diego, CA, USA)		6	linkage (repo file/function names indicate a four-bar finger linkage; not stated on the vendor page itself)				30 touch sensor values streamed (FSR-based per API and URDF 'no_fsr' variant); per-finger split not stated	BLE, I2C, UART or RS485; rate not stated			no		URDF (left, right, large, small, no-FSR) plus MuJoCo and Isaac Sim paths in the vendor repo	sold

33 rows. 216 of 528 cells (40 percent) are values no source stated. A blank price is a hand with no public list price rather than a free hand.

B.2 Where each specification came from

hand	source quality	claim date	caveat recorded in the row	source URL	parsed source in the manifest	note	corpus method rows running it
agibot_omnihand_2025	unavailable	no source: both store.agibot.com URLs 404, fetched 2026-09-17			no manifest entry	papers/notes/agibot_omnihand_2025.md	1: clutterdexgrasp_2025
onex_neo_hand_2026	vendor page	vendor page, 2026-07-09			no manifest entry	papers/notes/onex_neo_hand_2026.md	0
leap_hand_v2_adv_2025	vendor page	undated project page; API repo at commit a0936196, fetched 2026-09-17	the page claims 21 DoF with 17 powered motors and the README calls the same hand 17-DOF	https://github.com/leap-hand/LEAP_Hand_V2_Adv_API	no manifest entry	papers/notes/leap_hand_v2_adv_2025.md	2: bidex_teleop_2024, dexwild_2025
paxini_dexh13_2024	press	press, CES 2026-01-07			no manifest entry	papers/notes/paxini_dexh13_2024.md	0
clone_robotics_hand_2024	press	press, 2025-11-24, on a video of 2025-11-15	'under 2 pounds' is an upper bound, and the source does not say whether the 500 W pump and 36 valves are inside it		no manifest entry	papers/notes/clone_robotics_hand_2024.md	0
sanctuary_phoenix_hand_2024	press	undated press release; trade reprint 2024-12-17			no manifest entry	papers/notes/sanctuary_phoenix_hand_2024.md	0
bidexhand_2025	paper	paper, 2025 (ICRA Dexterity workshop abstract)	the paper states 16 independently actuated DoF and 21 joints; the README states 15 servos driving 15 joints	https://github.com/wengmister/BiDexHand	5 pp, 89f39715, fetched 2026-09-17	papers/notes/bidexhand_2025.md	0
ruka_v2_2026	paper	paper, 2026	the AS5600 encoders are attachable and detachable and are read for calibration and measurement only, not in the control loop	https://github.com/ruka-hand/RUKA	20 pp, d71d6f90, fetched 2026-09-17	papers/notes/ruka_v2_2026.md	0
orca_hand_2025	paper	paper, 2025	the 19.6 N previously tabulated as a fingertip force is the newton equivalent of the 2 kg index-finger payload at a control-imposed current limit, and no source states it as a force	https://github.com/orcahand/orca_core	8 pp, f3967f2e, fetched 2026-09-17	papers/notes/orca_hand_2025.md	0
leap_hand_2023	paper	paper, 2023		https://github.com/leap-hand/LEAP_Hand_API	11 pp, ac08a57e, fetched 2026-09-17	papers/notes/leap_hand_2023.md	11: bidex_teleop_2024, bidexhd_2024, cross_embodiment_world_models_2025, dextrac_2024, dexndm_2025, dexterous_functional_grasping_2023, dextrack_2025, dexwild_2025, dreureka_2024, unidex_2026, videodex_2022
ruka_2025	paper	paper, 2025		https://github.com/ruka-hand/RUKA	13 pp, 3a19161f, fetched 2026-09-17	papers/notes/ruka_2025.md	0
dexhand_open_source_2023	project page and GitHub README	page posts 2023-08-08 to 2023-10-01; no hardware release date	no joint count is given anywhere, so no DoF figure can be quoted; the servo count is 16 finger and thumb micro-servos plus 2 wrist servos, with a third optional. The \$300 is 'additional	https://github.com/TheRobotStudio/V1.0-Dexhand	no manifest entry	papers/notes/dexhand_open_source_2023.md	0

hand	source quality	claim date	caveat recorded in the row	source URL	parsed source in the manifest	note	corpus method rows running it
			total cost of components', excluding printing and the wrist servos				
ilda_hand_2021	paper	paper, accepted 2021-11-05		https://www.nature.com/articles/s41467-021-27261-0.pdf	13 pp, ed0e19c2, fetched 2026-09-17	papers/notes/ilda_hand_2021.md	0
pisa_iit_softhand_2014	paper	paper, IJRR 2014		https://www.centropiaggio.unipi.it/sites/default/files/PisaIIT_Soft_Hand_0.pdf	14 pp, 9c9cd17f, fetched 2026-09-17	papers/notes/pisa_iit_softhand_2014.md	0
faive_hand_2023	paper	paper, 2023	the row's open-hardware status is unconfirmed: the repo has LICENSE and LICENSE-NVIDIA files whose contents were not captured, and no open-hardware claim appears in the sources	https://github.com/sri-ethz/faive_gym_oss	7 pp, 13feb753, fetched 2026-09-17	papers/notes/faive_hand_2023.md	1: graspxl_2024
boston_dynamics_atlas_hand_2026	press	trade press, 2026-01-05 (CES); the Boston Dynamics blog of the same date has no hand content			no manifest entry	papers/notes/boston_dynamics_atlas_hand_2026.md	0
daxo_muscle_v0_2025	third-party tracker	third-party catalogue post, 2026-09-16; daxo-robotics.com returned HTTP 404	120 actuators in 750 g is about 6 g per actuator including structure, tendons, routing and skin; no DoF number is stated, and a redundant tendon architecture with no rigid joints has fewer kinematic DoF than actuators		no manifest entry	papers/notes/daxo_muscle_v0_2025.md	0
figure_03_hand_2025	third-party tracker	undated tracker page, fetched 2026-09-17; figure.ai launch page returned HTTP 404	the tracker's 'Degrees of freedom, hands / 20' sits beside 'Number of fingers / 10', so the 20 is most likely the pair and about 10 per hand; the 16-per-hand figure in circulation has no reachable source		no manifest entry	papers/notes/figure_03_hand_2025.md	0
tesla_optimus_hand_2025	press	Gen 2 video 2023-12-13; V3 patents filed 2024-10, relayed by press 2026-04	the 22 DoF in circulation is a Teslarati article's arithmetic of four DoF on each of five fingers plus two at the wrist, not a Tesla statement; Gen 2's own figure is 11 DoF. This row's actuation cell describes V3 and its tactile cell describes Gen 2	https://image-pubs.uspto.gov/dirsearch-public/print/downloadPdf/9505134	20 pp, 93664ecb, fetched 2026-09-17	papers/notes/tesla_optimus_hand_2025.md	0
xiaomi_cyberone_hand_2026	press	press, 2026-03-30			no manifest entry	papers/notes/xiaomi_cyberone_hand_2026.md	0
shadow_dexterous_hand_2005	datasheet	specification dated 2024-12-04	the four distal joints are coupled and not independently controllable, so the '24 DoF' quoted in learning papers is the joint count	https://shadowrobot.com/wp-content/uploads/2024/12/shadow_dexterous_hand_e_technical_specification_20241204.pdf	11 pp, 66541d02, fetched 2026-09-17	papers/notes/shadow_dexterous_hand_2005.md	21: asyidx_2024, bidexgrasp_2026, bimangrasp_2024, cross_embodiment_world_models_2025, dexdeform_2023, dexman_2025, dp3_2024, eureka_2023, graspxl_2024, maniptrans_2025, mjpc_2022, objdex_2024, openai_dexterity_2018,

hand	source quality	claim date	caveat recorded in the row	source URL	parsed source in the manifest	note	corpus method rows running it
							openai_rubiks_cube_2019, pddm_2019, pgdm_2023, pianomime_2024, resdex_2024, unidex_2026, unidexgrasp_2023, unidexgrasp_pp_2023
proception_prohand_2026	press	launch post and press, 2026-06-29	the evidence is a shipping announcement, not a datasheet: no weight, force, actuated-DoF count or control rate is stated anywhere		no manifest entry	papers/notes/proception_prohand_2026.md	0
sharpa_wave_2026	vendor page	undated page, copyright 2026, fetched 2026-09-17	the page footnotes 'Specifications may vary between products', and no column of the spec table is labelled		no manifest entry	papers/notes/sharpa_wave_2026.md	6: dexteleop0_2026, egoscale_2026, metis_2025, poise_2026, simtoolreal_2026, teledexter_2026
tesollo_dg5f_2024	vendor page	undated page, copyright 2026, download links dated 2026-08		https://en.tesollo.com/wp-content/uploads/2025/02/TESO-LLO-DG-5F-catalogue-En-2.pdf	no manifest entry	papers/notes/tesollo_dg5f_2024.md	0
unitree_dex5_2025	datasheet	undated page, copyright 2016-2025, fetched 2026-09-17		https://github.com/unitreerobotics/unitree_sdk2	no manifest entry	papers/notes/unitree_dex5_2025.md	0
wuji_hand_2025	vendor page	undated page, copyright 2026, fetched 2026-09-17	the page footnotes that these figures 'reflect the current Beta1 product and the SPEC will continue to iterate'		no manifest entry	papers/notes/wuji_hand_2025.md	2: toporetaret_2026, unidex_2026
allegro_hand_v4_2016	driver repo	driver repo, versions 1.0-4.0 undated, fetched 2026-09-17	allegrohand.com/v4 returned HTTP 404 and the Wonik wiki timed out; every cell here comes from the ROS driver repository, and weight, joint torque, payload and price have no reachable source	https://github.com/simlabrobotics/allegro_hand_ros_v4	no manifest entry	papers/notes/allegro_hand_v4_2016.md	35: anyrotate_2024, anyteleop_2023, asymdex_2024, castro_sap_contact_2021, cross_embodiment_world_models_2025, cyberdemo_2024, dexmachina_2025, dexpbt_2023, dexpilot_2020, dexplore_2025, dexpoint_2022, dexteritygen_2025, dexterous_handover_2025, dextrack_2025, dextreme_2022, dime_2022, dp3_2024, dynamic_handover_2023, eureka_2023, geometric_retargeting_2025, graspxl_2024, holo_dex_2022, hora_2022, hudor_2024, human2sim2robot_2025, maniptrans_2025, pang_global_planning_2022, penspin_2024, robot_synesthesia_2023, rotateit_2023, rotating_without_seeing_2023, twisting_lids_2024, unidex_2026,

hand	source quality	claim date	caveat recorded in the row	source URL	parsed source in the manifest	note	corpus method rows running it
							viserdex_2026, wm_dex_human_videos_2025
inspire_rh56dfx_2023	vendor page	undated page, fetched 2026-09-17	the vendor's own page reads 'Degrees of freedom 6, Numbers of joints 12'; the joints column here prints the joint count, so the vendor's DoF figure is the actuated-DoF column		no manifest entry	papers/notes/inspire_rh56dfx_2023.md	19: ace_teleop_2024, articulated_tools_inhand_2025, being_h05_2026, being_h0_2025, bidexgrasp_2026, dextramachina_2025, dextramimicgen_2024, dexplore_2025, dexumi_2025, dexvla_2025, humanoid_policy_human_policy_2025, humanoid_sim2real_recipe_2025, humanplus_2024, maniptrans_2025, metis_2025, okami_2024, omnih2o_2024, open_television_2024, unidex_2026
robotera_xhand1_2024	third-party tracker	undated tracker page, fetched 2026-09-17; the vendor page was not fetched	the 15 N in circulation comes from a survey table, not from the tracker or the vendor; the tracker states no fingertip force		no manifest entry	papers/notes/robotera_xhand1_2024.md	8: cross_embodiment_world_models_2025, deximit_2026, dextramachina_2025, dexora_2026, dexumi_2025, dp3_2024, force_grasp_sim2real_2026, unidex_2026
shadow_dex_ee_2024	vendor page	undated page, copyright 2026, fetched 2026-09-17			no manifest entry	papers/notes/shadow_dex_ee_2024.md	1: demostart_2024
brainco_revo2_2025	datasheet	undated page, fetched 2026-09-17			no manifest entry	papers/notes/brainco_revo2_2025.md	1: bidexgrasp_2026
linkerbot_l20_2025	vendor page	undated page, fetched 2026-09-17		https://github.com/linkerbotai/linker_hand_sdk	no manifest entry	papers/notes/linkerbot_l20_2025.md	1: being_h05_2026
psyonic_ability_hand_2021	vendor page	undated page, fetched 2026-09-17	the four-bar finger linkage rests on repository file names and the five fingers on URDF meshes; the vendor page states neither	https://github.com/psyonicinc/ability-hand-api	no manifest entry	papers/notes/psyonic_ability_hand_2021.md	7: ace_teleop_2024, asymdex_2024, bunny_visionpro_2024, cross_embodiment_world_models_2025, dextramachina_2025, hato_visuotactile_2024, unidex_2026

33 rows. 31 of 264 cells (11 percent) are values no source stated. `parsed source in the manifest` gives page count, the first eight characters of the sha256 of the stored copy, and the date it was fetched. A hand with no manifest entry was recorded from a page that could not be stored, and its row should be read as a claim with no retrievable source.

Appendix C. Per-paper reward term extraction

Table 5 marks nine recurring term families across the in-hand reorientation methods. This appendix is the extraction underneath it, over every method row that records a reward at all. `terms stated` counts the terms the paper itself names. `where the reward was read` gives the section of the paper or the file in the released code that the note quotes. A row with a `paper/code class` is one where the two disagree, and C.2 prints the disagreement in full. Cells and entries reproduce the text stored in the row, with the source's own punctuation and wording.

C.1 The extraction

method	yr	paradigm	terms stated	term names as the paper names them	where the reward was read	in Table 5	code released	paper/code class	confidence
openai_dexterity_2018	2018	RL	3		Method, reward (Sec 4.2 / App C.1)	yes	no		
openai_rubiks_cube_2019	2019	RL, distillation	3		Method, reward (Sec. 6.1)	yes	no		
pddm_2019	2019	MPC					yes	contradiction	high
dexpoint_2022	2022	RL	4	reach, contact, lift, action penalty			yes	contradiction	high
dextreme_2022	2022	RL	6		Reward block, paper Table 2 against allegro_hand_dextreme.py and the two DR yamls	yes	yes	contradiction	high
dexvip_2022	2022	RL	4				no		
hora_2022	2022	RL, distillation	5		Block C, paper Sec 3.1/App. C against configs/task/AllegroHand Hora.yaml	yes	yes	version-skew	high
mjpc_2022	2022	MPC	3				yes		
visual_dexterity_2022	2022	RL, distillation	7	c1_sparse_task, c2_dense_task, c3_fingertip_dist, c4_energy, c5_push_away_penalty, c6_table_contact, c7_finger_height	Method, Table S1 against dexenv/envs/rewards.py and dexenv/conf/task/dclaw.yaml	yes	yes	contradiction	high
aloha_act_2023	2023	BC					yes	internal-inconsistency	high
artigrasp_2023	2023	RL	4	r_p, r_c, r_reg, r_task			yes	parse-limitation	low
dexpbt_2023	2023	RL	4	r_reach, r_pick, r_targ, r_vel	Block C, paper Sec III-C/Table II against allegro_kuka_base.py and AllegroKuka.yaml	yes	yes	contradiction	high
dexterous_functional_grasping_2023	2023	RL	2	r_threshold, r_hand-obj			no		
diffusion_policy_2023	2023	BC, diffusion					yes	parse-limitation	high
dynamic_handover_2023	2023	RL	3				no		
eureka_2023	2023	RL			Method, App. G.1/G.2 LLM-authored rewards; generated code not checked into the repo	yes	yes	parse-limitation	low
pgdm_2023	2023	RL	2				no		
physhoi_2023	2023	RL	4	body motion reward, object motion reward, ig reward, cg reward			yes	contradiction	high
robot_synesthesia_2023	2023	RL, distillation	6		Method, reward (Sec IV.A.3 eq. 1); no code released	yes	no		

method	yr	paradigm	terms stated	term names as the paper names them	where the reward was read	in Table 5	code released	paper/code class	confidence
rotateit_2023	2023	RL, distillation	6		Block C (paper Sec 3.1, App. B); no code released	yes			
rotating_without_seeing_2023	2023	RL	6	r_rot, r_vel, r_fall, r_work, r_torque, r_dist	Reward block (Sec IV.A.3, App. D Eq. 9); no code released	yes			
unidexgrasp_2023	2023	grasp-synthesis, RL, distillation	4	goal, reach, lift, move			yes	contradiction	high
unidexgrasp_pp_2023	2023	RL, distillation	4	r_reach, r_lift, r_move, r_bonus			yes	parse-limitation	low
anyrotate_2024	2024	RL, distillation	10	keypoint distance reward, rotation reward, goal bonus reward, good contact reward, bad contact penalty, angular velocity penalty, pose penalty, work penalty, torque penalty, termination penalty	Block C (App. B.1 weights); no code released	yes	no		
asymdex_2024	2024	RL	4				yes	code-absent	medium
bidexhd_2024	2024	RL, distillation	4	r_appro, r_lift, r_bonus, r_track			yes	parse-limitation	high
bimangrasp_2024	2024	grasp-synthesis, diffusion	7				no		
demostart_2024	2024	RL, distillation	1		Method, reward (App. VII-A.1); no code released	yes	no		
dexmimicgen_2024	2024	data-collection, BC	1				yes	code-absent	high
dp3_2024	2024	BC, diffusion					yes	internal-inconsistency	medium
dreureka_2024	2024	RL	4	ang_z_vel_reward, lin_vel_penalty, object_fall_penalty, deviation_penalty	Method, App. B1 Prompt 13; cube-rotation code absent from the released repo	yes	yes	code-absent	high
graspxl_2024	2024	RL	8	r_dis, r_v, r_ω, r_m, r_c, r_f, r_reg, r_anatomy			yes	parse-limitation	low
hudor_2024	2024	RL	1				no		
humanplus_2024	2024	RL, BC	8				yes	parse-limitation	medium
objdex_2024	2024	BC, RL, distillation	3	rotation, translation, joint angle			no		
omnigrasp_2024	2024	RL, distillation	3	r_approach, r_pre-grasp, r_obj			yes	parse-limitation	medium
omnih2o_2024	2024	RL, distillation	24				yes	contradiction	medium
open_television_2024	2024	teleop-system, BC					yes	code-absent	low
penspin_2024	2024	RL, distillation	7		Method, paper Table 4 against penspin/tasks/allegro_hands_hora.py and	yes	yes	contradiction	medium

method	yr	paradigm	terms stated	term names as the paper names them	where the reward was read	in Table 5	code released	paper/code class	confidence
					configs/task/AllegroHand Hora.yaml				
pi0_2024	2024	VLA, flow					yes	version-skew	high
pianotime_2024	2024	RL+demo, distillation	2				yes	contradiction	high
resdex_2024	2024	RL, distillation	7	task, proposal, pose, reach, lift, move, bonus			yes	parse-limitation	medium
twisting_lids_2024	2024	RL	5				no		
articulated_tools_inhand_2025	2025	RL, BC, distillation	8	r_pos, r_quat, r_goal, r_timer, r_inc, r_contact, r_slip, r_act			no		
being_h0_2025	2025	VLA, BC					yes	code-absent	high
clutterdexgrasp_2025	2025	RL, diffusion, distillation	4	r_pos, r_neg, r_grasp, r_force			no		
cross_embodiment_world_models_2025	2025	world-model, MPC	2				no		
dexgraspvla_2025	2025	VLA, diffusion	1				yes	parse-limitation	high
dexmachina_2025	2025	RL	4				yes	internal-inconsistency	low
dexman_2025	2025	RL, trajopt	3				no		
dexndm_2025	2025	RL, distillation	7		Reward block (Sec. 3.1, App. A.1); no code released	yes	no		
dexplore_2025	2025	RL, distillation	7	r_j^h, r_r^h, r_p^o, r_r^o, r_d, r_c, r_energy			yes		
dexremoe_2025	2025	RL	5	csuccess, cdist, crot, cω, ca	Method, reward (Sec. III-B, Table II); no code released	yes	no		
dexteritygen_2025	2025	RL, diffusion			Method, reward/objective: the RL reward behind the pretrained controller is not quoted in the source	yes	no		
dexterous_handover_2025	2025	RL	3	r_BASE, contact_term (c-w_c), r_MANIP					
dextrack_2025	2025	RL+demo	5	r_o,p, r_o,q, r_wrist, r_finger, r_affinity			yes	parse-limitation	low
dexvla_2025	2025	VLA, diffusion	2				yes	parse-limitation	high
dydexhandover_2025	2025	RL	5	R_dist, R_obj, R_contact, P_action, P_hand			no		
geometric_retargeting_2025	2025	teleop-system	5	Ldir, Lcover, Lflat, Lpinch, Lcol			yes		
groot_n16_2025	2025	VLA					yes	version-skew	high
groot_n1_2025	2025	VLA, flow					yes	version-skew	high
human2sim2robot_2025	2025	RL	1	r_obj			yes		

method	yr	paradigm	terms stated	term names as the paper names them	where the reward was read	in Table 5	code released	paper/code class	confidence
humanoid_policy_hum an_policy_2025	2025	BC	2				yes		
humanoid_sim2real_r ecipe_2025	2025	RL, distillation	2				no		
maniptrans_2025	2025	RL, BC, diffusion	5	r_wrist, r_finger, r_smooth, r_object, r_contact			yes	parse-limitation	low
metis_2025	2025	VLA	2	Lar, Laction			no		
pi05_2025	2025	VLA, flow					yes	code-absent	high
pistar06_2025	2025	VLA, RL, flow					yes	code-absent	high
wm_dex_human_videos _2025	2025	world-model, MPC	2				no		
being_h05_2026	2026	VLA, flow					yes	code-absent	high
dexteleop0_2026	2026	teleop-system, MPC	3				no		
force_grasp_sim2rea l_2026	2026	RL	7	R_torque, R_Force, R_diff, R_outer, terminal- state penalty, R_action, R_vel	Method, reward (Sec 3.5.1, Table 2); no weights survive the parse; no code released	yes	no		
poise_2026	2026	RL	4	pose reaching, goal completion, grasp maintenance, drop and actuation regularization	Method, reward (Sec. IV- D); weights lost to the parse; no code released	yes	no		
simtoolreal_2026	2026	RL	5	r_smooth, r_approach, r_lift, r_goal, b_succ			no		
teledexter_2026	2026	RL, teleop-system			Reward block (Sec. 3.1); term formulas lost to the parse; no code released	yes	no		
toporetargt_2026	2026	trajopt, RL	4	object, link-position, joint-position, action- smoothness			no		
viserdex_2026	2026	RL, distillation	10	orientation tracking, success bonus, object dropped, object distance, object velocity, joint velocity, action magnitude, action rate, joint work, joint torques	Method, reward (App. Table IX); no code released	yes	no		

77 rows. 253 of 770 cells (32 percent) are values no source stated. A blank `terms stated` with a filled `term names` column is a paper that names its terms without numbering them. `in Table 5` marks the rows that are also in the body matrix, which covers in-hand reorientation only.

C.2 Paper against released code, in full

Each entry below is the disagreement text stored in the row, unedited. The class is what Sec. 5.8 and Sec. 8.1 count. `contradiction` means the paper states one value and the shipped code demonstrably states another. `parse-limitation` means this survey's own parse could not settle it and the accusation is withdrawn. `code-absent` means the described component is not in the released repository. `version-skew` means the repository is a later generation than the paper. `internal-inconsistency` means the paper disagrees with itself and no code is implicated.

contradiction, 10 rows.

- `pddm_2019` (high). Table 2 states obs-dim 46 for In-hand Reorientation while the released `cube_env.py` code sums to 39; the Baoding reward code includes an extra `-10*wrist_too_high` term absent from Table 2's printed formula.
- `dexpoint_2022` (high). The released code's reward adds several terms absent from the paper's four-term Eq. 5 (a lift-threshold bonus, a target-distance term, a rotation bonus, and an IK controller-tracking penalty) and reshapes the reach/lift terms into inverse-distance and clipped forms rather than the paper's plain distance/height-difference formulas.
- `dextreme_2022` (high). Action Delta Penalty weight is -0.25 in paper Table 2 but -0.2 in the ADR yaml and -0.01 in the ManualDR yaml; the Joint Velocity Penalty in code normalises velocity by $(\text{max_velocity} - \text{vel_tolerance})$ unlike the paper's stated formula; code has a `timeout_rew` term absent from the paper's reward table; Appendix Table 12 states critic learning rate $5e-4$ and KL threshold 0.16, vs body text/code values of $5e-5$ and 0.016.
- `visual_dexterity_2022` (high). Table S1 states 32000 teacher training environments, but the released config sets `alg.num_envs` to 8000 (parent config 16384); `fallDistance` differs across two shipped configs (0.24 vs 0.15, only the latter matching Table S1's threshold); the paper's Eq 8 penultimate-joint penalty ($c7=-2$) does not appear anywhere in the released reward code; the config carries a `dead distRewardScale=-10.0` key never used in `compute_reward`; and the paper's table-friction lower bound (0.05) differs by a factor of 10 from the code's randomized lower bound (0.005).
- `dexpbt_2023` (high). Paper presents the reward as 4 mutually exclusive stage terms (`r_reach`, `r_pick`, `r_targ`, `-r_vel`), but code's `compute_kuka_reward` sums 8 named components, one of which (`hand_delta_penalty`) is multiplied by 0 and disabled; there is no single `r_vel` term in code, instead separate kuka/allegro action penalties whose exact formula is not shown. Also, the paper reports zero experiments with domain randomization, yet the shipped `AllegroKuka.yaml` already carries a fully specified DR schedule (disabled via `randomize: False`). Review: R3 adversarial review: the disabled-randomisation half is withdrawn, since the note finds it consistent with the paper; the zeroed reward term stands
- `physhoi_2023` (high). The released code hardcodes the body position-velocity error and the object rotation/rotation-velocity errors to zero in `compute_humanoid_reward`, so despite Table 4 listing nonzero $\lambda^{\text{or}}=0.1/\lambda^{\text{orv}}=0.01$ weights for GRAB, the trained reward never actually tracks object orientation (position-only in practice).
- `unidexgrasp_2023` (high). The paper describes a four-term weighted reward (`r_goal` + `r_reach` + `r_lift` + `r_move` via Table 7's omega weights) but the released `compute_hand_reward` implements a different threshold-gated torch.where cascade with distinct hardcoded coefficients that do not map one-to-one onto the paper's weights.
- `omnih2o_2024` (medium). stumble weight -0.00125 (paper) vs -1250 (code); max-feet-height sign/magnitude differ (+1000 paper vs -2500 code, a penalty not a bonus); paper's $\exp(-c*//.//)$ form vs code's $\exp(-\text{err}^2/\text{sigma})$; curriculum level-down threshold 40 (paper) vs 50 (code) Review: R3 adversarial review: plausible alternative reading; held at medium confidence pending a direct code read
- `penspin_2024` (medium). The released code disables the paper's tactile observation channel and zeroes the described disturbance-force domain randomization, and the code's reward scale-key names (e.g. `rotate_reward_scale`, `pencil_z_dist_penalty_scale`) do not 1:1 name-match the paper's Table 4 weight list, though matched values agree. Review: R3 adversarial review: plausible alternative reading; held at medium confidence pending a direct code read
- `pianomime_2024` (high). Paper's Table 3 states 2 weighted reward terms (Key Press 2/3, Mimic 1/3), but the released code sums roughly 5 unweighted terms (key press doubled, sustain, energy and fingering hardcoded to return 0, forearm-

collision) plus a separately-added mimic wrapper term.

internal-inconsistency, 3 rows.

- `aloha_act_2023` (high). Algorithm 1 pseudocode states $L_{\text{reconst}} = \text{MSE}$, but Sec.IV.C's prose explicitly states L1 loss is used instead; the algorithm box and the implementation text disagree; code/md captures only function signatures, not bodies, for `policy.py`'s loss implementation
- `dp3_2024` (medium). the paper's prose states the network predicts the noise added to the data, but the shipped default config trains with `prediction_type: sample` (predicting the denoised action a^0 directly), not epsilon; the paper only qualifies this later in the same section (Fig. 7 ablates both).
- `dexmachina_2025` (low). paper describes a plain weighted sum $\lambda_{\text{taskr_task}} + \lambda_{\text{imir_imi}} + \lambda_{\text{bcr_bc}} + \lambda_{\text{conr_con}}$ with unspecified weights; code implements a multiplicative task term with per-component beta decay, an unmentioned 0.1 force-penalty term, and curriculum-driven decay of auxiliary weights not described as such in the paper Review: R3 adversarial review: the multiplicative form credited to the code is printed in the paper itself

version-skew, 4 rows.

- `hora_2022` (high). The parsed code commit's joint-noise range, disabled default disturbance force, and default cube object type differ from the paper's stated values; the code README itself says to use tag v0.0.1, not the parsed commit, to reproduce paper numbers.
- `pi0_2024` (high). The released repo (`openpi`, commit 215abfb2) has evolved past this paper: it also documents `pi0-FAST` (autoregressive) and `pi0.5` checkpoints/configs not described in this paper, which describes only the flow-matching `pi0` and the non-VLM `pi0-small` ablation.
- `groot_n16_2025` (high). code/md is the current Isaac-GR00T main branch (N1.7 generation, commit 51d4c89f), not the N1.6 checkpoint; the repo names a separate `n1d6` branch for N1.6 that was not fetched. The page states N1.6's backbone is "an internal NVIDIA Cosmos-2B VLM variant," but the repo's own changelog states the Eagle backbone (`nvidia/Eagle-Block2A-2B-v2`) was used through N1.6 and only replaced by Cosmos-Reason2-2B in N1.7 -- whether the page's Cosmos-2B variant is a third distinct model or the page is describing N1.7 under an N1.6 headline is not resolvable from what was fetched. Review: The page attributes a Cosmos backbone to N1.6 while the repository changelog says N1.6 used Eagle and N1.7 made the switch, so the repository is a later generation than the page describes.
- `groot_n1_2025` (high). The parsed code repo (Isaac-GR00T, commit 51d4c89f) is a later N1.7 generation, not the paper's GR00T-N1-2B: VLM backbone changed from the paper's Eagle-2 to Cosmos-Reason2-2B (via Qwen3-VL); action horizon expanded from the paper's stated $H=16$ to 40; state/action dimensions expanded from 29 to 132; DiT layers changed from 32 to 16; embodiment scope widened to include Unitree G1, AgiBot G1, and a 'YAM' arm not described in the paper.

code-absent, 8 rows.

- `asymdex_2024` (medium). Code contains a GraspAndPlace-task reward variant (`AllegroHandDualArmGraspAndPlaceAsymDex.py`) with exponential hand-distance and block-alignment terms not one of the four BiDexHands tasks described in the paper body.
- `dexmimicgen_2024` (high). The environment code includes an unused 'reward_shaping' branch (if `self.reward_shaping: pass`) that is never described or used in the paper text; only the sparse binary completion reward is used, as the paper states.
- `dreureka_2024` (high). The released DrEureka repo (code/md) contains only `forward_locomotion/` and `globe_walking/` trees; the LEAP-hand cube-rotation training/deployment code and the quoted Prompt-13 reward are the LEAP-hand authors' own code, not included in the released repo, so the cube-rotation reward cannot be cross-checked against a code file.
- `open_television_2024` (low). The paper's prose states 25k training iterations, $\text{lr } 5\text{e-}5$, batch size 45, but the code repo's example training command instead shows 50000 epochs and an explicit `kl_weight=10` not mentioned anywhere in the

paper text. Review: R3 adversarial review: a behaviour-cloning paper with no reward; the charge compared paper iterations to a README example's epochs

- **being_h0_2025** (high). The actual downstream action-head module referenced in §4.3 (learnable query tokens, proprioceptive projector f_p , regression head f_r) is not present in the parsed code; the repo's own TODO list marks 'Training code and scripts' as unreleased, so the action-chunk-length=16 value and exact proprioception-vector fields come only from inference/eval CLI examples, not a verifiable implementation.
- **pi05_2025** (high). The paper describes a hybrid discrete-FAST-token-pretrain-then-flow-matching-post-train recipe (Eq. 1, α : 0 to 10), but the released repo states plainly it 'currently only support[s] the flow matching head for both pi0.5 training and inference.'
- **pistar06_2025** (high). The general openpi framework (pi0/pi0.5/FAST, flow-matching action head, ALOHA/DROID/LIBERO adapters) is released at commit 215abfb2, but the RECAP-specific pieces described in the paper -- the advantage-conditioning input, the distributional value function, the RECAP training loop (Algorithm 1), and the espresso/laundry/box-assembly checkpoints and datasets -- are not present in the parsed repo.
- **being_h05_2026** (high). Released code (commit e12ac44f) exposes only benchmark/inference/config scaffolding (policy.py, beingh_policy.py, dataset-transform configs). no MoT/MoF model definition, no rectified-flow action-expert source, and no config for any dexterous-hand embodiment; the only post-train configs present (libero_robotcasa.yaml, libero_all.yaml, robocasa_human.yaml, so101_example.yaml) are all gripper-only, so the dexterous-hand claims cannot be cross-checked against released code.

parse-limitation, 13 rows.

- **artigrasp_2023** (low). paper's two regulariser weights ($w_{rh}=0.5$, $w_{ro}=0.2$) correspond to five separate velocity penalties in code (-0.5,-0.2,-0.5,-0.5,-0.3); the fingertip weight 12.0 and λ 5.0 in Table 6 do not appear in the yaml; reward weights differ between curriculum phase-1 and phase-2 configs though Table 6 reports only one set Review: R3 adversarial review: the reward body is C++ and absent from the parse; the weight split is the paper's documented curriculum
- **diffusion_policy_2023** (high). the paper's core loss equations (DDPM training loss, EBM/InfoNCE, score-matching) are unrecovered images in the parsed text, and code/md is a signature-only API map, so compute_loss's body could not be verified from code; only the exact shipped noise-scheduler config (DDPMScheduler, 100 train steps) is confirmed.
- **eureka_2023** (low). The released repo's default eureka/cfg/config.yaml ships iteration:1, sample:3 (a fast-test default), not the paper's reported experimental configuration of $N=5$ iterations $\times$ $K=16$ samples $\times$ 5 independent runs, and LLM-generated reward code itself is never checked into the repo (written at runtime to a gitignored outputs/ directory). Review: R3 adversarial review: the repo README documents the paper's defaults; the cited config block is a fast-test preset
- **unidexgrasp_pp_2023** (low). The paper's four named reward weights (0.5, 0.1, 2, 10) numerically match literals in the released compute_hand_reward, but the code never labels them with the paper's variable names, and the paper's own display equations were dropped from the PDF, so the correspondence is inferred rather than confirmed by either source.
- **bidexhd_2024** (high). Table 5 weights (w_r , w_t , w_{1-w4} , $\lambda_w=0.12$, $\lambda_{ft}=0.48$) cannot be matched to code reward lines because the code md truncates the reward function before stage-2/final sum; λ_w and λ_{ft} numerically match the code's grasp-condition thresholds (0.12, 0.12x4 fingers) rather than confirmed reward weights
- **graspxl_2024** (low). The code splits the paper's single $r_{reg} = -w_h//u_h//^2 - w_o//u_o//^2$ into four separately-weighted terms (wrist vel, wrist qvel, object vel, object qvel), and the floating-phase object-velocity coefficient (-1.5) does not match the paper's stated $w_o=0.1$, so the regularization weighting cannot be reconciled 1:1 with Table 9. Review: R3 adversarial review: the environment source is not in the parse and the note says the reward cannot be verified
- **humanplus_2024** (medium). code implements a superset of legged_gym reward terms (lin_vel_z, ang_vel_xy, orientation, base_height, torques, dof_vel, dof_acc, action_rate, collision, termination, dof/torque limits, tracking_lin/ang_vel, feet_air_time, stumble, stand_still, feet_contact_forces, target_jt) beyond the 8 terms documented in the paper's Table 1

- `omnigrasp_2024` (medium). The code's `compute_pregrasp_reward_time` hard-codes `w_pos`, `w_rot` = 0.9, 0.1, silently overriding the `rwd_specs` weight arguments passed into the function; separately, the paper's pre-grasp reward weights (`w_hp`, `w_hr`) are not resolvable from either the paper text or Appendix Table 7.
- `resdex_2024` (medium). The base-policy and residual-policy task files implement different variants of `compute_hand_reward` (whether finger/hand-distance terms are included alongside `goal_hand_rew`/`hand_up`/`bonus` differs), a difference the paper's Sec 4.1/4.4 text does not describe, and the released configs' env count (4096 default) and iteration counts (2500/5000) do not match the paper's stated 11,000 envs and 5,000/20,000 iterations.
- `dexgrasvla_2025` (high). Hardware-interfacing code is explicitly withheld ("Due to intellectual property constraints, we are unable to open-source the hardware-related code"); only function signatures for the controller loss (`compute_loss`, `noise_assignment`) are present in `code/md`, not their bodies, so the MSE/Immiscible-Diffusion loss cannot be verified against code.
- `dextrack_2025` (low). The paper gives one canonical reward-weight table (Table 3), but the released repo's task-config YAMLs contain several divergent, unreconciled reward-coefficient sets across AllegroHandTracking/... TrackingGeneralist/other variants, and it is not identifiable from the parsed source which config produced the paper's headline Table 1 numbers.
- `dexvla_2025` (high). `code/md` only carries class/function signatures, not bodies, for the diffusion head (`modeling_scaledp.py`) and no config file with the α weight was captured; the paper's diffusion loss and denoising process cannot be verified against code bodies. only the multi-head ScaleDP signature and a DDIM fp32 patch are available, so none of the paper's tables can be reproduced from what was parsed.
- `maniptrans_2025` (low). paper's PPO learning rate ($5e-4$) and env count (4096) differ from the shipped README/yaml defaults ($2e-4$, 8192) Review: R3 adversarial review: the cited values are a README override and an unused fallback; the repo's own config matches the paper

Appendix D. The existing surveys, and what each covers

Fourteen corpus entries are themselves surveys or engine-comparison studies. Table 10 sets them on one set of columns. The columns record what each work covers, not how well. Section 1 states what this survey adds over them and does not repeat the comparison here.

`an_dexil_survey_2025` is the closest in subject, covering imitation learning for multi-fingered hands by learning family, end-effector class and demonstration source. It gives reinforcement learning no taxonomy, treats bimanual work as a single "multi-agent" subsection, compares no simulator, and defines no evaluation metric. `welte_iil_survey_2025` finds only seven dexterous works that use interactive imitation learning and carries a fifteen-hand commercial table, with no bimanual section, no benchmark table and no contact modelling. `zhao_sim2real_survey_2020` supplies the standard sim-to-real split and predates GPU-parallel simulation. `firoozi_foundation_models_2023` has zero occurrences of bimanual, tactile or in-hand.

`bai_unified_manip_survey_2025` spans all of manipulation across 212 pages. Its Sec. 4.3 on dexterous manipulation runs about 720 words, the third longest of its ten task subsections behind grasping and quadrupedal manipulation, which is a real treatment and not a passing mention. The difference is elsewhere. Its Sec. 1.2 lists dexterous manipulation among the topics that "existing surveys" cover from "narrower perspectives" and defers it to two of them.

`zhao_dexhand_survey_2026` is the most recent hand-centred survey, with a 29-hand anatomy table and a task-by-paradigm taxonomy. Its Sec. IV-C names the reference-versus-rollout split that Section 1 credits it with, and its Sec. III-F is the one bimanual subsection, 16 cited works with no coordination analysis.

`nine_physics_engines_review_2024` is the predecessor closest to the engine comparison, and it reviews Brax, Chrono, Gazebo, MuJoCo, ODE, PhysX, PyBullet, Unity and Webots for reinforcement learning research. It scores each on documentation, model and environment creation, URDF and MJCF support, and readiness for multi-agent work. It runs no benchmark of its own and says so in its Sec. V, that implementing the same scenarios across nine engines "goes beyond the

scope of this paper". Its running bodies are ant-and-humanoid RL benchmarks rather than hands, and it discusses no timestep, no friction model, no contact formulation and no penetration.

Table 10. Existing surveys and what each covers

survey	yr	scope	taxonomy used	bimanual covered	hardware covered	evaluation covered	gaps it names
isaac_sim_2026	2026	one simulator's ecosystem and application domains, reviewed rather than measured	qualitative capability matrix, Table 1, over simulators	one cited GR00T task called bimanual, no hand, DoF or number	no hand named anywhere in paper or code parse	none. The paper runs no experiment of its own	computational cost, configuration complexity, learning curve
zhao_dexhand_survey_2026	2026	dexterous hands end to end: hardware anatomy, methods, datasets, directions	five task categories, each split by learning paradigm. Hardware is split by actuation, transmission and perception	one subsection, III-F, 16 cited works, no coordination analysis	Table I, 29 hands, 12 columns, secondary values	names two layers: physical plausibility including penetration before execution, success during it. No threshold, method or count	hardware feasibility, perception fusion, learning beyond benchmark-centric optimisation, industrialisation, absent evaluation standards
an_dexil_survey_2025	2025	imitation learning for multi-fingered end-effectors	IL family by end-effector class by demonstration source. RL gets none	one subsection, II.E, framed as multi-agent	hands named in prose, no hand table	no metric defined and no trial count. Calls for protocols, proposes none	contact dynamics in engines, data-collection standards, cross-hand transfer, failure datasets, end-effector morphology
bai_unified_manip_survey_2025	2025	all of robot manipulation. Dexterous manipulation is one of ten task subsections, about 720 words, and Sec. 1.2 defers it to other surveys	high-level planning, action modelling, actuation control, plus a bottleneck taxonomy of data and generalisation	bimanual means two arms. One dual-hand mention in the paper	hands named, no DoF or actuation table	success rate and checkpoint selection, six lines. Never says how success is judged for a dexterous task	no scaling law, sim-to-real for contact-rich tasks, fragmented datasets, reliability as important as success
welte_iil_survey_2025	2025	interactive imitation learning, seven dexterous works found	IIL feedback type, plus a keyword bibliometric of 326 papers	three mentions in 687 lines, no section	Table 1, 15 commercial hands, manufacturer figures	no metric defined, no benchmark table	tactile feedback, long-horizon tasks, generalisation, human-feedback interface
nine_physics_engines_rev_2024	2024	nine physics engines for RL research, scored on documentation and usability	13-axis feature and usability matrix, Table II, plus citation-count popularity	none. MARL readiness is the multi-agent axis	none. Ant and humanoid RL bodies are the running examples	no benchmark of its own. Throughput claims are second-hand	no cross-engine MARL performance comparison exists in the literature
contact_models_comparison_2023	2023	LCP, CCP, RaiSim and NCP contact models re-implemented in one framework and ranked	contact model by solver, with the physical property each one violates	none	one Allegro hand as a benchmark system, a ball dropped into it	NCP criterion, self-consistency against a 1e-5 s reference, iteration cost	no fully satisfactory contact model, and gradients through simulation artifacts are unexplored
firoozi_foundation_models_2023	2023	foundation models in robot decision-making, perception and embodied AI	background, robotics, and robotics-adjacent papers, then by application	none, zero occurrences	none. Parallel-jaw end effectors throughout	none defined. Benchmarking appears as a reproducibility problem	data scarcity, variability, uncertainty, safety, real-time inference, and simulators that neglect contact physics
zhao_sim2real_survey_2020	2020	sim-to-real transfer in deep RL, eight pages, 21 works tabulated	zero-shot, system identification, domain randomisation, domain adaptation, learning with disturbances, simulator choice	none, zero occurrences	two cited hand works, no hand table	no metric defined, no trial count, no success rate	domain randomisation has no formal account, and domain adaptation assumes matched feature spaces
piazza_century_2019	2019	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>
physics_engine_comparison_2015	2015	five engines on one shared model: speed, self-consistency, conservation, grasp stability	none. Four test systems, one comparison per test	none	one 35-DoF rig modelled on the Shadow Hand	largest timestep that holds a grasp, and a speed-accuracy Pareto curve	restricted feature subset by design, and the authors are MuJoCo's developers
roa_suarez_grasp_quality_2015	2015	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>
ma_dollar_dexterity_2011	2011	<i>PDF fetched 2026-09-18 after the note was written, no note read from it</i>	<i>PDF fetched 2026-09-18 after the note was written, no note read from it</i>	<i>PDF fetched 2026-09-18 after the note was written, no note read from it</i>	<i>PDF fetched 2026-09-18 after the note was written, no note read from it</i>	<i>PDF fetched 2026-09-18 after the note was written, no note read from it</i>	<i>PDF fetched 2026-09-18 after the note was written, no note read from it</i>
okamura_overview_2000	2000	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>	<i>source not obtained</i>

14 rows, one per corpus entry of class `survey`, 3 of which could not be obtained and are entered as such, and 1 of which was fetched too late to be read into a note. Every filled cell is read from `papers/notes/<key>.md` and is checked against a quotation from that note by `tools/make_survey_table.py`. A cell whose evidence has gone from the note prints empty rather than printing an unchecked claim. The columns record what each work covers, not how well, and a blank cell in `bimanual covered` or `hardware covered` is a scope decision by its authors rather than a failure.

Table 10's last rows carry the cost of the corpus. `okamura_overview_2000`, `piazza_century_2019` and `roa_suarez_grasp_quality_2015` are behind publisher paywalls with no author-hosted copy found on 2026-09-18, and `bicchi_hands_2000` is in the same position with no row at all. `ma_dollar_dexterity_2011` is a different case. The fetch that failed when its note was written succeeded afterwards, so a seven-page PDF is on disk with a recorded hash, and no note has been read from it. All five are cited by metadata only and nothing in this survey describes their contents. The open chapter `bicchi_grasping_chapter_2001` overlaps the paywalled Bicchi paper without being identical to it, so it is quoted in its own right.