

Learning Dexterous Manipulation: Hands, Simulators, Training, and Evaluation

Luai Abuelsamen

luai_abuelsamen@berkeley.edu

Abstract—A dexterous hand can change an object’s pose without setting it down. This survey covers how learning-based systems acquire that ability: hand design and tactile sensing, contact simulation, policy training, bimanual control, and evaluation. Its evidence base is 218 structured records, every claim traced to a parsed source or a repository at a named revision, so it reports what the literature states and what its artefacts show. Three findings recur. First, experiments concentrate on four established hands, while eight hands that can be bought or built from published designs appear in no method paper here. Second, no closed-loop policy in the corpus reports interpenetration for the rollouts of its own trained policy, although simulators already in the corpus compute the quantity. Third, nine of the 62 methods with inspectable code state a training objective in the paper that differs from the one in the repository; each comparison names a public file and revision, and their authors were written to before posting and asked to correct it. The survey reruns and ranks nothing, reporting what the present literature can support and proposing an evaluation protocol for the missing quantities. Because an absent value may reflect this survey’s extraction as well as its source, every coverage count is a lower bound.

Index Terms—Dexterous manipulation, multi-fingered hands, bimanual manipulation, reinforcement learning, imitation learning, physics simulation, contact modelling, benchmarking, reproducibility.

I. INTRODUCTION

A hand is dexterous when it can change an object’s pose without putting the object down. That property, and not the finger count, is what separates a hand from a gripper. Bicchi [1] draws the same line in Sec. 1.1, between restraining an object and “manipulating objects with fingers, in contrast to manipulation with the robot arm”. Restraint is a static question about whether the contacts prevent motion. In-hand manipulation is a dynamic question about whether contacts can be broken and remade while the object stays held. An et al. [2] put the same idea in their abstract as the ability “to skillfully control, reorient, and manipulate objects through precise, coordinated finger movements and adaptive force modulation”. Both definitions place the work in the fingers rather than in the arm.

What follows is a survey in the ordinary sense, covering hands, simulators, training, bimanual control and evaluation, and it is built so that a second question can be asked of it. Fourteen entries in this corpus are themselves surveys of this

field. What this one adds to them is a habit rather than a subject: every claim is carried back to something a reader can open, a parsed source, a vendor’s own specification, or a released file at a named revision. That habit is also what turned up the third finding, since no prior work in robotics we are aware of reads a field’s released code against the rewards its own papers describe.

The analytic theory answered the static question and stalled on the dynamic one. Form closure has a first-order test on the grasp matrix and known contact counts, four in the plane and seven in three dimensions for any polyhedron, per [1] Sec. 1.3.1. Force closure adds the wrench balance and the hand Jacobian. Neither delivers what a controller needs. The same chapter states in Sec. 1.5 that “force closure does not guarantee stability”, and its Sec. 1.8 names the reason the theory could not be pushed further, which is that “the nonsmooth nature of grasp dynamics, because of the unilateral constraints on displacements and forces, has made a thorough analysis very difficult”. That verdict is one architect’s, on one chapter, and this survey did not survey the tradition it judges. Learned control did not solve those modelling problems. It went around them by sampling a simulator instead of solving a model, and by scoring a rollout instead of certifying a configuration. Of the 112 method papers in this corpus, 53 train with reinforcement learning and 35 run in Isaac Gym, against 7 on its successors Isaac Lab and Isaac Sim. The routes drawn in Fig. 1 converge on that recipe: almost everything ends in a policy trained in a GPU simulator and distilled to vision.

The survey has three main findings. The first concerns the link between papers and their code. Of the 112 method papers, 62 released code that could be compared with the paper. Thirty-eight comparisons required an explanatory note, and nine contain a direct conflict between a published value and the corresponding public implementation. PhysHOI [3], for example, assigns weights to object-rotation errors in its paper but sets those errors to zero in the public reward implementation. Its reported success criterion uses position alone, so it does not reveal the difference. This is a comparison between two public documents, not a claim about which code produced the reported experiments. Section V-F explains the classification, and Table III provides the repository, revision, file, and values for all nine cases. It also records seven rows withdrawn during review and two narrowed, and explains how corrections will be handled. Fig. 9 draws the narrowing.

The second finding concerns physical plausibility. How far a hand passes into the object it is holding is not a number the policies in this corpus report. A method row is one method

Correspondence: luai_abuelsamen@berkeley.edu. Every claim in this paper traces to a structured note taken from a parsed source, and every count is recomputed from those notes at build time. The corpus, the notes and the build that produce them are available from the author on request, and will be deposited with a persistent identifier.




Fig. 1: The field on one page, and the map of this survey: nodes sized by corpus papers, edges drawn only for routes the corpus contains, each column head naming the section that covers it.

paper read into this survey’s structured record, and it is the unit its counts are taken over. Eleven of the 96 method rows whose notes settle the question address interpenetration at all, seven of the eleven do it outside a closed-loop policy in a grasp synthesiser, a trajectory optimiser or a contact model, and we found none that reports a penetration number for its own trained policy’s rollouts. The claim is about learned closed-loop control and not about the field. Grasp synthesis and hand-object reconstruction have reported penetration depth and intersection volume as comparative columns for years, and four rows of this corpus do it: OakInk [4] scores a dataset split on penetration depth, solid intersection volume and simulation displacement, BiDexGrasp [5] prints penetration depth beside a prior method’s, BimanGrasp [6] fails any grasp whose total penetration exceeds 1.5 mm, and TopoRetarget [7] reports a maximum penetration and a share of frames past 2 mm against a baseline retargeter. Every one of those numbers scores a pose or a reference trajectory rather than the behaviour a trained policy produced, and it is the rollout that is missing. The obstacle is not the engines. NVIDIA’s own IsaacGymEnvs repository already computes a per-environment maximum interpenetration depth in Warp and gates the policy update on a 1 mm threshold. Section IV-B has the file and the lines.

The third finding concerns hardware coverage. Nineteen of the 33 hand rows in Table V and Table VI appear in no method row, and 8 of those can be bought today or built from published designs. Thirty-five method rows run on the Allegro, whose weight, joint torque, payload and price have no reachable source, because its product page returns HTTP 404 and everything Table V confirms about it comes from its ROS driver.

These findings build on a substantial literature in reproducibil-

ity and evaluation. Collberg and Proebsting [8] examined 601 papers in computer systems research for whether the code behind them could be obtained and built at all. Xu et al. [9] align 48 bioinformatics projects with their publications at sentence-to-function granularity under expert annotation, and SciCoQA [10] collects 92 real paper-code discrepancies, mined from issue trackers and reproducibility reports, into a benchmark for detecting such discrepancies automatically. In reinforcement learning the phenomenon itself is a known result. Engstrom et al. [11] show that code-level optimisations present only in the implementation account for most of PPO’s reported gain over TRPO, and Meta-World+ [12] finds undocumented changes accumulated across one benchmark’s own versions, which make comparisons between those versions unfair. Both establish it on a single codebase.

The closest relative to the audit here is Knox et al. [13], who review nineteen reinforcement-learning publications on autonomous driving, characterise the reward functions of ten of them exhaustively in a standard form, apply eight sanity checks and report near-universal flaws in reward design. Their ground truth for what each reward was is the authors, obtained through correspondence with them rather than by reading a released repository, and what they establish is that published reward descriptions are incomplete: one of the ten described its reward, discount factor, termination conditions and timestep thoroughly. Reading the code instead needs no correspondence and supports a different claim, which is that where code exists it sometimes contradicts the description. Raff [14] took the opposite ground truth on purpose, reimplementing 255 papers from their text alone and never opening the authors’ code, which is what makes the choice of arbiter a position rather than an accident.

So the claim here is narrow. We are aware of no prior work in robotics, and none in dexterous manipulation, that reads a field’s released reward implementations against the rewards its own papers describe. Those seven works are cited from outside this corpus and enter none of its counts. The exposure of the claim belongs beside the count above: the evidence is a repository at a fetched commit, every hash recorded in the code manifest, and a repository at a commit is evidence about that repository rather than about the run that produced a paper’s numbers, because the commit may postdate, precede or diverge from it. Hora [15] states the problem in its own README, which directs a reader to tag v0.0.1 and not to the commit parsed here to reproduce the published numbers. Every comparison this survey has revised, seven rows withdrawn from the count and two narrowed within it, is recorded in the row it concerns for the same reason: those revisions are the evidence that the nine left standing were checked rather than counted.

Fourteen corpus entries are themselves surveys or engine-comparison studies. Table I sets all fourteen on one set of columns. Four of the fourteen could not be obtained, or were fetched too late to read, and are entered as such. Two of the three things this survey adds are columns none of the fourteen fills. One is Table VII, which takes the engines [22] scored on documentation and usability and adds contact model, solver, iteration count, default timestep and penetration exposure, conditioned on what a hand does to a solver. Erez et al. [24] and Le Lidec et al. [23] do measure engines, on five engines and on four contact formulations, and neither surveys the field those engines are used in. The other is two hands on one object as its own problem, which none of the four field surveys gives more than a subsection.

A cell is a level and not a score: ✓ where the work gives the topic a structure of its own, a section, a table or a measurement it defines; ○ where the topic is there without one, a subsection with no analysis, a handful of mentions, or a metric named and never defined; × where it is absent or put outside the work’s scope. Each level coarsens one sentence read from `papers/notes/<key>.md`, and it is printed only while the quotation behind that sentence is still in the note, so a glyph cannot outlive its evidence. The supporting sentences remain in the reproducibility materials and `paper/tables/table10_surveys.md`. Three of the fourteen sources could not be obtained and one is on disk and has never been read, which is why their cells are empty; they are grouped as such rather than scored, because this survey knows nothing of them beyond a title and a DOI. The *gaps named* column is uniform by design: every predecessor names gaps, and what none of them attaches to a gap is a measurement.

The third addition is the penetration measurement, on an axis a predecessor had already named. Zhao et al. [16] state in their Sec. IV-C that the field assesses “at least two layers of performance: the quality of grasps or poses prior to execution, and the performance of policies or generators during downstream execution”, and names “physical plausibility, including penetration” as the most common criterion for the first layer. That is the reference-versus-rollout split, named before this survey measured it, and the idea that contact quality

belongs on the evaluation axis is Zhao’s. What Sec. IV-C does not give is a measurement method or a count, and this survey supplies those two. It does not supply a threshold. The 2 mm figure the field uses is taken from [7] with no independent justification, the captured human grasps in [29] sit above it at 3.25 mm, and Section VII-D states plainly that 2 mm is a simulator convention rather than a physical bound.

The scope is single-hand and bimanual multi-fingered manipulation. Parallel-jaw manipulation enters only as a comparison, which matters because 12 of the 112 method rows name a parallel-jaw gripper among their own embodiments. Locomotion is excluded. Prosthetics are excluded except where a hand crosses over into robot use, as the Psyonic Ability Hand does. The corpus holds 221 bibliography entries, of which 218 carry a structured row read from a note. Those 218 rows are 112 method papers, 33 hands, 15 simulators, 15 datasets, 14 benchmarks, 14 surveys, 8 tactile sensors and 7 evaluation protocols. “Method row” throughout means one of the 112, and every headline count in this survey has one of those eight classes as its denominator.

Two limits apply to every number here. This is a corpus of the learned era, which is a selection effect and not a judgement: 138 of the 221 entries are dated 2024 or later, 16 predate 2018, and two of the 112 method rows predate 2018, so any statement about a trend over time is a statement about 2022 onward. The analytic tradition is represented by one readable chapter rather than surveyed, and the planning line that took up the dynamic question directly, finger gaiting and rolling-contact manipulation and regrasp planning, is not in the corpus at all. Every coverage statistic is also a floor rather than a rate, because a value this survey failed to extract is indistinguishable from a value the paper never reported, and every such miss converts a reporting paper into a silent one.

Fig. 1 puts the field on one page. What follows works outward from the task: the six families and what makes each hard, then the hands and their makers, then the simulators and the contact models underneath them, then training, then two hands on one object as a problem of its own. Section VII is the longest, and it proposes an evaluation frame rather than a leaderboard. Each of those sections closes on the gap it owns, with the evidence and the experiment that would settle it; Section VIII is the list of those claims in one place, one sentence each, and says what to do about them, addressed to someone publishing, running experiments or buying a hand. Appendix A describes the method and the underlying corpus. Appendix B carries the three reference tables used in Sections III and IV. The complete discrepancy catalogue and the statistical derivations behind the proposed protocol are supplied with the reproducibility materials rather than repeated in the article.

II. A TAXONOMY OF THE PROBLEM

A. Task families

The six families below are read off the rows rather than imposed on them, and the last paragraph here says what they miss. Grasping is the largest. Fifty-seven of the 112 method rows carry the grasp label, the labels are not exclusive, and one

TABLE I: What 14 predecessor surveys cover: ✓ covered, ○ partly, × not. 24 of 84 coverage cells are empty.

survey	yr	dexterous scope	taxonomy	bimanual	hardware	evaluation	gaps named
Surveys of the field							
Zhao et al. [16]	2026	✓	✓	○	✓	○	✓
An et al. [2]	2025	✓	○	○	○	○	✓
Bai et al. [17]	2025	○	✓	○	○	○	✓
Welte et al. [18]	2025	○	✓	○	✓	×	✓
Firoozi et al. [19]	2023	×	○	×	×	○	✓
Zhao et al. [20]	2020	×	✓	×	○	×	✓
Simulator and contact-model studies							
NVIDIA Isaac Sim [21]	2026	×	○	○	×	×	✓
Kaup et al. [22]	2024	×	✓	×	×	○	✓
Le Lidec et al. [23]	2023	×	✓	×	○	✓	✓
Erez et al. [24]	2015	×	×	×	○	✓	✓
Sources not obtained, or on disk and never read							
Piazza et al. [25]	2019	—	—	—	—	—	—
Roa and Suárez [26]	2015	—	—	—	—	—	—
Ma and Dollar [27]	2011	—	—	—	—	—	—
Okamura et al. [28]	2000	—	—	—	—	—	—

paper can sit in several families. Success is a lift that survives a hold, and the thresholds differ by more than an order of magnitude. DexGraspVLA [30] requires the object “held 10 cm above the table for 20 s”, while Omnigrasp [31] requires it “held at least 0.5 s in simulation”. What makes grasping hard at scale is the continuum of starting configurations. UniDexGrasp++ [32] states it in Sec. 4.3, that “we are dealing with an infinite number of tasks considering the initial object pose can change continuously”.

Functional and tool use is second with 45 rows. It is the family where a stable grasp can still be the wrong answer. Dexterous Functional Grasping [33] gives the case in Sec. 2.1, that “grabbing a hammer from the head or handle are both equally valid ways of using it”, and only one of the two lets the tool be used. Success is defined against the tool’s function, and the field has no shared way to state it.

In-hand reorientation has 31 rows and the most settled criteria, because the community inherited one number. OpenAI [34] declared the goal achieved below 0.4 rad of orientation error, DeXtreme [35] keeps 0.4 rad at test time against a 0.1 rad training tolerance, and Eureka [36] counts consecutive successes at 0.1 rad. A shared tolerance is not a shared protocol, because the stopping rule differs. Visual Dexterity [37] measures error “when the controller predicts it has reached the goal and stops”, which makes the policy a judge of its own trial.

Tracking a human reference has 12 rows and inverts the problem. The target is a trajectory rather than an endpoint, so success is a per-frame error band. ManipTrans [38] requires all four of 30 degrees of rotation, 3 cm of translation, 8 cm of joint error and 6 cm of fingertip error. The difficulty is that the reference came from a human hand and was never dynamically feasible for the robot. DexMachina [39] warns in App. B.4 that scoring by timesteps inside the thresholds makes “the results highly sensitive to the threshold values”.

Bimanual coordination has 43 rows and handover has 10. Handover is the smallest family and the one whose failure has a single moment, because the giver must release only after the

receiver has the object. Dynamic Handover [40] reports a hit rate above its success rate and blames the gap in Sec. 5.4 on “occasional challenges encountered during the grasping phase of the catcher”. Not every paper here learns both sides. In [41] only the receiver is learned and the giver is a scripted arm.

Those six families do not cover the corpus. Twenty-four of the 112 method rows carry a label from outside them and eight carry no label from the six at all. Twenty fall in a catch-all class, mostly generalist policies evaluated on a task suite rather than on a dexterous task family, among them [42], [43], [44], [45] and [46]. Three adjacent families are named here rather than absorbed. Locomanipulation is [47], [48] and [49], where the base is not fixed and the gravity argument below changes character. Piano playing is [50], [51] and [52], discussed in Section VI, where success is a per-timestep F1 against a MIDI score rather than an object pose. Deformable manipulation is [53], whose object carries its own state and its own physics. Ferrari and Canny’s grasp-quality measure [54] carries no task label, because it is a measure rather than a task. Table II’s six rows are the families with enough papers to compare, and not a partition of the corpus.

B. Sources of difficulty

Contact is non-smooth, and that is a property of the problem rather than of any solver. Bicchi [1] states in Sec. 1.2 that contact constraints are unilateral, and that losing a contact “involves an abrupt change of the structure of the model under consideration”. Its Sec. 1.4 gives the sharper version, that a rod sliding on rough ground has configurations with no consistent solution and configurations with more than one. Simulators inherit the difficulty. The re-implementation study [23] concludes in Sec. V that “there is no fully satisfactory approach at the moment, as all existing solutions compromise either accuracy, robustness, or efficiency”. Pang et al. [55] treat non-smoothness as the thing to remove, curating every contact pair in Sec. III-D so that contact points and normals change smoothly.

The hand has more actuators than the object has degrees of freedom and is still short of authority over it. The object moves only through contacts, and the contacts are what cannot be commanded. Bicchi [1] states in Sec. 1.4 that in multifingered grasps “the number of independent contact forces is much larger than the number of actuators. Thus, from a controllability standpoint, not all the contact forces are controllable”.

Vision is occluded by the hand doing the work. Bai et al. [17] name it in Sec. 4.3, that “occlusion hampers object tracking”. DexPoint [56] keeps vision and adds imagined hand points to the point cloud. Yin et al. [57] remove vision entirely and rotate objects from 16 binary touch sensors over the palm, links and fingertips.

Gravity direction changes the task rather than scaling it. Visual Dexterity [37] states that reorientation with the hand below the object “is much easier”, because “with a downward-facing hand, the hand must manipulate the object while simultaneously counteracting gravity”. AnyRotate [58] makes the direction a randomised variable by “randomly initializing hand orientations between episodes”, and its Sec. 5.3 measures the cost, with performance dropping progressively from palm up and palm down, through base up and base down, to thumb up and thumb down.

The second hand, and three counts. Three counts describe two hands and they measure different things. Fifty-three of the 112 method rows record two hands on the robot, which is all the corpus’s two-hand flag claims. Forty-three carry the task label for bimanual coordination, the narrower claim that coordinating the hands is the task. Section VI narrows again, to the 28 rows whose notes place a learned closed-loop controller on two multi-fingered hands, and its opening paragraph names every exclusion that takes the 53 down to the 28. That 28 is the denominator for every architecture count in this survey. Every bimanual claim here names which of the three it uses.

What the second hand adds. Four things genuinely change when the second hand arrives. Contact stays non-smooth, occlusion stays, and gravity stays the same problem.

Role asymmetry is one. AsymDex [59] assigns a dominant hand with full finger and wrist control and a facilitating hand with 6-DoF base pose only, so that “the facilitating hand repositions and reorients one object, while the dominant hand performs complex manipulations”.

Two hands on one object also close a kinematic chain through the object. In [1] Sec. 1.4 the hand and object dynamics are separate and are “linked through the n rigid-body contact constraints”. A second hand adds a second constraint set on the same object rather than a second independent problem. That is where physical plausibility fails first. BimanGrasp [6] is the corpus row that measures it, rejecting a synthesised grasp when “total penetrations exceeded 1.5 mm” and reporting in Sec. IV-C that “penetration remains the primary cause of grasp failure”.

Two-arm collision does not exist for one hand at all. DyDexHandover [60] resets an episode when “any unintended arm contact with the environment or self-collisions” occurs. Bunny-VisionPro [61] pays for it in the controller, where adding a sphere-approximated self-collision cost raises motion-control time from 0.74 ms to 7.85 ms. Bidex [62] solves it in hardware,

mounting the hands so the human arm and the teacher arm “are perpendicular to each other and do not collide”. DexMimicGen [63] states plainly that it does not handle inter-arm collision at all.

And the action space doubles. Bi-DexHands [64] runs 20 tasks on two Shadow Hands and finds single-agent PPO beating multi-agent RL on most of them, offering in Sec. 5.2 that “PPO algorithm is able to use all observations for training the policy, while MARL can only use partial observations”. AsymDex [59] attacks the dimensionality from the other side, halving the observation and action dimension through its role split and a frame relative to the facilitating hand’s object.

Two cells say the field has no agreed criterion, and both are in families where a second body is involved. A criterion naming a distance, a duration and a trial count can be re-run by someone else. A rubric cannot, and Section VII takes up what follows from that.

III. HANDS AND WHO MAKES THEM

A. Design axes


No degree-of-freedom, force, weight or price figure in Table V or Table VI, both set in Appendix B, was measured by anyone outside the maker. Six rows are the exception, from peer-reviewed papers with stated protocols: ILDA, Pisa/IIT, ORCA, RUKA, LEAP and BiDexHand. Every other claim is dated at its source. Table VI prints the date beside each announced hand, and the repository tabulation carries a claim-date column for the hands that can be bought, running from Shadow’s specification of 4 December 2024 to vendor pages fetched on 17 September 2026.

The degree-of-freedom count is the first number a vendor states and the least comparable one. Shadow’s specification of 4 December 2024 reads “20 actuated DOF and a further 4 under-actuated movements for a total of 24 joints” [69]. Table V prints joints first and actuated DoF second, and the gap is the informative number. Where a vendor’s own DoF figure differs from the joint count, the joint count is what the table prints, as with Inspire’s “Degrees of freedom 6, Numbers of joints 12” [70]. An actuator count is neither of those two, and the repository tabulation keeps it in a column of its own. DexHand states no joint count at all, only 16 finger micro-servos and two wrist servos, so 18 servos is the only count quotable for it [71]. Daxo states 120 actuators, no DoF figure, and a tendon structure with no rigid joints, which has far fewer kinematic degrees of freedom than actuators [72].

The actuation ratio is the real decision, and Fig. 2 sets the three answers side by side. The Pisa/IIT SoftHand drives 19 joints from one motor and its authors state the cost plainly: “No in-hand dexterous manipulation is required for this prototype” [73]. Tendon drive moves the motors off the fingers and the mass follows them. Shadow’s hand plus forearm weighs 4.3 kg, against 1.1 kg for the ILDA hand alone with every motor in its palm [69], [74]. ILDA states an 18 kg payload against Shadow’s 4 kg in a power grasp, and Shadow states no fingertip force at all. What tendons cost is state estimation. The Faive Hand estimated joint angles from tendon length through a Kalman filter and its cube reorientation failed on hardware,

TABLE II: Task families against the properties that define success.

task family	what is held	what moves	what gravity does	what counts as failure	shortest honest success criterion
In-hand reorientation	one rigid object, held by the fingers for the whole episode	the object's orientation in the palm frame, while contacts break and remake	a disturbance whose direction is set by hand orientation, and a load to counteract with the palm down, [37]	the object leaves the hand, e.g. the fall reset at goal distance 0.24 m in [35]	orientation error below a stated tolerance under a stated stopping rule, 0.4 rad in [34] and [35], 0.1 rad in [36]
Grasping	nothing at the start, the object rests on a support	the hand closes, then the object is lifted clear	the test itself, since the grasp exists to resist it	no lift, or a drop inside the hold window	a stated lift height held for a stated duration, 10 cm for 20 s in [30], 10 cm to episode end in [65], 0.5 s in [31]
Functional and tool use	a tool, held in the pose its function requires rather than the most stable one	the tool's own joint, or the tool acting on a second object	secondary, the function dominates	a stable grasp in the wrong place, [33]	no agreed criterion. Criteria range from never dropping the object, [33], to a four-point rubric, [66]. The one mechanical threshold is holding the tool joint past a commanded angle for the episode, [67]
Tracking a human reference	whatever the human held in the captured clip	hand and object follow a recorded trajectory under physics	full, and the reference never obeyed it, so the reference is not dynamically feasible	per-frame error leaves the band, or the object drops	all stated per-frame errors below threshold for the whole clip, 30 deg and 3 cm and 8 cm joint and 6 cm fingertip in [38]. Thresholds are not comparable across papers, and DexMachina [39] calls the results "highly sensitive to the threshold values"
Bimanual coordination	one object by two hands, or one object per hand	both hands, the object, and the relative pose of the hands	can be carried by one hand while the other acts, [59]	the object drops, the arms collide, or the hands work against each other	no agreed criterion. Task-specific thresholds dominate, e.g. 0.035 m block-to-cup distance in [59]. The one criterion with a physical-plausibility gate is 2.0 s of hold with total penetration under 1.5 mm, [6]
Handover and in-hand transfer	by the giver at the start, by the receiver at the end	the object crosses between hands, released only after it is regained	unopposed during the exchange, and total in a thrown transfer, [40]	a drop at the moment of release, and the catcher's grasp phase, [40] Sec. 5.4	the receiver holds the object and moves it clear, more than 10 cm from the first hand in [68], or holds it to episode end in [60]

**Fig. 2:** Three answers to where the motor goes, and what each one costs. A coupling bar saves an actuator by tying two joints, so the fingertip reaches a curve through the reachable region rather than the region itself.

which its authors blamed on “poor joint angle measurement from the EKFs, especially when there is contact” [75]. Ruka-v2 adds detachable AS5600 encoders to measure that problem rather than close the loop on it, and those encoders put its open-loop linear joint-to-motor map at 8.26 degrees of error over seven joints [76].

Direct drive trades torque density for transparency. LEAP’s Dynamixel joints resist 19.5 N in a pull-out test against the Allegro Hand’s 8.5 N, at 595 g for the LEAP hand [77], and the Allegro’s own weight has no reachable source. Linkage drive is the argument against the forearm. ILDA puts 15 motors, drivers and ball screws inside the palm and reports 34 N at the

fingertip in the bent pose and 28 N stretched [74]. Linkages buy coupling for free, as in BiDexHand’s four-bar that drives each DIP off its PIP, whose paper claims 16 actuated DoF and 21 joints while its README describes 15 servos driving 15 joints [78].

Fingertip force is the number a hand buyer reads first, and across these rows it is six different measurements. Table V puts the quantity beside the figure: pull-out resistance for LEAP’s 19.5 N, pinch for RUKA’s 2.74 N and BrainCo’s 15 N, fingertip normal force under a 1 cm indenter for Unitree’s 10 N, a calibrated five-trial fingertip force for BiDexHand’s 2.14 N, a peak vendor claim for 1X’s 45 N, and two unlabelled spec columns for Sharpa’s “20 N | 12 N”. BrainCo’s whole-fist grip of 50 N is a separate figure and is no longer in the same column as a pinch. ORCA’s 19.6 N was none of these. It is the newton equivalent of a 2 kg index-finger payload at a fixed 600 mA motor current, so its force cell is empty and the figure is recorded as a payload instead [79]. Payload is no more comparable than force, which is why Table V does not print it. Nine of its twenty hands state one, in the repository tabulation’s payload column, and no two state it under the same test: Shadow 4 kg in a power grasp with no protocol given, Tesollo 2.5 kg pinching and 10 kg enveloping as rated figures, Unitree 3.5 kg palm-down and 4.5 kg palm-left on a 5 cm round object, ORCA 10.5 kg on four fingers and 2 kg on the index alone at a fixed motor current, BiDexHand 4.54 kg lifted, and RUKA 6.0 kg as the weight at which its joint-angle error passes 15 degrees.

Compliance is repairability seen twice over, and both are mechanical fuses: the Pisa/IIT SoftHand’s rolling-contact joints return to assembly after over-extension and ORCA’s “poppable” pin joints dislocate rather than break [73], [79]. ORCA also




Fig. 3: Which hands the literature runs on, banded by how the hand can be got, with every tabulated hand shown and those no paper uses as named slots at zero. A paper naming two hands counts under both.

reports tendon drive as maintenance: “Prolonged use requires manual re-tensioning to maintain performance”. What fails is rarely the motors. Its durability run found silicone skin degrading on two fingertips after about 2,000 to 4,000 grasp cycles and sensor wires snapping on three after about 4,500 to 7,000. The sensing wore out an order of magnitude sooner than the hand.

Two properties matter before any of the above, and neither is a column of Table V: the repository tabulation prints a control rate and a URDF-or-MJCF column for every row, and this one does not. The first is the rate a loop can be closed at, which spans 250 Hz on a Tesollo to 1 kHz on a Shadow, a Unitree or a Wuji, with the Allegro at 333 Hz on its own CAN clock and Inspire stating no rate at all. LEAP’s 500 Hz is a serial query ceiling and its own sim-to-real policy runs at 20 Hz, which answer different questions [77]. A reader deploying torque control is bitten by that first. The second is whether a URDF or MJCF exists and who ships it, which is the bridge to Section IV and to this survey’s recommendation to pick a hand a simulator already carries. Shadow, Sharpa, Wuji, the Ability Hand, RUKA, Faive and the Allegro driver ship one. LEAP is the warning: its paper says the URDF is released and the released API repository contains none [77].

B. Hands in use, and what they cost

The two oldest designs in Table V carry 49 of the 103 method rows that name a hand at all, and seven rows use both. Neither design’s date is confirmed by its own sources, so 2005 and 2016 are the bibliography’s. Fig. 3 counts, per hand, the method papers whose own experiments use it. Of 112 method rows, 103 name a hand. The Allegro accounts for 35, Shadow for 21, the Inspire RH56 family for 19, a parallel-jaw gripper for 12 and LEAP for 12. Seventy-five of the 103 name an Allegro, a Shadow or Adroit model, LEAP or an Inspire. What no row states is how big any of them is. Neither hand table carries a dimension column, because no source behind them gives a hand’s length or its width: the specification a reader could hold against their own hand is the one these sources omit, while the degree-of-freedom count, which all of them give, is compatible with a hand of any size at all. Mass is the nearest thing on record and it is not the same quantity.

The Allegro’s position is the uncomfortable part. Its product page at allegrohand.com/v4 returned HTTP 404 while the Wonik wiki timed out [80]. Everything Table V confirms about the most-used hand in the corpus comes from its ROS driver: 16 joints, four fingers, a torque interface, a 333 Hz CAN clock, 12 V, no tactile sensing. Its weight, joint torque, payload and price have no reachable source. The field’s most common platform cannot be specified from a page a reader can open.

Concentration would matter less if the hand did not move the result, and it does. On the same simulated cube rotation, LEAP reaches 0.2288 rad/s against the Allegro’s 0.0828 rad/s [77], and RUKA reports a 2.74 N pinch against the Allegro’s 1.60 N under the same three-trial pinch test [81]. A method compared only on Allegro hardware is compared at one point in a space where a single axis moves the headline number two or three times over. Inspire, XHand and Sharpa take 29 of the 103 rows between them, four rows use two of the three, and none is earlier than 2024. Eight take none at all: ORCA, RUKA, Ruka-v2, BiDexHand, DexHand, the Proception ProHand, the Tesollo DG-5F and the Unitree Dex5.

Open hardware and the collapse in cost. Six rows in Table V are open hardware, and LEAP Hand V2 in Table VI is a seventh. Five of the seven state a dollar cost: \$2,000 for LEAP, \$3,000 for LEAP Hand V2, \$1,500 for Ruka-v2, \$1,300 for RUKA and \$300 for DexHand. ORCA states a material cost below 2,000 CHF that the price column leaves unconverted, and BiDexHand states none [79], [78]. The Faive Hand states neither a cost nor a licence that could be read, so it is not counted as open hardware here [75]. Two of the stated bases need saying. DexHand’s \$300 is “additional total cost of components”, excluding the printing and the wrist servos [71], and ORCA’s own figure sits against Ruka-v2’s table listing ORCA at about \$3.5K [76].

The expensive end of the collapse is secondhand throughout. The only six-figure numbers on disk are RUKA’s comparison table at \$100,000 for a Shadow Hand and Faive’s “steep price tag of 110k GBP” [81], [75]. Shadow’s own page says to discuss pricing and Table V’s price cell for it is empty. The collapse is real at the cheap end and secondhand at the expensive one, and it has barely moved the literature. Twelve

of the 103 hand-naming method rows use an open-hardware hand, and all twelve are LEAP.

What the cheap hands give up is sensing. LEAP has none and names touch sensors as future work [77], RUKA states its design “lacks tactile sensing” [81], and BiDexHand’s sources never mention touch [78]. ORCA is the exception, with binary FSR fingertips whose threshold was measured as low as 0.05 N on a fresh fingertip, against the sensor’s rated 0.29 N, and 6.38 N on a degraded one [79].

C. Tactile sensing

DIGIT set the cost floor. It is 20 by 27 by 18 mm, weighs about 20 g, streams 640x480 at 60 fps, and its paper states a “total estimated manufacturing cost is approximately 15 USD per sensor ... when manufactured in a batch of 1000” [82]. Durability was as much the contribution as price: its gel degraded 0.3 percent over 15 abrasion passes, against 805 and 918 percent for the two gels compared with it.

Digit 360 is the same lineage at a different operating point, claiming about 8.3 million taxels, spatial features to 7 μ m, normal and shear force resolution of 1.01 mN and 1.27 mN, and on-device processing that cuts event-to-action latency from 6 ms to 1.2 ms [83]. It is also not for sale, and no corpus paper’s parsed text names it.

XELA’s uSkin is the non-camera alternative and the one that bolts onto hands the field already uses. Its curved fingertip kit for the Allegro V4 and LEAP carries 30 three-axis sensing points, a full Allegro integration reaches 368, resolution is 0.1 gram-force, and the kit runs at 275 Hz [84]. The 500 Hz figure often quoted belongs to the product family, not the curved fingertip. One corpus paper names it, [85], whose self-supervised encoders are pre-trained on 462.7k tactile images and beat matched end-to-end models by 95.1 percent when both see 33 to 50 percent of the labels. Its own bead-maze policies never complete the maze on the real robot.

The usage number is the one to keep, and it needs its inclusion rule stated: a sensor physically on the hand, read by the deployed policy. Eight of 112 method rows meet it: [58], [67], [86], [87], [68], [88], [89] and [57]. The rule excludes [90], whose 20 binary contacts are simulated on an Allegro without tactile hardware and whose released configuration disables tactile input; it also excludes [91] and [92] for the same reason. Sixty-three of the 112 method papers use the word tactile somewhere in their parsed text and 35 carry it in their structured note. Against the eight that meet the rule above, either number is the gap worth quoting.

Almost none of the eight uses a high-resolution sensor. Yin et al. [57] remove vision entirely and rotate objects from 16 binary touch sensors over the palm, links and fingertips, deployed zero-shot to a real Allegro, and Robot Synesthesia uses 16 force-sensing resistors read as binary [88]. Two papers exclude touch deliberately, HORA reporting rotation “even without the usage of vision and tactile sensing” [15], and ORCA’s authors dropping it from their reinforcement learning “due to the additional complexity involved in accurately modeling them” [79]. Taxel counts have risen by three orders of magnitude while the policies consuming them have stayed at binary contact.

D. The catalogue against the literature

Three classes of evidence. Behind Table VI’s rows sit three kinds of evidence, and conflating them is how a DoF figure with no source ends up in a survey. Vendor prose carrying numbers is the strongest, as on 1X’s page of 9 July 2026 [93]. Video is second and carries none. Press or bibliography assertion is third, as with Tesla’s V3 hand, known here only through a paraphrase of patents because the USPTO PDF parsed empty [94]. Table VI blanks the cells resting on the third class and footnotes who did the arithmetic.

A survey can go one step past recording that a claim is unverified, which is to say which claims are implausible on their face. Daxo’s 120 actuators in 750 g is about 6 g per actuator including structure, tendons, routing and skin [72]. Clone’s 27 degrees of freedom under 2 pounds excludes a 500 W pump the source does not confirm is excluded [95]. Figure 03’s “Degrees of freedom, hands | 20” sits on the same tracker page as “Number of fingers | 10”, so it is almost certainly the pair [96]. Tesla’s repeated 22 is one article’s arithmetic of four DoF on each of five fingers plus two at the wrist, and Gen 2’s own figure was 11 [94].

Scepticism belongs to the evidence class, not to which table a row lands in. Sharpa’s 22 of 22, Wuji’s 20 of 20 and Tesollo’s 20 of 20 are vendor claims about unmeasured hardware and they sit in Table V, where Sharpa’s page footnotes “Specifications may vary between products” and Wuji’s says the spec “will continue to iterate” on a Beta1 product [97], [98]. Even the best-specified vendor page in the corpus leaves its fingertip-force columns unlabelled.

What is sold against what is published on. The band at zero in Fig. 3 carries the finding. None of the nine company-announced hands in Table VI appears in a single method row whose own experiments use it. Tesla, Figure, 1X, Sanctuary, Boston Dynamics, Xiaomi, Clone, Daxo and PaXini account for zero of the 112 method papers’ experiments. The nearest thing to a counterexample is [99], a Figure blog post claiming a 35-DoF whole-upper-body action space at 200 Hz that includes individual finger control. It never names the hand, gives no per-hand DoF count, and reports no success rate or trial count for any task. It is the maker describing its own unreleased hand, which is the evidence class the finding is about. The four research prototypes in Table VI are in a different position, since the Faive Hand and LEAP Hand v2 Advanced account for three method rows between them [65], [62].

Table VI’s emptiness is measurable and part of the same finding. Over the eight specification columns it prints, from DoF to price, 58 percent of its cells are values no source stated, against 42 percent for the hands that can be bought; both captions state their own count. Over the twelve specification columns the repository tabulation carries, the same split is 61 to 39. The hands with the highest advertised DoF counts have the least specification behind them.

The one hand that crosses the gap crosses it because its maker published rather than announced. ByteDexter reaches a method row through a ByteDance technical report stating 21 DoF per hand, 16 piezoresistive fingertip channels in the action vector and real-robot success rates [100], not through a

product page. That mechanism is available to every vendor in Table VI and none has used it.

The gap runs the other way too, and the last column of Table V shows it. Eight hands that can be bought or built from published designs take zero method rows each: the Unitree Dex5 with 94 pressure sensors on its P variant [101], the fully actuated Tesollo DG-5F [102], the 22-DoF Proception ProHand [103], ORCA, RUKA, Ruka-v2, BiDexHand and DexHand. Cheapness is established only for RUKA, Ruka-v2 and DexHand, because Table V’s price cell is empty for Unitree, Tesollo, the ProHand, ORCA and BiDexHand, and the only Dex5 price on disk is Ruka-v2’s secondhand “~\$25K” [76]. A reader choosing a hand on this corpus’s evidence has four options with published precedent behind them: the Allegro at 35 method rows, a Shadow at 21 and 4 more for its Adroit model, the Inspire RH56 family at 19, and LEAP at 12, of which only LEAP has both a published design and a stated price. Below those four the record thins fast: eight rows for the XHand, seven for the Ability Hand, six for Sharpa, two each for Wuji and LEAP Hand v2, one each for five more, and none at all for the other 19 hands of Table V and Table VI. Those 19 sit in the tables on their specifications alone, and no method row in this corpus runs on any of them.

Model availability. Nineteen of the 33 hands in Table V and Table VI appear in no method row, and 11 of those are neither sold nor open and appear in none for that reason, so 33 is not the denominator for a software-lag claim: the eight hands of the paragraph above are. The rule that decides used from unused is one regular expression per hand against the method rows’ own hand field, in `tools/hand_usage.py`, so the partition can be recomputed rather than argued about. The lag runs through the tooling as well as the papers. Seven of the fifteen simulator rows ship a first-party hand model at all, six of those hands are real, Brax’s being a synthetic claw, and every one of the six is an Allegro or a Shadow; only four of the 33 tabulated hands appear in any engine. Bench2Dex [104] meanwhile compares 12 hands and DexVerse [105] six without stating a DoF count for any. A useful next step is a conformance suite for hand models: one URDF or MJCF per hand, with fixed joint-limit, mass and collision checks and a published pass or fail per engine.

IV. SIMULATORS AND THE PHYSICS UNDERNEATH

A. Requirements of a dexterous simulation

The clearest demonstration that hands are the hard case came from a benchmark that was not about hands. Erez et al. built a 35-DOF arm modelled on the Shadow Hand, closed it around a capsule with fixed spring-dampers, and asked four of the five engines for the largest timestep at which the object was still in the hand. MuJoCo held the grasp at 16 ms, PhysX at 2 ms, ODE at 0.25 ms and Bullet at 0.03 ms, a spread of a factor of 500 [24, Sec. IV-D]. Havok is the fifth and was excluded for want of a working PD controller. Three caveats travel with that spread. The engines run a deliberately restricted common model, hinge joints with sphere and capsule geometry, no boxes and no meshes (Sec. II). The authors call it “an open question whether this model system can be tuned to work better in the




Fig. 4: The geometry the solver uses against the geometry the renderer draws. A single convex hull fills the concavity with material that exists only in the collision model, and holds the fingertip out of a contact the visual mesh allows.

gaming engines” (Sec. IV-D). The timesteps are log-spaced and good only to a factor of two, and the authors wrote MuJoCo and disclose it.

A grasp is hard for nameable reasons. It is many contacts at once, all persistent, all near stiction, on an object much lighter than the mechanism holding it. Persistence and multiplicity make the problem hyperstatic, and Le Lidec et al. show that per-contact solvers of the projected Gauss-Seidel family, and RaiSim’s, then inject spurious jamming forces at stiction that vanish only once the object slides. The mass ratio makes it ill-conditioned, and in their stacked-cube test at a 10^3 to 10^{-3} kg ratio those same methods fail to converge. Global methods with proximal regularisation stay robust in both cases [23, Sec. IV-A]. What degrades a truncated solve is therefore conditioning and redundancy, not stiffness, and a grasp supplies both by construction. Locomotion is forgiving by comparison. On flat ground the contact model and solver choice “hardly affects” their quadruped’s tracked base velocity, and only on rough, slippery terrain do RaiSim and CCP deviate. A walking robot makes and breaks a few contacts against ground far heavier than itself. A hand does neither.

Speed enters by the same door. MuJoCo Playground reports that contact time scales with the number of possible contacts rather than the active ones, because JAX requires static shapes, which is why its tasks override the bound on contact points and the bound on geometry pairs by hand [106, Sec. VI]. Every override printed in that paper is a locomotion port. A hand’s bound would have to be set the same way, for the worst case, but no hand configuration is printed.

Fig. 5 sets out the stages of one simulation step. Three of them make overlap, and they do not answer to the same knob. In an engine that enforces non-penetration at the velocity level, integration turns any residual approach velocity into overlap of order v times Δt ; an engine that enforces the gap at the next configuration carries no such term, which is why Dojo’s hard-contact NCP keeps its feet above the floor at every timestep it was tested at. Constraint assembly fixes the compliance a loaded contact then rests at. A truncated solver leaves a residual that grows with conditioning. Which of the three dominates in a grasp is not measured anywhere in this corpus, and the three are not ordered here. A fourth thing is not on the figure, because it is not a stage of the step and not a source of overlap: a mismatch between the geometry the solver uses and the

geometry the renderer draws, which runs in both directions and is drawn in Fig. 4.

B. Contact models and solvers

Table VII, set in Appendix B, is the engine-by-engine comparison, one row per simulator, built only from what a note confirmed. Le Lidec et al. supply the taxonomy that organises it, checking each formulation against the Signorini condition, Coulomb’s law, and the maximum dissipation principle. Linear complementarity, the family of Bullet, ODE, PhysX and DART, satisfies Signorini alone, because linearising the friction cone to a pyramid biases friction toward its corners. The cone complementarity problem satisfies the other two but relaxes Signorini, so contact acts at a distance of size $\Delta t \cdot \mu \cdot \|c_T\|$. The full nonlinear problem satisfies all three and is non-convex, by Table II of [23]. A contact model is not the algorithm that solves it. PhysX linearises the cone, which places its model in the LCP family, and it does not solve an LCP. Its Temporal Gauss-Seidel scheme folds substepping into the Gauss-Seidel sweep and exposes position and velocity iteration counts separately [107, Sec. 3].

MuJoCo is inside that taxonomy rather than outside it. Le Lidec et al. classify it as CCP-MuJoCo, a convex relaxation solved by a Newton method on the primal QCQP, in the same family as CCP-Drake, the SAP-style scheme, by Table III of [23]. The relaxation makes two errors of opposite sign in one engine. A loaded contact carries a violation, and a sliding contact carries force at a positive gap. Drake’s SAP inherits the same pair, and its gliding effect at distance $\varphi \approx \delta t \cdot \mu \cdot \|v_t\|$ “unfortunately does not go away as $\delta t \rightarrow 0$ ” [108, Sec. V-A]. Le Lidec et al. call that compliance a “numerical trick designed to circumvent the issues due to hyper-staticity or ill-conditioning at the cost of impairing the simulation”.

The depth a loaded contact carries is a chosen number. In a regularised formulation the steady-state violation at a contact is the normal load times the compliance. It is zero at zero load and it grows with the load carried. MuJoCo drives that violation coordinate with a critically damped stabiliser parameterised by ε and κ , and for an object resting under gravity the steady-state depth has a closed form independent of the object’s mass [109, Sec. V]. The closed form did not survive the parse of that paper, so the cancellation is quoted and the algebra is not. Neither MuJoCo paper states why mass cancels, and the explanation this survey offers is its own inference rather than a cited one: the regulariser is scaled by the inverse effective inertia at the contact, so compliance falls as $1/m$ exactly as the gravity load rises as m . The nearest support in a parsed source is for a different solver of the same engine, MuJoCo’s diagonal solver, a “mass-aware spring-damper” that uses the diagonal of the A matrix to keep contacts critically damped [110, Sec. II-E]. This is not a penalty spring, and depth is not always non-zero. The impulse solves a regularised convex program over the whole contact set rather than a per-contact function of the gap, both MuJoCo papers reject spring-dampers by name, and the 2012 ball-drop figure is captioned “there is no penetration” [110, Fig. 2].

The first of two concrete measurements comes from the other end. Dojo solves a hard-contact nonlinear complementarity

problem with an exact second-order friction cone, by an interior-point method converging within 15 iterations on the three robots of its convergence study. Table II of [111] drops an Atlas humanoid and reports foot-floor penetration against the timestep. MuJoCo penetrates -28 mm at $\Delta t = 0.01$ s and -46 mm at $\Delta t = 0.001$ s, while Dojo stays above the floor at every step tested [111, Sec. V-A]. The MuJoCo column is not a trend. A ten-times-smaller step produces more overlap, which no timestep-independent stabiliser does. Either that configuration ties the compliance to Δt , or the quantity is an impact transient on a drop. Neither reading is a steady-state grasp depth.

Drake’s SAP quotes 2.5×10^{-5} m at $\delta t = 10^{-2}$ s and 2.5×10^{-7} m at $\delta t = 10^{-3}$ s, three and five orders of magnitude below Dojo’s two MuJoCo cells at the same steps, on an engine that is also compliant [108, Sec. V-B]. Those two figures are analytical bounds for a single point mass at rest on a plane under that paper’s near-rigid stiffness rule, not measured depths, and they are not a Drake grasp-penetration number: a point-mass bound and a humanoid drop transient differ in load, effective inertia and regime, so the distance between them is not a measurement of anything. What the pair of engines does show is that in a compliant formulation the depth follows from a stiffness that someone chose.

Table V of [111] times 1000 steps of forward simulation with gradients, at a matched $\Delta t = 0.01$ s, for engines that are not all computing gradients. MuJoCo is fastest on every system, 0.335 ± 0.001 s against Dojo’s 1.159 ± 0.077 s on a Franka Panda, and the authors call the comparison difficult because Dojo is stable at five times the step size (Sec. VI-B). What the regularisation buys, on Le Lidec’s reading quoted above, is conditioning on a hyperstatic problem, and the depth it costs is tuned separately. Whether that is also what holds Erez’s grasp at 16 ms is a question his data do not answer. His planar chain is contact-free, so its numbers speak to the coordinate formulation and not to contact: MuJoCo runs at 243.2 kHz there against Bullet’s 22.8 and PhysX’s 6.4, and Bullet’s articulated Featherstone mode at 81.4 kHz beats every Cartesian-coordinate engine [24, Sec. IV-B]. Those are throughput in evaluations per second, not a largest-stable-timestep result, and the articulated Bullet mode was never run on the grasp test at all, being usable only in tests without contact, which that paper’s own appendix states of its Featherstone implementation. Joint coordinates explain the contact-free speed advantage. The grasp timestep is a contact result, and the paper attributes it to nothing: it reports that the other engines go unstable and “effectively simulate a different physics model which can no longer hold the object” (Sec. IV-D), without an experiment that separates the coordinate formulation from the soft absorption of penetration.

The GPU era moved the compliance knob rather than removing it. ComFree-Sim resolves contact in closed form in the dual cone of the friction cone, so penetration becomes an explicit tuning parameter and the paper reports it. On a drop test of convex primitives about 5 cm across at $dt = 0.002$ s, MuJoCo Warp penetrates 1.7 ± 4.9 mm and ComFree-Sim runs from 3.9 ± 6.9 mm at $k_{\text{user}} = 0.1$ down to 0.9 ± 1.5 mm at $k_{\text{user}} = 0.5$, though its prose describes the opposite direction from its own table [112, Sec. IV-A]. Its impedance acts on the signed gap through the “identical” default hyperparameter API




Fig. 5: One simulation step, and where engines differ: the stages every rigid-body engine runs, with the choices that separate the engines used for hands under each; dashed arrows mark the three stages that overlap.

as MuJoCo’s solver, so the knob it is credited with exposing is one MuJoCo users already set. Those are 5 cm primitives under their own weight, not a fingertip loaded by a grasp, and depth scales with the normal load. No paper in this corpus publishes the fingertip number. The baseline it measures against is MuJoCo Warp, which Newton builds on and which MuJoCo Playground names as its intended replacement for JAX [106, Sec. VI].

That brings Table VII’s most important column, and the claim it does not support. The column records whether a parsed source reported a penetration depth, not what an engine can compute. Two of the 15 engines report one, Dojo and ComFree-Sim, and both are engines whose paper is about contact accuracy. Four are recorded as not reporting it: Brax, Isaac Gym, Isaac Lab and Orbit. Nine rows are blank. The word “penetration” appears nowhere in the Isaac Gym paper, and Isaac Lab’s contact sensor reports force, duration and an average contact point with no contact-quality metric. The Isaac family carries 42 of the 112 method papers in this corpus.

The tooling exists and the number is still not recorded. NVIDIA’s own IsaacGymEnvs repository computes interpenetration depth in simulation. Its IndustReal tasks load plug and socket meshes into Warp, sample points on one, query them against the other, and reduce to a per-environment maximum interpenetration distance (code/md/isaacgym_2021.md, lines 8809 to 8862). The policy update is gated on that number. Environments are split on whether their maximum stays under a threshold, and the reward of those that survive falls as the maximum approaches it. The configured limit is 1 mm (lines 3630 and 3753). That is a shipped Isaac Gym task measuring simulated interpenetration per environment at a millimetre threshold, during RL, and acting on it. Table VII records Isaac Gym as not exposing penetration, and so does its row in the corpus, which the code parse in that same corpus contradicts. Tactile Genesis makes the point from the other side, shipping penetration depth as a sensor on an analytic SDF backend and a BVH backend [113, App. A.1], while Table VII leaves the Genesis cell blank. The depth is computable from the poses and the meshes in a few lines of Warp, and the field’s own benchmark repository already does it. Interpenetration in a dexterous rollout is a setting these engine papers do not document and a measurement no method row in this corpus

reports.

A policy will exploit what is not measured. DexTrack’s configurations cap PhysX’s separation velocity at 10.0 or 1000.0 depending on the task variant, with no explanation in the paper or in a configuration comment [114]. The same parameter appears across unrelated stock IsaacGymEnvs tasks at 5.0, 10.0, 100.0 and 1000.0, at five places in the same parsed file, which is the signature of a stock template rather than of a per-variant choice. The parameter caps the rate at which the solver pushes overlapping bodies apart, so it sets how long an overlap persists and how violently it is undone, not how deep the overlap gets. So the one knob here that bears on interpenetration behaviour carries stock values and takes two of them across variants of one method. That pattern is what the corpus shows; why the values were chosen is not in the record, and matching numbers cannot establish it. DexTrack’s paper defines a maximum hand-object penetration depth, applies it only to its input kinematic references, and presents tolerance of “severe hand-object penetrations” as evidence of robustness (App. B). TopoRetarget [7] is the one corpus method that reports the number carefully, and it reports it on retargeted references rather than on a rollout, which Section VII takes up.

The rest of Table VII is largely empty, and the emptiness is a result. Forty of the 105 cells in its seven property columns, 38 percent, are values no parsed source stated, which is what its caption counts. The repository tabulation is emptier still, because it carries four columns this one does not: 69 of the 150 cells outside its engine column, 46 percent. No engine paper states a default physics timestep. Three report one for a named experiment, and the timestep column reports those experiment settings. Isaac Gym’s cell is its Shadow Hand step, from the only per-task timestep table any engine paper here publishes, which runs 1/120 s for Shadow Hand and Allegro, 1/200 s for ANYmal and TriFinger and 1/60 s for Franka [107, Table 4]. MuJoCo’s 0.01 s is the 27-DoF humanoid test’s step and ComFree-Sim’s 0.002 s is its benchmark step, against a stated stability limit near 0.02 s. Seven of the fifteen engines ship any dexterous hand at all. Two of the columns only the repository tabulation carries are worth the trip: four engines state a solver iteration count there, and two state a licence, which is the question a reader choosing an engine actually has.

C. GPU-parallel engines and their throughput

The GPU-parallel turn. Isaac Gym set the pattern. Physics, observations, rewards and actions stay on the GPU, and PhysX resolves contacts with the Temporal Gauss-Seidel sweep described above. Its per-task timesteps are published, which is rare: the Shadow Hand runs a 1/120 s physics step under a 1/60 s control step, or 1/20 s in the OpenAI variant. The result that reorganised the field is that reproducing OpenAI’s Shadow Hand cube reorientation took under an hour on one A100, against 30 hours on 6144 CPU cores and 8 V100s [107, Sec. 6.4.1]. Thirty-five of the 112 method papers in this corpus run on it.

Orbit and Isaac Lab moved the stack to PhysX 5, and the dexterous offering is thinner than the predecessor’s: the first-party suite is lifting, grasping and reorienting with the KUKA Allegro hand, while Shadow and Ability hands appear only through third-party grasp evaluation [115, Sec. 7.2.3]. MJX and MuJoCo Playground took the other route, putting MuJoCo’s own solver on the GPU through JAX at the cost described above. MuJoCo Warp ports the same solref and solimp soft-constraint rows to NVIDIA Warp, and its README states that the legacy PGS solver and the noslip pass are unsupported and that differentiability “is not yet available” [116]. GPU MuJoCo is not a differentiable MuJoCo. Genesis is the broadest, pairing a rigid constraint solver with FEM, MPM, PBD and SPH solvers and a differentiable backward path, and it ships a Shadow Hand. Its README quotes no throughput number at all, so the large FPS figures that circulate for Genesis trace to nothing in [117].

Newton changes how these names should be read. It is explicitly multi-solver: SolverMuJoCo wraps MuJoCo Warp, SolverKamino is a proximal-ADMM and dual-variational-inequality contact solver, and XPBD and VBD do position-based constraint projection [118]. Its contact model is solver-dependent, not a property of the engine, which is how Table VII records it. Isaac Lab already offers a Newton–MuJoCo Warp backend in its hands demo, and its roadmap announces Newton integration. An experiment reported as Isaac Lab may be running PhysX 5 with TGS, or MuJoCo’s soft constraint rows under Warp, and those two make different contact errors. The name also fails to fix the physics inside one engine. MuJoCo’s convex solver is a family, interior point or projected Newton, conjugate gradient or Gauss-Seidel [110, Sec. II-D], and MuJoCo Warp supports neither PGS nor the noslip pass. PhysX 4 and PhysX 5 differ in whether non-convex rigid bodies get SDF collision, which is a contact-geometry difference between Isaac Gym and Isaac Lab under one vendor name [119], [115]. Naming a simulator no longer names its physics. Papers should report the backend and the solver beside the framework.


Reported throughput, and why the numbers do not compare. Four headline figures measure four different quantities. Isaac Gym reports parallel environment steps per second with physics, observations and rewards on device: 150,000 for the Shadow Hand at 16,384 environments on one A100. Isaac Lab reports frames per second in training, which includes the learning update: over 900,000 for the DextrAH teacher

task at 16,384 environments on eight RTX Pro 6000 GPUs, with no single-GPU dexterous number anywhere in the paper. ManiSkill3 reports “up to 30,000+ FPS” with RGBD and segmentation on one RTX 4090, measured over 1000 random actions with reward and termination removed from the timing, and does not state the environment count behind the figure [120]. Orbit reports a 125,000 FPS physics-only ceiling on an RTX 3090, also with no environment count [119, Sec. VII].

The one number reported cleanly enough to reuse is MuJoCo Playground’s, in PPO steps per second on a single A100 over five seeds. LeapCubeReorient runs at $76,354 \pm 143$ and PandaRobotiqPushCube at $487,341 \pm 4,346$. Same hardware, same measurement, same codebase, a factor of 6.4 between two environments. It is not a factor from the task in isolation. Playground tunes solver iterations, line-search iterations, timestep and contact bounds per environment, with values as far apart as one and four solver iterations, by Table III of [106], and it does not print the two configurations side by side. That omission is what makes the two figures incomparable. The same confound sits inside Isaac Gym’s paper on one A100: 700,000 environment steps per second for Ant, 200,000 for Humanoid, 150,000 for the Shadow Hand. A training-loop FPS is also mostly not a measurement of physics. Playground breaks the fractional cost down on an RTX 4090: for CartpoleBalance, physics is 0.02, rendering 0.06, inference 0.01 and the policy update 0.91, and the policy update still dominates on the Franka task.

Cross-framework comparisons add further free parameters. ManiSkill2’s PickCube table takes the best result over 16 to 512 environments for each system [121, Table 1a], giving ManiSkill2 with a render server 2487 ± 24 FPS at its optimum of 64 environments against Isaac Gym’s 865 ± 35 at its optimum of 512. The env-count tuning is the second problem here. The first is that the two systems are not the same kind of thing: ManiSkill2 runs rigid-body physics on CPU worker processes behind a shared GPU render server, Isaac Gym runs physics on the GPU, and its own paper says so plainly. The measured quantity includes 128×128 rendering at 500 Hz simulation and 20 Hz control for both, so this is a visual sample-collection loop rather than two physics engines. Its authors add the fidelity caveat themselves, and Playground is equally explicit that its cross-simulator plot borrows its Isaac Lab and ManiSkill3 numbers from the ManiSkill3 paper. Several sources give no number at all: MuJoCo Warp’s README points to an external nightly dashboard, and Newton’s and Genesis’s READMEs contain no FPS or speedup anywhere.

The failure reaches past the engine papers into the methods. Of the 47 corpus method papers that name a GPU-batched simulator, 19 state no environment count anywhere, and Isaac Lab’s own paper states no physics timestep, no decimation and no solver iteration count for any task. A throughput figure is interpretable only with the environment count, the GPU, the timesteps, the solver iteration budget, and a statement of whether rendering and the learning update sit inside the measurement. Almost nobody reports all five.



d is the depth by which the two solids overlap: the quantity every other engine in Sec. IV works to remove

the same d , thresholded, is the contact flag:
on above 0.5 mm, off below 0.2 mm



The inversion. Everywhere else in this survey the overlap is the error. For a compliant tactile sensor it is the signal, so the quantity a policy elsewhere is penalised for producing is the one this sensor is calibrated to read.

Fig. 6: A compliant tactile sensor reads the overlap depth. The tactile image, the distributed force and the binary contact flag are all functions of the one quantity every other engine in this section works to drive to zero.

D. Tactile simulation and the sim-to-real gap

The three tactile simulators in this corpus calibrate against three different things, and none of them is a manipulation outcome. TACTO is a rendering layer over a host engine, by default PyBullet’s rigid contact model. It reads post-solve link poses and the engine’s reported normal force and maps that force to gel-mesh deformation at the rendering level, so it contributes no contact physics of its own. Its only sim-to-real number is a tactile pose-estimation task, at 1.66 ± 0.16 mm with colour-jitter augmentation against 0.76 ± 0.07 mm for a model trained on 128 real datapoints, by Table II of [122].

Taxim is example-based rather than simulated, with an optical model calibrated from 50 real indentations, and it beats TACTO on every optical-similarity metric against real images. Its marker-motion model is calibrated against an ANSYS FEM, and the errors are 1.00×10^{-2} mm real against FEM, 1.02×10^{-2} mm real against Taxim, and 3.96×10^{-3} mm FEM against Taxim. Taxim tracks its own reference more tightly than that reference tracks reality, so its residual is bounded below by the calibration target. The authors also state that they simulate quasi-static contact only, with slip left as future work [123, Sec. V]. Slip is what a hand needs.

Tactile Genesis inverts the survey’s thread. It implements two penetration-depth backends on the collision geometry, an analytic SDF query and a BVH raycast, then binarises the result with Schmitt hysteresis at 5×10^{-4} m on and 2×10^{-4} m off. Here the penetration depth is the sensor signal, so the quantity every other engine treats as an error sets this sensor’s operating point, which is Fig. 6. Against a real GelSight it reports relative marker RMSE of 0.329 in dilation and 0.174 in shear, against HydroShear’s 0.403 and 0.217, with each simulator tuned to match the real image first [113, Fig. 2]. Its own sim-to-real check is a matched success count rather than a fidelity measurement. The real XHand1 SDK reports a per-taxel raw pressure field as well as an aggregate contact pressure, but documents no taxel positions or response characteristics, so the raw field cannot be registered to the simulated probe

layout and only the aggregate is comparable (App. C). The undocumented calibration is the barrier, not a missing signal.

The sim-to-real gap for hands. Most of what is written about this gap is attribution without measurement. PenSpin asserts that the pure physics gap “cannot be bridged by extensive domain randomization alone” while reporting no experiment that isolates it [90]. MuJoCo Playground attributes its LEAP hand failures to physical flex in low-cost hardware and names more accurate collision geometries as the fix, without measuring either. DeXtreme lists four candidate causes for its shortfall, including a malfunctioning Allegro thumb used in most trials, and disambiguates none by experiment [35].

What has been measured is narrower, and Section VII collects the transfer numbers. The one result that links a quantified model error to transfer is a humanoid recipe paper, which ranks its system identification runs by dynamics-model MSE and pairs each with real success over 10 trials: the lowest-MSE model grasps 8 out of 10, the median-MSE model 3 out of 10, the highest-MSE model 0 out of 10 [124, Table 1]. Where a cause has been pinned down elsewhere it is usually perception or actuation, not contact. OpenAI’s pose estimator has 3.12 mm error on rendered images and 9.27 ± 4.02 mm on 992 real ones, and PDDM reports a camera tracker with 5 mm average error and 20 ms latency as the unmodelled source in its real numbers [125, App. C].

The contact side stays unmeasured, and the one paper that looks at it is usually read backwards. DeXtreme’s real-to-sim replay interpenetrated because the replayed poses carried the pose estimator’s error, not because the contact model failed. The paper says so: “there is still some sim-to-real gap in pose estimation. This is manifested when we played back the real states in sim (real-to-sim) with physics enabled, which sometimes resulted in interpenetrations. Therefore, we were not able to easily calibrate physics parameters of the cube” [35, Sec. 5]. A replayed trajectory is a placement, so the overlap it shows bounds the state estimate rather than the physics. That is why it could not calibrate the cube, and it is why calibrating contact against hardware still has no worked example for a hand. The only direct measurement of simulator fidelity against hardware in this corpus is Dojo’s, an average final-position gap of about 0.5 cm over 5 box-pushing trials. It is a parallel-jaw arm pushing a box, and there is no equivalent number for a hand.

V. TRAINING A DEXTEROUS POLICY

A. Sources of supervision

A dexterous policy is defined by what supervises it, and four sources are in use across the corpus’s 112 method rows: a reward function supervises reinforcement learning, a human demonstration supervises imitation, a human reference trajectory supervises a physics-based tracker, which is imitation with a simulator in the loop, and nothing supervises a model-based planner, which is handed a cost and a model instead.

The labels do not partition the corpus. Of the 112 method rows, 53 are tagged reinforcement learning, 27 behaviour cloning, 23 distillation, 18 generalist or vision-language-action, 14 teleoperation systems, 14 diffusion, 10 data collection, 8

flow matching, 6 reinforcement learning from demonstrations, 6 trajectory optimisation, 5 model-predictive control, 4 grasp synthesis and 2 world models. The tags sum to far more than 112 because most methods published since 2024 sit on two branches at once. Fig. 7 draws the tree and the cross-links. The repository carries the row-by-row version of the same thing, one row per method, and that is what to scan when looking for work comparable to your own.

What the deployed policy looks like once training is done matters more than the name a paper gives itself, and on that axis the field has converged hard. Section V-E shows why.

What the tabulation does not record. The emptiest columns of the method tabulation are the ones a reader most needs: only 31 rows state an environment count, only 70 state how many real trials are behind the headline number, and only 39 state how many unseen objects were tested. The trial and unseen-object figures are the audited ones, after Section VII-A recovered 15 trial counts and 7 unseen-object counts that the notes carried and the extraction had dropped. A mostly empty row is not a weak method, but it is one that cannot be compared with any other row here. All 112 rows carry the columns this survey reads, and Appendix A says where they are and how to read them.

B. Reinforcement learning

The standard recipe. Fifty-nine of the 112 method rows learn from a reward. Forty-eight name PPO as the algorithm, and the next most frequent, DAPG, appears four times. Forty-three run their own experiments in a GPU-parallel simulator from the Isaac family or Genesis, and 36 do both. Only 31 rows state an environment count, with a median of 8192 and a maximum of 64000. Twenty-three rows distil a privileged teacher into a deployable student, and 16 combine all three of PPO, a GPU simulator and distillation. That 16-row intersection is the recipe as it is actually practised.

The hand is almost always an Allegro. It appears in 35 of the 112 method rows and in 17 of the 31 reorientation rows, ahead of Shadow at 21 and 6 and LEAP at 12 and 4.

Privilege is handled two ways. OpenAI [34] gives the value network object and target orientation, joint angles, joint velocities and object velocities the policy never sees, with a footnote recording that current object orientation was left out of the policy inputs by accident, and DeXtreme [35] uses the same asymmetric critic with a 2048-unit LSTM against the actor’s 1024. Hora [15] instead compresses nine privileged object properties into an eight-dimensional vector regressed from 30 steps of proprioception, and Visual Dexterity [37] distils twice, from a state teacher to a synthetic point cloud and then to a rendered one, for a fivefold speedup.

Domain randomisation is the part of the recipe with the least discipline. OpenAI [126] made it adaptive, pushing each range boundary out when performance at that boundary exceeds 20 successes and pulling it in below 10, and DeXtreme [35] reproduced the mechanism with a 256-sample queue and 40 percent of environments dedicated to boundary evaluation. Both publish the discovered ranges. Below that standard, two shipped configurations disagree with their own papers about what was

randomised: Hora [15] states a joint-noise range of $U(0, 0.005)$, while its public configuration uses 0.02; its README warns that this revision does not reproduce the paper. PenSpin [90] describes random disturbance forces, while the only public configuration that names them disables them. DexPBT [127] is consistent: its configuration disables randomisation, as the paper says.

Reward engineering. The repository puts the 21 in-hand reorientation methods against nine recurring term families and marks each cell by where the term was found: in the paper, in the code, in both, or in the code with every shipped configuration setting its weight to zero.

The families are not equally popular. Nineteen of 21 methods have a goal or rotation tracking term, which is the task. Thirteen penalise effort as torque, work or joint velocity, 13 pay a sparse success bonus, 11 penalise object velocity and 10 penalise a drop. Eight penalise deviation of the hand from a canonical grasp pose, a family the plan for that matrix did not anticipate and which had to be added. Six penalise action rate or magnitude, five reward closing the distance from fingertips to the object, and four carry a contact or force mark at all.

Three marks an earlier draft recorded as *code* now read *code* (0), because in each the term is in the released code with every shipped configuration setting its weight to zero: DeXtreme [35]’s timeout reward, DexPBT [127]’s fall penalty, and PenSpin [90]’s action penalty. Marking them *code* would tell a reader the code optimises something the paper does not state, and leaving them unmarked would hide a term that is in the file. The repository carries the config key and the zero beside each of the three marks.

Only three of the four contact marks are the finding. AnyRotate [58] scores good and bad fingertip contacts, Lin et al. [128] reward a friction-cone wrench margin, and Zhao et al. [129] track a commanded grasp force. The fourth, [37], penalises the *object* touching the table, a task-shaping term against using the table as a third finger rather than a hand-object term, and its table-contact flag defaults to off, with no shipped config setting it true. Three of 21 in-hand reorientation methods, then, put a hand-object contact or force quantity in the reward, and none puts interpenetration in it. Li et al. [130] penalise interpenetration with a differentiable signed-distance term, but during offline reference construction, not in the policy’s reward. Across all 112 method rows, 85 do not address penetration at all, 5 constrain it, 3 measure it, 3 penalise it and 16 say nothing either way. The physical quality of the contact is not something this literature optimises.

Term counts range from one to ten. DemoStart [131] gives a single binary terminal reward of 1 and lets a curriculum over demonstration states do the shaping. AnyRotate [58] and ViserDex [132] each publish ten weighted terms. ViserDex is the cleanest specification of the 21, with every term, weight and equation in one appendix table, and a statement that the same weights are used for every object.

Weights are frequently unrecoverable. Three of the 21 rows give no numeric weight that survives PDF conversion, because the equation blocks are images, which is a limit of this survey and is recorded as such. Robot Synesthesia [88] is a worse case. Its six coefficients are described only as “tuned hyper-



Eight further rows carry a trajectory-optimisation or MPC tag inside a learned pipeline and are counted on the branch that learns. The double-dashed edge is the one overlap that is a theorem rather than a pipeline: pang_global_planning_2022 proves that a policy gradient and an analytic log-barrier relaxation compute the same local model of contact.

Fig. 7: What supervises a dexterous policy, subdivided to where the 112 method papers actually differ; each leaf names two of its corpus keys and counts the rest, bars scaled by row count, dashed curves marking overlap.

parameters” and no number is printed anywhere in the paper, so the reward it reports cannot be reconstructed from the paper. DexReMoE [133] names an angular-velocity penalty weight in prose with no entry in its hyperparameter table, and DexTrack [114] names an affinity reward that is missing from its weight table.

Curricula, populations and machine-written rewards. Curricula in this literature relax physics or tighten tolerances. DexPBT [127] tightens a success tolerance from 0.075 to 0.01 by a factor of 0.9 every 3000 environment steps once three successes are logged, and ManipTrans [38] starts at zero gravity and high friction and restores both while narrowing a fingertip threshold from 6 cm to 4 cm. DexMachina [39] attaches virtual spring-damper controllers to the object, lets them do the task at first, then decays their gains to zero. The two that ablate the curriculum report it is load-bearing: RotateIt [89] sets its rotation-penalty weight to zero and ramps it to

–0.1, because applying it from the start makes the policy learn only to hold the object still, and ViserDex [132], whose performance-based curriculum runs over regularisation weight, action latency and time between successes, fails completely without the regularisation component.

Population-based training appears once. DexPBT [127] runs populations of 8, 16 and 32 agents, splits them 30/40/30, mutates the middle and replaces the bottom with mutated copies of the top, each float hyperparameter multiplied or divided by a factor drawn from $U(1.1, 1.5)$ with probability 0.2. It reports 30 hours on a single V100 for a five-billion-transition single-arm run, and 0.32 trillion environment steps for its largest population.

Eureka [36] has GPT-4 write the reward function directly, constrained to return a total and a dictionary of named components, and feeds per-component statistics back between iterations. Its appendix gives an example whose reward

correlates at -0.26 with the human-written one and still scores 1.45 on the human-normalised metric, the strongest argument in the corpus that hand-tuned weights are not a ceiling. DrEureka [134] moves the safety requirement into the prompt instead of a penalty term, asking in words for a cube rotating at about 0.25 radians per second with fingers penalised for leaving their initial pose. Neither can be checked against a file. Eureka’s generated rewards are written at run time to a gitignored directory and never committed, so the reward behind any reported number cannot be recovered from the repository, though its README documents the paper’s budget of five iterations and sixteen samples as the default and marks the shipped one-iteration block a fast-test preset. DrEureka’s repository contains only the locomotion and globe-walking trees, so its LEAP-hand reward has no code at all.

Where reinforcement learning stalls. Reinforcement learning stalls on objectives that are not learnable as stated. AnyRotate [58] found the angular-velocity objective unlearnable in the multi-axis setting and replaced it with a moving keypoint target. RotateIt [89] reports that its multi-axis policy “does not converge when training with only reinforcement learning” and adds an imitation loss against single-axis oracles. Eureka [36] cannot spin a pen from scratch and needs a pretrain-then-finetune split, with the from-scratch ablation failing to complete one cycle. Section V-C’s contact-graph term exists because of the general shape of the problem: contact moves the object off the reference, so return goes down, so the policy learns not to touch the object.

It stalls on geometry too. Hora [15] fails on objects under 4 cm across because the fingers collide with each other. What it does not stall on is publication: the real-trial counts behind these headlines are small enough that Section VII-A treats them as the reporting problem they are.

C. Learning from human data

Teleoperation and retargeting. Thirty corpus rows describe a system for getting human motion onto a robot hand; the repository lists them with the retargeting objective compressed to one clause each.

Retargeting splits four ways. The dominant family matches task-space vectors between the human and robot hands. DexPilot [135] set the pattern with a weighted squared distance between fingertip vectors, switching weights of 1, 200 and 400 by whether a pair is within threshold, a minimum-separation constant of 3 cm between primary fingers, and a regulariser toward the open hand, solved online with SLSQP; AnyTeleop [136], Open-TeleVision [137], Bunny-VisionPro [61] and ACE [138] all use the same formulation. The second copies joint angles where the kinematics allow it, as Holo-Dex [139] does for index, middle and ring while solving the thumb by inverse kinematics and discarding the pinky the Allegro does not have. The third learns the map offline: Geometric Retargeting [140] trains a per-finger network in three to five minutes against motion-direction, coverage, flatness, pinch and collision criteria, then runs at 1 kHz against the 60 to 100 Hz of online optimisation, raising configuration-space coverage from 38 to 90 percent and one-time grasp success from 55 to 87.5 percent.

The fourth removes retargeting from the loop: DexUMI [87] builds an exoskeleton whose fingertip workspace is optimised to match the target hand, then regresses encoder values straight to motor values.

What that tabulation does not carry is the point. Only 12 of the 30 rows record a latency at all, three of those give no number, and only 7 state a rig cost. Where costs are stated they are low and falling, from [139]’s \$399 headset through the \$600 exoskeleton of [138] and the \$600 haptic glove of [141] to the \$4,000 backpack of [142] and the \$6,000 rig of [62]. Latency is not one quantity: DexPilot [135]’s “about one second” and Holo-Dex [139]’s “under 100 milliseconds” are end to end, while Geometric Retargeting [140]’s 1 kHz and AnyTeleop [136]’s 9 to 10 ms retargeting are single stages and several cells give a control rate. That tabulation’s latency column mixes the two, so its spread is not a range. Only 13 rows report how many trajectories were collected, and only 7 report hours.

Where throughput is measured, human collection beats teleoperation by a wide margin. Qiu et al. [143] time a grasp demonstration at 3.79 seconds for a bare human and 19.72 seconds through a teleoperated humanoid, and a pour at 4.81 against 37.31 seconds. It attributes the gap to retargeting latency and to the workspace of a 7-DoF arm rather than to anything about the hand.

Imitation architectures. Four architectures cover the corpus. Action chunking came from [144], which predicts a chunk of joint targets with a CVAE and combines overlapping chunks by exponentially weighted ensembling, reaching 80 to 90 percent on fine bimanual tasks from about 50 demonstrations each. Diffusion came from Diffusion Policy [145], which denoises an action sequence and reports a 46.9 percent average success improvement over LSTM-GMM, IBC and BET across 15 tasks. 3D Diffusion Policy [146] swaps the image encoder for a small point-cloud encoder and reports 74.4 percent against 59.8 percent across 72 simulated tasks, and 85.0 against 35.0 percent on four real tasks from 40 demonstrations each. Flow matching is the 2024-onward default for large models and carries 8 of the 112 rows. Discrete action tokens are the fourth, used by [45] and [147], and are the only one of the four that needs no continuous head at all.

Data generation is a separate lever and is underrated. DexMimicGen [63] turns 60 human source demonstrations into 21,000 simulated ones across 9 tasks and 3 embodiments, and a real Fourier GR1 with two Inspire hands reaches 90 percent from 40 generated demonstrations against 0 percent from the 4 source ones. Dex1B [148] iterates optimisation, a CVAE proposal model and a simulator filter into roughly a billion synthetic grasps, and its CVAE baseline beats the prior best by 22 points.

Human video without a robot. Seven rows train a dexterous-hand policy from video with no teleoperation at any stage. DexMV [149] retargets 700 self-recorded demonstrations by matching palm-to-fingertip task-space vectors and estimating actions by inverse dynamics, VideoDex [150] mines 965 trajectories from EpicKitchens to pretrain a neural dynamic policy, and DexVIP [151] goes further into the reward, clustering 715 curated HowTo100M frames into one consensus grasp pose

per object category and adding it as a reward term. OKAMI [152], Lum et al. [153] and Guzey et al. [154] work from a single video each, and Goswami et al. [155] pretrain a world model on 829 hours of EgoDex.

The embodiment gap. Three papers measure the gap directly rather than asserting it. OKAMI [152] runs the same reference-plan and warping pipeline with and without human body and hand retargeting, and the version without it loses 75.0, 42.1, 82.0 and 74.0 points across four task and setting pairs. That is the largest clean embodiment-gap number in the corpus.


Qiu et al. [143] ablate two mitigations separately on the same vertical-grasp task over 10 trials. The full method scores 4 out of 10. Removing the interpolation that slows human trajectories by a factor of four drops it to 1 out of 10. Removing the unified 54-dimensional state space while keeping the slowdown drops it to 0 out of 10, because the policy is handed a shortcut for telling the two embodiments apart. EgoZero [156] reports the sharpest version of the same effect. Removing 3D augmentation drops every one of its seven tasks to 0 out of 15, and substituting monocular depth for triangulated depth does the same, because the best metric depth models it tested carry more than 5 cm of error. It also quantifies what its hand-pose pipeline costs, at 1 to 2 cm of action-label error.

Chen et al. [157] supply the counterexample. Their method retargets only the wrist and lets reinforcement learning discover finger motion, and their ablation that adds a fingertip-matching reward “does not yield benefits and even leads to lower performance”. Richer human correspondence is not uniformly better, and this is the one result in the corpus that says so with an ablation.

Tracking a human reference with physics. Twelve method rows sit in the human-reference tracking family. Eight track a human hand on an object and are covered here; the other four sit at the edges, Lum et al. [153] tracking only the object’s trajectory, DexMan [158] retargeting bimanual video onto a full humanoid, and HumanPlus [47] and OmniH2O [48] tracking whole-body motion. A reference is retargeted onto the robot and a policy trained to make the simulated body follow it: the reference supplies the shaping reward engineering would otherwise have to invent, and the simulator the physical consistency pure imitation lacks.

PhysHOI [3] is the origin of the reward form. It multiplies a body term, an object term, an interaction-graph term and a contact-graph term, and reaches 95.4 percent success on GRAB against 27.0 percent for a DeepMimic baseline. The contact-graph term exists to stop the policy learning not to touch the object. Omnigrasp [31] replaces the action space with a pretrained latent motion prior and reports 94.6 percent grasp success and 84.8 percent trajectory success on GRAB, while stating outright that it omits penetration metrics.

DexTrack [114] adds an imitation loss to PPO and mines its own demonstrations through a homotopy search over easier neighbouring trajectories, reaching 46.70 and 65.48 percent on GRAB at loose and strict thresholds against 38.58 and 54.82 for the best baseline. It has a penetration-depth formula and applies it only to input references, never to its own rollouts, and presents tolerance of “severe hand-object penetrations” as robustness. ManipTrans [38] freezes a generalist hand imitator



The gap. Joint-angle error is a sum over joints and carries no term for the object, so it cannot see the right-hand panel at all. Errors small enough to pass at every joint compose along the finger into a fingertip displacement larger than a thin object is thick, and the pose is scored as good with the object inside the finger. Sec. V’s one method that measures penetration measures it on the reference, before the policy runs.

Fig. 8: A retargeted pose can be scored good and still be inside the object. The joint-angle error is a sum over joints with no term for the object; a few degrees at each joint compose along the finger into a fingertip displacement larger than a thin object is thick.

and trains a per-task residual on top, reaching 58.1 percent single-hand and 39.5 percent bimanual success on OakInk-V2. Its stance on contact is to raise the friction coefficient above the real value rather than model skin deformation, and its real deployment is open-loop replay.

DexMachina [39] moves to Genesis with 12,000 environments and six hand URDFs, and drives the object with virtual controllers that decay to zero. Its re-implementation of [38]’s curriculum does not beat no curriculum on its own setup, which is a direct disagreement between two papers worth quoting in both directions. Dexplore [92] drops the explicit retargeting stage and treats the reference as soft guidance, with termination thresholds derived from the failure rate, reaching 87.7 percent on GRAB with an Inspire hand.

TopoRetarget [7] is the only corpus method that quantifies penetration as an evaluation quantity, reporting maximum depth and the fraction of frames past 2 mm, and constraining it during retargeting with a 1 mm soft tolerance and a 30 mm hard bound. It reaches 87.5 percent on its own pen-spinning set against 46.9 for the best baseline. It does not re-measure penetration after the tracking policy runs, so the property it constrains is a property of the reference and not of the behaviour. Fig. 8 draws why the usual retargeting objective cannot see the failure at all: joint-angle error is summed over joints and carries no term for the object, so errors small enough to pass at every joint compose along the finger into a fingertip displacement larger than a thin object is thick, and the pose is scored as good with the object inside the finger.

That is the pattern across all eight, and it is the reference-versus-rollout split in its sharpest form. Penetration is handled at the reference, if at all, and never at the rollout. Chen et al. [157] complete the set with the only real-robot numbers among them, from 100 percent on a microwave and a laptop down to 41.2 percent on a ketchup bottle over 20 trials each.

Cross-hand retargeting. Human data does not port across hands, and the map is usually left unstated. Fifty-three method rows use human data and name 44 distinct hand strings between them, and twenty of those appear among the 30 teleoperation systems with a stated retargeting objective. The

other 33 never say how the human motion reached the hand, and the stated objectives do not converge, from a fingertip keypoint-vector energy in [136] to a per-joint regression from exoskeleton encoders in [87] to no retargeting at all in [156]. A shared benchmark could make these mappings comparable. TopoRetarget’s format [7] is a useful starting point: it reports contact precision, contact alignment, maximum penetration and the share of frames past 2 mm, per hand and per dataset.

D. Generalist and vision-language-action policies

The finding is the size of the hand. Eighteen rows carry the generalist tag. Fourteen of them settle whether the reported evaluation ran on a multi-fingered hand. Eight of those fourteen did, six report no hand result at all, and the remaining four never say, GR00T N1.6 [49] naming no end effector anywhere on its page.

Five of the eight state a hand size, and that comparison has to be made in actuated degrees of freedom rather than joints, for the reason Section III opens with. Four are six, the Inspire RH56DFX among them actuating six of its twelve joints, and one is twelve, [159]’s XHAND. Their median is 6. Two of the four rows that never settle the hand question do state a size, and both are large: 21 for [100]’s ByteDexter V2 and 22 for [160]’s Sharpa Wave. Of the 59 reward-learning rows, 48 state a count, 38 of those are 16 or above, and their median is 16: an Allegro and a LEAP actuate 16, and a Shadow actuates 20 of its 24 joints. So no row that settles the question reaches the band the reinforcement-learning literature of Section V-B works in, and the two rows that reach it on paper are the two that never say whether the hand was in the evaluation. An earlier version of this claim counted eleven rows as evaluating on a hand, which mixed the settled eight with rows the notes leave open.

The six with no hand result are parallel-jaw throughout (π_0 [42], whose released code encodes each ALOHA gripper as one scalar, $\pi_{0.5}$ [43], $\pi_{0.6}^*$ [44], OpenVLA [45], Helix [99], which claims individual-finger control of a 35-DoF upper body with no success rate or trial count for any task, and Gemini Robotics [161], whose Apollo hand appears in a figure caption with no number attached). Its successor [46] closed that gap with Apollo progress scores, not success rates, of 0.74 down to 0.62 across five generalisation axes. DexVLA [66] is the borderline case the other way: a hand in one real-robot task family, grippers in the other five and in all 91 pretraining tasks.

The released code also drifts. GR00T N1 [162]’s repository is a later N1.7 generation with a different backbone, a different action horizon and embodiments the paper does not describe, and the $\pi_{0.5}$ repository [43] supports only the flow-matching head, not the hybrid recipe the paper describes.

Dexterous hands in generalist policies. The gap is size rather than absence. The generalist policies that do touch a hand run it at roughly a third of the actuation the reinforcement-learning literature assumes, and the rows that never settle the question are the smaller half of the same gap. The standard generalist suite, with the hand’s actuated degrees of freedom and its vendor printed beside the number, should include at least one dexterous-hand task.

E. Model-based baselines and hybrids

Three corpus methods solve dexterous tasks without learning a policy, and they are the control group the rest of this section lacks.

PDDM [125] learns an ensemble of dynamics models and plans through it with an MPPI-style optimiser, replanning every step. It needs 1 to 2 hours of data in simulation and 2 to 4 hours on a real 24-DoF Shadow Hand. Predictive Sampling [163] removes the learned model too and samples ten rollouts per step through MuJoCo itself, planning in 1 to 20 milliseconds. It reorients a cube with a Shadow Hand in real time from scratch, and reports no success criterion, no trial count and no real-robot result.

The work of Pang et al. [55] is the most substantial of the three. It proves that the randomised smoothing implicit in reinforcement learning and an analytic log-barrier relaxation compute the same local linear model of contact, then uses the analytic version inside a trajectory optimiser and an RRT. Allegro in-hand rotation takes 19.59 seconds to optimise and its plate-pickup task 117.16 seconds of planning on a 16-core CPU, against the GPU-days of Section V-B, and it is the only method here that imposes non-penetration as a hard constraint. Its own limitation is the honest part: its 3D systems transfer to hardware far worse than its 2D ones, because the quasi-dynamic assumption breaks and planned grasps miss contacts under a second-order solver.

Hybrids, and where the field has converged. The recurring shape is reinforcement learning in simulation distilled into a policy that looks like an imitation policy, and it appears in four variants. The first distils a privileged teacher into a vision student inside one paper, which is [15], [37], [89], [88] and [132]. The second uses reinforcement learning as a demonstration factory: DexTrack [114] mines demonstrations with per-trajectory RL and trains a generalist on them, ManipTrans [38] does the same to build a 3.3K-episode dataset, and DexterityGen [164] trains a diffusion controller on ten billion simulated grasp-to-grasp transitions and projects a human teleoperator’s coarse command onto it, taking four reorientation tasks from 0 out of 20 under raw teleoperation to 12, 13, 10 and 9 out of 20. The third generates data without reinforcement learning at all, which is [63], [148] and [165]; Dex1B’s row is classed as a dataset, so Fig. 7’s leaf, which counts method rows, holds the other two. The fourth wraps an RL policy in a residual, which is [166] at 88.8 percent over 3,200 objects in 12 GPU-hours.


In the variants that dominate 2025 and 2026 work, the reward is no longer the objective of the deployed policy. It is the objective of the process that produced the deployed policy’s training data. Reward engineering has moved upstream into data curation, and every reward-shaping pathology in Section V-B now reaches the shipped policy through a dataset rather than a gradient.

F. Paper against released code

Thirty-eight of the 112 method rows record a disagreement between a paper and the code it released, and all 38 released code, so they sit inside the 62 rows that released anything

TABLE III: The 9 rows where a paper and its released repository state different values, at the commit fetched, with confidence.

method	conf.	repository, fetched commit	file in it, and where	what the paper prints	what that file contains
PDDM [125]	high	google-research/pddm 06b88cdebaf	pddm/envs/cube/cube_ env.py _get_obs	observation dimension 46 for in-hand reorientation, Table 2	the six fields <code>_get_obs</code> concatenates sum to 39 on the code's own inline comments
DexPoint [56]	high	yzqin/ dexpoint-release 17f1e238bb	dexpoint/env/rl_env/ relocate_env.py AllegroRelocateREnv. get_reward	a four-term reward, Eq. 5	<code>get_reward</code> sums nine, among them 1.0 / (0.06 + <code>finger_ object_dist</code>) and <code>controller_penalty</code>
DeXtreme [35]	high	isaac-sim/ IsaacGymEnvs aeed298638	isaacgymenvs/cfg/task/ AllegroHandDextremeADR. yaml actionDeltaPenalty Scale	action-delta penalty weight -0.25, Table 2	<code>actionDeltaPenaltyScale</code> : -0.2, and -0.01 in the ManualDR yaml beside it
Visual Dexterity [37]	high	Improbable-AI/dexenv ad9634e9d2	dexenv/conf/dclaw.yaml alg.num_envs	32,000 teacher training environments, Table S1	<code>num_envs</code> : 8000, and 16384 in the parent <code>hand_default.yaml</code>
DexPBT [127]	high	NVIDIA-Omniverse/ IsaacGymEnvs aeed298638	isaacgymenvs/tasks/ allegro_kuka/allegro_ kuka_base.py compute_kuka_reward	four staged reward terms, Sec. III-C	eight summed components, <code>hand_ delta_penalty</code> multiplied by 0 with the comment currently disabled
PhysHOI [3]	high	wyhuai/PhysHOI 6095c605e2	physhoi/env/tasks/ physhoi.py compute_humanoid_ reward	object-rotation weights 0.1 and 0.01 for GRAB, Table 4	<code>eor</code> and <code>eorv</code> set to <code>torch. zeros_like(ep)</code> , the computation commented out on the same lines
UniDexGrasp [167]	high	PKU-EPIC/UniDexGrasp 36c9bfcf7c	dexgrasp_policy/ dexgrasp/tasks/shadow_ hand_grasp.py compute_hand_reward, goal_cond branch	seven named weights from 0.1 to 10, Table 7	a <code>torch.where</code> cascade on literals -0.5, -1.0, 0.9, 0.1, 0.2, none of the seven present
PenSpin [90]	medium	HaozhiQi/penspin 5035c52dc9	configs/task/Allegro HandHora.yaml forceScale	a disturbance force of 0.2 times object mass at probability 0.25, Table 8	<code>forceScale</code> : 0.0 and random <code>ForceProbScalar</code> : 0.0, the only shipped config naming either five summed terms; <code>_compute_ energy_reward</code> ends return 0 and <code>_compute_fingering_ reward</code> return 0.0
PianoMime [52]	high	sNiper-Qian/ pianomime c4abefac8d	single_task/piano_ with_shadow_hands_res. py _set_rewards	two weighted terms, 2/3 and 1/3, Table 3	

**Fig. 9:** Papers and their own released repositories state different things, and the funnel ends at the 9 the survey stands behind as contradictions. Table III names them, with the file and both values.

(Fig. 9). They are not one kind of thing. Nine are contradictions, where paper and code state different values or different terms. Thirteen are limits of this survey's own parse, where the body or config that would settle the question was never recovered and the row says so. Eight released code without the described component in it, four are version skew against a later repository, and four are a paper disagreeing with itself. The fourth version skew is [49], which ships a main branch one generation later than the checkpoint its page describes.

Nine is the number to quote, eight at high confidence and one, PenSpin [90], held at medium. Nine of 62 is 15 percent, and it is a floor rather than a rate: the census covers method rows only, so RoboPianist [50], whose row is a benchmark,

does not enter it although it sums five reward terms against the three its Table 2 documents. It is also a screened count rather than a first pass. Sixteen rows were drawn at first and seven have since left the class, six under adversarial re-reading and one, OmniH2O [48], once writing to its authors sent someone back to the evidence; two more were narrowed and stayed, and each revision is recorded in the row it concerns beside the comparison it revises. That screening is evidence that the nine were checked, not a bound on how many a fuller reading would find. Among the 21 reorientation methods scored term by term, seven released a repository and four of those state something different from their paper, the other three being a missing environment, a later generation and a default-value question their own READMEs settle.

PhysHOI [3] is the clearest example. Its paper assigns nonzero weights to object rotation and rotation velocity on the GRAB dataset. In the public implementation, both errors are set to zero and the calculations are commented out. The paper does specify zero weights for the separate BallPlay task, but the implementation does not condition this choice on the dataset. Its reported success criterion uses position alone, so that score cannot expose the difference. The claim here is limited to this conflict between the paper and the public revision listed in Table III; it does not establish which code produced the published experiments.

The other cases vary in importance. DeXtreme [35] gives one action-change penalty in the paper and a different value in each of two public configurations. PianoMime [52] describes two reward terms, while its implementation sums five and forces two of them to return zero. Table III gives the exact

file, revision, and values for these and the remaining six cases, where they can be checked without interrupting the argument with source-code identifiers.

Scope and correction. The authors of the nine works were written to on 19 September 2026, before this survey was posted, and asked whether the reading was right and whether they wanted the wording changed; each was given until 10 October 2026, and told that silence would be recorded as silence rather than as agreement. The comparison is nonetheless limited to two public artefacts: the paper and the repository revision named in Table III. It makes no claim about intent or about the private code used for the reported experiments: a fetched commit is not the code that produced a paper’s numbers. It may predate, postdate, or differ from that code. Hora [15] says exactly that about itself by directing readers to a tagged release rather than its default branch. The letters are in `outreach/`, with the address each went to and its provenance in `outreach/RECIPIENTS.md`, and replies are recorded in the row they concern. Every comparison is correctable in public through the address in the author block; supporting evidence changes the record, the derived counts, and the next version of the text.

The classification distinguishes a term that is absent, present with zero weight, or assigned a different weight. In 13 cases the parsed repository cannot settle the comparison, and those cases are recorded as limitations rather than conflicts. The reproducibility materials preserve all 38 records, including the file, both values, confidence, and review history.

Review and correction. This is a result about publishing practice in robot learning rather than about dexterous manipulation in particular: the dexterous corpus is its sample, not its subject. Nine is a lower bound because 46 method papers released no code to inspect and four comparisons remain unsettled. It is also the result of a conservative review. The first draft identified sixteen conflicts; seven have since left the count and two more were narrowed and remain in it. Six failed under adversarial rereading, and a seventh was reclassified when its apparent weight mismatch proved consistent with a curriculum factor. PenSpin [90], the sole medium-confidence case retained, concerns a disabled disturbance rather than the tactile input claimed in an earlier draft. Each change and its reason is preserved with the underlying record. Future papers can avoid most of this ambiguity by generating the reward table directly from the configuration used for training and citing the corresponding revision.

VI. BIMANUAL DEXTEROUS MANIPULATION

A warning first. Fifty-three corpus method rows carry the two-hand flag, and that flag says only that the robot has two end effectors. Thirteen of the 53 put no dexterous hand on the robot at all. ALOHA, RDT-1B, EgoMimic and H-RDT [144], [168], [169], [170] and UMI [171] are parallel-jaw throughout; π_0 , $\pi_{0.5}$ and $\pi_{0.6}^*$ [42], [43], [44] and Diffusion Policy [145] name no hand and run grippers on every embodiment; Gemini Robotics [161] and Gemini Robotics 1.5 [46] report every bimanual number on grippers and show five-fingered hands only qualitatively; and Helix [99] and GR00T N1.6 [49] name

no end effector anywhere on the page. Three more run a gripper on one embodiment and a hand on another: ACE and 3D Diffusion Policy [138], [146] and Open-TeleVision [137], whose Unitree H1 carries 6-DoF Inspire hands and whose Fourier GR-1 carries a 1-DoF jaw.

The denominator for this section is 28: the method rows whose notes place a learned closed-loop controller on two multi-fingered hands. It is the 53 less those 16; less two static grasp-synthesis methods, BiManGrasp [6] and BiDexGrasp [5]; less three systems with no learned policy, Castro et al. and DexTeleop-0 [108], [86] and Pang et al. [55]; and less four rows where no learned policy holds both hands: OmniH2O [48], whose fingers are mapped open-loop from the Vision Pro and sit outside its 19-DoF policy; OKAMI [152], whose headline pipeline is open-loop retargeting with a learned policy only in a side experiment; DexDeform [53], a skill model refined by trajectory optimisation rather than a closed-loop controller; and Omnigrasp [31], a simulated human body with no bimanual task. Frau-Alfaro et al. [41] never enter, because their row records one hand. Benchmarks and datasets are outside the 28 by class, Bi-DexHands and Bench2Dex [64], [104] and RoboPianist [50] being benchmarks and RP1M [51] and HumanoidGen [172] datasets; they are quoted here as evidence and never counted. Every paper below runs two multi-fingered hands unless said otherwise.

A. Coupling, roles and collision

When both hands hold the same object, the object closes a kinematic loop between them. Neither hand can move without changing what the other must do. BiDexGrasp [5] reports that the coupled objective “often yields imbalanced solutions, where one hand dominates stability while the other contributes marginally.” ArtiGrasp [173] reports simulation speed scaling “roughly quadratically with the number of contacts,” and trains each hand alone before pairing them.

Role asymmetry is an assumption most of these systems make silently. BiDexHD [174] states it outright: “we assume the robot to be right-handed by default, i.e., the left hand handles the target object and the right hand handles the tool,” and Lin et al. [175] bake the same split into their reward, putting reference contact keypoints on the bottle base for the left fingertips and on the lid for the right. The bias is in the data first: TACO [176] recruited 14 right-handed subjects and measures the right hand moving consistently faster.

The arms collide, and the corpus handles this by construction rather than by control. RoboPianist [50] ships a forearm-forearm collision term that its own reward table never lists. Two papers name arm collision as unsolved. DexMimicGen [63] “does not explicitly handle collisions,” and DexMan [158] says its control parameterisation “overlooks full arm posture,” which it needs to avoid them.

The action space doubles and the observation grows faster. Bi-DexHands [64] uses a 52-dimensional action for 18 of its 20 tasks against observations of 398 to 446 dimensions. GR-Dexter [100] emits an 88-dimensional chunk of arm joints, end-effector poses, hand joints and fingertips. What the larger vector does not carry is any statement of who does what. The

usual answer is one scalar reward summed over both sides, and the cost is a compounding failure rate. ManipTrans [38] scores success only if both hands succeed, and its OakInk-V2 success rate falls from 58.1 percent on single-hand sequences to 39.5 percent on bimanual ones.

B. Coordination architectures

Fig. 10 sets the four mechanisms side by side. Every row of the 28 carries exactly one architecture, read from its note: one policy over both hands, a network per hand, leader and follower, or a note that does not say. The fourth panel is not a fourth architecture but the relative frame, and it is the same single row as the third.

Twenty-one of the 28 put one policy over both hands, which is three quarters of the set: the trackers and tracking-adjacent methods [39], [158], [38], [157] and [47]; the reinforcement-learning tasks [127], [175], [36], [52] and [124]; the demonstration pipelines [62], [142], [177], [63], [68] and [143]; and the generalist policies [100], [162], [159], [147] and [160]. In every one of them a single network takes a concatenated two-hand observation and emits a two-hand action. Two rows do not say which they are, Bunny-VisionPro [61] and DexImit [165], whose notes describe the rig and the data pipeline but never the policy’s own decomposition. The concentration is not the outcome of a comparison that was won.

Four of the 28 give each hand its own network, and the two papers that compare the choice disagree. Bi-DexHands [64], a benchmark row and so outside the 28, ships the MARL baselines and finds PPO over the full observation beats HAPPO and MAPPO “in most cases,” with the gap narrowing on tasks that need both hands, because PPO “can use all observations” where MARL sees only part. BiDexHD [174] concludes the opposite. Its independent PPO teachers score 74.59 percent stage-two tracking rate on trained tasks against 53.88 for a centralised policy over both observations. Both are simulation only, on different tasks and hands, so neither settles it. ArtiGrasp [173] trains one PPO policy per hand, and Dynamic Handover [40] and DyDexHandover [60] use MAPPO with one agent per arm-hand system.

One of the 28 assigns explicit leader and follower roles, and the same one is the only policy expressed in a relative frame. AsymDex [59] is that paper, and the only corpus paper that puts either mechanism inside a policy. It gives the dominant hand full finger and wrist control, restricts the facilitating hand to a 6-DoF base pose, and writes the dominant hand and the object in a frame attached to the object the facilitating hand holds. That cuts the observation from 176 dimensions to 88 and the action from 52 to 26. It ablates the two mechanisms separately, which no other corpus paper does. On Block in cup the full method scores 0.7701 over five seeds, against 0.1086 for relative frames without asymmetry and 0.0164 for asymmetry without them. Twist Lid transfers zero-shot at 18 of 20 real trials. The real rig pairs a 16-DoF Allegro with a 6-DoF Ability Hand because only one Allegro was available, so the roles are confounded with the hardware. Two other rows use a relative quantity without giving either hand a role: DexWild [177] appends the inter-hand pose to its observation,

and DexMimicGen [63] applies one shared SE(3) transform to both arms’ source segments when generating data, which is the mechanism offline rather than in a policy.

Evidence still needed. One architecture carries three quarters of the field, and the choice behind it has been compared twice, with opposite outcomes, and ablated once. A decisive comparison would run one handover task on three architectures with the same hand and the same seeds, reporting separate giver and receiver returns rather than a single shared one.

C. Benchmarks, datasets and the shared axes

Two purpose-built bimanual dexterous suites exist in the corpus, four years apart.

Bi-DexHands [64] is 20 tasks on two Shadow Hands in Isaac Gym, ordered by the infant age at which humans acquire the skill, at 2048 environments and a reported 30,000-plus FPS. Its measurement discipline is weaker than its coverage. It reports reward and normalised score, never a success rate. Its only success flag in code tests the object-to-goal distance against 3 cm, which ignores orientation and exists only in the four catching tasks. Any success rate later work attributes to Bi-DexHands comes from that flag or its own definition.

Bench2Dex [104] is the more instrumented of the two. It runs 26 long-horizon tasks in Isaac Lab across 12 arm-and-hand embodiments, with roughly 1.3K teleoperated demonstrations in eight modalities, including a ray-cast tactile image that maps all 12 hand geometries into one representation. Success needs the terminal predicate held for 0.5 s. Matched, GR00T N1.5 leads at 631 of 1300 rollouts, 48.5 percent, against ACT at 29.5, $\pi_{0.5}$ at 27.3 and Diffusion Policy at 12.9. Under combined perturbation those first two compress to 19.8 and 19.7. Tasks are not crossed with embodiments, so no contrast isolates the hand, and the authors state the hand and tactile models are not calibrated against matching physical hardware.

Ten corpus rows are bimanual datasets, and five are two-hand interaction capture. ARCTIC [178] gives 339 mocap sequences of 11 one-DoF articulated objects held in both hands, and deliberately labels over-shooting vertices as in contact, so its ground truth records interpenetration as contact; Fig. 11 is the chain that produces it. TACO [176] gives 2.5K bimanual tool-use sequences fitted markerlessly under an attraction loss and a penetration loss the paper says trade against each other. OakInk-V2 [179] is the only corpus dataset publishing a penetration number for its own annotations, a mean depth of 0.25 cm over its 627 long-horizon sequences. HOT3D [180] gives 833 egocentric minutes with mocap-grade poses for both hands and up to six objects, and GigaHands [181] gives 2,034 markerless minutes from 56 subjects, neither with a penetration metric. These five feed the tracking work here: DexMachina [39] and ArtiGrasp [173] from ARCTIC, BiDexHD [174] from TACO, ManipTrans [38] from OakInk-V2 and DexMan [158] from both.

Methods on the same axes. The method tabulation carries the bimanual papers on the same columns as every other method here, and three of those columns are worth reading together.

Real trials. AsymDex [59] reports 20 per task, HATO [68] 10 per condition, DexMimicGen [63] 20 on a Fourier GR1





Fig. 10: Four mechanisms for controlling two dexterous hands, counted over the 28 corpus rows with a learned closed-loop controller on two multi-fingered hands, each read from its note. The panels are not a partition: the fourth is the same single row as the third, and the two rows whose notes never say which architecture they use appear in neither.



The annotation is as good as the fit and no better, and a penetration is recorded as a contact rather than as an error. Of the five bimanual capture sets in Sec. VI, one publishes a penetration number for its own annotations.

Fig. 11: Where a capture set’s contact annotation comes from. Nothing in the chain is measured at a surface, and a vertex that has gone through the surface is written down as a contact rather than as an error.

and BiDexGrasp [5] 260 across 30 objects. ManipTrans [38] reaches real Inspire hands by open-loop replay with no trial count and no success rate. Seven of the 28 have no real robot at all: [173], [174], [39], [158], [127], [60] and [52]. The two piano rows outside the denominator, RoboPianist [50] and RP1M [51], are simulation-only as well.

Penetration, and Fig. 12 is why it has to be reported beside a success rate rather than instead of one. Most bimanual rows read “not addressed.” BimanGrasp [6] gates on it, failing any grasp whose total penetration across hand-object, self and inter-hand exceeds 1.5 mm, and reports that “penetration remains the primary cause of grasp failure.” BiDexGrasp [5] reports it as a number, 0.15 to 0.20 cm maximum depth for its own grasps. DexImit [165] carries a named hand-hand penetration term in its grasp-synthesis objective. DexMachina [39] resolves it once, replaying retargeted joints against a fixed object “to eliminate object penetrations” and measuring nothing during rollout. ArtiGrasp [173] declines the metric outright, because its baselines “include a physics simulation which exhibits no interpenetration.”

Comparability. BimanGrasp [6] reports 54.03 percent success in Isaac Gym at friction 3. BiDexGrasp [5] re-runs the same




Fig. 12: Two failures of synthesised two-hand grasps, each enlarged at the contact: (A) a finger inside the object, (C) a finger inside the other hand. Both look plausible whole [6].

Fig. 9 of Shao and Xiao, arXiv:2411.15903, CC BY 4.0; panels (A) and (C) of Fig. 9.

grasps in MuJoCo at friction 0.6 and gets 26.80 percent, at 1.52 cm penetration depth. Neither number transfers, and Section V-C’s re-implementation result is the same lesson on the training side.

D. Handover, and what transfers from one hand

Handover is the one bimanual task where the hands are unambiguously asymmetric, because one gives and one receives. All three corpus handover papers use a single shared reward across giver and receiver. None defines separate objectives for the two roles.

Dynamic Handover [40] states one formula, $r = r_{\text{dis}} + r_{\text{linvel}} + r_{\text{torque}}$, and never says whether thrower and catcher receive different decompositions of it. DyDexHandover [60] is explicit that they do not: “Both hands share aligned objectives, forming a fully cooperative relationship,” with one weighted sum feeding both agents. Only the thrower gets an extra KL regulariser toward a policy pretrained on human throws, which constrains style and not role.

The third case is weaker still. In [41] only the receiver is learned. The giver is a UR5e with an Allegro hand that “holds the object without moving during the whole episode.” There is one agent, one reward and no second policy, its corpus row accordingly records one hand, and it is one of the rows the 28 excludes. The 94 percent often attached to this paper needs its conditions. It is Total Success, which counts “Indetermination” cases where the simulator failed to resolve collisions and the object clipped through the giver’s hand, on the short prism, in simulation, over 100 episodes, with no real robot in the paper.

Nothing in the handover literature measures contact quality. Contact appears only as a positive signal, a boolean per-phalange touch in [41] and a boolean contact reward in [60]. The two systems that do handover well sidestep the problem. HATO [68] reaches 10 of 10 on a real slippery handover by imitating teleoperation with fingertip touch sensing and no physics model. Lin et al. [124] stage their reward with a discrete variable switching which hand’s fingertips are scored, and still calls handover its hardest task at 52.5 percent real success.

What transfers from single-hand work. The training recipe transfers intact. PPO with thousands of parallel environments, an asymmetric critic on privileged state, domain randomisation and distillation into a vision policy is the same in [175] and [124] as in Section V-B. Reward forms transfer literally, and the contact-goal reward $r_{\text{contact}} = \sum 1/(1 + \alpha d)$ appears in both.

Multi-task learning does not transfer, and the bimanual benchmarks show it clearest. Bi-DexHands [64] finds multi-task PPO and ProMP fail on MT4, MT20, ML4 and ML20. RoboPianist [50] finds F1 on the shared training song drops “from roughly 0.7 F1 for 1 song to almost 0 F1 for 16 songs,” with no positive transfer “regardless of the size of the pre-training tasks.”

Single-hand evaluation does not transfer either. Success criteria compound, because ManipTrans [38] fails a trajectory “if either hand fails to meet these conditions.” And object-centric metrics cannot say which hand caused the failure. BiDexHD [174] scores the steps where both objects track their references, and DexMachina [39] averages ADD across object parts before an AUC. Both say whether the task worked and nothing about which hand lost it.

One failure mode has no single-hand counterpart. Two hands can penetrate each other, and the corpus almost never looks. Three papers carry an inter-hand penetration term, and all three are static grasp synthesis: [6], [5] and [165]. Not one learned bimanual controller in the corpus measures or penalises hand-hand penetration during a rollout. PianoMime [52] ships a flag that turns inter-hand collision off in the physics.

VII. EVALUATION, AND A PROTOCOL FOR COMPARISON

A. Reporting practice

Of the 112 method rows in the corpus, 89 report a real-robot experiment, 22 do not and one row is unsettled, which is 80 percent of the 111 the note settled. Among those 89, 70 state how many real trials produced the headline number, 79 percent of them. The 89 is the denominator that belongs to




Fig. 13: What gets reported, and what does not: six shares of the 112 corpus method rows, each bar drawn against the denominator the text gives, with the two rarest reporting practices marked in colour.

this statistic: the 22 rows with no real robot cannot state a real trial count, and counting them as silent turns a definitional impossibility into a reporting failure. Thirty-nine rows state a count of unseen test objects, 35 percent. Ninety-eight state how a rollout is scored, 88 percent. Sixty-two released code and 46 did not, with four rows unsettled, 57 percent of the 108 the note settled. Fig. 13 draws these six shares, each against the denominator that belongs to it. Every bar is a lower bound, for the reason the next paragraph sets out. Two of the six are marked, because they are the pair a reader needs in order to compare any two methods: without a count of unseen objects a success rate has no generalisation denominator, and without contact or penetration handling it says nothing about whether the hand passed through the object. They are also the two the field reports least.

These are counts of what this survey captured, not of what papers reported, and every one is a floor. A structured row holds a scalar. A paper that reports ten trials on each of nine tasks, or a scoring rubric instead of a threshold, or a count spread over four tables, has nothing the extraction can reduce to one integer, so it produces a null, and a null is then indistinguishable from a paper that said nothing. The bias runs one way: every miss converts a reporting paper into a silent one, and the survey’s argument is that the field reports badly, so the artefact flatters the argument. Section V-F makes the same disclosure about the paper/code count, subtracting the 13 disagreements that are limitations of this survey’s own parsing before declaring which number to quote, and the coverage statistics above need it more.

The nulls behind those six shares were therefore audited by hand against the notes they came from, and every number above is post-audit. Appendix A gives what that audit recovered and what it could not: the counts moved by tens of rows, 19 trial counts and 4 criteria remain genuinely unsettled, and every bar in Fig. 13 is still a lower bound.

The remaining gap is the one that matters. Nineteen of the 89 papers with a real robot never say how many times they ran it. A percentage with no denominator cannot be given an interval, so it cannot be compared with anything.

Where the denominator is stated it is small, and it is not one quantity. Some stored counts are per-cell, meaning ten trials on each task, or twenty per condition, or five per object. Others are grand totals over every cell. The rows now record which of the two each value is, and the two distributions are quoted separately. The 39 per-cell counts run from 5 to 100 with a median of 15 and quartiles at 10 and 20. The modal cell is




Fig. 14: Where a denominator is stated it is small: the 39 per-cell real-trial counts, binned by fives, median 15 marked. The 24 grand totals are a separate quantity, counted beneath the axis rather than pooled in.

10 trials, in 16 rows, then 20, in 12. The 24 grand totals run from 12 to 750 with a median of 110. Pooling the two gives a median of 20 and a range of 5 to 1287, and that pooled figure is the one an earlier draft of this section quoted. It describes nothing, because Hora [15]’s 240 and Visual Dexterity [37]’s 20 are experiments of comparable size recorded on different bases. Seven further counts could not be assigned a basis at all. Fig. 14 is the per-cell distribution.

The per-cell count is the one a Wilson interval attaches to. For s successes in n trials, with $\hat{p} = s/n$ and $z = z_{1-\alpha/2} = 1.96$ at the 95 percent level, the bounds of that interval are

$$\frac{2n\hat{p} + z^2 \pm z\sqrt{4n\hat{p}(1-\hat{p}) + z^2}}{2(n + z^2)}, \quad (1)$$

which, unlike the normal approximation, stays inside $[0, 1]$ and does not collapse to zero width at $\hat{p} = 0$ or $\hat{p} = 1$. At 20 trials a reported 60 percent carries a 95 percent Wilson interval of 39 to 78 percent, and a reported 80 percent an interval of 58 to 92 percent. Overlapping intervals are not a test, so the point is made with one: 12 of 20 against 16 of 20 is $z = 1.38$, $p = 0.17$, and the difference is not established. Two methods separated by 20 points at 20 trials, the larger of the two modal cell sizes, are not separated at all.

The share of papers stating a count has risen, from 39 of the 65 rows before 2025, 60 percent, to 31 of the 47 rows from 2025 and 2026, 66 percent. The counts themselves have not. The median stated count is 20 in 2024, 22.5 in 2025 and 20 in 2026. Before 2024 the per-year medians rest on one to seven observations and should not be read as a trend.

The denominators also sit on different hardware. The 103 method rows that name their own hand give 78 distinct hand strings between them, and this survey applies no normalisation to those strings, so 78 is a count of strings and not of hand designs. Matching on the string, Allegro appears in 35, Shadow in 21, Inspire in 19 and LEAP in 12. A success rate on a 16-degree-of-freedom Allegro and a success rate on a 6-actuator Inspire hand are not measurements of the same thing.

Nor are the criteria. DexVerse [105] counts PickCube a success when the cube is “lifted at least 0.20 m above its resetting height”. Bench2Dex [104] requires its terminal predicate to hold for a continuous dwell time of 0.5 s, to reject transient contacts. The Colosseum [182] counts an episode

successful “if the model completes the task fully”. DexTrack [114] reports every success rate as a pair under two threshold sets, which on GRAB gives 46.70 and 65.48 percent for the same rollouts. Of the 14 rows that still state no criterion, ten have no success predicate at all. They report radians rotated or time-to-fall and never define a success, which is a fact about the paper rather than a gap in this survey. And four are unsettled by the note. Of the 19 criteria the audit recovered, all but one are a rubric, a staged partial credit, a judgement by eye or a deferral to a benchmark’s own definition rather than a threshold, and exactly one, [44], states a verbatim numeric one; Appendix A counts the four kinds. A rubric is a milder failure than silence and a worse one than a threshold, because it is reproducible inside a lab and not across two.

The axes that matter. Seven quantities dissociate in the published data, so they have to be reported separately.

Task success. Binary success discards the difference between near-misses and inaction. Snyder et al. [183] put it plainly: “a policy that completes 90% of the task is clearly better than a policy that is frozen the whole time, yet their success rates would be identically 0%”. In [184] policy C scores 17 percent overall on the pancake task while picking up the spatula and flipping the pancake in 23 of 23 attempts.

Robustness to perturbation. The Colosseum [182] measures a 30 to 50 percent success drop under single perturbation factors and at least 75 percent under all 14 together. In [104] GR00T N1.5 leads the matched condition at 48.5 percent and falls to 19.8 percent under combined shift, while $\pi_{0.5}$ goes from 27.3 to 19.7 and retains the most at 72.1 percent. The ranking at the anchor is not the ranking under shift.

Generalisation to unseen objects. Only 39 rows state a count and the median is 11 objects.

Physical plausibility of the contact. Eleven of the 96 rows whose contact handling the note settled address it, 11 percent, with 16 rows unknown. Section VII-B takes them apart.

Sample and wall-clock cost. Thirty-one of 112 rows state a parallel environment count and 18 a simulated episode count. RoboPianist [50] is the exception, at 5 million samples per song and roughly 5 hours per run on four Tesla K80 GPUs.

Real-robot transfer. Simulated rank order is not real rank order. AutoEval [185] scores Open- π_0 on put-eggplant-in-sink at 6 of 50 in SIMPLER and 47 of 50 on the real WidowX. Badithela et al. [186] state the limit directly: “the simulation-to-real gap precludes rigorous statistical inferences about real-world outcomes from simulation results alone”.

Reproducibility. 62 rows released code that could be parsed against the paper, and 38 of the 112 rows record a disagreement of some kind between the paper and that code. All 38 released code, so the raw rate among code-releasing rows is 61 percent. That raw rate is not the finding, because the 38 are not one thing. Section V-F classifies them: 9 contradictions, 13 limitations of this survey’s own parsing, 8 components never released, 4 version skews and 4 inconsistencies internal to a paper. Only the contradictions are a finding about the work rather than about this survey, so 9 of 62 code-releasing rows, which is 15 percent, is the figure this section and Table IV use. PhysHOI [3] is the clearest of the 9. It lists a non-zero object-orientation weight for GRAB in its Table 4, and the reward function in the

112 method rows in the corpus

96 of them whose note settles the question

11 of those 96 address interpenetration in any form

4 of those 11 inside a closed-loop policy



report a penetration number for the rollouts of their own trained policy

Fig. 15: Interpenetration handling among all 112 method rows, the settled 96 shown against the light track behind them, split by closed-loop policy versus reference-only scoring.

repository at the commit this survey fetched sets that orientation error to zero, so in that file the object is tracked in position only.

B. Physical plausibility

The quantity most specific to a hand is the one closed-loop policies do not record, and Fig. 15 is the funnel that narrows to zero. Eleven method rows handle interpenetration in any form: three penalise it, three measure it, five constrain it, 11 percent of the settled rows. The denominator is 96, not 112, because the contact-handling field is null for 16 rows, and a null there means the note did not settle the question, not that the paper ignored penetration.

That eleven is not a claim that penetration goes unmeasured in general, and reading it that way would be wrong. Outside closed-loop control the quantity is a standard comparative column, and has been one for years in grasp synthesis and in hand-object reconstruction. Four rows of this corpus show the practice. BiDexGrasp [5] prints a penetration depth beside a prior method’s, BimanGrasp [6] fails any grasp that exceeds 1.5 mm of total penetration, TopoRetarget [7] reports a maximum penetration and a share of frames past 2 mm against a baseline retargeter, and OakInk [4] scores a dataset split on penetration depth, solid intersection volume and simulation displacement. The finding is narrower than the field and concerns learned closed-loop control: all eleven measure at the reference rather than at the rollout, seven of them scoring a pose, a trajectory or a contact model before anything executes and the four that are closed-loop policies scoring the references they were given, and we found none that reports the measurement for rollouts of its own trained policy.

The null result was audited. Earlier audits showed that the extraction could miss reported values, so 25 of the 85 apparently silent methods were drawn with a fixed random seed and checked again against their papers and repositories. None reports penetration depth, intersection volume, or another quantitative plausibility measure on its own rollouts. With zero recoveries, a one-sided 95 percent finite-population bound allows at most 8 hidden cases among the 85. The reported count of eleven is therefore a lower bound and nineteen an upper bound.

GraspXL [65] is the closest borderline case: human raters considered interpenetration as one component of a broader realism score. We treat that as a judgement rather than a penetration measurement. Other near misses concern self-collision, trajectory constraints, or a solver’s separation setting rather than measured hand–object overlap during a rollout. The sample, search terms, and result for every audited row are included with the corpus; Appendix A gives the full procedure.

Where in the pipeline those eleven act is the reference-versus-rollout split that Section I takes from [16]. A reference is a pose or a trajectory scored before execution, and a rollout is what the trained policy actually did. What this section supplies on that axis is a measurement method and a count, and it does not supply a threshold. The count is the eleven of 96 above, with four closed-loop policies inside it and none we found reporting a number for its own rollouts, which is Fig. 15. The method is the plausibility row of Table IV: maximum and mean penetration depth over the evaluation rollouts, on a dense surface sample, computed by code that never entered the reward or the termination rule. The threshold is borrowed, and Section VII-D says from where and why it does not bind. Six of the eleven are grasp synthesisers or trajectory optimisers, namely [5], [6], [165], [55], [7] and [167], and Castro et al. [108] give a contact model rather than a controller. That leaves four closed-loop policies in the whole corpus: [187], [39], [114] and [130].

The reason the number is four is a measurement trap. A quantity a policy optimises cannot also judge it, because the policy learns the measure rather than the property the measure stands for. PhysHOI [3] documents the failure in the clean direction. Its kinematic imitation reward was maximised by not touching the object at all, because contact perturbed the reference trajectory: “the policy may learn not to touch the object and falls into a local optimal”. A contact-graph reward closed the hole, and that reward reads a force threshold rather than geometry, so it constrains contact presence and says nothing about penetration depth.

DexTrack [114] shows the trap in the other direction. It defines a maximum hand-object penetration depth over all frames in its Appendix B, and applies it only to the input kinematic references, as one component of a reference-quality score. No penetration number appears for any of its own rollouts, in simulation or on the LEAP hand. Tolerance of the failure is then reported as a result: “Despite severe hand-object penetrations in Figure 4c and Figure 4a, the hand still interacts effectively with the object, highlighting the resilience of our tracking controller”.

TopoRetarget [7] is the strongest case in the corpus and still stops one step short on the same reference-versus-rollout line. It constrains penetration during retargeting with a 1 mm soft tolerance and a 30 mm hard bound, and it reports two numbers on 25 ContactPose grasps: a maximum penetration of 1.07 mm and 0.00 percent of frames above 2 mm, against 22.22 mm and 96 percent of frames for its GeoRT baseline. Then a PPO controller tracks those references, and its four reward terms and its five termination criteria govern object pose, link position, joint error and action smoothness, never penetration. The constrained quantity is the reference, and the rollout is

not re-measured.

Definitions are not shared either. GRAB [29] estimates contact by proximity, because “contact cannot be directly observed”, with a 4.5 mm tolerance, and reports that “‘Use’ grasps have 3.25 ± 0.68 mm average penetration”, without saying whether 3.25 mm is a maximum, a mean or a median. OakInk [4] supplies the fullest published vocabulary: penetration depth, solid intersection volume and simulation displacement. Its Table 3 scores the GRAB GrabNet split at 2.53 cm penetration depth, against GRAB’s own 3.25 mm. The two differ by a factor of about eight, and neither source states its distance function precisely enough to reconcile them, and the two are not scored over the same grasps either, since GRAB’s figure is over its own captured “use” grasps and OakInk’s is a model’s output on the GrabNet split, so a difference of population and a difference of definition are confounded in the same ratio.

The analytic tradition scored a grasp without simulating it, and Ferrari and Canny’s measure [54] and Roa and Suárez’s review [26] are its wrench-space reference points. Neither could be obtained. The first DOI fetch returned HTTP 202 and the second a Springer JavaScript interstitial, so both are cited by metadata only and no definition here rests on their contents. The learned literature has not replaced that tradition with anything it measures on its own rollouts. Penetration is a quantity between meshes, so it needs a simulator or a mesh reconstruction, and the protocol below treats it as a simulation-only axis.

Reporting penetration. The gap is specific to learned closed-loop control rather than general, and it is a choice rather than a capability. Evaluations should report maximum and mean penetration depth over the evaluation rollouts, on a dense surface sample, computed by a measure the policy never optimised.

C. Statistical practice, and a proposed protocol

Kress-Gazit et al. [184] give the field’s reference protocol, and it prescribes process, not numbers. Write the success criteria before the run and have someone other than their author score the runs. Match initial conditions with image overlays. Interleave the policies blind within one session. Report counts rather than percentages, alongside the initial conditions and the failure modes. Use a posterior over the Bernoulli parameter instead of a point estimate. It prescribes no minimum trial count anywhere and no frequentist confidence-interval width anywhere. Its own example report uses 10 initial conditions with two runs each, 20 evaluations per policy, and its worked case shows what 20 buys. Pancake success of 15 of 18 against 11 of 17, nominally 83 against 65 percent, leaves a 0.11 posterior probability that the worse policy is actually better. At 150 of 180 against 110 of 170 the same rates separate.

The TRI LBM Team [188] supplies the missing numbers by fiat rather than derivation: 50 rollouts per task per policy per condition on hardware, 200 in simulation, blind, with randomised policy order inside per-initial-condition bundles. It replaces confidence intervals with Beta-posterior violins and corrects its pairwise tests under Bonferroni, “unless otherwise noted”. Its own warning is the strongest sentence in this

literature: “there is significant risk that many robotics papers are measuring statistical noise due to insufficient statistical power”.

The rest fall short of their own advice. RoboArena [189] runs 612 double-blind pairwise comparisons across seven institutions and 4284 rollouts, and reports no confidence intervals and no per-policy trial counts. The Colosseum [182] evaluates 235 test sets at 25 episodes each with “one training seed and one evaluation seed” per baseline and gives no intervals in simulation. AutoEval [185] runs 50 trials per policy per task and calls ± 10 percent “the natural variance of robot evaluations”, without giving the formula behind the intervals it plots.

Two papers do give usable numbers. Badithela et al. [186] pair real and simulated trials and de-biases the simulation with a rectifier, saving 20 to 25 percent of hardware trials at paired correlations of roughly 0.6 to 0.7, and nothing at all at a correlation near zero. The 25 percent figure is the better of its two reported settings, the DP case at $\rho = 0.702$. Its decision rule is exact: combining helps only when the rectifier variance is below the variance of the real evaluations. Snyder et al. [183] give the largest saving. On the LBM rubrics their sequential test on graded scores cuts simulated evaluation by about 70 percent and hardware by about 45 percent, 286 trials against a nominal 500, with per-task decisions landing in 12 to 36 paired trials. On RoboArena’s data a 30-point gap on continuous progress scores reaches significance in 18 trials, while a 20-point gap on binary success needs about 80.

A proposed protocol. Every count below is printed by `tools/make_eval_tables.py` in its derivation mode, and each axis is derived for the statistic that axis actually reports: a single rate takes a Wilson half-width, a matched comparison takes McNemar, a ratio takes the standard error of the log ratio, a correlation takes the Fisher-z interval.

The reproducibility materials give the full derivation; its conclusions are as follows, and every count in Table IV is one of these. A single reported rate takes 100 real rollouts per cell, because 93 is where the worst-case Wilson half-width of (1) falls to 10 points and 10 points is the coarsest width at which a rate is worth printing. A comparison is a different question and a harder one. The design is paired, so the count follows McNemar and depends on the discordance rate rather than on the two rates alone, and separating 50 from 70 percent at 80 percent power takes 57 matched pairs per arm at the assumed discordance of 0.3, which is 114 rollouts against the 186 two independent arms would need. Simulated cells take 200 episodes for a 6.9-point half-width. Perturbation axes are screened at 40 per axis, which ranks the axes and cannot establish that any one of them hurt, and certified at 101 per arm. Unseen objects resample the object rather than the trial, at 20 objects and 5 trials each, and a 10-point claim about an object distribution needs about 93 objects. The transfer axis takes the 100 real trials paired to 100 simulated, and about 457 pairs to separate two correlations 0.1 apart. One hundred is a cap and not a bill, because a sequential test on a graded score reached its decision in 12 to 36 paired trials in [183], and a cell that stops early reports a confidence sequence rather than a Wilson interval, never both. The appendix also keeps the counts an earlier draft of this section quoted and this one

withdrew, since the test they were computed for was the wrong one.

D. Feasibility and limitations

The proposed protocol has not been run. In the twelve-method comparison assembled for this survey, no method supplied all the measurements needed for a complete protocol row. DeXtreme [35], for example, reports a success criterion and a confidence interval, but uses ten trials. This limits retrospective comparison; it does not imply that the reported experiments failed their own stated objectives.

What would have to be true. The protocol is expensive. For each policy and task it calls for 57 matched anchor trials and 100 unseen-object trials. Comparing two policies on three tasks therefore requires 942 real rollouts, or about 16 hours at one minute per rollout before repairs and scoring. Pairing saves roughly a fifth relative to an independent design because each pair still requires two rollouts. AutoEval [185] completed about 850 WidowX episodes in 24 hours with three interventions and regular cooling pauses; a tendon-driven hand is likely to stretch the same workload across several days.

That bill is large but not unprecedented. AutoEval [185] records that evaluating OpenVLA against its baselines took more than 2500 rollouts and more than 100 hours of human labour across three institutions, and the TRI LBM Team [188] analysed about 1800 real rollouts across nine hardware stations. What is unprecedented is paying it for a single dexterous-hand paper, where the modal per-cell count is 10 trials and the median per-cell count 15, and where the median stated count has been 20 or 22.5 in each of the last three years.

Four things would have to change. Reviewers would have to reward 57 matched trials on one task over 20 unmatched trials on five, and nothing in the corpus suggests that is happening. The loop would have to be automated, and AutoEval [185] shows it is buildable for a gripper at one to three hours of setup, while stating that it supports binary success only and no robustness axes. Somebody would have to run the penetration measure on their own rollouts, which is a choice rather than a capability: Section IV-B shows the depth is computable from the poses and the meshes in a few lines of Warp, and IsaacGymEnvs already ships a task that does it every step. An earlier draft of this section made the engine the barrier, and that claim is withdrawn, because the released code refutes it. And the comparison would have to be sequential, because the savings in [183] are the only reason 100 is a cap rather than a cost.

The proposal has four limits. It has not been run, and its sample sizes come from power and interval calculations rather than dexterous-hand experiments. The worst-case calculation assumes a success rate of 0.5, while the paired design assumes a discordance rate of 0.3 that the literature does not report. Its perturbations come from parallel-jaw benchmarks [182], [190], where contact is simpler. Finally, the 2 mm penetration threshold follows [7]; it is a simulator convention, not a physical limit. Captured human grasps in GRAB [29] average 3.25 mm under the same kind of measurement. Establishing a meaningful threshold will require measurements across hands, engines, and accepted behaviours.

The methodology literature has no dexterous hand in it. Seven corpus rows are evaluation protocols, a small denominator, and not one of them uses a dexterous hand: Badithela et al. [186] run a parallel-jaw gripper and the other six state no hand. Everything proposed above is therefore assembled from work on grippers and on whole-arm tasks. A modest first test would run the protocol in Table IV once, on one in-hand reorientation task with one 16-DoF hand, and release the rollouts and per-trial outcomes for independent analysis.

VIII. CONCLUSION

Progress in dexterous manipulation now depends as much on verification as on new methods. The evidence in this survey supports seven conclusions.

- 1) Nine of the 62 method rows that released parseable code state, in a named file at a named commit, something other than the value their paper prints, and seven further rows have been withdrawn from that count since its first draft, with two more narrowed. Table III gives each of the nine as a repository, a commit, a file and two values, and all of their authors were written to before this was posted: Section V-F says that beside the finding, with the letters in `outreach/` and the route by which a disputed case is corrected (Section V-F).
- 2) No closed-loop policy in the corpus reports interpenetration for the rollouts of its own trained policy, and every one of the eleven rows that handles penetration at all sits on the reference side of the reference-versus-rollout split (Section VII-B, Fig. 15).
- 3) The evaluation-methodology literature this survey builds its protocol on contains no dexterous hand at all: of its seven corpus rows, one runs a parallel-jaw gripper and the other six state no hand (Section VII-D, Table IV).
- 4) The generalist and vision-language-action policies that do evaluate on a multi-fingered hand run it at a median of 6 actuated degrees of freedom, against 16 across the reinforcement-learning rows that state a count, and the rows reaching the larger band on paper never say whether the hand was in the evaluation (Section V-D).
- 5) Eight hands that can be bought today or built from published designs take zero method rows between them, while the corpus's experiments concentrate on four designs (Section III-D, Fig. 3).
- 6) Twenty-one of the 28 rows that put a learned closed-loop controller on two multi-fingered hands run one policy over a concatenated two-hand observation, a choice compared twice with opposite outcomes and ablated once (Section VI-B, Fig. 10).
- 7) Human data does not port across hands, and the map is usually unstated: 33 of the 53 method rows that use human data never say how the human motion reached the robot hand (Section V-C).

Six concern gaps in the literature. The first concerns publishing practice; the corrections made during this review are recorded beside the result in Section V-F.

Three of those claims carry a case worth remembering. A reward table in a paper is a claim about a document, not about a

TABLE IV: The proposed evaluation protocol, and this survey’s own contribution. No cell is empty: every cell is a proposal.

axis	what is measured	minimum trials	reported with it
Task success	fraction of episodes meeting a criterion written before the run, and a graded rubric score in $[0, 1]$	100 real and 200 sim per policy, task and condition for an absolute rate at a 10-point interval; a comparison alone needs 57 matched pairs per arm	raw counts, the criterion verbatim, the rubric, the initial-condition protocol, and an interval whose kind is named
Robustness	success under each perturbation axis, and the ratio to the unperturbed anchor	40 per axis per policy to rank the axes; 101 per axis and 101 at the anchor before any claim that a named axis hurt	both absolute rates with their Wilson intervals, the axis ranges, and the ratio without an interval unless the certifying count was run
Unseen objects	mean over held-out objects of per-object success	20 objects at 5 trials to screen; 93 objects for a 10-point claim	the object list, the per-object rates, a bootstrap interval over objects, and the within-object trial count
Physical plausibility	maximum and mean hand-object penetration depth per rollout, and the fraction of frames above 2 mm	100 rollouts, simulation only	an interval over rollouts, the sample density, the threshold, the solver’s depenetration settings, and at least one rendered rollout
Sample and wall-clock cost	environment steps and wall-clock to the checkpoint that produced the headline number	3 training seeds	the range over seeds, and the statement that 3 seeds is a range and not an interval
Real-robot transfer	paired real and simulated outcomes on matched initial conditions, their correlation, and the rectifier variance against the real variance	the 100 real trials paired to 100 sim; about 457 pairs to separate two correlations 0.1 apart	the correlation with its Fisher-z interval, the variance ratio with its bootstrap interval, and which decision the count supports
Reproducibility	the released config and checkpoints against the paper’s stated objective	not a trial count	the file and line of every disagreement, or an explicit statement that none was found

run, and PhysHOI [3] is the instance to keep in mind: the term its table weights at 0.1 is set to zero in the file at the commit this survey fetched, and the success criterion the paper reports is itself position-only. Contact is what separates a hand from a gripper, and the obstacle to measuring it is not the engines. IsaacGymEnvs ships a task that computes a per-environment maximum interpenetration depth against meshes and gates the policy update on a 1 mm threshold, and Tactile Genesis [113] offers penetration depth on Genesis geometry as a sensor. The tooling sits in the field’s own benchmark repository and the number is still not reported. DexTrack [114] has the formula and points it at its inputs. And hardware and software have come apart on a narrower claim than the hand count first suggests, because 11 of the 19 hands that take no method row are neither sold nor open and take none for that reason, which is why the fifth claim is eight hands and not nineteen.

What this survey cannot establish is which method is better than which. It re-runs nothing, and Section VII argues that the published numbers do not compare. Six works are cited by metadata only, and a seventh, Ma and Dollar 2011, is on disk but unread; no claim rests on any of them. Every coverage statistic here counts what this survey’s extraction captured rather than what the literature reported. Each is a floor and not a rate, because every miss converts a reporting paper into a silent one.

Several practical changes require little additional computation. Authors can generate reward tables from the configurations used for the reported run, cite the exact revision, state the trial count and success criterion, and release per-trial outcomes. Evaluations can report penetration depth over policy rollouts as a diagnostic, without optimizing the policy against that measure. TopoRetarget [7] shows how informative the result can be. Hardware choices also deserve a broader base: the Allegro, Shadow, Inspire, and LEAP hands account for most of the published use in this corpus, while every other listed hand

appears in eight method papers or fewer. Better reporting would make results easier to reproduce; broader hardware coverage would show which results survive a change of embodiment.

APPENDIX A METHOD

The survey is written from a corpus that lives on disk. No claim about a paper is made from memory. Each claim traces to a note, each note to a parsed markdown file, and each markdown file to a PDF or repository with a recorded hash or commit.

Selection. Six topic-specific bibliographies were assembled in parallel (simulators and physics, hands and vendors, reinforcement learning, imitation and human data, bimanual, benchmarks and evaluation), seeded with the canonical works in each area and extended by search up to 2026-09-17. Every arXiv identifier was checked by fetching the abstract page and matching the title, and entries that could not be checked that way are marked. The six lists were merged with deduplication on arXiv identifier and normalised title, giving 221 entries.

Seven further bibliography entries are not part of that corpus and are not counted anywhere in this survey. They are the prior audits and case studies Section I positions this survey against, they sit in a bibliography of their own as prior work from outside the corpus, and none of them carries a structured row. No source for them was parsed either, so each is quoted only from its abstract and its stated method.

Acquisition. PDFs were downloaded from arXiv or, for work without a preprint, from the publisher or vendor page recorded in the bibliography. Vendor pages for hands without any paper were fetched as HTML and converted to markdown, so that a specification quoted in this survey is quoted from a stored copy of the page and dated. Repositories were shallow-cloned and parsed, then deleted. What remains is one markdown per

repository holding the README, a pruned file tree, the task and reward configuration files, and the bodies of reward and observation functions.

Parsing. PDFs were converted with `pymupdf4llm`, falling back to raw text extraction when the layout parse returned too little. The corpus manifest records, per entry, the source URL, page count, sha256 of the PDF and the converter used. Two consequences matter for reading this survey. Equations rendered as images do not survive conversion, so where a reward weight exists only inside a figure it is recorded as unreadable rather than guessed. Tables are flattened, so a number taken from a table is quoted with the table it came from and, where the flattening is ambiguous, the ambiguity is stated.

Notes. Each entry was read into a structured note under a fixed template: embodiment, learning method, objective, contact handling, evaluation, reproducibility, stated limitations, and quotable claims. Notes were written only from the parsed files. Where the source is silent the note says “not stated”, and nothing was inferred from the reviewer’s prior knowledge of the work. For method papers the reward or loss was quoted from the paper and, separately, from the released code, so that disagreements between the two are visible rather than smoothed over. Those disagreements turned out to be common enough to become a finding in their own right. The reading into notes was done with language-model assistance, from the parsed files only. The template’s conventions are the safeguard against what that can get wrong: a quotation rather than a paraphrase, “not stated” rather than an inference, and a named source beside every value, so a note can be checked against the file it came from. Two fields carrying a headline claim were re-read against the sources afterwards, the trial counts and the penetration field, and both audits are below; the coverage statistics are reported as floors for the same reason.

Sources that could not be obtained. Six works are cited by metadata only and no claim in this survey rests on their contents. Five are behind publisher paywalls with no author-hosted copy found on 2026-09-18: Okamura et al. 2000, Bicchi 2000 (IEEE T-RO), Piazza et al. 2019, Roa and Suarez 2015, and Butterfass et al. 2001 on DLR-Hand II. The sixth is Hwangbo et al. 2018 on RaiSim, whose hosted PDF returned no body. The freely circulating PDF often taken for Bicchi 2000 is a different work, a book chapter, and is listed separately.

Ma and Dollar 2011 is a seventh case with a different cause. The fetch that failed when its note was written succeeded on 2026-09-18, so a seven-page PDF and its hash are in the manifest, but no note has been read from it. It is cited by metadata only for that reason and not because the source is unavailable.

Three of the 221 bibliography entries carry no structured row. Two are the paywalled Bicchi 2000 and DLR-Hand II entries. The third is [1], which was read into a note and quoted throughout but is a book chapter rather than a work with an embodiment, a method or a result to record in a row.

The reference list of this edition is shorter than the corpus, and the two numbers are different quantities. It prints 205 entries: the 198 corpus entries that some sentence, table or figure of this paper cites, plus the 7 prior-work entries from outside the corpus. The other 23 corpus entries carry a row and

a note and are counted in every statistic here, but no passage in the paper names them, so they have nothing to be cited from and do not appear in the list. A reader counting the reference list should get 205, and a reader counting the corpus should get 221.

What this method cannot do. Mention counts over the corpus are counts of mentions, not of use: a related-work sentence counts the same as an experiment. Vendor specifications are manufacturer claims and are labelled as such throughout, and where a page has since gone offline the note says so. The corpus is large but not exhaustive, and selection by search favours work that is indexed, in English, and posted as a preprint. It also favours recent work. Of the 221 bibliography entries, 138 are dated 2024 or later and 16 predate 2018, so this is a corpus of the learned era and any claim here about a trend over time is a claim about 2022 onward.

Every coverage statistic in this survey measures what this survey’s extraction captured, not what the literature reported. A structured row holds a scalar. A paper that reports a per-task count, a rubric, or a total spread across several tables produces a null, and a null is then counted as silence. The bias runs one way. Every miss converts a reporting paper into a silent one, so the field is made to look worse at reporting than it is. The size of the effect was measured on the statistic the survey leads with. Of the 34 method rows that had a real robot and no recorded trial count, 15 carried a count in plain text in their own note, dropped because the paper reports it per task and the field takes a single integer. Those 15 have since been re-extracted, which moved the stated-trial-count row from 55 to 70. The success criterion and the unseen-object count were audited the same way, rising from 79 to 98 and from 32 to 39. What those 19 recovered criteria say is mostly not a threshold: eleven score by rubric or staged partial credit, five judge binary completion against a task description by eye, two defer to a benchmark’s own definition, and exactly one, [44], states a verbatim numeric threshold. The audit counted only cases where the note itself carried the number, so a count the note also missed is still uncounted, and 19 trial counts, 8 unseen-object evaluations whose object count the note never gives, and 4 criteria remain genuinely unsettled. Fifteen of the recovered trial counts are named in the rows: π_0 [42] at ten trials per task, RDT-1B [168] at 139 across seven tasks, UMI [171] at 260, $\pi_{0.6}^*$ [44] at 750, Gemini Robotics [161] at twenty per task, and ten more, which is 44 percent of the audited nulls and moved the “never says” figure from 34 of 89 down to 19. Read every coverage statistic in this survey as a floor rather than as a rate, and read the bars in Fig. 13 the same way.

The penetration field, audited the same way. The same mechanism reaches the penetration, code-release and failure-mode fields. The penetration field has since been audited by hand, because this survey’s second finding is a null over it and a null is worth what the search behind it is worth. The population is the 85 method rows whose contact-handling field records that the work does not address penetration. It excludes the 16 nulls, which are already counted as unsettled rather than as silence, so recovering one would not move the finding. Twenty-five of the 85, 29 percent, were drawn with

`random.Random(20260919).sample` over the sorted population, so the draw is fixed and can be redrawn. Each sampled row was read again in `papers/md/`, in its OCR recovery where one exists, and in `code/md/` where a repository was parsed — and not in `papers/notes/`, because the note is the artefact under suspicion: a field is wrong exactly when the source says something the note did not carry. A regular expression over the whole source collected every occurrence of penetration, interpenetration, intersection, intersection volume, solid intersection, simulation displacement, contact consistency, physical plausibility, signed distance and contact depth, and of the words that are mistaken for them, and every hit was read in its context. A row counted as a recovery if its own source reported a measurement of any of those quantities on that method’s own rollouts. A statement that penetration occurs, a citation whose title says physically plausible, a collision-avoidance constraint on a reference trajectory and a solver setting did not count, and the three recurring near misses of that kind are listed in Section VII-B. Recoveries: none of the 25, against the 20 to 45 percent the three earlier audits recovered on the fields above. The one-sided 95 percent bound, computed hypergeometrically over the finite population, is therefore 8 of the 85, and Section VII-B states it beside the claim and names the one borderline case that a reader might count differently. Why this field held where the others did not is visible in the per-row results: the recovered trial counts and criteria were numbers present in the source and dropped by a field shaped to hold a scalar, whereas a penetration number is absent from the source altogether. The earlier audits measured a defect in this schema; this one looked for an absence in the literature. `reviews/penetration_audit.md` holds the seed, the sample and a result per row beside the sentence it rests on, and `tools/audit_penetration.py` redraws the sample and re-runs the sweep. The code-release and failure-mode fields have still not been audited this way, and their counts stay floors.

Where the tabulation lives. Most of this survey’s tabulation is not printed in this paper. The 112 method rows, the 21 reward-term objectives scored against their released code inside a 77-row extraction whose ten columns hold 253 cells no source stated, the 30 teleoperation systems, the 33 hand specifications with the source sentence behind each one, the 15 contact-model and solver descriptions verbatim, and the sentence behind every cell of the predecessor-survey comparison: all of it is reference material, consulted a row at a time and never read through. Printed, it ran to more pages than the argument it supports. It is in the repository instead.

Read it there in one of three forms. `corpus/rows/` holds one JSON record per work, with every field at full length and no cell cut; this is the form to query, and it is the form every number in this paper is computed from at build time. `paper/tables/` holds the same content as markdown tables, one file per table, for reading rather than querying. `papers/notes/` holds the structured note each row was read from, with the quotation that supports each field, so that a cell can be traced to a sentence in a parsed source. The generators are in `tools/`: every table in this paper and in the repository is written by a script there from those records, so a reader who disagrees

with a classification can change it and rebuild rather than argue with a PDF.

The corpus is not deposited yet, and this paper prints no link to it for that reason. It is available from the author at the address in the author block, and the archived release will carry a persistent identifier, which will be the copy to cite for any number quoted here. A survey that argues for checkable artefacts should point at its artefact rather than reprint it, and a table a reader cannot recompute is worth less than the record it was compressed from.

Three tables are exceptions and are printed, because the prose reads individual cells out of them and does not stand up without them: the two hand tables and the engine table, in Appendix B. The proposed evaluation protocol appears in the body alongside its justification. Table I is the comparison to the fourteen predecessor surveys, which is this survey’s positioning and is therefore set where a reader meets it rather than in the repository. It prints one of three levels per cell; the sentence supporting each level remains in the reproducibility materials.

Reproduced figures. One figure in this edition is not this survey’s own artwork. It is a render reproduced from the paper it discusses, rendered from the stored PDF in the corpus, and it states its source in its own caption, keyed to the corpus entry it came from and to the page it was taken from. An earlier draft reproduced eighteen. A plate is kept only where it carries something a drawing cannot, and only where this paper has the right to print it: a contact defect invisible at the scale of a grasp and unmistakable at the scale of the contact, on a source that is CC BY 4.0. The other seventeen were redrawn in the paper’s own style or dropped, and what each one was standing in for is said in the comment at the head of the figure that replaced it, or in the paragraph that absorbed it. Two of those seventeen went last and for a reason worth stating, because it is a limit on what this paper shows rather than a tidying of its figures. Both would have needed an author’s permission, since arXiv’s non-exclusive licence to distribute grants a third party nothing, and this paper is posted without asking anyone for anything. One of them argued from geometry and was redrawn, as Fig. 8. The other photographed five research hands to scale beside a human hand and a ruler, and it was dropped rather than redrawn: the honest redraw is a to-scale outline annotated in millimetres, and no source in the corpus states a hand’s length or its width, so every envelope in such a drawing would have been invented. Section III now states that absence instead of illustrating around it. The one surviving plate is cleared: its source carries CC BY 4.0, read directly from the licence page on the date the register records, and its caption prints the attribution the licence asks for, naming the creator, the work, the licence with a link and the change made to the plate. No permission request and no author email is outstanding, and nothing in this paper waits on a reply. The register that records, per figure, the licence found, the URL it was read from and the position it closed at, is `tex/PERMISSIONS.md`. The build does not enforce that register.

APPENDIX B

THE HARDWARE AND ENGINE TABLES

Sections III and IV argue over individual cells of three tables: which column is emptiest, which cell is blank and why, which figure rests on a press release. Those three are set here rather than in the body, because a reader consults them rather than reading them through. Every other tabulation in this survey is in the repository, and Appendix A says where.

Nothing in these tables is a measurement by this survey. Every specification is a manufacturer claim, a figure in a paper, or a figure in a press article. Table V and Table VI split the 33 hand rows by whether the hand can be bought or built, and Table VII carries the 15 engines. The controlled vocabularies in all three compress a sentence the corpus stores in full: *corpus/rows/* holds the actuation, tactile, force-kind, contact-model and solver text at its own length, with the source’s own punctuation and wording, and *papers/notes/* holds the quotation each was read from.

Each caption states how many of its own cells no source filled. Two conventions behind those counts are worth stating once. A blank price is a hand with no public list price rather than a free hand; and a hand with no entry in the manifest was recorded from a page that could not be stored, so its row should be read as a claim with no retrievable source.

Table V holds the hands that can be had: sold, open-source, or with no release status stated. A row whose joint count no source states sorts last rather than being ranked by a number nobody stated. Its empty-cell count covers the eight specification columns, DoF through price; the hand, maker, open-hardware, status and usage columns are left out of it because they are never blank and counting them would dilute the finding. *force N* is not a ranking: read *force kind* first, because pull-out resistance, pinch force, a fingertip normal force under an indenter and an unlabelled vendor number are different measurements. *uses* counts the method rows whose own experiments run on that hand.

Table VI carries the columns of Table V plus the date of the claim and the class of artefact it was read off, over the same eight specification columns. Its empty share is far higher, and that is the finding of the table rather than a defect of it: an unreleased hand is announced without a specification. *evidence* is *datasheet* for a paper, *datasheet* or *vendor specification* page, *video* for a demonstration reel, *press* for a press release or a third-party tracker. *claim date* prints *undated* where the source page carries no date of its own; a fetch date is not a claim date.

Table VII prints, for each engine, the contact formulation it names, whether it exposes interpenetration depth and how many first-party hand models it ships; its empty-cell count is over the seven property columns. *contact model* and *solver* are controlled vocabularies: *soft* for a compliant or regularised contact, *convex relaxation* for a convexified formulation, *complementarity* for an LCP, *hard NCP* for a nonlinear complementarity problem, *impulse* for a velocity-level update, and *rigid, model unstated* where the source describes collision geometry but names no contact law. The sentence each was read from runs to several hundred characters; *corpus/rows/*

carries each one verbatim, because Section IV’s argument turns on its wording. *penetration exposed* is whether the engine makes interpenetration depth readable by a user’s code without patching it. *hands* counts first-party hand models in the parsed snapshot, and they are: *Brax* synthetic 4-fingered claw hand; *Genesis* Shadow Hand; *Isaac Gym* Shadow Dexterous Hand, *Allegro* Hand, *TriFinger*; *Isaac Lab* KUKA *Allegro* hand, *ShadowHand*, *AbilityHand*, *Fourier* GR1-T2 dex hands, *Unitree* *Inspire* hand, *Unitree* trihand; *ManiSkill3* *Allegro* Hand, *Ability* Hand, *Inspire* Hand, *Delto* 3-finger hand, *TriFinger*; *Newton* *Allegro* Hand; *Orbit* *Allegro* hand.

REFERENCES

- [1] Bicchi, “Robotic Grasping and Manipulation (book chapter),” in *RAMSETE: Articulated and Mobile Robotics*, Springer LNCIS. Springer, 2001.
- [2] An *et al.*, “Dexterous Manipulation through Imitation Learning: A Survey,” arXiv:2504.03515, 2025.
- [3] Wang *et al.*, “PhysHOI: Physics-Based Imitation of Dynamic Human-Object Interaction,” arXiv:2312.04393, 2023.
- [4] Yang *et al.*, “OakLink: A Large-scale Knowledge Repository for Understanding Hand-Object Interaction,” in *CVPR*, 2022.
- [5] Lin *et al.*, “BiDexGrasp: Coordinated Bimanual Dexterous Grasps across Object Geometries and Sizes,” arXiv:2604.06589, 2026.
- [6] Shao *et al.*, “Bimanual Grasp Synthesis for Dexterous Robot Hands,” *IEEE RA-L*, 2024.
- [7] Wu *et al.*, “TopoRetarget: Interaction-Preserving Retargeting for Dexterous Manipulation,” arXiv:2606.16272, 2026.
- [8] Collberg and Proebsting, “Repeatability in Computer Systems Research,” *Communications of the ACM*, vol. 59, no. 3, 2016, related work outside this survey’s corpus; no source parsed, no structured row, and no count rests on it.
- [9] Xu *et al.*, “Do Papers Tell the Whole Story? A Benchmark and Framework for Uncovering Hidden Implementation Gaps in Bioinformatics,” arXiv:2603.22018, 2026, related work outside this survey’s corpus; no source parsed, no structured row, and no count rests on it.
- [10] Baumgärtner and Gurevych, “SciCoQA: Quality Assurance for Scientific Paper-Code Alignment,” arXiv:2601.12910, 2026, related work outside this survey’s corpus; no source parsed, no structured row, and no count rests on it.
- [11] Engstrom, Ilyas, Santurkar, Tsipras, Janoos, Rudolph, and Madry, “Implementation Matters in Deep Policy Gradients: A Case Study on PPO and TRPO,” in *ICLR*, 2020, related work outside this survey’s corpus; no source parsed, no structured row, and no count rests on it.
- [12] McLean *et al.*, “Meta-World+: An Improved, Standardized, RL Benchmark,” in *NeurIPS Datasets and Benchmarks*, 2025, related work outside this survey’s corpus; no source parsed, no structured row, and no count rests on it.
- [13] Knox, Allievi, Banzhaf, Schmitt, and Stone, “Reward (Mis)design for Autonomous Driving,” *Artificial Intelligence*, vol. 316, 2023, related work outside this survey’s corpus; no source parsed, no structured row, and no count rests on it.
- [14] Raff, “A Step Toward Quantifying Independently Reproducible Machine Learning Research,” in *NeurIPS*, 2019, related work outside this survey’s corpus; no source parsed, no structured row, and no count rests on it.
- [15] Qi *et al.*, “In-Hand Object Rotation via Rapid Motor Adaptation,” in *CoRL*, 2022.
- [16] Zhao *et al.*, “Towards Robotic Dexterous Hand Intelligence: A Survey,” arXiv:2605.13925, 2026.
- [17] Bai *et al.*, “Towards a Unified Understanding of Robot Manipulation: A Comprehensive Survey,” arXiv:2510.10903, 2025.
- [18] Welte *et al.*, “Interactive Imitation Learning for Dexterous Robotic Manipulation: Challenges and Perspectives – A Survey,” *Frontiers in Robotics and AI*, 2025.
- [19] Firoozi *et al.*, “Foundation Models in Robotics: Applications, Challenges, and the Future,” *IJRR*, 2023.
- [20] Zhao *et al.*, “Sim-to-Real Transfer in Deep Reinforcement Learning for Robotics: a Survey,” in *IEEE SSCI*, 2020.
- [21] Gao *et al.*, “NVIDIA Isaac Sim: Enabling Scalable, GPU-Accelerated Simulation for Robotics,” arXiv:2606.03551, 2026.
- [22] Kaup *et al.*, “A Review of Nine Physics Engines for Reinforcement Learning Research,” arXiv:2407.08590, 2024.

TABLE V: Hands that can be obtained, sorted by joint count. 68 of 160 specification cells are empty.

hand	maker	DoF	act. DoF	actuation	mass g	force N	force kind	tactile	price USD	open HW	status	uses
Shadow Dexterous Hand [69]	Shadow Robot Co.	24	20	tendon	4 300	–	–	optional	–	no	sold	21
Proception ProHand 1.0 [103]	Proception Inc	22	–	tendon	–	–	–	unstated	–	no	sold	0
SharpaWave 22-DoF [97]	Sharpa Robotics	22	22	–	1 300	20 / 12	grip	unstated	–	no	sold	6
BiDexHand [78]	Northwestern Univ.	21	16	tendon	–	2.14	fingertip normal	none	–	yes	open-source	0
Ruka-v2 [76]	NYU and NYU Shanghai	20	16	tendon	–	–	–	optional	1 500	yes	open-source	0
TESOLLO DG-5F [102]	Tesollo Inc.	20	20	direct drive	1 763	–	–	none	–	no	sold	0
Unitree Dex5-1 [101]	Unitree Robotics	20	16	tendon	1 100	10	fingertip normal	array	–	no	sold	0
Wuji Hand [98]	Wuji Tech.	20	20	direct drive	–	–	–	fingertip	–	no	sold	2
ORCA [79]	ETH Zurich	17	17	tendon	1 200	–	–	fingertip	–	yes	open-source	0
Allegro Hand V4 [80]	Wonik Robotics Co. Ltd.; earlier versions...	16	16	–	–	–	–	none	–	no	sold	35
LEAP Hand [77]	CMU	16	16	direct drive	595	19.5	pull-out	none	2 000	yes	open-source	12
RUKA [81]	NYU	15	11	tendon	–	2.74	pinch	none	1 300	yes	open-source	0
Inspire Robots RH56DFX [70]	Beijing Inspire Robots Tech. Co., Ltd.	12	6	–	540	10	vendor, no protocol	none	–	no	sold	19
ROBOTERA XHAND1 [191]	ROBOTERA	12	12	geared	1 100	–	–	unstated	14 000	no	sold	8
DEX-EE [192]	Shadow Robot Co., in collaboration with...	12	–	–	4 100	–	–	unstated	–	no	sold	1
BrainCo Revo 2 [193]	BrainCo Inc.	11	6	–	383	15	pinch	unstated	–	no	sold	1
AgiBot OmniHand 2025 [194]	–	–	–	–	–	–	–	–	–	–	unstated	1
DexHand [71]	IoT Design Shop	–	–	tendon	–	–	–	none	300	yes	open-source	0
LinkerBot Linker Hand [195]	LinkerBot	–	–	–	–	–	–	unstated	–	no	sold	1
PSYONIC Ability Hand [196]	PSYONIC	–	6	linkage	–	–	–	unstated	–	no	sold	7

TABLE VI: Hands announced, prototyped or held internally, which cannot be obtained. 60 of 104 specification cells are empty.

hand	maker	DoF	act. DoF	actuation	mass g	force N	force kind	tactile	price USD	open HW	status	uses	claim date	evidence
Clone Robotics [95]	Clone Robotics	27	–	fluidic muscle	<910	–	–	array	–	no	internal-only	0	2025-11-24	video
1X NEO [93]	1X Tech.	25	25	tendon	–	45	vendor, no protocol	unstated	–	no	announced	0	2026-07-09	datasheet
LEAP Hand V2 [197]	CMU	21	17	tendon	–	–	–	none	3 000	yes	announced	2	undated	datasheet
Sanctuary AI Phoenix [198]	Sanctuary AI	21	–	fluidic muscle	–	–	–	unstated	–	no	internal-only	0	2024-12-17	press
ILDA hand [74]	Korea Univ.	20	15	linkage	1 100	34	fingertip normal	array	–	no	prototype	0	2021-11-05	datasheet
Pisa/IIT SoftHand [73]	Centro E. Piaggio, Univ. Pisa and IIT	19	1	tendon	–	–	–	none	–	no	prototype	0	2014	datasheet
Faive Hand [75]	ETH Zurich	16	11	tendon	1 100	–	–	none	–	–	prototype	1	2023	datasheet
Boston Dynamics [199]	Boston Dynamics	–	–	geared	–	–	–	unstated	–	no	prototype	0	2026-01-05	press
Daxo Robotics Muscle [72]	Daxo Robotics	–	–	tendon	750	–	–	–	–	–	prototype	0	2026-09-16	press
Figure 03 [96]	Figure AI	–	–	–	–	–	–	unstated	–	no	prototype	0	undated	press
PaXini DexH13 [200]	PaXini Tech	–	–	–	–	–	–	array	–	–	announced	0	2026-01-07	press
Tesla Optimus [94]	Tesla	–	–	tendon	–	–	–	unstated	–	no	prototype	0	2023-12-13	video
Xiaomi [201]	Xiaomi	–	–	tendon	–	–	–	unstated	–	no	prototype	0	2026-03-30	press

TABLE VII: 15 simulators and physics engines, by contact formulation and penetration exposure. 40 of 105 property cells are empty.

engine	contact model	solver	diff.	GPU	dt (s)	penetration exposed	hands
Brax [202]	convex relaxation	explicit/spring	yes	yes	–	no	1
ComFree-Sim [112]	complementarity-free	closed form	–	yes	0.002	yes	–
Dojo [111]	hard NCP	interior point	yes	no	–	yes	–
Genesis [117]	soft	Newton	yes	yes	–	–	1
Isaac Gym [107]	convex relaxation	PGS	no	yes	0.008333	no	3
Isaac Lab [115]	rigid, model unstated	Gauss-Seidel	no	yes	–	no	4
ManiSkill3 [120]	convex relaxation	PGS	–	yes	–	–	5
MuJoCo [110]	complementarity	Newton	no	no	–	–	–
MuJoCo convex [109]	complementarity-free	closed form	no	no	0.01	–	–
MuJoCo Warp [116]	soft	PGS	no	yes	–	–	–
Newton [118]	soft	primal-dual	–	yes	–	–	1
Orbit [119]	convex relaxation	PGS	no	yes	–	no	1
PyBullet [203]	–	–	–	yes	–	–	–
RaiSim [204]	–	–	–	–	–	–	–
SAPIEN [205]	convex relaxation	PGS	–	no	–	–	–

- [23] Le Lidec *et al.*, “Contact Models in Robotics: a Comparative Analysis,” *IEEE T-RO* 2024, 2023.
- [24] Erez *et al.*, “Simulation tools for model-based robotics: Comparison of Bullet, Havok, MuJoCo, ODE and PhysX,” in *ICRA*, 2015.
- [25] Piazza *et al.*, “A Century of Robotic Hands,” *Annual Review of Control, Robotics, and Autonomous Systems*, 2019, full text not obtained; cited by metadata only.
- [26] Roa and Suárez, “Grasp quality measures: review and performance,” *Autonomous Robots*, 2015, full text not obtained; cited by metadata only.
- [27] Ma and Dollar, “On dexterity and dexterous manipulation,” in *ICAR*, 2011.
- [28] Okamura, Smith, and Cutkosky, “An overview of dexterous manipulation,” in *ICRA*, 2000, full text not obtained; cited by metadata only.
- [29] Taheri *et al.*, “GRAB: A Dataset of Whole-Body Human Grasping of Objects,” in *ECCV*, 2020.
- [30] Zhong *et al.*, “DexGraspVLA: A Vision-Language-Action Framework Towards General Dexterous Grasping,” arXiv:2502.20900, 2025.
- [31] Luo *et al.*, “Omnigrasp: Grasping Diverse Objects with Simulated Humanoids,” in *NeurIPS*, 2024.
- [32] Wan *et al.*, “UniDexGrasp++: Improving Dexterous Grasping Policy Learning via Geometry-aware Curriculum and Iterative Generalist-Specialist Learning,” in *ICCV*, 2023.
- [33] Agarwal *et al.*, “Dexterous Functional Grasping,” in *CoRL*, 2023.
- [34] OpenAI *et al.*, “Learning Dexterous In-Hand Manipulation,” *IJRR* 2020, 2018.
- [35] Handa *et al.*, “DeXTreme: Transfer of Agile In-hand Manipulation from Simulation to Reality,” in *ICRA* 2023, 2022.
- [36] Ma *et al.*, “Eureka: Human-Level Reward Design via Coding Large Language Models,” in *ICLR* 2024, 2023.
- [37] Chen *et al.*, “Visual Dexterity: In-Hand Reorientation of Novel and Complex Object Shapes,” *Science Robotics* 2023, 2022.
- [38] Li *et al.*, “ManipTrans: Efficient Dexterous Bimanual Manipulation Transfer via Residual Learning,” in *CVPR*, 2025.
- [39] Mandi *et al.*, “DexMachina: Functional Retargeting for Bimanual Dexterous Manipulation,” arXiv:2505.24853, 2025.
- [40] Huang *et al.*, “Dynamic Handover: Throw and Catch with Bimanual Hands,” in *CoRL*, 2023.
- [41] Frau-Alfaro *et al.*, “Learning Dexterous Object Handover,” in *RO-MAN*, 2025.
- [42] Black *et al.*, “pi0: A Vision-Language-Action Flow Model for General Robot Control,” arXiv:2410.24164, 2024.
- [43] Physical Intelligence, “pi0.5: a Vision-Language-Action Model with Open-World Generalization,” arXiv:2504.16054, 2025.
- [44] —, “pi*0.6: a VLA That Learns From Experience,” arXiv:2511.14759, 2025.
- [45] Kim *et al.*, “OpenVLA: An Open-Source Vision-Language-Action Model,” in *CoRL*, 2024.
- [46] Gemini Robotics Team, “Gemini Robotics 1.5: Pushing the Frontier of Generalist Robots with Advanced Embodied Reasoning, Thinking, and Motion Transfer,” arXiv:2510.03342, 2025.
- [47] Fu *et al.*, “HumanPlus: Humanoid Shadowing and Imitation from Humans,” in *CoRL*, 2024.
- [48] He *et al.*, “OmniH2O: Universal and Dexterous Human-to-Humanoid Whole-Body Teleoperation and Learning,” in *CoRL*, 2024.
- [49] NVIDIA, “GR00T N1.6: An Improved Open Foundation Model for Generalist Humanoid Robots,” https://research.nvidia.com/labs/gear/gr00t-n1_6/, 2025, company press release, December 2025, research.nvidia.com, fetched 2026-09-17.
- [50] Zakka *et al.*, “RoboPianist: Dexterous Piano Playing with Deep Reinforcement Learning,” in *CoRL*, 2023.
- [51] Zhao *et al.*, “RPIM: A Large-Scale Motion Dataset for Piano Playing with Bi-Manual Dexterous Robot Hands,” in *CoRL*, 2024.
- [52] Qian *et al.*, “PianoMime: Learning a Generalist, Dexterous Piano Player from Internet Demonstrations,” in *CoRL*, 2024.
- [53] Li *et al.*, “DexDeform: Dexterous Deformable Object Manipulation with Human Demonstrations and Differentiable Physics,” in *ICLR*, 2023.
- [54] Ferrari and Canny, “Planning optimal grasps,” in *ICRA*, 1992.
- [55] Pang *et al.*, “Global Planning for Contact-Rich Manipulation via Local Smoothing of Quasi-dynamic Contact Models,” *IEEE T-RO* 2023, 2022.
- [56] Qin *et al.*, “DexPoint: Generalizable Point Cloud Reinforcement Learning for Sim-to-Real Dexterous Manipulation,” in *CoRL*, 2022.
- [57] Yin *et al.*, “Rotating without Seeing: Towards In-hand Dexterity through Touch,” in *RSS*, 2023.
- [58] Yang *et al.*, “AnyRotate: Gravity-Invariant In-Hand Object Rotation with Sim-to-Real Touch,” in *CoRL*, 2024.
- [59] —, “AsymDex: Asymmetry and Relative Coordinates for RL-based Bimanual Dexterity,” in *RSS 2025 WCBM workshop (spotlight)*, 2024.
- [60] Zhou *et al.*, “DyDexHandover: Human-like Bimanual Dynamic Dexterous Handover using RGB-only Perception,” arXiv:2509.17350, 2025.
- [61] Ding *et al.*, “Bunny-VisionPro: Real-Time Bimanual Dexterous Teleoperation for Imitation Learning,” arXiv:2407.03162, 2024.
- [62] Shaw *et al.*, “Bimanual Dexterity for Complex Tasks,” in *CoRL*, 2024.
- [63] Jiang *et al.*, “DexMimicGen: Automated Data Generation for Bimanual Dexterous Manipulation via Imitation Learning,” in *ICRA* 2025, 2024.
- [64] Chen *et al.*, “Towards Human-Level Bimanual Dexterous Manipulation with Reinforcement Learning,” in *NeurIPS Datasets and Benchmarks*, 2022.
- [65] Zhang *et al.*, “GraspXL: Generating Grasping Motions for Diverse Objects at Scale,” in *ECCV*, 2024.
- [66] Wen *et al.*, “DexVLA: Vision-Language Model with Plug-In Diffusion Expert for General Robot Control,” arXiv:2502.05855, 2025.
- [67] Atar *et al.*, “In-Hand Manipulation of Articulated Tools with Dexterous Robot Hands with Sim-to-Real Transfer,” arXiv:2509.23075, 2025.
- [68] Lin *et al.*, “Learning Visuotactile Skills with Two Multifingered Hands,” in *ICRA* 2025, 2024.
- [69] Shadow Robot Company, “Shadow Dexterous Hand (technical specification, Dec 2024),” https://shadowrobot.com/wp-content/uploads/2024/12/shadow_dexterous_hand_e_technical_specification_20241204.pdf, 2005, vendor product page, shadowrobot.com, fetched 2026-09-17.
- [70] Inspire Robots, “Inspire Robots RH56DFX dexterous hand (RH56 series, FTP tactile variants),” <https://en.inspire-robots.com/product/rh56dfx/>, 2023, vendor product page, en.inspire-robots.com, fetched 2026-09-17.
- [71] Knight (The Robot Studio) and Trevor Blank, “DexHand: An Open Source Dexterous Humanoid Robot Hand,” <https://www.dexhand.org/>, 2023, open-hardware project page, dexhand.org, fetched 2026-09-17.
- [72] Daxo Robotics, “Daxo Robotics Muscle V0 (120-actuator compliant tendon hand),” <https://www.daxo-robotics.com/>, 2025, company announcement, hand not released, daxo-robotics.com, fetched 2026-09-17.
- [73] Catalano *et al.*, “Adaptive synergies for the design and control of the Pisa/IIT SoftHand,” *IJRR*, vol. 33, no. 5, 2014.
- [74] Kim *et al.*, “Integrated linkage-driven dexterous anthropomorphic robotic hand,” *Nature Communications*, vol. 12, 2021.
- [75] Toshimitsu *et al.*, “Getting the Ball Rolling: Learning a Dexterous Policy for a Biomimetic Tendon-Driven Hand with Rolling Contact Joints,” in *IEEE Humanoids*, 2023.
- [76] Liu *et al.*, “Ruka-v2: Tendon Driven Open-Source Dexterous Hand with Wrist and Abduction for Robot Learning,” arXiv:2603.26660, 2026.
- [77] Shaw *et al.*, “LEAP Hand: Low-Cost, Efficient, and Anthropomorphic Hand for Robot Learning,” in *RSS*, 2023.
- [78] Weng, “BiDexHand: Design and Evaluation of an Open-Source 16-DoF Biomimetic Dexterous Hand,” in *ICRA Dexterity Workshop*, 2025.
- [79] Christoph *et al.*, “ORCA: An Open-Source, Reliable, Cost-Effective, Anthropomorphic Robotic Hand for Uninterrupted Dexterous Task Learning,” in *IROS*, 2025.
- [80] Wonik Robotics, “Allegro Hand V4 (Wonik Robotics),” <https://www.allegrohand.com/v4/>, 2016, vendor product page, allegrohand.com, fetched 2026-09-17.
- [81] Zorin *et al.*, “RUKA: Rethinking the Design of Humanoid Hands with Learning,” in *RSS*, 2025.
- [82] Lambeta *et al.*, “DIGIT: A Novel Design for a Low-Cost Compact High-Resolution Tactile Sensor with Application to In-Hand Manipulation,” *IEEE RA-L*, 2020.
- [83] —, “Digitizing Touch with an Artificial Multimodal Fingertip,” arXiv:2411.02479, 2024.
- [84] XELA Robotics, “XELA uSkin magnetic 3-axis tactile skin (Allegro Hand fingertips),” [https://www.hannovermesse.de/apollo/hannover_messe_2022/obs/Binary/A1167836/XELA%20Robotics%20-%20uSkin%20Curved%20\(A4\).pdf](https://www.hannovermesse.de/apollo/hannover_messe_2022/obs/Binary/A1167836/XELA%20Robotics%20-%20uSkin%20Curved%20(A4).pdf), 2020, vendor product page, hannovermesse.de, fetched 2026-09-17.
- [85] Higuera *et al.*, “Sparsh: Self-supervised touch representations for vision-based tactile sensing,” in *CoRL*, 2024.
- [86] Liu *et al.*, “DexTeleop-0: Force-Aware Bimanual Dexterous Teleoperation with Ego-Centric Perception towards Shared Autonomy,” arXiv:2606.23431, 2026.
- [87] Xu *et al.*, “DexUMI: Using Human Hand as the Universal Manipulation Interface for Dexterous Manipulation,” in *CoRL*, 2025.
- [88] Yuan *et al.*, “Robot Synesthesia: In-Hand Manipulation with Visuotactile Sensing,” in *ICRA* 2024, 2023.
- [89] Qi *et al.*, “General In-Hand Object Rotation with Vision and Touch,” in *CoRL*, 2023.
- [90] Wang *et al.*, “Lessons from Learning to Spin ‘Pens,’” in *CoRL*, 2024.

- [91] Liu *et al.*, “DexNDM: Closing the Reality Gap for Dexterous In-Hand Rotation via Joint-Wise Neural Dynamics Model,” arXiv:2510.08556, 2025.
- [92] Xu *et al.*, “Dexplore: Scalable Neural Control for Dexterous Manipulation from Reference-Scoped Exploration,” in *CoRL*, 2025.
- [93] 1X Technologies, “NEO’s Hands: An API to the Physical World (1X tendon-driven hand),” <https://www.1x.tech/discover/neos-hands>, 2026, company announcement, 1x.tech, fetched 2026-09-17.
- [94] Tesla, “Tesla Optimus hands: Gen 2 (11 DoF, in-hand actuators, 2023) and Gen 3 / V3 (22 DoF, forearm-actuated tendons, patents 2025-2026),” <https://image-ppubs.uspto.gov/dirsearch-public/print/downloadPdf/9505134>, 2025, company announcement and USPTO patent filings, hand not released, image-ppubs.uspto.gov, fetched 2026-09-17.
- [95] Clone Robotics, “Clone Robotics hydraulic Myofiber hand (27 DoF),” <https://www.theresarobotforthat.com/clones-robot-hand-has-artificial-muscles/>, 2024, company announcement, hand not released, thesarobotforthat.com, fetched 2026-09-17.
- [96] Figure AI, “Figure 03 hands (palm camera, 3 g fingertip sensitivity),” <https://www.figure.ai/news/introducing-figure-03>, 2025, company announcement, hand not released, figure.ai, fetched 2026-09-17.
- [97] Sharpa Robotics, “SharpaWave 22-DoF dexterous hand with Dynamic Tactile Array,” <https://www.sharpa.com/pages/wave>, 2026, vendor product page, sharpa.com, fetched 2026-09-17.
- [98] Wuji Technology, “Wuji Hand / Wuji Hand 2: 20-DoF direct-drive dexterous hand,” <https://docs.wuji.tech/docs/en/wuji-hand/latest/overview/>, 2025, vendor product page, docs.wuji.tech, fetched 2026-09-17.
- [99] Figure AI, “Helix: A Vision-Language-Action Model for Generalist Humanoid Control,” <https://www.figure.ai/news/helix>, 2025, company blog post, February 2025, figure.ai, fetched 2026-09-17.
- [100] Wen *et al.*, “GR-Dexter Technical Report,” arXiv:2512.24210, 2025.
- [101] Unitree Robotics, “Unitree Dex5-1 / Dex5-1P dexterous hand,” <https://www.unitree.com/Dex5-1>, 2025, vendor product page, unitree.com, fetched 2026-09-17.
- [102] TESOLLO, “TESOLLO DG-5F / DG-5F-M / DG-5F-S humanoid robot hand,” <https://en.tesollo.com/wp-content/uploads/2025/02/TESOLLO-DG-5F-catalogue-En-2.pdf>, 2024, vendor product page, en.tesollo.com, fetched 2026-09-17.
- [103] Proception (YC), “Proception ProHand 1.0 (22-DoF tendon-driven hand),” <https://www.ycombinator.com/launches/R78-proception-the-dexterous-robot-hand-we-wish-existed>, 2026, company announcement, ycombinator.com, fetched 2026-09-17.
- [104] Yang *et al.*, “Bench2Dex: Benchmarking Visuo-Tactile Bimanual Dexterous Manipulation Across Dexterous Hands,” arXiv:2609.15726, 2026.
- [105] Yao *et al.*, “DexVerse: A Modular Benchmark for Multi-Task, Multi-Embodiment Dexterous Manipulation,” arXiv:2607.08751, 2026.
- [106] Zakka *et al.*, “MuJoCo Playground,” arXiv:2502.08844, 2025.
- [107] Makovychuk *et al.*, “Isaac Gym: High Performance GPU-Based Physics Simulation For Robot Learning,” in *NeurIPS Datasets and Benchmarks*, 2021.
- [108] Castro *et al.*, “An Unconstrained Convex Formulation of Compliant Contact,” *IEEE T-RO* 2022, 2021.
- [109] Todorov, “Convex and analytically-invertible dynamics with contacts and constraints: Theory and implementation in MuJoCo,” in *ICRA*, 2014.
- [110] Todorov *et al.*, “MuJoCo: A physics engine for model-based control,” in *IROS*, 2012.
- [111] Howell *et al.*, “Dojo: A Differentiable Physics Engine for Robotics,” arXiv:2203.00806, 2022.
- [112] Borse *et al.*, “ComFree-Sim: A GPU-Parallelized Analytical Contact Physics Engine for Scalable Contact-Rich Robotics Simulation and Control,” arXiv:2603.12185, 2026.
- [113] Chung *et al.*, “Tactile Genesis: Exploring Tactile Sensors at Scale for Learning Dexterous Tasks,” arXiv:2606.22332, 2026.
- [114] Liu *et al.*, “DexTrack: Towards Generalizable Neural Tracking Control for Dexterous Manipulation from Human References,” in *ICLR*, 2025.
- [115] Mittal *et al.*, “Isaac Lab: A GPU-Accelerated Simulation Framework for Multi-Modal Robot Learning,” arXiv:2511.04831, 2025.
- [116] Google DeepMind and NVIDIA, “MuJoCo Warp: GPU-optimized version of the MuJoCo physics simulator,” https://github.com/google-deepmind/mujoco_warp, 2025, software release announcement, NVIDIA GTC 2025, github.com, no fetch date recorded.
- [117] Genesis Authors (Zhou *et al.*), “Genesis: A Generative and Universal Physics Engine for Robotics and Beyond,” <https://github.com/Genesis-Embodied-AI/Genesis>, 2024, software release page, github.com, no fetch date recorded.
- [118] NVIDIA, Google DeepMind, and Disney Research, “Newton: An open-source, GPU-accelerated physics simulation engine built upon NVIDIA Warp,” <https://github.com/newton-physics/newton>, 2025, software release announcement, Linux Foundation, github.com, no fetch date recorded.
- [119] Mittal *et al.*, “Orbit: A Unified Simulation Framework for Interactive Robot Learning Environments,” *IEEE RA-L*, 2023.
- [120] Tao *et al.*, “ManiSkill3: GPU Parallelized Robotics Simulation and Rendering for Generalizable Embodied AI,” in *RSS* 2025, 2024.
- [121] Gu *et al.*, “ManiSkill2: A Unified Benchmark for Generalizable Manipulation Skills,” in *ICLR*, 2023.
- [122] Wang *et al.*, “TACTO: A Fast, Flexible, and Open-source Simulator for High-Resolution Vision-based Tactile Sensors,” *IEEE RA-L* 2022, 2020.
- [123] Si and Yuan, “Taxim: An Example-based Simulation Model for GelSight Tactile Sensors,” *IEEE RA-L* 2022, 2021.
- [124] Lin *et al.*, “Sim-to-Real Reinforcement Learning for Vision-Based Dexterous Manipulation on Humanoids,” in *CoRL*, 2025.
- [125] Nagabandi *et al.*, “Deep Dynamics Models for Learning Dexterous Manipulation,” in *CoRL*, 2019.
- [126] OpenAI *et al.*, “Solving Rubik’s Cube with a Robot Hand,” arXiv:1910.07113, 2019.
- [127] Petrenko *et al.*, “DexPBT: Scaling up Dexterous Manipulation for Hand-Arm Systems with Population Based Training,” in *RSS*, 2023.
- [128] Lin *et al.*, “Learning In-Hand Object Reaching to General 6D Poses,” arXiv:2609.13761, 2026.
- [129] Zhao *et al.*, “Closing the Reality Gap: Zero-Shot Sim-to-Real Deployment for Dexterous Force-Based Grasping and Manipulation,” arXiv:2601.02778, 2026.
- [130] Li *et al.*, “Towards Human-level Dexterous Teleoperation,” arXiv:2607.11481, 2026.
- [131] Bauza *et al.*, “DemoStart: Demonstration-led auto-curriculum applied to sim-to-real with multi-fingered robots,” arXiv:2409.06613, 2024.
- [132] Bhardwaj *et al.*, “ViserDex: Visual Sim-to-Real for Robust Dexterous In-hand Reorientation,” arXiv:2604.11138, 2026.
- [133] Wan *et al.*, “DexReMoE: In-hand Reorientation of General Object via Mixtures of Experts,” arXiv:2508.01695, 2025.
- [134] Ma *et al.*, “DrEureka: Language Model Guided Sim-To-Real Transfer,” in *RSS*, 2024.
- [135] Handa *et al.*, “DexPilot: Vision Based Teleoperation of Dexterous Robotic Hand-Arm System,” in *ICRA*, 2020.
- [136] Qin *et al.*, “AnyTeleop: A General Vision-Based Dexterous Robot Arm-Hand Teleoperation System,” in *RSS*, 2023.
- [137] Cheng *et al.*, “Open-TeleVision: Teleoperation with Immersive Active Visual Feedback,” in *CoRL*, 2024.
- [138] Yang *et al.*, “ACE: A Cross-Platform Visual-Exoskeletons System for Low-Cost Dexterous Teleoperation,” in *CoRL*, 2024.
- [139] Arunachalam *et al.*, “Holo-Dex: Teaching Dexterity with Immersive Mixed Reality,” in *ICRA* 2023, 2022.
- [140] Yin *et al.*, “Geometric Retargeting: A Principled, Ultrafast Neural Hand Retargeting Algorithm,” arXiv:2503.07541, 2025.
- [141] Zhang *et al.*, “DOGlove: Dexterous Manipulation with a Low-Cost Open-Source Haptic Force Feedback Glove,” in *RSS*, 2025.
- [142] Wang *et al.*, “DexCap: Scalable and Portable Mocap Data Collection System for Dexterous Manipulation,” in *RSS*, 2024.
- [143] Qiu *et al.*, “Humanoid Policy ~ Human Policy,” in *CoRL*, 2025.
- [144] Zhao *et al.*, “Learning Fine-Grained Bimanual Manipulation with Low-Cost Hardware,” in *RSS*, 2023.
- [145] Chi *et al.*, “Diffusion Policy: Visuomotor Policy Learning via Action Diffusion,” in *RSS*, 2023.
- [146] Ze *et al.*, “3D Diffusion Policy: Generalizable Visuomotor Policy Learning via Simple 3D Representations,” in *RSS*, 2024.
- [147] Fu *et al.*, “METIS: Multi-Source Egocentric Training for Integrated Dexterous Vision-Language-Action Model,” arXiv:2511.17366, 2025.
- [148] Ye *et al.*, “Dex1B: Learning with 1B Demonstrations for Dexterous Manipulation,” in *RSS*, 2025.
- [149] Qin *et al.*, “DexMV: Imitation Learning for Dexterous Manipulation from Human Videos,” in *ECCV* 2022, 2021.
- [150] Shaw *et al.*, “VideoDex: Learning Dexterity from Internet Videos,” in *CoRL*, 2022.
- [151] Mandikal and Grauman, “DexVIP: Learning Dexterous Grasping with Human Hand Pose Priors from Video,” in *CoRL* 2021, 2022.
- [152] Li *et al.*, “OKAMI: Teaching Humanoid Robots Manipulation Skills through Single Video Imitation,” in *CoRL*, 2024.
- [153] Lum *et al.*, “Crossing the Human-Robot Embodiment Gap with Sim-to-Real RL using One Human Demonstration,” arXiv:2504.12609, 2025.
- [154] Guzey *et al.*, “Bridging the Human to Robot Dexterity Gap through Object-Oriented Rewards,” in *ICRA* 2025, 2024.

- [155] Goswami *et al.*, “World Models for Learning Dexterous Hand-Object Interactions from Human Videos,” arXiv:2512.13644, 2025.
- [156] Liu *et al.*, “EgoZero: Robot Learning from Smart Glasses,” in *CoRL*, 2025.
- [157] Chen *et al.*, “Object-Centric Dexterous Manipulation from Human Motion Data,” in *CoRL*, 2024.
- [158] Hsieh *et al.*, “DexMan: Learning Bimanual Dexterous Manipulation from Human and Generated Videos,” arXiv:2510.08475, 2025.
- [159] Zhang *et al.*, “Dexora: Open-source VLA for High-DoF Bimanual Dexterity,” arXiv:2605.18722, 2026.
- [160] Zheng *et al.*, “EgoScale: Scaling Dexterous Manipulation with Diverse Egocentric Human Data,” arXiv:2602.16710, 2026.
- [161] Gemini Robotics Team, “Gemini Robotics: Bringing AI into the Physical World,” arXiv:2503.20020, 2025.
- [162] NVIDIA, “GR00T N1: An Open Foundation Model for Generalist Humanoid Robots,” arXiv:2503.14734, 2025.
- [163] Howell *et al.*, “Predictive Sampling: Real-time Behaviour Synthesis with MuJoCo,” arXiv:2212.00541, 2022.
- [164] Yin *et al.*, “DexterityGen: Foundation Controller for Unprecedented Dexterity,” in *RSS*, 2025.
- [165] Mu *et al.*, “DexImit: Learning Bimanual Dexterous Manipulation from Monocular Human Videos,” arXiv:2602.10105, 2026.
- [166] Huang *et al.*, “Efficient Residual Learning with Mixture-of-Experts for Universal Dexterous Grasping,” in *ICLR 2025*, 2024.
- [167] Xu *et al.*, “UniDexGrasp: Universal Robotic Dexterous Grasping via Learning Diverse Proposal Generation and Goal-Conditioned Policy,” in *CVPR*, 2023.
- [168] Liu *et al.*, “RDT-1B: a Diffusion Foundation Model for Bimanual Manipulation,” in *ICLR 2025*, 2024.
- [169] Kareer *et al.*, “EgoMimic: Scaling Imitation Learning via Egocentric Video,” in *ICRA 2025*, 2024.
- [170] Bi *et al.*, “H-RDT: Human Manipulation Enhanced Bimanual Robotic Manipulation,” arXiv:2507.23523, 2025.
- [171] Chi *et al.*, “Universal Manipulation Interface: In-The-Wild Robot Teaching Without In-The-Wild Robots,” in *RSS*, 2024.
- [172] Jing *et al.*, “HumanoidGen: Data Generation for Bimanual Dexterous Manipulation via LLM Reasoning,” arXiv:2507.00833, 2025.
- [173] Zhang *et al.*, “ArtiGrasp: Physically Plausible Synthesis of Bi-Manual Dexterous Grasping and Articulation,” in *3DV 2024*, 2023.
- [174] Zhou *et al.*, “Learning Diverse Bimanual Dexterous Manipulation Skills from Human Demonstrations,” in *AAAI 2026 (oral)*, 2024.
- [175] Lin *et al.*, “Twisting Lids Off with Two Hands,” in *CoRL*, 2024.
- [176] Liu *et al.*, “TACO: Benchmarking Generalizable Bimanual Tool-ACTION-Object Understanding,” in *CVPR*, 2024.
- [177] Tao *et al.*, “DexWild: Dexterous Human Interactions for In-the-Wild Robot Policies,” in *RSS*, 2025.
- [178] Fan *et al.*, “ARCTIC: A Dataset for Dexterous Bimanual Hand-Object Manipulation,” in *CVPR 2023*, 2022.
- [179] Zhan *et al.*, “OAKINK2: A Dataset of Bimanual Hands-Object Manipulation in Complex Task Completion,” in *CVPR*, 2024.
- [180] Banerjee *et al.*, “HOT3D: Hand and Object Tracking in 3D from Egocentric Multi-View Videos,” in *CVPR 2025*, 2024.
- [181] Fu *et al.*, “GigaHands: A Massive Annotated Dataset of Bimanual Hand Activities,” in *CVPR 2025 (highlight)*, 2024.
- [182] Pumacay *et al.*, “THE COLOSSEUM: A Benchmark for Evaluating Generalization for Robotic Manipulation,” in *RSS*, 2024.
- [183] Snyder *et al.*, “Beyond Binary Success: Sample-Efficient and Statistically Rigorous Robot Policy Comparison,” arXiv:2603.13616, 2026.
- [184] Kress-Gazit *et al.*, “Robot Learning as an Empirical Science: Best Practices for Policy Evaluation,” arXiv:2409.09491, 2024.
- [185] Zhou *et al.*, “AutoEval: Autonomous Evaluation of Generalist Robot Manipulation Policies in the Real World,” arXiv:2503.24278, 2025.
- [186] Badithela *et al.*, “Reliable and Scalable Robot Policy Evaluation with Imperfect Simulators,” arXiv:2510.04354, 2025.
- [187] Chen *et al.*, “ClutterDexGrasp: A Sim-to-Real System for General Dexterous Grasping in Cluttered Scenes,” in *CoRL*, 2025.
- [188] TRI LBM Team, “A Careful Examination of Large Behavior Models for Multitask Dexterous Manipulation,” arXiv:2507.05331, 2025.
- [189] Atreya *et al.*, “RoboArena: Distributed Real-World Evaluation of Generalist Robot Policies,” arXiv:2506.18123, 2025.
- [190] Li *et al.*, “Evaluating Real-World Robot Manipulation Policies in Simulation,” in *CoRL*, 2024.
- [191] ROBOTERA, “ROBOTERA XHAND1 dexterous hand,” <https://humanoid.guide/product/xhand1/>, 2024, vendor product page, humanoid.guide, fetched 2026-09-17.
- [192] Shadow Robot Company and Google DeepMind, “DEX-EE: Shadow Robot hand developed with Google DeepMind,” https://shadowrobot.com/dex-ee_series/, 2024, vendor product page, shadowrobot.com, fetched 2026-09-17.
- [193] BrainCo, “BrainCo Revo 2 dexterous hand,” <https://www.brainco-hz.com/docs/revolimb-hand/en/revo2/parameters.html>, 2025, vendor product page, brainco-hz.com, fetched 2026-09-17.
- [194] AgiBot (Zhiyuan), “AgiBot OmniHand 2025 / OmniHand Pro 2025,” <https://store.agibot.com/products/omnihand-2025>, 2025, vendor product page, store.agibot.com, fetched 2026-09-17.
- [195] LinkerBot, “LinkerBot Linker Hand L20 (and O6),” <https://openlab.io/products/linker-hand-linkerbot-l20-21>, 2025, vendor product page, openlab.io, fetched 2026-09-17.
- [196] PSYONIC, “PSYONIC Ability Hand,” <https://www.psyonic.io/robots>, 2021, vendor product page, psyonic.io, fetched 2026-09-17.
- [197] Shaw and Pathak, “LEAP Hand V2 Advanced: Dexterous, Low-Cost Hybrid Rigid-Soft Hand for Robot Learning,” in *IEEE Humanoids*, 2025.
- [198] Sanctuary AI, “Sanctuary AI Phoenix hydraulic hand (21 DoF),” <https://sanctuary.ai/news/sanctuary-ai-demonstrates-in-hand-manipulation-capabilities-for-improved-general-purpose-robot-dexterity/>, 2024, company announcement, hand not released, sanctuary.ai, fetched 2026-09-17.
- [199] Boston Dynamics, “Boston Dynamics electric Atlas hands (7-DoF three-finger, production four-digit gripper),” <https://bostondynamics.com/blog/boston-dynamics-unveils-new-atlas-robot-to-revolutionize-industry/>, 2026, company announcement, hand not released, bostondynamics.com, fetched 2026-09-17.
- [200] PaXini Tech, “PaXini DexH13 tactile dexterous hand and multi-dimensional tactile sensors,” <https://humanoid.guide/product/dexh13/>, 2024, vendor product page, humanoid.guide, fetched 2026-09-17.
- [201] Xiaomi, “Xiaomi next-gen CyberOne bionic hand (22-27 DoF, liquid-cooled forearm),” <https://eu.36kr.com/en/p/3947315943472512>, 2026, company announcement, hand not released, eu.36kr.com, fetched 2026-09-17.
- [202] Freeman *et al.*, “Brax - A Differentiable Physics Engine for Large Scale Rigid Body Simulation,” in *NeurIPS Datasets and Benchmarks*, 2021.
- [203] Coumans and Bai, “PyBullet, a Python module for physics simulation for games, robotics and machine learning,” <https://github.com/bulletphysics/bullet3>, 2016, software documentation page, github.com, no fetch date recorded.
- [204] Hwangbo *et al.*, “Per-Contact Iteration Method for Solving Contact Dynamics,” *IEEE RA-L*, 2018, full text not obtained; cited by metadata only.
- [205] Xiang *et al.*, “SAPIEN: A SimulATED Part-based Interactive ENvironment,” in *CVPR*, 2020.